
AWS Summit
AWS技术峰会2015 · 上海
亚马逊
aws

Intro to Amazon EC2 Container Service

Aaron Kao

Sr. Product Marketing Manager

Agenda

Why Containers?

- Container Use Patterns
- Production Challenges

Cluster Management

Amazon EC2 Container Service

Demo

What Containers?

What are Containers?

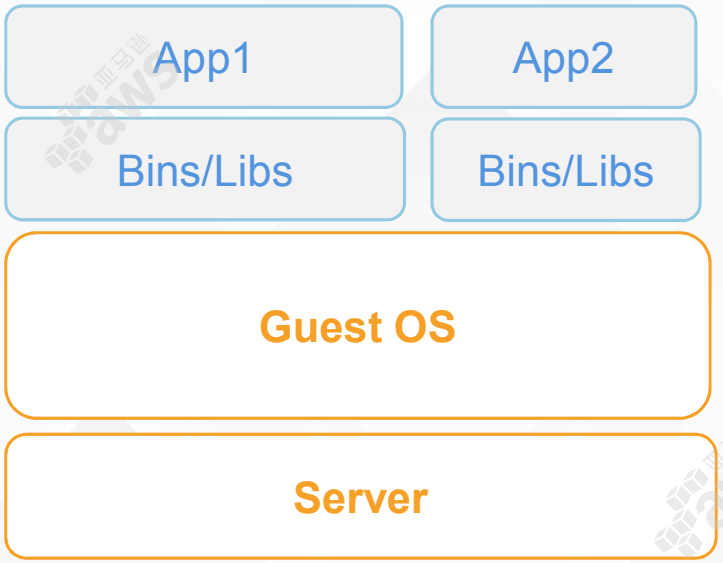

OS virtualization

Process isolation

Images

Automation

Container advantages

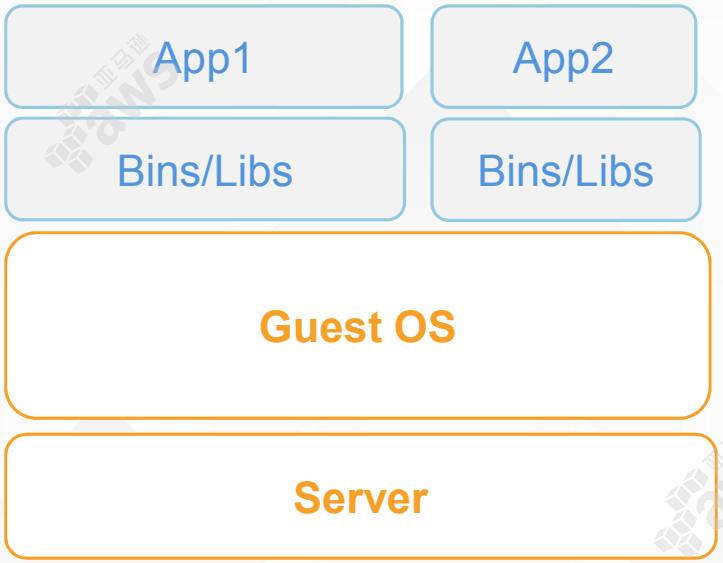

Portable

Flexible

Fast

Efficient

Services evolve to microservices

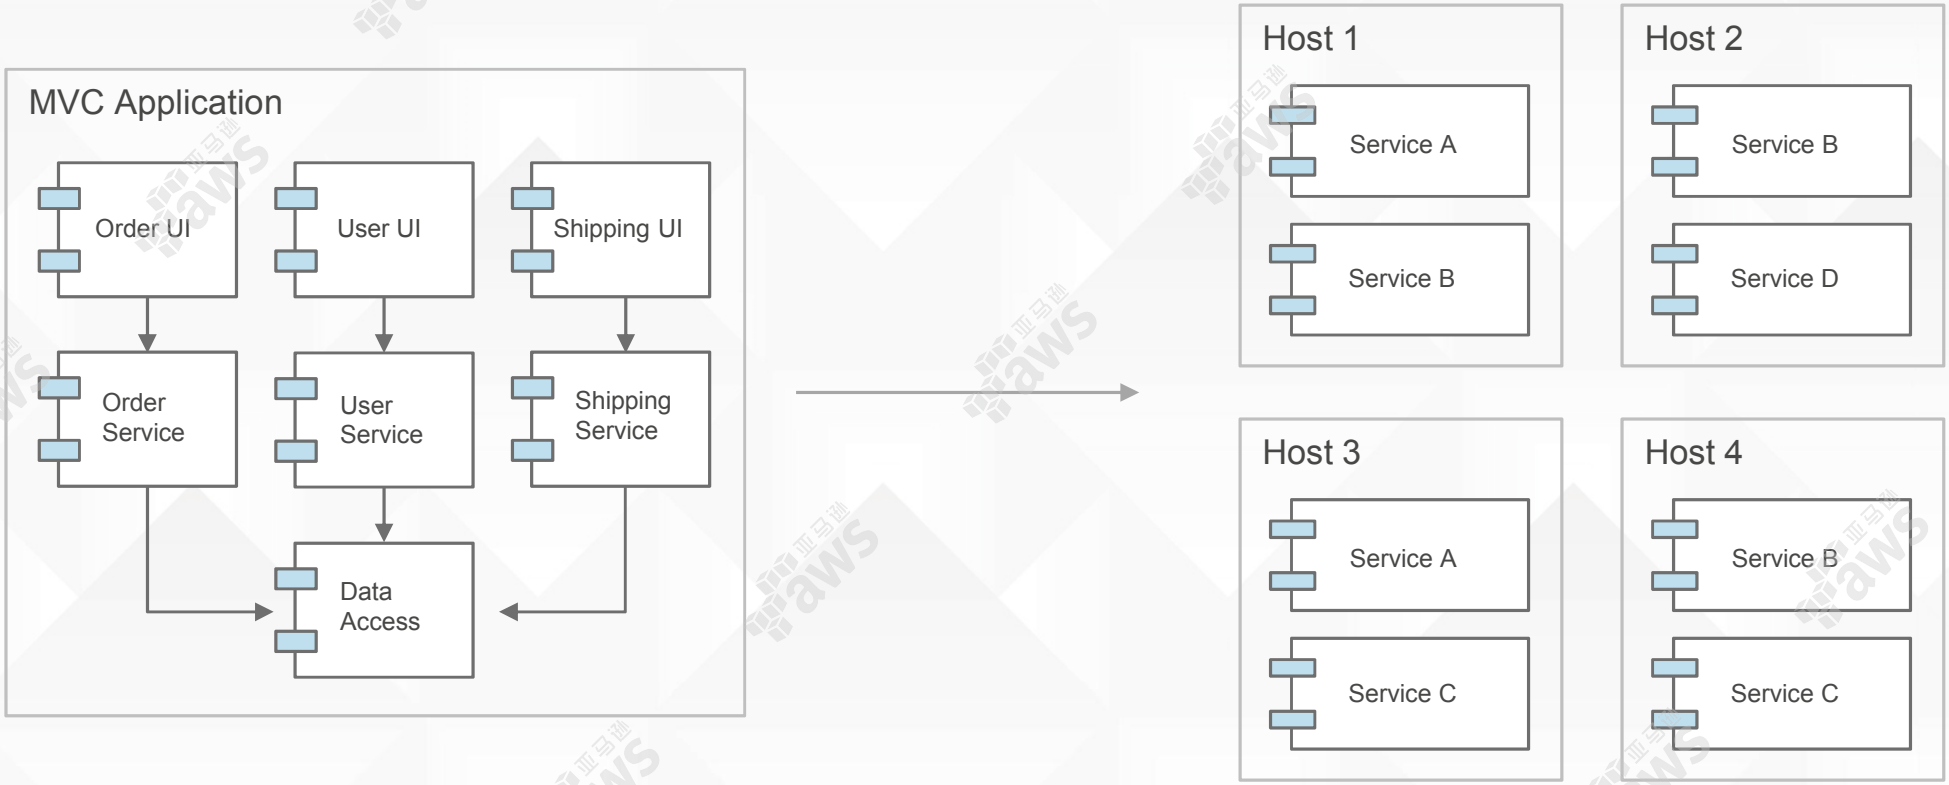

Containers are natural for services

Simple to model

Any app, any language

Image is the version

Test & deploy same artifact

Stateless servers decrease change risk

Scheduling


7a
8a
9a
10a
Pick up dry cleaning
11a
12p
Eat lunch
1p
2p
Bike to the store
3p
4p
Relax
5p
6p
7p

Scheduling one resource is straightforward

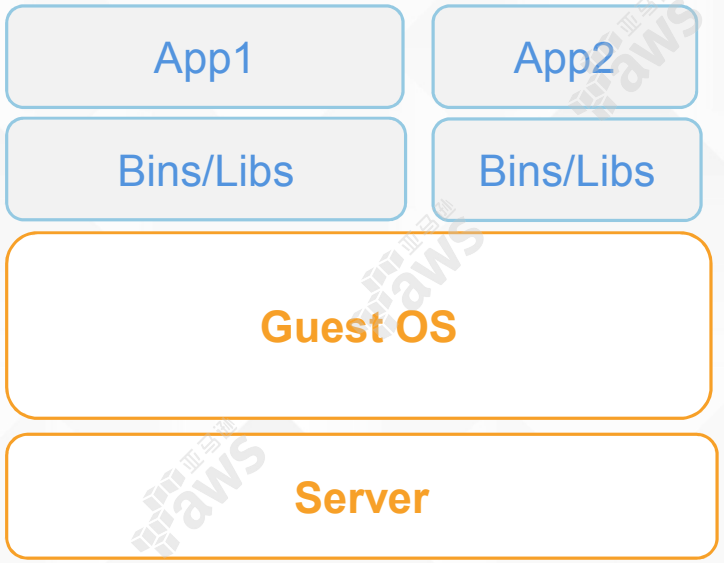

Scheduling a cluster is hard

Scheduling 101

Know your constraints

Find resources that meet the constraints

Request a resource

Confirm the resource

Cluster Management

Cluster Management

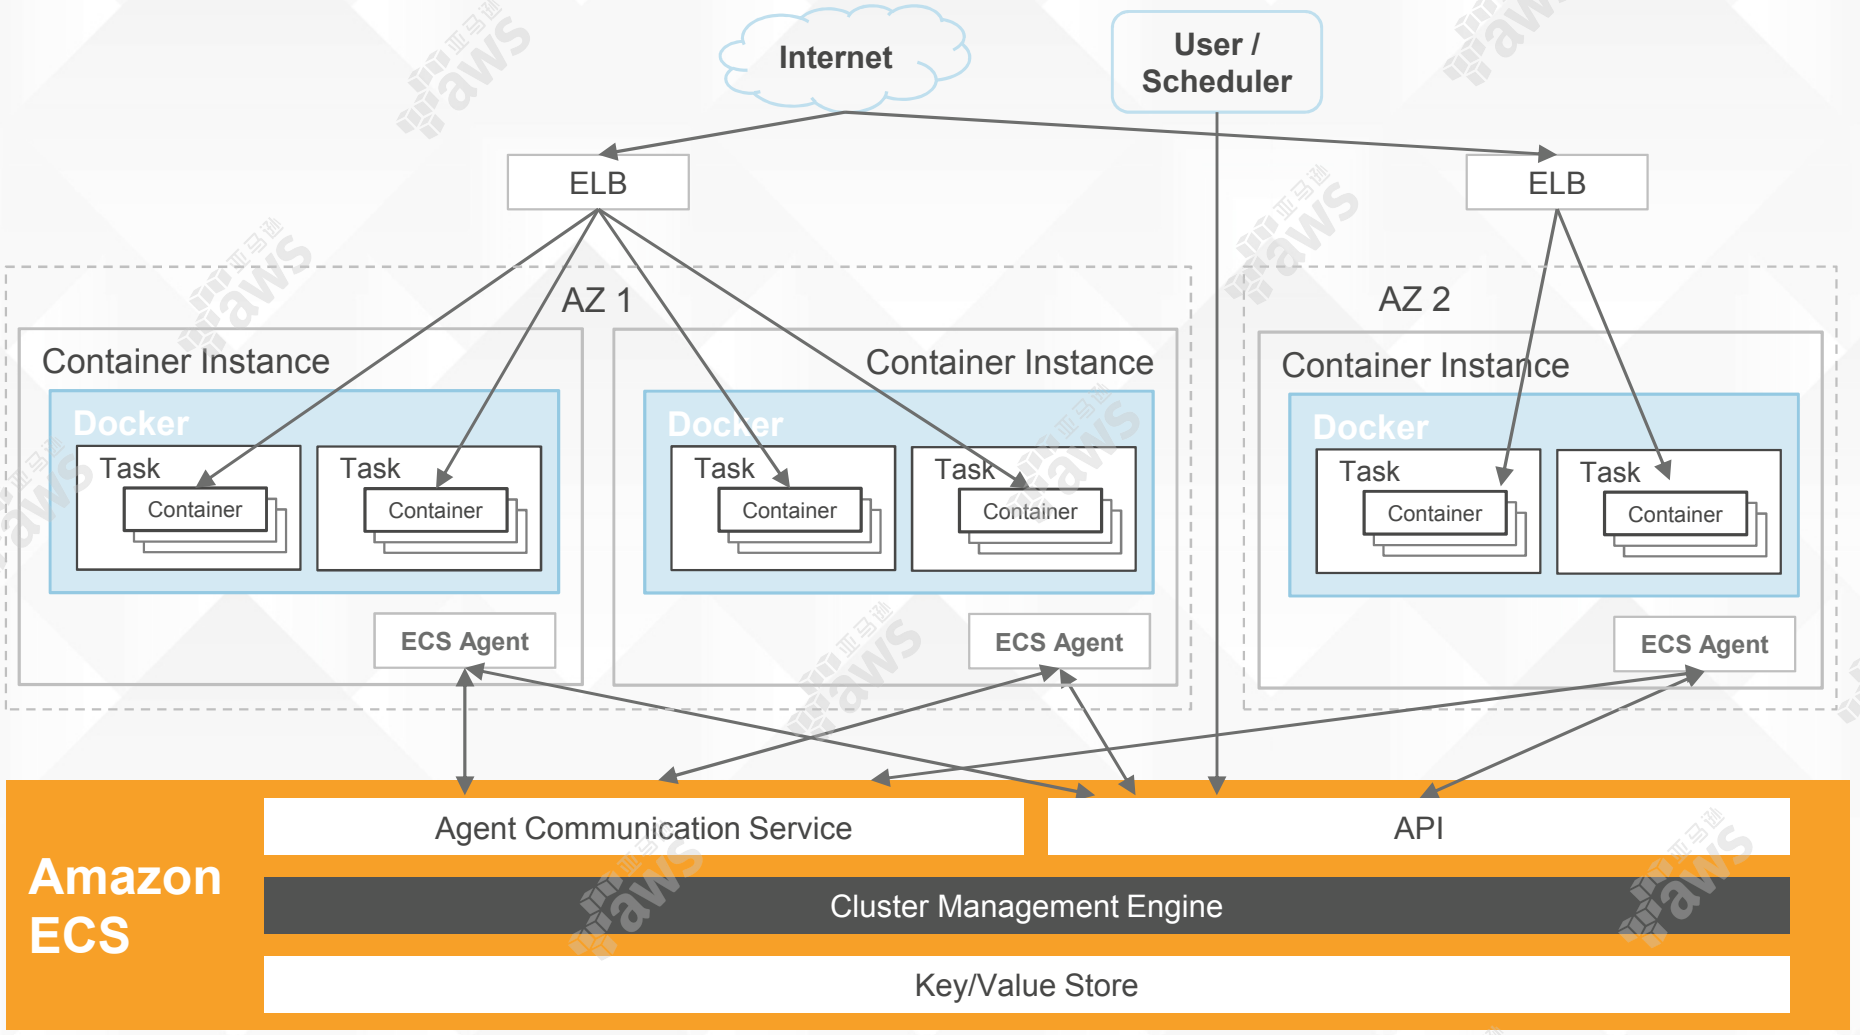

Cluster Management with Amazon ECS

Management of followers via ECS Agent

Dispatching of sub-tasks to proper location

Cluster state inspection

http://amzn.to/1jlHvnU

Cluster Management under the Hood

Paxos-based transactional journal based data store

Writes are committed as transaction in the journal with order-based ID. The current value is the sum of all transactions made as recorded by the journal.

Reads are simply a snapshot in time of the journal. For a write to succeed, the write proposed must be the latest transaction since the last read.

http://bit.ly/1M9gGiv

Amazon EC2 Container Service

Easily Manage Clusters for Any Scale

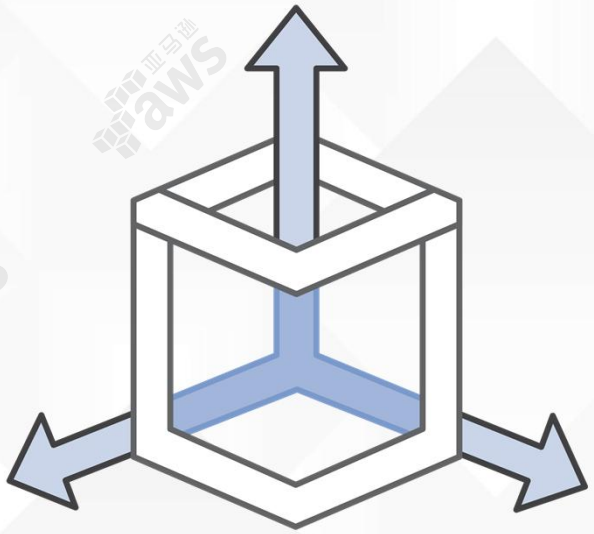

Nothing to run

Complete state

Control and monitoring

Scale

Scalable

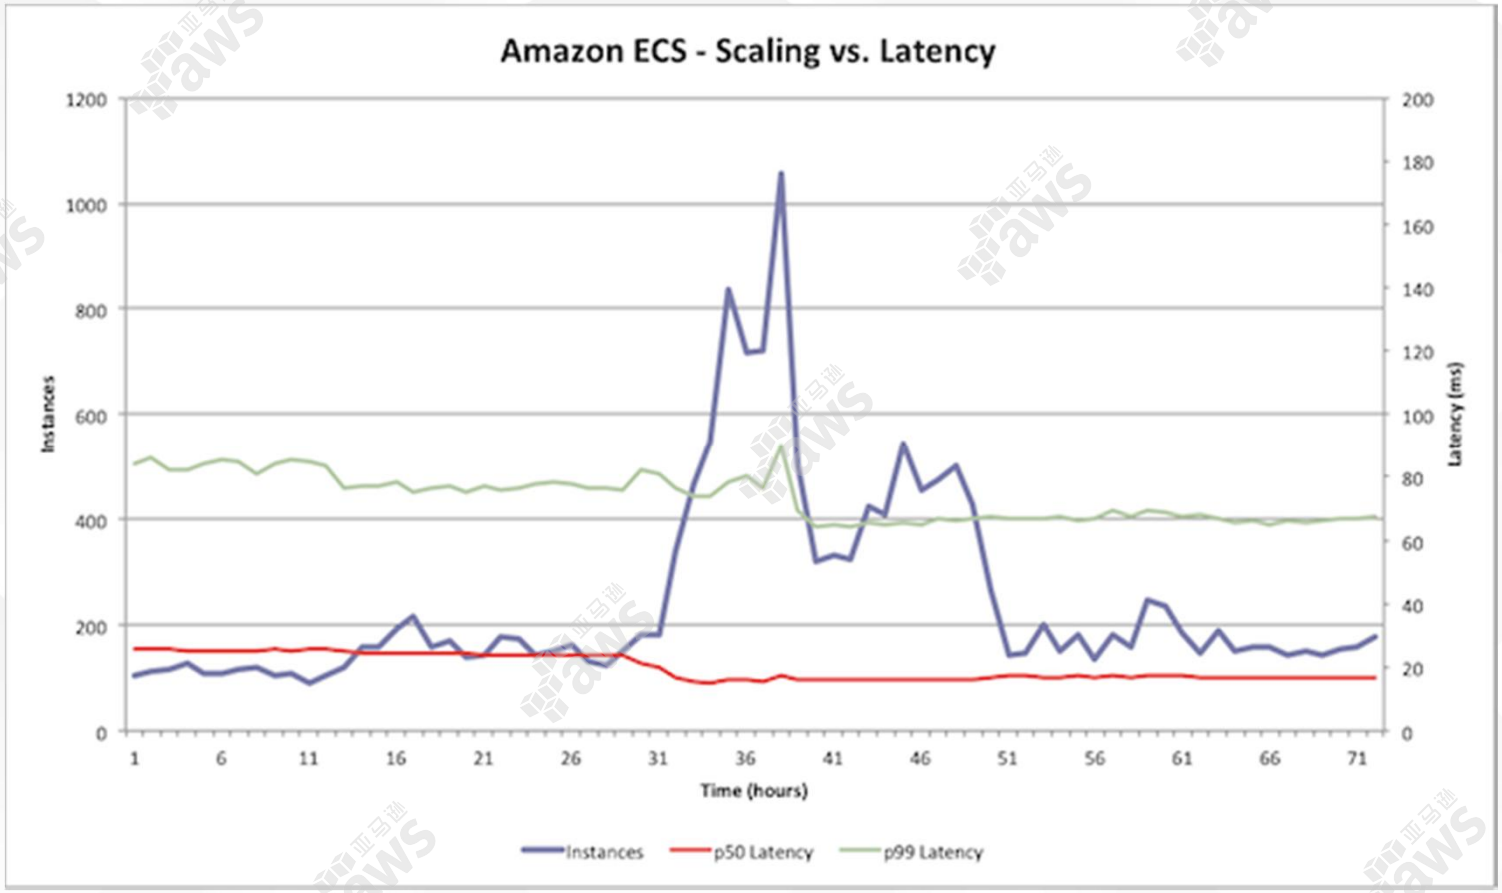
Amazon ECS - Scaling vs. Latency
Instances
Latency (ms)
Time (hours)
1200
1000
800
600
400
200
0
200
180
160
140
120
100
80
60
40
20
0
1
6
11
16
21
26
31
36
41
46
51
56
61
66
71
Instances
p50 Latency
p99 Latency

Flexible Container Placement

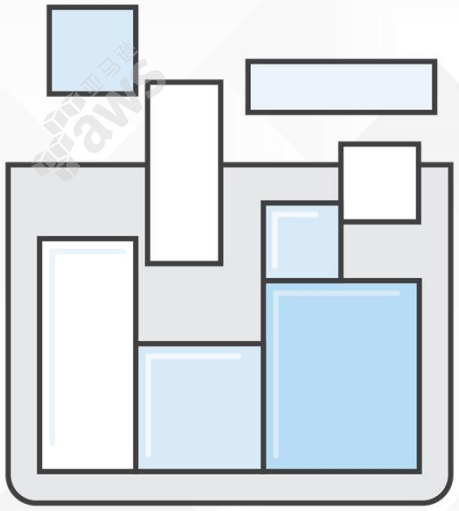

Applications

Batch jobs

Multiple schedulers

Designed for use with other AWS services

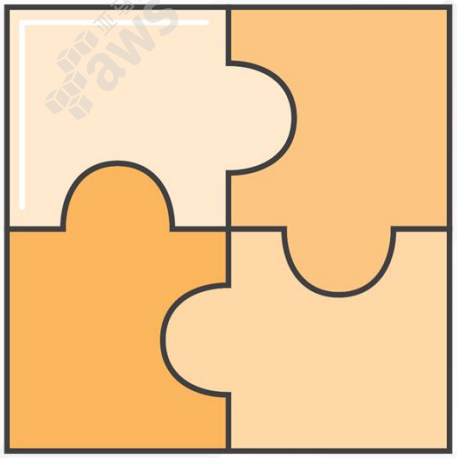

- Elastic Load Balancing
- Amazon Elastic Block Store
- Amazon Virtual Private Cloud
- AWS Identity and Access Management
- AWS CloudTrail

Extensible

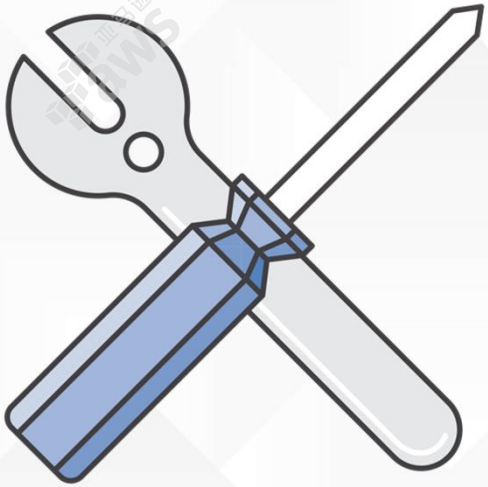

Comprehensive APIs

Open source agent

Custom schedulers

Key Components: Container Instances

Amazon EC2 instances

Docker daemon

Amazon ECS agent

Key Components: Clusters

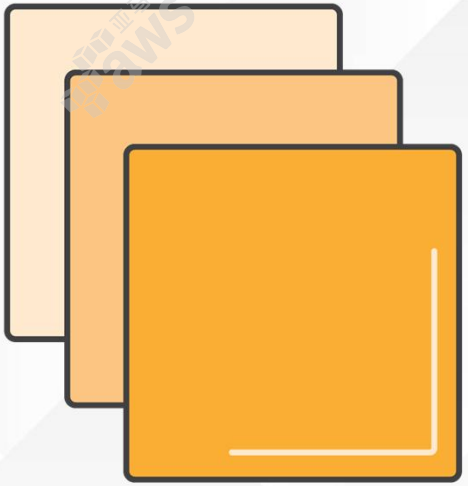

Regional

Resource pool

Grouping of Container Instances

Start empty, dynamically scalable

Key Components: Task Definitions

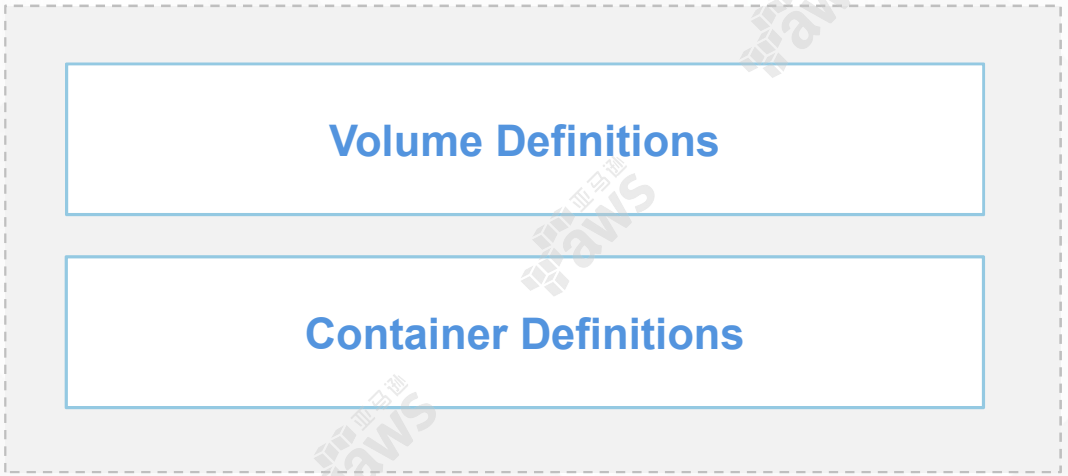

Key Components: Task Definitions

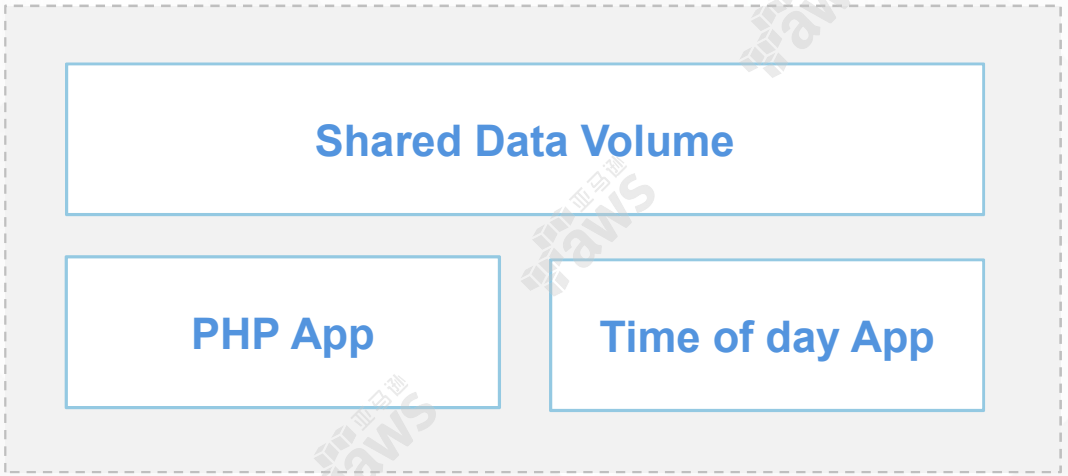

Key Components: Task Definitions

```
{
            "environment": [],
            "name": "simple-demo",
            "image": "my-demo",
            "cpu": 10,
            "memory": 500,
            "portMappings": [
                {
                    "containerPort": 80,
                    "hostPort": 80
                }
            ],
            "mountPoints": [
                {
                    "sourceVolume": "my-vol",
                    "containerPath": "/var/www/my-vol"
                }
            ],
            "entryPoint": [
                "/usr/sbin/apache2",
                "-D",
                "FOREGROUND"
            ],
            "essential": true
        },
```

```
{
            "name": "busybox",
            "image": "busybox",
            "cpu": 10,
            "memory": 500,
            "volumesFrom": [
            {
              "sourceContainer": "simple-demo"
            }
            ],
            "entryPoint": [
                "sh",
                "-c"
            ],
            "command": [
                "/bin/sh -c \"while true; do /bin/date >
/var/www/my-vol/date; sleep 1; done\""
            ],
            "essential": false
        }
```

Key Components: Task Definitions

```
{
        "environment": [],
        "name": "simple-demo",
        "image": "my-demo",
        "cpu": 10,
        "memory": 500,
        "portMappings": [
            {
                "containerPort": 80,
                "hostPort": 80
            }
        ],
        "mountPoints": [
            {
                "sourceVolume": "my-vol",
                "containerPath": "/var/www/my-vol"
            }
        ],
        "entryPoint": [
            "/usr/sbin/apache2",
            "-D",
            "FOREGROUND"
        ],
        "essential": true
    },
```

10 CPU Units (1024 is full CPU), 500 Megabytes of Memory

Expose port 80 in container to port 80 on host

Create and mount volumes

Essential to our Task

Key Components: Task Definitions

From Docker Hub

Mount volume from other container

Command to exec

```
{
            "name": "busybox",
            "image": "busybox",
            "cpu": 10,
            "memory": 500,
            "volumesFrom": [
            {
              "sourceContainer": "simple-demo"
            }
            ],
            "entryPoint": [
                "sh",
                "-c"
            ],
            "command": [
                "/bin/sh -c \"while true; do /bin/date >
/var/www/my-vol/date; sleep 1; done\""
            ],
            "essential": false
        }
```

Key Components: Tasks

Unit of work

Grouping of related Containers

Run on Container Instances

Key Components: Tasks

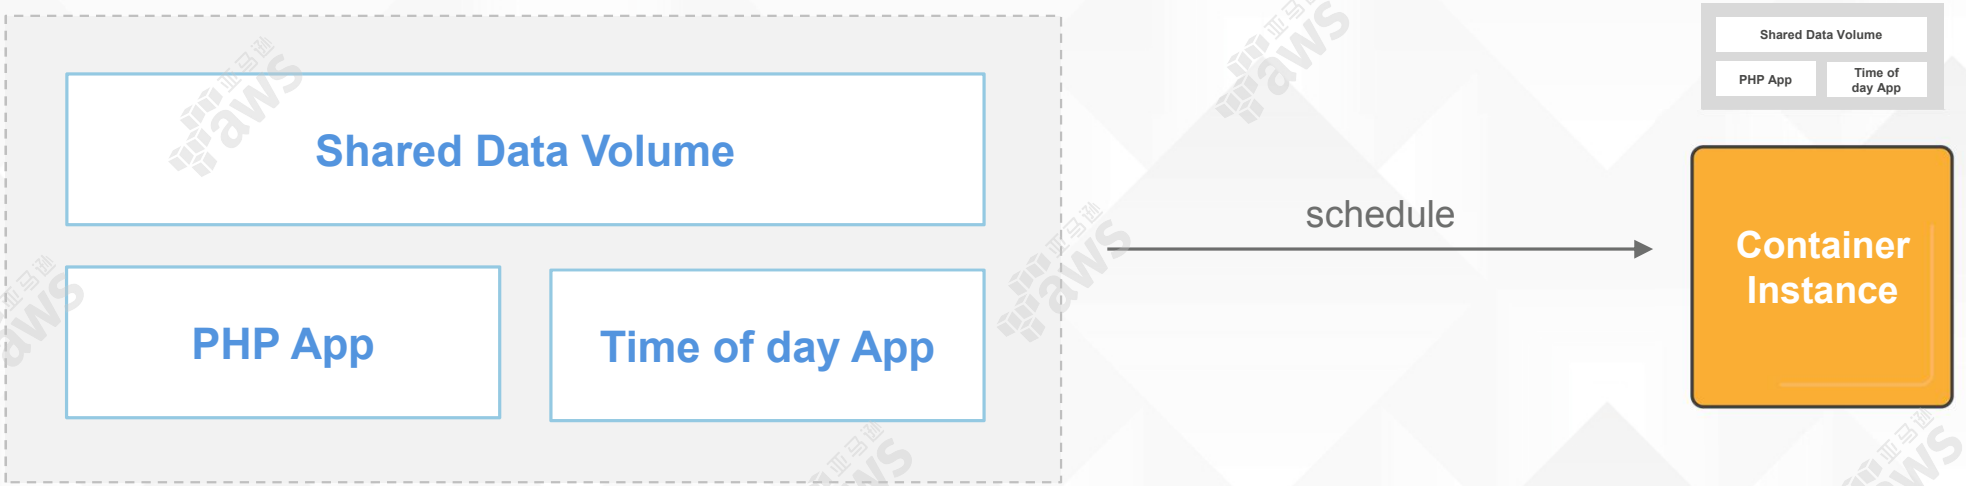

Running Services

Run a task

Good for short-lived containers, e.g. batch jobs

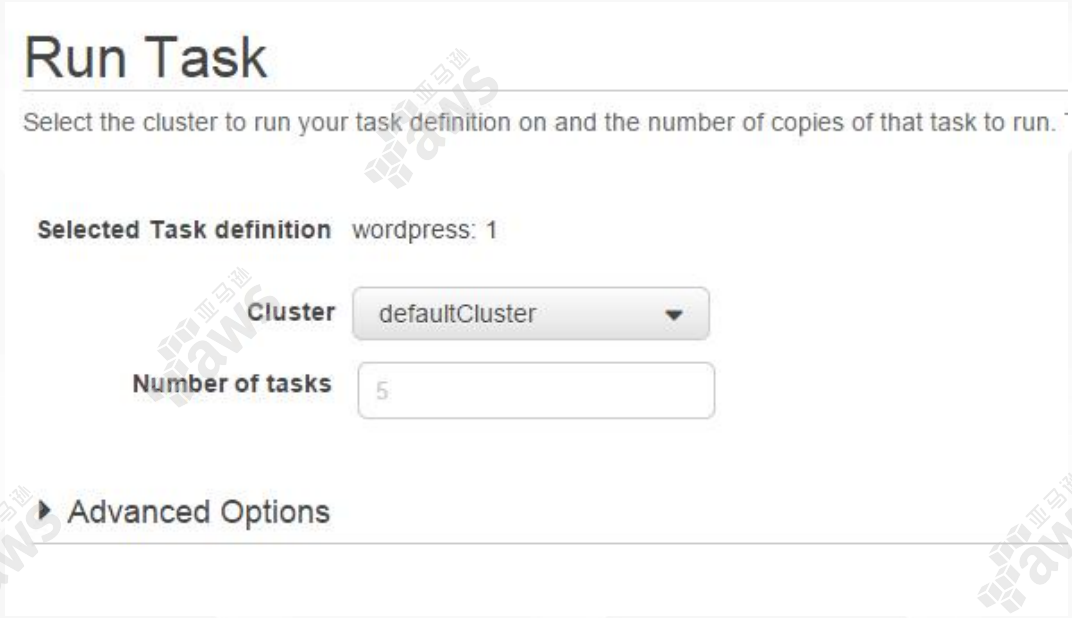

Create a Service

Good for long-running applications and services

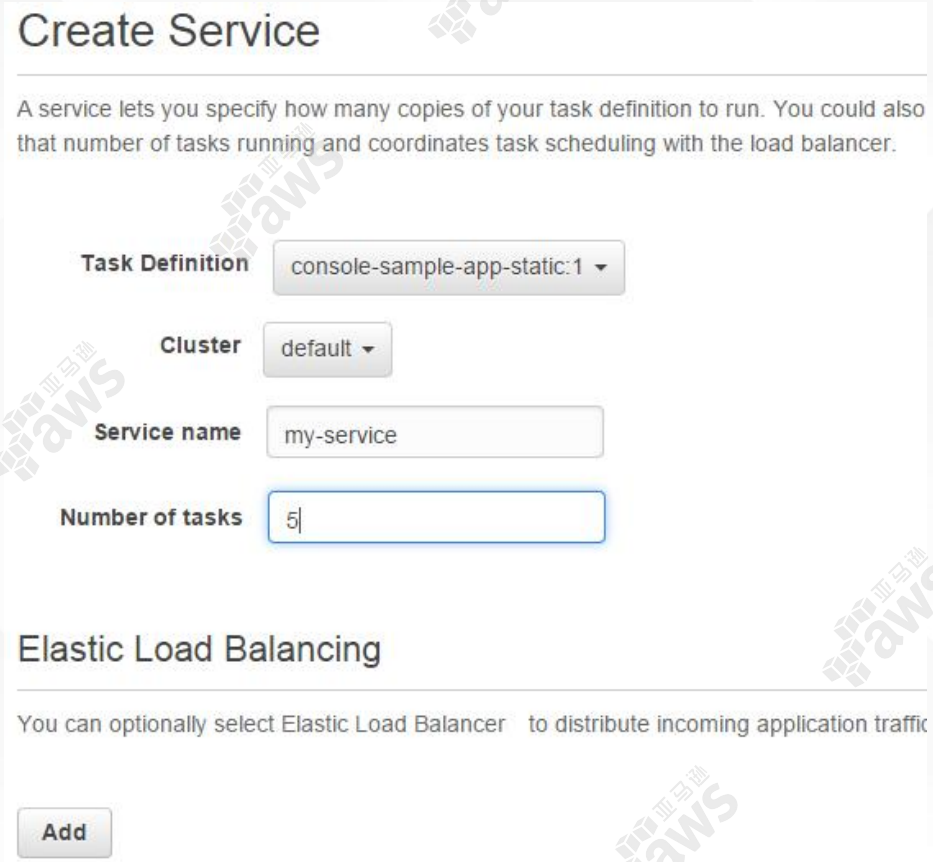

Create Service

Load Balance traffic across containers

Automatically recover unhealthy containers

Discover services

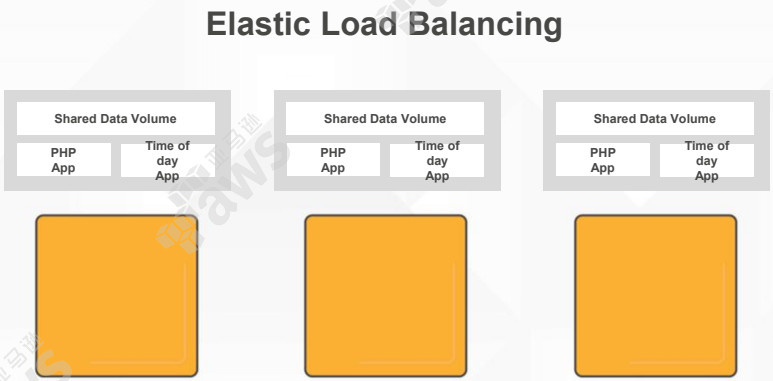

Update Service (cont.)

Scale up

Scale down

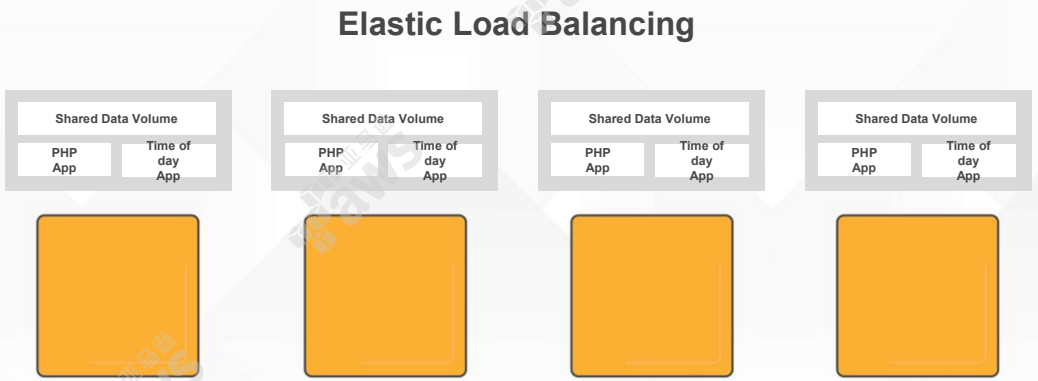

Update Service (cont.)

Deploy new version

Drain connections

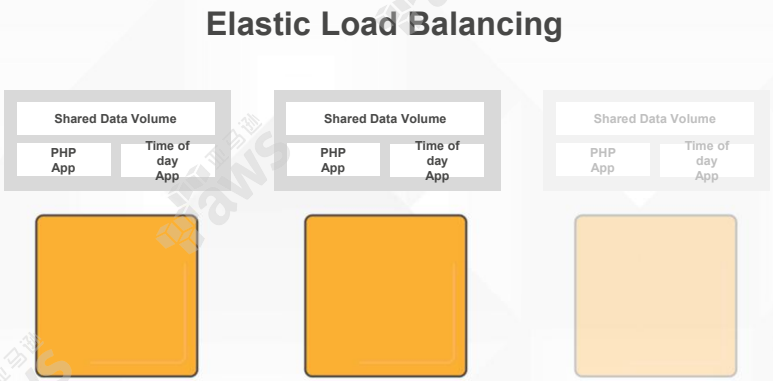

Update Service (cont.)

Deploy new version

Drain connections

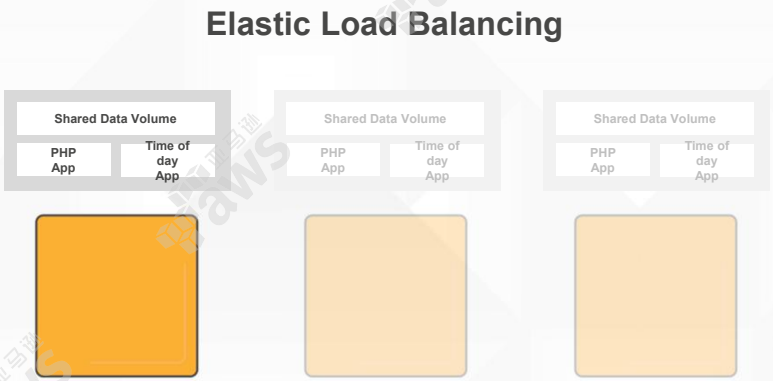

Update Service (cont.)

Deploy new version

Drain connections

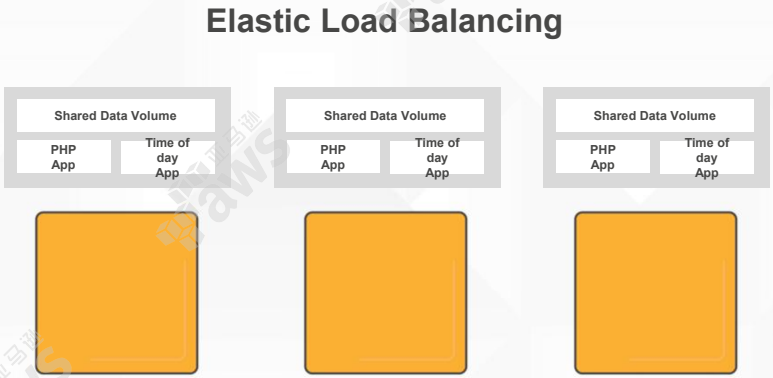

Demo

Thank You