



# Oracle技术嘉年华

Oracle Technology Carnival 2015

稳健•高效•云端 - 数据技术最佳实践

## 算法为王的 SQL优化

怀晓明



# 目录

- Who am I?
- 应用困境
- 影响数据库性能的因素
- 从一个复杂的案例开始
- 一个看似简单的案例
- SQL优化的本质
- 以一个“简单”的案例结束
- 总结

# Who am I ?

- Lastwinner@itpub.net，资深数据库开发版主
- 《剑破冰山——Oracle开发艺术》、《DBA手记2》合著者
- 善于搜索的Troubleshooter

# 应用困境

- 现状
  - 随着 society 对 IT 的需求不断扩展与扩张，IT 应用呈爆发式增长
  - 大量的需求尤其是新需求，需要大量的开发者来完成
  - 于是，“码农”应运而生，而且是大量产生，这使得开发人员的总体水平越来越低
  - 然而，这是 IT 发展过程中的正常现象
  - 每个系统的数据量越来越大，随之而来的问题当然就是应用大量的性能问题

## 应用困境

- 大多数公司的解决办法
  - 升级硬件
  - 拓展架构
- 预想的实施结果
  - 解决性能问题
- 然而现实总是残酷的
  - 硬件升级太昂贵，架构拓展成本也不低
  - 即便钱不是问题，一年后呢？两年后呢？ .....
  - 老板不是傻瓜，无底洞的硬件投入，不可持续发展

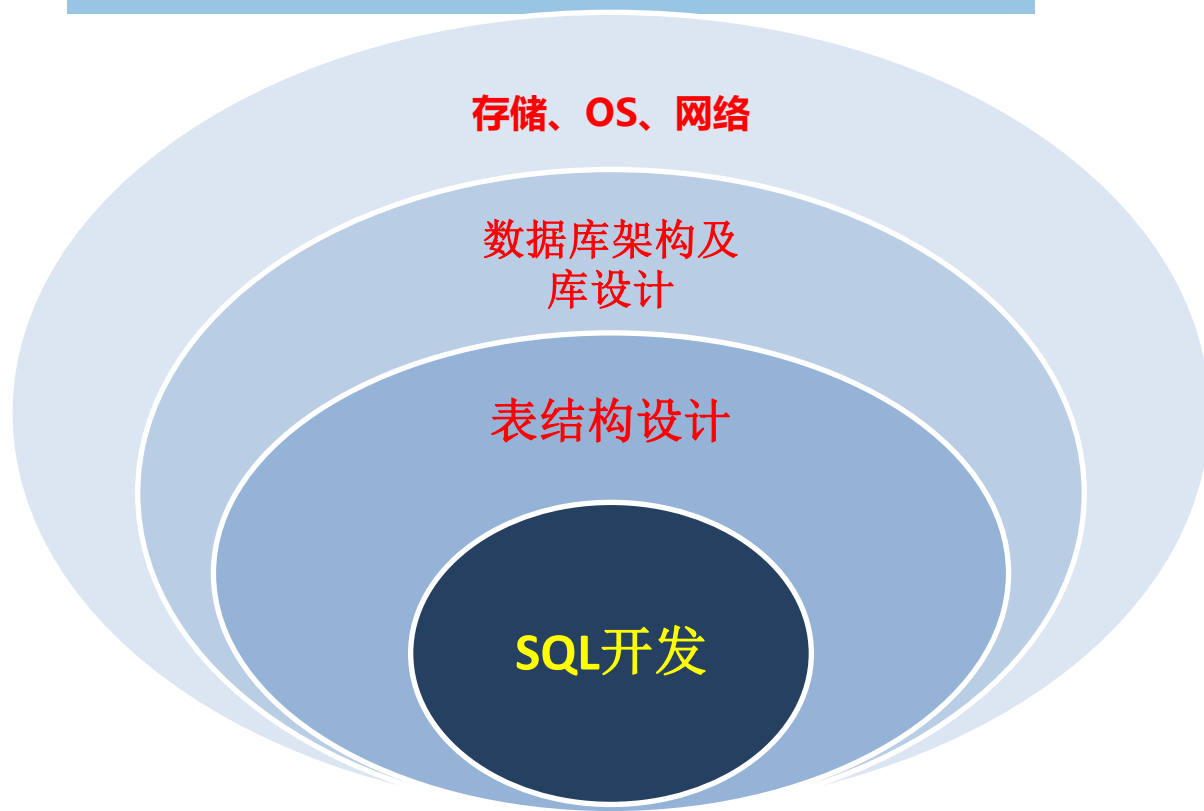
## 应用困境

- 如何真正的解决应用性能问题？
  - 战略：治病要除根
  - 战术：程序代码优化，尤其是SQL优化

# 影响数据库性能的因素

- 框架
- 表
- 索引
- 统计信息
- 执行计划
- SQL写法

# 影响数据库性能的因素—框架





# 影响数据库性能的因素一表

- 大表没分区
- 表没主键
- 外键无索引
- 主子表设计不合理
- .....

# 影响数据库性能的因素—索引

- 未建索引
- 索引无效
- 索引不可见
- 索引不合理
- 索引过于冗余
- .....

# 影响数据库性能的因素—统计信息

- 缺失
- 过期
- 不准
- .....

# 影响数据库性能的因素—执行计划

- 执行计划改变
- 执行计划不稳定（有多个）
- 执行计划非最佳
  - stats
  - CBO
  - hint
  - .....
- 执行计划错误
- .....

# 影响数据库性能的因素—SQL写法

- 未使用绑定变量
- 隐式转换
- SQL写得不够好
- SQL写得不太好
- SQL写得比较差
- SQL写得很差劲
- SQL写得很糟糕
- SQL写得.....

# 影响数据库性能的因素—SQL写法

- 写得不好的SQL如何优化？
- 当然是靠改写！

# 从一个复杂的案例开始—SQL 文本

- ```
select sii.P_Name processName, sii.PID processId, sii.VID versionId, (nvl(max(t.ciid), 0)+nvl(max(queueMember.ciid), 0)+nvl(max(s.ciid), 0)+nvl(max(abs.ciid), 0)) totalTask from (select SI.*, IP.P_Name as P_Name from (select Current_Process_Id as PID, Current_Version_Id as VID, Step_Id as SID from step_info where step_status in (1, 6)) SI left join Installed_Process IP on SI.PID = IP.Process_Id and SI.VID = IP.Version_Id ) sii left join (select Current_Process_Id, Current_Version_Id, count(v2.Incident_Id) ciid from step_info v2 where getTaskPerformer(Belong_to, Assign_to, Escalation_to)=7785 and Current_Process_Id>0 and Current_Version_Id>0 AND step_status in (1, 6) group by Current_Process_Id, Current_Version_Id) t on sii.pid=t.Current_Process_Id and sii.vid=t.Current_Version_Id left join (select Current_Process_Id, Current_Version_Id, count(v2.Current_Process_Id) ciid from step_info v2, Queue_Member q where (v2.TaskFlag='Q' and v2.urlid=q.urlid and q.User_Num= 7785) and QPICKUP <> 'Y' and Current_Process_Id>0 and Current_Version_Id>0 AND step_status in (1, 6) group by Current_Process_Id, Current_Version_Id) queueMember on sii.pid=queueMember.Current_Process_Id and sii.vid=queueMember.Current_Version_Id left join (select Current_Process_Id, Current_Version_Id, count(v2.Incident_Id) ciid from step_info v2 where step_status=4 and getTaskPerformer(belong_to, assign_to, escalation_to)=7785 and Current_Process_Id>0 and Current_Version_Id>0 group by Current_Process_Id, Current_Version_Id) s on sii.pid=s.Current_Process_Id and sii.vid=s.Current_Version_Id left join (select Current_Process_Id, Current_Version_Id, count(v2.Incident_Id) ciid from step_info v2 where assign_to > 0 and belong_to =7785 and belong_to= assign_by and escalation_to is null and Current_Process_Id>0 and Current_Version_Id>0 AND step_status in (1, 6) group by Current_Process_Id, Current_Version_Id) abs on sii.pid=abs.Current_Process_Id and sii.vid=abs.Current_Version_Id where (t.ciid is not null or queueMember.ciid is not null or s.ciid is not null or abs.ciid is not null) group by sii.pid, sii.vid, sii.P_Name order by nlsort(sii.P_Name, 'NLS_SORT=CHINESE_PINYIN_M')
```

## 从一个复杂的案例开始—相关信息

- 多个执行计划

| # | Plan Hash Value | Total Elapsed Time(ms) | Executions |
|---|-----------------|------------------------|------------|
| 1 | 4029382908      | 32,855,784             | 9,958      |
| 2 | 191117238       | 10,052                 | 2          |

- 执行计划信息

| Stat Name         | Statement Total | Per Execution | % Snap Total |
|-------------------|-----------------|---------------|--------------|
| Elapsed Time (ms) | 32,855,784      | 3,299.44      | 9.17         |
| CPU Time (ms)     | 5,204,766       | 522.67        | 2.90         |
| Executions        | 9,958           |               |              |
| Buffer Gets       | 535,692,906     | 53,795.23     | 6.42         |
| Disk Reads        | 1,802           | 0.18          | 0.00         |
| Parse Calls       | 9,958           | 1.00          | 0.01         |
| Rows              | 149,220         | 14.98         |              |



| Id | Operation                   | Name                  | Rows | Bytes | Cost (%CPU) | Time     |
|----|-----------------------------|-----------------------|------|-------|-------------|----------|
| 0  | SELECT STATEMENT            |                       |      |       | 10211 (100) |          |
| 1  | SORT ORDER BY               |                       | 42   | 4830  | 10211 (2)   | 00:02:03 |
| 2  | HASH GROUP BY               |                       | 42   | 4830  | 10211 (2)   | 00:02:03 |
| 3  | HASH JOIN RIGHT OUTER       |                       | 640K | 70M   | 10151 (1)   | 00:02:02 |
| 4  | INDEX FULL SCAN             | INSTALLED_PROCESS_IND | 196  | 4900  | 1 (0)       | 00:00:01 |
| 5  | FILTER                      |                       |      |       |             |          |
| 6  | HASH JOIN RIGHT OUTER       |                       | 640K | 55M   | 10146 (1)   | 00:02:02 |
| 7  | VIEW                        |                       | 4877 | 97540 | 5951 (1)    | 00:01:12 |
| 8  | HASH GROUP BY               |                       | 4877 | 147K  | 5951 (1)    | 00:01:12 |
| 9  | NESTED LOOPS                |                       | 4877 | 147K  | 5949 (1)    | 00:01:12 |
| 10 | INDEX RANGE SCAN            | QM_IND1               | 4877 | 53647 | 149 (0)     | 00:00:02 |
| 11 | TABLE ACCESS BY INDEX ROWID | STEP_INFO             | 1    | 20    | 3 (0)       | 00:00:01 |
| 12 | INDEX UNIQUE SCAN           | PK_STEP_INFO          | 1    |       | 2 (0)       | 00:00:01 |
| 13 | HASH JOIN RIGHT OUTER       |                       | 312K | 20M   | 4177 (1)    | 00:00:51 |
| 14 | VIEW                        |                       | 81   | 1620  | 83 (2)      | 00:00:01 |
| 15 | HASH GROUP BY               |                       | 81   | 2916  | 83 (2)      | 00:00:01 |
| 16 | TABLE ACCESS BY INDEX ROWID | STEP_INFO             | 81   | 2916  | 82 (0)      | 00:00:01 |
| 17 | INDEX RANGE SCAN            | STEPINFO_IND2         | 81   |       | 7 (0)       | 00:00:01 |
| 18 | HASH JOIN RIGHT OUTER       |                       | 312K | 14M   | 4091 (1)    | 00:00:50 |
| 19 | VIEW                        |                       | 5    | 100   | 13 (8)      | 00:00:01 |
| 20 | HASH GROUP BY               |                       | 5    | 180   | 13 (8)      | 00:00:01 |
| 21 | TABLE ACCESS BY INDEX ROWID | STEP_INFO             | 5    | 180   | 12 (0)      | 00:00:01 |
| 22 | INDEX RANGE SCAN            | STEPINFO_IND2         | 5    |       | 7 (0)       | 00:00:01 |
| 23 | HASH JOIN RIGHT OUTER       |                       | 312K | 9151K | 4076 (1)    | 00:00:49 |
| 24 | VIEW                        |                       | 1    | 20    | 5 (20)      | 00:00:01 |
| 25 | HASH GROUP BY               |                       | 1    | 38    | 5 (20)      | 00:00:01 |
| 26 | TABLE ACCESS BY INDEX ROWID | STEP_INFO             | 1    | 38    | 4 (0)       | 00:00:01 |
| 27 | INDEX RANGE SCAN            | STEPINFO_IND6         | 1    |       | 3 (0)       | 00:00:01 |
| 28 | INDEX FAST FULL SCAN        | STEPINFO_IND1         | 312K | 3050K | 4069 (1)    | 00:00:49 |

## 从一个复杂的案例开始—格式化

```
SELECT sii.P_Name processName, sii.PID processId, sii.VID versionId,
       (NVL (MAX (t.ciid), 0) + NVL (MAX (queueMember.ciid), 0) + NVL (MAX (s.ciid), 0)
        + NVL (MAX (ABS.ciid), 0)) totalTask
FROM (SELECT SI.*, IP.P_Name AS P_Name
      FROM (SELECT Current_Process_Id AS PID, Current_Version_Id AS VID, Step_ID AS SID
            FROM step_info
            WHERE step_status IN (1, 6)) SI
      LEFT JOIN Installed_Process IP ON SI.PID = IP.Process_ID AND SI.VID = IP.Version_ID) sii
LEFT JOIN
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
  FROM step_info v2
  WHERE      getTaskPerformer (Belong_to, Assign_to, Escalation_to) = 7785
            AND Current_Process_Id > 0
            AND Current_Version_Id > 0
            AND step_status IN (1, 6)
 GROUP BY Current_Process_Id, Current_Version_Id) t
ON sii.pid = t.Current_Process_Id AND sii.vid = t.Current_Version_Id
LEFT JOIN
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Current_Process_Id) ciid
  FROM step_info v2, Queue_Member q
  WHERE      (v2.TaskFlag = 'Q' AND v2.urlid = q.urlid AND q.User_Num = 7785)
            AND QPICKUP <> 'Y'
            AND Current_Process_Id > 0
            AND Current_Version_Id > 0
            AND step_status IN (1, 6)
 GROUP BY Current_Process_Id, Current_Version_Id) queueMember
ON sii.pid = queueMember.Current_Process_Id AND sii.vid = queueMember.Current_Version_Id
LEFT JOIN
```

## 从一个复杂的案例开始—格式化

```
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
  FROM step_info v2
 WHERE      step_status = 4
           AND getTaskPerformer (belong_to, assign_to, escalation_to) = 7785
           AND Current_Process_Id > 0
           AND Current_Version_Id > 0
 GROUP BY Current_Process_Id, Current_Version_Id) s
 ON sii.pid = s.Current_Process_Id AND sii.vid = s.Current_Version_Id
LEFT JOIN
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
  FROM step_info v2
 WHERE      assign_to > 0
           AND belong_to = 7785
           AND belong_to = assign_by
           AND escalation_to IS NULL
           AND Current_Process_Id > 0
           AND Current_Version_Id > 0
           AND step_status IN (1, 6)
 GROUP BY Current_Process_Id, Current_Version_Id) ABS
 WHERE (t.ciid IS NOT NULL OR queueMember.ciid IS NOT NULL OR s.ciid IS NOT NULL OR ABS.ciid IS NOT NULL)
GROUP BY sii.pid, sii.vid, sii.P_Name
ORDER BY NLSSORT (sii.P_Name, 'NLS_SORT=SCHINESE_PINYIN_M')
```

## 从一个复杂的案例开始—分析

```
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
   FROM step_info v2
  WHERE      getTaskPerformer (Belong_to, Assign_to, Escalation_to) = 7785
            AND Current_Process_Id > 0
            AND Current_Version_Id > 0
            AND step_status IN (1, 6)
 GROUP BY Current_Process_Id, Current_Version_Id) t
```

```
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
   FROM step_info v2
  WHERE      step_status = 4
            AND getTaskPerformer (belong_to, assign_to, escalation_to) = 7785
            AND Current_Process_Id > 0
            AND Current_Version_Id > 0
 GROUP BY Current_Process_Id, Current_Version_Id) s
```

## 从一个复杂的案例开始—查询合并

```
WITH tmp_a
AS ( SELECT step_status, Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
      FROM step_info
      WHERE      getTaskPerformer (Belong_to, Assign_to, Escalation_to) = 7785
              AND Current_Process_Id > 0
              AND Current_Version_Id > 0
              AND step_status IN (1, 4, 6)
      GROUP BY Current_Process_Id, Current_Version_Id, step_status),
t
AS ( SELECT Current_Process_Id, Current_Version_Id, SUM (ciid) ciid
      FROM tmp_a
      WHERE step_status IN (1, 6)
      GROUP BY Current_Process_Id, Current_Version_Id),
s
AS (SELECT Current_Process_Id, Current_Version_Id, ciid
      FROM tmp_a
      WHERE step_status = 4)
```

## 从一个复杂的案例开始—查询合并

```
WITH tmp_a AS ( SELECT step_status, Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
                  FROM step_info
                  WHERE      getTaskPerformer (Belong_to, Assign_to, Escalation_to) = 7785
                        AND Current_Process_Id > 0 AND Current_Version_Id > 0 AND step_status IN (1, 4, 6)
                  GROUP BY Current_Process_Id, Current_Version_Id, step_status),
t AS (SELECT Current_Process_Id, Current_Version_Id, SUM (ciid) ciid
      FROM tmp_a WHERE step_status IN (1, 6)
      GROUP BY Current_Process_Id, Current_Version_Id),
s AS (SELECT Current_Process_Id, Current_Version_Id, ciid
      FROM tmp_a WHERE step_status = 4),
tmp_b AS (SELECT Current_Process_Id, Current_Version_Id, Step_ID, Incident_Id, TaskFlag,
                urlid, QPICKUP, /*虽然条件中没有指明此字段来源, 但可以猜出应该属于step_info表*/
                assign_to, assign_by, belong_to, escalation_to
      FROM step_info WHERE step_status IN (1, 6)),
sii AS (SELECT si.*, ip.p_name
        FROM tmp_b si LEFT JOIN Installed_Process IP
        ON SI.Current_Process_Id = IP.Process_ID AND SI.Current_Version_Id = IP.Version_ID),
queueMember AS ( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Current_Process_Id) ciid
                 FROM tmp_b v2, Queue_Member q
                 WHERE (v2.TaskFlag = 'Q' AND v2.urlid = q.urlid AND q.User_Num = 7785)
                      AND QPICKUP <> 'Y'
                      AND Current_Process_Id > 0
                      AND Current_Version_Id > 0
                 GROUP BY Current_Process_Id, Current_Version_Id),
/*ABS是oracle内置的函数, 命名时应避免, 所以此处改为了v_abs*/
v_abs AS ( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
          FROM tmp_b
          WHERE      assign_to > 0
```

## 从一个复杂的案例开始—小结

- 其实，单纯从SQL文本上看，就能做出如上改写
- 本次改写涉及到：
  - 查询合并
  - 查询结果分发（行列转换）

## 一个看似简单的案例—SQL文本

```
SELECT DEPART_ID DEPART_ID, B.CUC_CHNL_CODE CUC_DEPART_CODE
FROM DEPART A, CHANNEL B
WHERE A.DEPART_ID = B.RSRV_STR9
AND B.CUC_CHNL_CODE IS NOT NULL
UNION
SELECT DEPART_ID DEPART_ID, CUC_DEPART_CODE
FROM DEPART
WHERE CUC_DEPART_CODE IN
    (SELECT A.CUC_DEPART_CODE
     FROM DEPART A
     WHERE A.CUC_DEPART_CODE IS NOT NULL
     AND NOT EXISTS (SELECT 1
                     FROM CHANNEL B
                     WHERE B.RSRV_STR9 = A.DEPART_ID))
MINUS
SELECT B.CUC_CHNL_CODE
FROM DEPART A, CHANNEL B
WHERE A.DEPART_ID = B.RSRV_STR9
AND B.CUC_CHNL_CODE IS NOT NULL)
ORDER BY DEPART_ID
```



# 一个看似简单的案例—相关信息

- 表所占空间

|                |    |
|----------------|----|
| – SEGMENT_NAME | MB |
| – DEPART       | 20 |
| – CHANNEL      | 8  |

- 记录数以及据此计算出表所占block总大小

|              |          |             |     |
|--------------|----------|-------------|-----|
| – TABLE_NAME | NUM_ROWS | AVG_ROW_LEN | MB  |
| – DEPART     | 28810    | 154         | 4.2 |
| – CHANNEL    | 28390    | 237         | 6.4 |

- SQL问题

- SQL在指定运行时间窗口内没执行完，而作为某业务关键路径上的SQL，其执行时间平常都没超过10秒
- 如何优化？

## 一个看似简单的案例—改写方案

```
WITH t AS
(SELECT DEPART_ID DEPART_ID, B.CUC_CHNL_CODE CUC_DEPART_CODE
 FROM DEPART A, CHANNEL B
 WHERE A.DEPART_ID = B.RSRV_STR9
 AND B.CUC_CHNL_CODE IS NOT NULL)
SELECT * FROM t
UNION
SELECT DEPART_ID DEPART_ID, CUC_DEPART_CODE
 FROM DEPART
 WHERE CUC_DEPART_CODE IN
 (SELECT A.CUC_DEPART_CODE
  FROM DEPART A
  WHERE A.CUC_DEPART_CODE IS NOT NULL
  AND NOT EXISTS (SELECT 1
                   FROM CHANNEL B
                   WHERE B.RSRV_STR9 = A.DEPART_ID)
  MINUS
  SELECT B.CUC_CHNL_CODE
  FROM t)
ORDER BY DEPART_ID
```

## 一个看似简单的案例—思路

- 如何做出更好的改写？
- 读懂业务！算法为王！！

## 一个看似简单的案例—SQL文本副本

```
SELECT DEPART_ID DEPART_ID, B.CUC_CHNL_CODE CUC_DEPART_CODE
FROM DEPART A, CHANNEL B
WHERE A.DEPART_ID = B.RSRV_STR9
AND B.CUC_CHNL_CODE IS NOT NULL
UNION
SELECT DEPART_ID DEPART_ID, CUC_DEPART_CODE
FROM DEPART
WHERE CUC_DEPART_CODE IN
    (SELECT A.CUC_DEPART_CODE
     FROM DEPART A
     WHERE A.CUC_DEPART_CODE IS NOT NULL
     AND NOT EXISTS (SELECT 1
                     FROM CHANNEL B
                     WHERE B.RSRV_STR9 = A.DEPART_ID))
MINUS
SELECT B.CUC_CHNL_CODE
FROM DEPART A, CHANNEL B
WHERE A.DEPART_ID = B.RSRV_STR9
AND B.CUC_CHNL_CODE IS NOT NULL)
ORDER BY DEPART_ID
```

## 一个看似简单的案例—更好的改写方案

```
WITH t AS
  (SELECT DEPART_ID DEPART_ID, A.CUC_DEPART_CODE, B.CUC_CHNL_CODE, B.RSRV_STR9
   FROM DEPART A, CHANNEL B
   WHERE A.DEPART_ID = B.RSRV_STR9(+))
SELECT DEPART_ID, CUC_CHNL_CODE CUC_DEPART_CODE
  FROM t
 WHERE CUC_CHNL_CODE IS NOT NULL
UNION
SELECT DEPART_ID, CUC_DEPART_CODE
  FROM t
 WHERE      CUC_DEPART_CODE IS NOT NULL
        AND RSRV_STR9 IS NULL
        AND CUC_DEPART_CODE NOT IN (SELECT CUC_CHNL_CODE
                                       FROM t
                                       WHERE CUC_CHNL_CODE IS NOT NULL)
ORDER BY DEPART_ID
```

# 一个看似简单的案例—改写方案

- 优化结果

- 如果你只能学一个SQL的优化技能，那么以下案例，你想学哪一个（均非缺索引）？
  - a. 一分钟到0.25秒
  - b. 9分钟到4.5分钟
  - c. 1小时18分到3秒
  - d. 400秒到40毫秒
  - e. 3.5秒到1秒以内（未有实际环境验证效率）

## 一个看似简单的案例一小结

- 查询合并
- 熟悉各种join的执行结果会是什么样
- 熟悉集合操作的结果会是什么样
- 读懂业务，算法为王！



# SQL优化的本质

- SQL优化，是数据库性能优化的重要一环
- 一般的，会将数据库性能优化当作就是SQL优化
- SQL优化要做哪些工作？
  - 包括SQL改写在内的SQL层面的优化，60%以上
  - 表和索引等对象的优化，25%左右
  - 模型优化（如主子表），7%左右
  - 其他，8%左右
- SQL优化更偏向业务层面，只有对业务理解透彻，才能做好SQL优化



# SQL 优化的本质

- 按目标定义，什么是SQL优化的本质？
  - 片面：SQL运行得越快越好
  - 深刻：平衡系统内的资源使用，充分发挥硬件资源，让大多数SQL运行得越快越好
  - 全面：平衡SQL优化服务成本与SQL优化效果期望，做到广义上的资源平衡，达到最优的性能目标

# SQL优化的本质

- SQL优化对企业的意义？
  - 从根源上提升系统效率，降低硬件设备损耗，延长硬件寿命
  - 延缓企业为解决性能问题而采购新设备的周期，等价于为企业省钱
  - 同样的硬件资源可以支撑更多的应用系统，等价于为企业省钱
- 成功案例：某客户系统，经过我方优化，虚拟化环境下的两节点的RAC，CPU从每节点12颗直接降低到每节点8颗，大大节约了宝贵的硬件资源。而这部分节省出的资源，客户又将其划分给其他系统使用，无异于节省了昂贵的硬件成本。

# SQL 优化的本质

- 如何找到能提供优秀的SQL优化的服务团队？



云和恩墨  
ENMOTECH

数据驱动 成就未来  
*Make Your Data Dance*

# Who am I ?

- 2012 年全国SQL大赛评委
- 2013 年全国SQL大赛评委

# SQL 优化的本质

- 如何加入如此优秀的SQL优化服务团队？
  - 熟悉SQL开发及数据库基本知识
  - 逻辑严密，思路清晰
  - 工作认真仔细，能吃苦
  - 工作年限不限，优秀的毕业生也可
  - HR@enmotech.com
- 技术方面你能获得什么提升？
  - 一年之后 能处理大部分初中级的SQL性能问题
  - 两年之后 能处理所有的初中级问题和部分高级的SQL性能问题
  - 三年之后 独立处理各种系统的SQL优化事宜

## 以一个“简单”的案例结束—SQL文本

```
SELECT sii.P_Name processName, sii.PID processId, sii.VID versionId,
       (NVL (MAX (t.ciid), 0) + NVL (MAX (queueMember.ciid), 0) + NVL (MAX (s.ciid), 0)
        + NVL (MAX (ABS.ciid), 0)) totalTask
FROM (SELECT SI.*, IP.P_Name AS P_Name
      FROM (SELECT Current_Process_Id AS PID, Current_Version_Id AS VID, Step_Id AS SID
            FROM step_info
            WHERE step_status IN (1, 6)) SI
      LEFT JOIN Installed_Process IP ON SI.PID = IP.Process_ID AND SI.VID = IP.Version_ID) sii
LEFT JOIN
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
  FROM step_info v2
  WHERE      getTaskPerformer (Belong_to, Assign_to, Escalation_to) = 7785
            AND Current_Process_Id > 0
            AND Current_Version_Id > 0
            AND step_status IN (1, 6)
  GROUP BY Current_Process_Id, Current_Version_Id) t
ON sii.pid = t.Current_Process_Id AND sii.vid = t.Current_Version_Id
LEFT JOIN
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Current_Process_Id) ciid
  FROM step_info v2, Queue_Member q
  WHERE      (v2.TaskFlag = 'Q' AND v2.urlid = q.urlid AND q.User_Num = 7785)
            AND QPICKUP <> 'Y'
            AND Current_Process_Id > 0
            AND Current_Version_Id > 0
            AND step_status IN (1, 6)
  GROUP BY Current_Process_Id, Current_Version_Id) queueMember
ON sii.pid = queueMember.Current_Process_Id AND sii.vid = queueMember.Current_Version_Id
LEFT JOIN
```

以一个“简单”的案例结束—优化思路

读懂业务，  
算法为王！

## 以一个“简单”的案例结束—优化结果

```
/* 实际上查询消耗最大的是在join操作，因此减少join操作，应能取得良好的效果
*/
SELECT sii.P_Name processName, sii.PID processId, sii.VID versionId,
       (NVL (MAX (t.ciid), 0) + NVL (MAX (queueMember.ciid), 0) + NVL (MAX (s.ciid), 0) + NVL (MAX (ABS.
FROM (SELECT SI.*, IP.P_Name AS P_Name
      FROM ( SELECT Current_Process_Id AS PID, Current_Version_Id AS VID, Step_ID AS SID
              FROM step_info
              WHERE step_status IN (1, 6)
              /*仅仅添加这一行*/
              GROUP BY Current_Process_Id, Current_Version_Id, Step_ID) SI
      LEFT JOIN Installed_Process IP ON SI.PID = IP.Process_ID AND SI.VID = IP.Version_ID)) sii
LEFT JOIN
( SELECT Current_Process_Id, Current_Version_Id, COUNT (v2.Incident_Id) ciid
  FROM step_info v2
  WHERE   getTaskPerformer (Belong_to, Assign_to, Escalation_to) = 7785
         AND Current_Process_Id > 0
         AND Current_Version_Id > 0
         AND step_status IN (1, 6)
 GROUP BY Current_Process_Id, Current_Version_Id) t
ON sii.pid = t.Current_Process_Id AND sii.vid = t.Current_Version_Id
```



## 以一个“简单”的案例结束一小结

- 四两拨千斤，有木有？

# 总结

- 程序=数据+算法
- 读懂业务，算法为王，写出极致SQL，达到SQL优化的极高境界！
- 对于开篇那个复杂的案例，如果你能独立完成全面的SQL优化工作（即不仅仅是我讲解过的方法的组合），不要犹豫，即刻发送简历，您有可能直接入职云和恩墨！

# 关注微信公众号 获取文档和更新



云和恩墨

(支持ORA错误自动查询)



恩墨学院




Oracle新闻



## z3 - SQL审核工具提升SQL质量

### • 独特的 SQL 视角

- SQL生命周期管理 - 捕获、分析、归档、形成SQL全周期管理平台



z3 - SQL审核



## zData - 高性能弹性分布式存储解决方案

### • 大数据整合与集中面临的平台压力

- 去“三”架构，通过Virtual SAN替代FC SAN
- 同样的成本获得 20倍+ 的IO性能
- 动态扩展，高性能的存储解决方案



zData - 分布式存储



## Announcing: BayMax 自动化巡检 即将开放云服务

- 自动化巡检 - 让DBA告别繁琐24x7最有价值的工作



BayMax自动化巡检

# 特别感谢 合作伙伴



云和恩墨  
ENMOTECH



Shannon Systems  
宝 存 科 技

ACOUG  
All China Oracle User Group  
中国 Oracle 用户组

恩墨学院  
ENMOEDU

Data Driven World  
EN·CORE  
恩核(北京)信息技术有限公司

OCMU  
OCM 联盟

SDOUG  
Shandong Oracle User Group  
山东 Oracle 用户组

ITPUB

ChinaUnix

IT168.com  
www.it168.com

BI 会议

Mellanox  
TECHNOLOGIES  
Connect. Accelerate. Outperform.™

乐导科技  
alidao.com

慕课网  
imooc.com



清华大学出版社

Broadview®  
www.broadview.com.cn

TURING  
图灵教育

华章科技

The background is a solid red color. In the top-left corner, there is a faint, light-red decorative floral element. In the bottom-right corner, there is a large, vibrant, multi-colored decorative floral element with shades of yellow, orange, pink, and blue. The word "THANKS" is centered in the middle of the image in a white, sans-serif font.

THANKS