

.conf2015

# Making the Most of the New Splunk Scheduler

Paul J. Lucas

Sr. Software Engineer, Splunk

splunk>

# Disclaimer

During the course of this presentation, we may make forward looking statements regarding future events or the expected performance of the company. We caution you that such statements reflect our current expectations and estimates based on factors currently known to us and that actual events or results could differ materially. For important factors that may cause actual results to differ from those contained in our forward-looking statements, please review our filings with the SEC. The forward-looking statements made in the this presentation are being made as of the time and date of its live presentation. If reviewed after its live presentation, this presentation may not contain current or accurate information. We do not assume any obligation to update any forward looking statements we may make.

In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. Splunk undertakes no obligation either to develop the features or functionality described or to include any such feature or functionality in a future release.

# Personal Introduction

- Paul J. Lucas, Sr. Software Engineer, Splunk
- On the Core Server Engineering Team
- Worked on Search Scheduler improvements to Splunk Enterprise
- Also worked on various parts of the Deployment Server
- Has been using C++ since the “cfront” days at AT&T Bell Labs

# Agenda

- **Scheduled Searches:**
  - A. Introduction
  - B. How Cron Works
  - C. Cron vs. Splunk Scheduler
- **Splunk Scheduler:**
  - A. Deferred vs. Skipped
  - B. Lag
  - C. `limits.conf` Settings
- **Splunk Scheduler Improvements:**
  - A. How the Splunk Scheduler Works
  - B. Priority Scoring
  - C. Schedule Windows
  - D. Variable `max_searches_perc`
  - E. Continuous Saved Searches “Catch-up”
- **Takeaways**
- **Splunk Enterprise 6.3 Test Results**

.conf2015

# Scheduled Searches

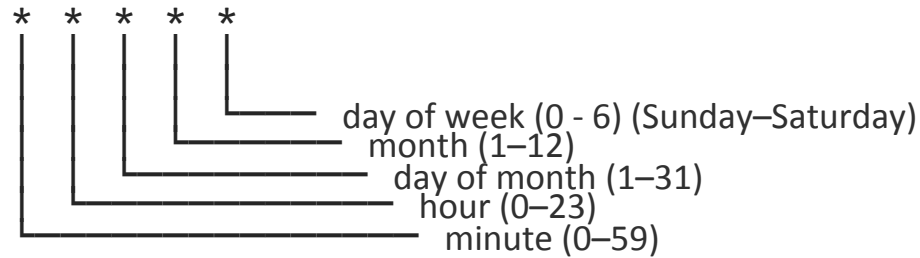
# Scheduled Searches: Introduction

- Splunk allows you to save your searches and run them on a schedule
- Scheduled searches can be used to trigger an alert action (possibly when a condition is met) or to speed-up dashboards
- An alert action is either sending an e-mail or running a script
- **Example:** index=\_internal source=\*splunkd.log\* error NOT debug

<b>Title</b>	Too many errors
<b>Trigger condition</b>	Number of Results
<b>Number of results is</b>	Greater than: 5
<b>in</b>	1 Minute

# Scheduled Searches: Introduction

- Scheduling is specified via a five-field cron string:



- All (\*), ranges (1–5), lists (1,8,15,22), and “every  $n$ ” ( $*/6$ ) allowed
- **Example:** `0 */6 1,15 * *` = every 6 hours on the hour on the 1st and 15th

# How Cron Works

1. For each cron entry, calculate the next run-time of the command
2. Place all commands in a priority queue by time
3. Enter main loop:
  - A. Examine the entry at the head of the queue
  - B. Calculate the delta between that entry's next run-time and *now*
  - C. If  $\text{delta} > 0$ , sleep for that period of time
  - D. Run the entry's command (in the background)
  - E. Calculate the *next* run-time of the command and place it back on the queue with that new time value



# Cron vs. Splunk Scheduler

## Cron

- No job quotas
- Entirely manual scheduling — have to skew searches by hand:

```
0 0 * * * command-1
15 0 * * * command-2
30 0 * * * command-3
45 0 * * * command-4
```

- Limited to a single machine

## Splunk Scheduler

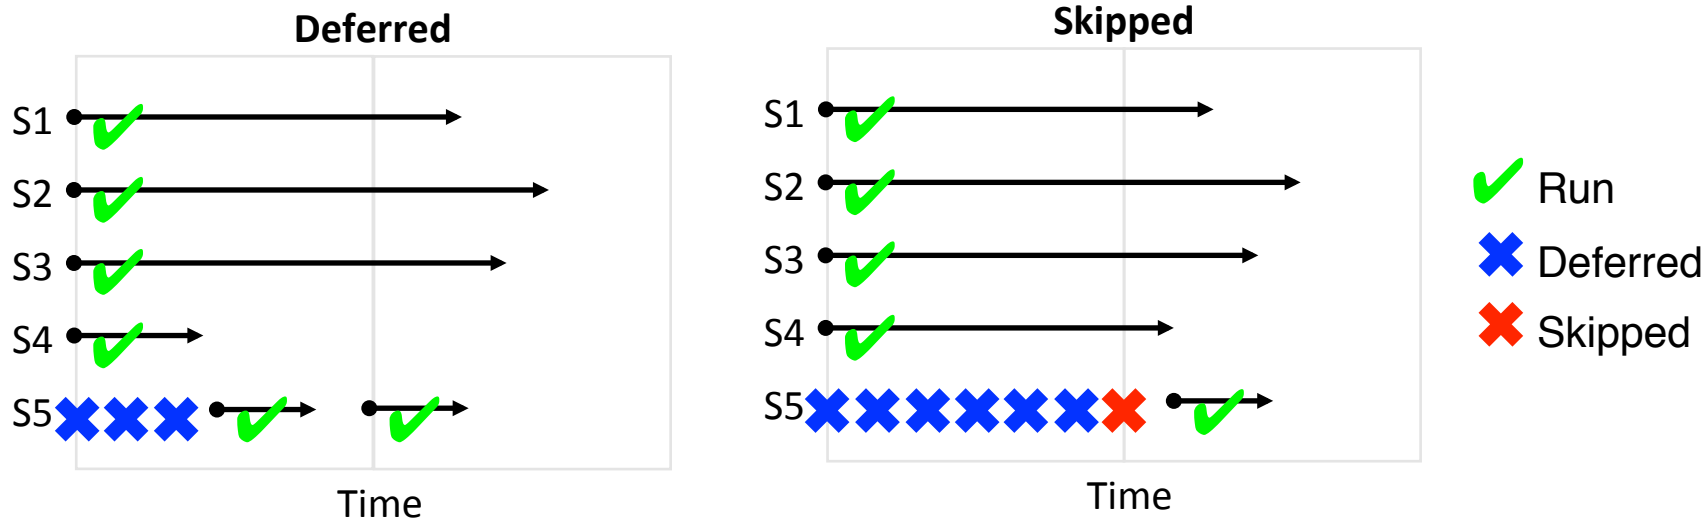
- Quotas: limit search concurrency — reserves CPU for other tasks
- Searches over quota are *deferred*, but implicitly retried repeatedly for the duration of their periods until either run or *skipped*
- Can distribute searches across a cluster of machines

.conf2015

# Splunk Scheduler


# Deferred vs. Skipped

- As mentioned, searches over quota are *deferred*, but are implicitly retried repeatedly for the duration of their periods until either run or *skipped*



# Lag

- “Lag” is the difference between a search’s *dispatch* time and its *scheduled time*



- Non-zero lag means scheduler is over-subscribed (at least temporarily)
- Causes delays and may cause skipping
- May be mitigated by *schedule windows* (new in 6.3 — more later)

# limits.conf Settings

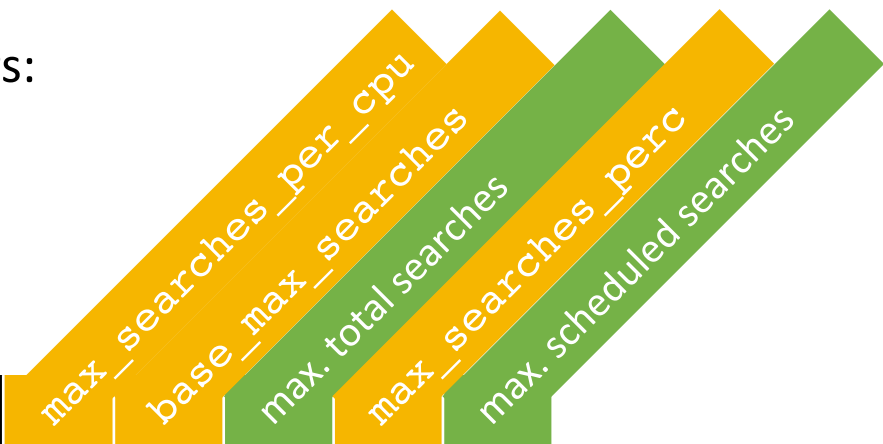
- **base\_max\_searches**: A constant added to max. total searches (default = 6)
- **max\_searches\_per\_cpu**: Maximum number of concurrent searches per CPU (default = 1)
- Given those, the total maximum number of concurrent searches allowed is:

$$\begin{aligned} \text{max. total searches} &= \text{max\_searches\_per\_cpu} \times \text{number of CPUs} \\ &+ \text{base\_max\_searches} \end{aligned}$$

- **max\_searches\_perc**: Maximum number of concurrent searches the scheduler can run as a percentage of max. total searches (default = 50)

# limits.conf Settings

- Some example numbers:



CPU	max_searches_per_cpu	base_max_searches	max. total searches	max_searches_perc	max. scheduled searches
1	1	6	7	50%	3
8	1	6	14	50%	7
64	1	6	70	50%	35

Legend:

- .conf setting
- Calculated

.conf2015

# Scheduler Improvements In Splunk Enterprise 6.3

splunk>

# How the Splunk Scheduler Works

1. For each search, calculate the next run-time of the search
2. Place all searches in a map<search\_id,next\_runtime>
3. Enter main loop:
  - A. For each search, if its next run-time  $\leq now$ , add it to the candidate search list
  - B. Randomly shuffle the candidate list
  - C. For each candidate search, calculate its *priority score*
  - D. Sort all candidate searches by priority score
  - E. For each candidate search, if it doesn't exceed quota, run it, calculate the *next* run-time of the search, and update the map



# Priority Scoring

- **Problem in 6.2:** Simple single-term priority scoring could result in saved search lag, skipping, and starvation (under CPU constraint)
- **Solution in 6.3:** Better multi-term priority scoring mitigates problems and improves performance by 25%

```
score(j) = next_runtime(j)
          + average_runtime(j) × priority_runtime_factor
          − skipped_count(j) × period(j) × priority_skipped_factor
          + schedule_window_adjustment(j)
```

# Schedule Windows

- **Problem in 6.2:** Scheduler can not distinguish between searches that (A) *really should* run at a specific time (just like cron) from those that (B) don't have to. This can cause lag or skipping
- **Solution in 6.3:** Give a *schedule window* to searches that don't have to run at specific times  
**Example:** For a given search, it's OK if it starts running sometime between midnight and 6am, but you don't really care when specifically
- A search with a window helps *other* searches
- Search windows *should not* be used for searches that run every minute
- Search windows *must* be less than a search's period

# Variable `max_searches_perc`

- **`max_searches_perc`**: Maximum number of concurrent searches the scheduler can run as a percentage of max. total searches (default = 50)
- **Problem in 6.2**: Constant value limits scheduler during off-peak times
- **Solution in 6.3**: Allow **`max_searches_perc`** to vary by time or day:

```
max_searches_perc = 50
```

```
# Allow value to be 75 anytime on weekends.
```

```
max_searches_perc.1 = 75
```

```
max_searches_perc.1.when = * * * * 0,6
```


```
# Allow value to be 90 between midnight and 5am.
```

```
max_searches_perc.2 = 90
```

```
max_searches_perc.2.when = * 0-5 * * *
```

# Continuous Scheduled Searches “Catch-up”

- **Problem in 6.2:** Continuous Scheduled Searches (CSSs) are missed due to Splunk downtime creating data gaps
- **Solution in 6.3:** By remembering last execution time, missed CSSs are run as soon as Splunk comes back up to fill in data gaps



.conf2015

# Splunk Enterprise 6.3 Test Results

splunk>

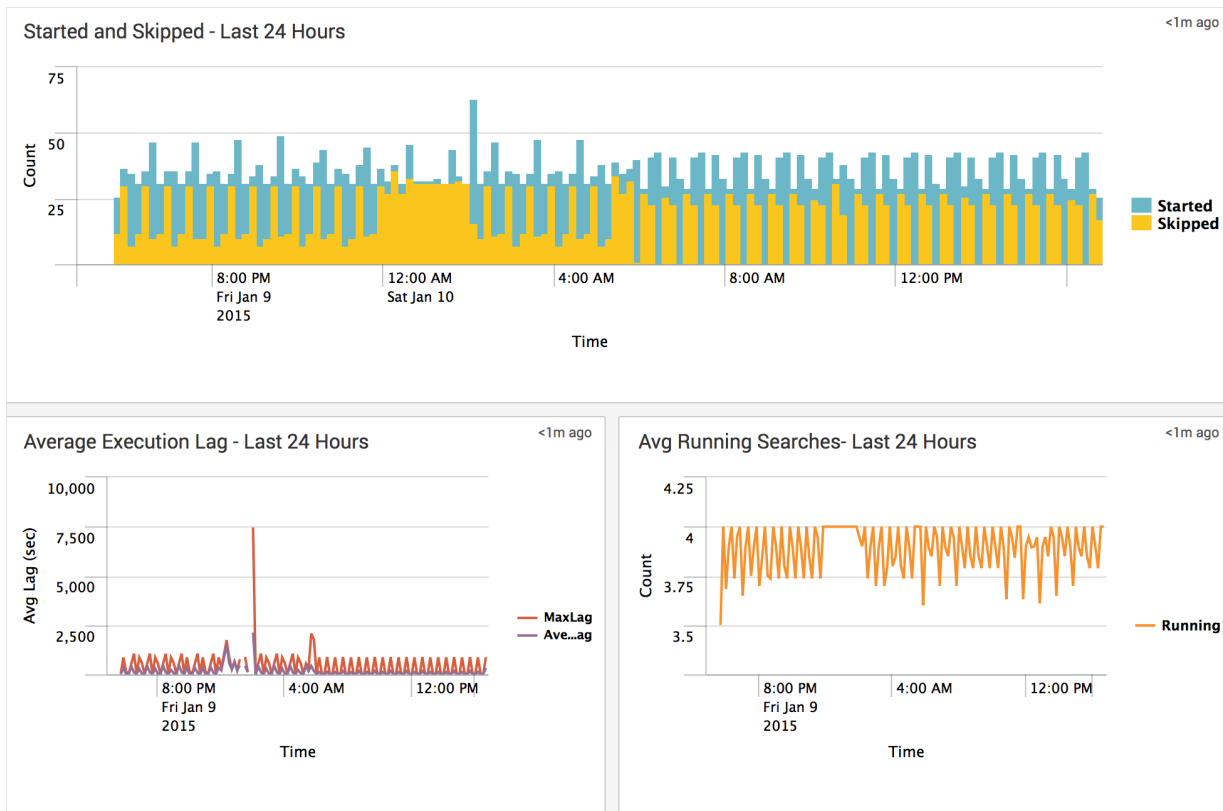
# Testing Conditions

- Stand-alone instance of Splunk (non-cluster)
- Given lots of searches and intentionally hamstrung by a low (15%) `max_searches_perc` to test performance under duress
- In production, for the same number of searches, you should have a higher `max_searches_perc`, a bigger CPU, or a search cluster

# Splunk 6.2 Test Results

## Things to notice:

- Started/skipped forms pattern: same searches are run in same order
- From 12–2am, many searches are skipped
- Average running searches is erratic



# Splunk 6.3 Test Results


## Things to notice:

- Started/skipped forms less of a pattern: the search order is being perturbed
- From 12–2am, fewer searches are skipped
- Average running searches is mostly constant






# Splunk 6.2 Test Results



## Things to notice:

- Number skipped gets worse in the order as the searches are in savedsearches.conf
- The “sample” searches, despite being very short, are almost always skipped

# Splunk 6.3 Test Results



## Things to notice:

- Number skipped is much more evenly distributed regardless of savedsearches.conf
- The “sample” searches are now almost never skipped (which is the sensible thing to do)

# Takeaways

- Splunk Enterprise 6.3 adds better *priority scoring* and *search windows* for much improved saved search scheduling by at least 25%
- For infrequent searches (hourly, daily, etc.) use *schedule windows*
- Use the built-in scheduler performance reports (under *Activity > System Activity > Scheduler*) to monitor performance: lots of skipped searches or high lag is bad
- If, despite tuning, you still have frequently skipped searches or high lag, then you probably need a bigger CPU or more nodes in your cluster

.conf2015

Questions



.conf2015

THANK YOU

splunk>