



构建可视化交互大数据查询平台

Admaster 向磊

大数据的交互式查询

1. 交互式查询很重要

我们不能要求所有人使用Linux命令行
减轻开发人员压力
需求即结果

2. 可选择余地很少

HUE, Qubole, phpHiveAdmin, Shib...

3. 难点

作业的提交
作业监控
集群监控
操作HDFS

为什么要开发交互式平台

一个最根本的原因是，不是所有人都会用Linux shell

不是所有人都能从shell操作Hadoop, Hive, Spark..., 而且这样不安全

让专业的人做专业的事，让产品经理自己去写无聊的SQL吧

一个数据分析作业在运行以前，可以在界面里先进行测试

界面化交互式平台可以解决以上问题

别人开发的东西难用，不会维护

最重要的事情提高数据分析工作的开发效率，干什么都要快

各平台对比

	HUE	phpHiveAdmin	Qubole
开源	Yes	Yes	No
开发语言	Python/Django	php/CodeIgniter	Python/java SDK
部署难度	Easy	Medium	AWS
功能	多/强大	Hive	Hive
二次开发	Python+Java	容易	Only SDK
Spark支持	Y	N	N
集群作业监控	Y	N	?
HDFS管理	Y	Half	?
可维护性	较难	矛盾	找Qubole
架构	B/S+C/S	B/S	Unknown

典型案例

一个典型的例子是我在那家涨停N次以后停牌的公司做phpHiveAdmin的时候不要问我关于那家公司股票的问题，我没有的

结果出乎意料

1. 极大的提高了在Hive上做数据分析的效率
2. 进而极大的提高了数据的产出效率
3. 开源后有N多国内外公司使用，并提交patch
4. 但是我仍然建议最好自己写一个
5. 这是一个公司从核心技术层面深入了解Hadoop和Spark的好机会
6. 将Spark/Hadoop/HBase/Hive整合为一个统一管理的可视化数据平台，大幅度降低分别管理和维护的成本。
7. 大幅度降低数据分析业务开发的周期和成本。

因何降低成本？

phpHiveAdmin

HQL查询

浏览HDFS

HiveQL模板

用户管理

历史查询

登出

clean_mac

default

mac

raw_mac

stranger

test

✕ 删除数据库

添加数据表

	表名
<input type="checkbox"/>	tb_users_history
<input type="checkbox"/>	sys_oui_s
<input type="checkbox"/>	sys_oui
<input type="checkbox"/>	radacct

结果集下载

7 Bytes

count(*)

全选 / 反选

把N多数据分析作业以众包形式分解，以提高分析作业的开发效率

phpHiveAdmin

HQL查询

浏览HDFS

HiveQL模板

用户管理

历史查询

登出

tb_users_history

sys_oui_s

sys_oui

radacct

openssid

返回 default

id	wlanusermac	wlanuserip	wlanapmac	wlanbssid	wlanssid	dev_id	dev_group
bigint	string	string	string	string	string	string	string
3	48:74:6e:8c:4f:49	192.168.78.126	84:82:f4:18:ef:e4	86:82:f4:18:ef:e5	WX-TTLY	1297	10035
4	48:74:6e:8c:4f:49	192.168.78.126	84:82:f4:18:ef:e4	86:82:f4:18:ef:e5	WX-TTLY	1297	10035

select * from default.tb_users_history wh

click to select

HQL 校验器

default.tb_users_history

MAP

REDUCE

2015-12-31 01:23:23,390 Stage-1 map = 92%, reduce = 25%, Cumulative CPU 28.39 sec
2015-12-31 01:23:20,245 Stage-1 map = 91%, reduce = 25%, Cumulative CPU 27.51 sec
2015-12-31 01:23:17,101 Stage-1 map = 87%, reduce = 25%, Cumulative CPU 25.41 sec
2015-12-31 01:23:13,971 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 23.45 sec
2015-12-31 01:23:06,678 Stage-1 map = 55%, reduce = 0%, Cumulative CPU 17.6 sec
2015-12-31 01:23:04,587 Stage-1 map = 51%, reduce = 0%, Cumulative CPU 16.41 sec
2015-12-31 01:23:03,532 Stage-1 map = 30%, reduce = 0%, Cumulative CPU 14.36 sec
2015-12-31 01:23:02,486 Stage-1 map = 16%, reduce = 0%, Cumulative CPU 6.33 sec
2015-12-31 01:22:50,937 Stage-1 map = 0%, reduce = 0%

Hadoop job information for Stage-1: number of mappers: 4; number of reducers: 1
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1446042486419_59074
Starting Job = job_1446042486419_59074, Tracking URL = http://rm.hadoop:8088/proxy/application_1446042486419_59074/
set mapreduce.job.reduces=
In order to set a constant number of reducers:
set hive.exec.reducers.max=

获取结果

提交

自己开发一个难吗？

关键在了解各种接口及对接口数据的获取

	HDFS	YARN/MR	Hive	Pig	HBase	Spark
Restful	Y	Y	N	N	Y	1.4 later
Thrift	N	N	Y	N	Y	SparkSQL
CLI	Y	Y	Y	Y	Y	Y

需要各接口的整合开发，首先以HDFS为例：

Hadoop提供基于HTTP/HTTPS的HDFS访问接口

- 同时支持对HDFS的读写操作
- 可以从程序或脚本中访问
- 可以用命令行工具如curl或wget来实现数据

文件的上传下载

- 启用webhdfs或httpfs
- webHDFS为内置，但不支持HA
- httpfs为外置tomcat，但支持HA
- 题外话，集群对拷也可以用webhdfs方式

REST = REpresentational State Transfer

Restful HDFS

通过浏览器或CURL访问HDFS

`curl -i -L "http://x.x.x.x:50070/webhdfs/v1/?OP=LISTSTATUS&user.name=hdfs"`

HTTP GET组内OP参数可以列目录，文件状态.....

HTTP PUT组内OP参数可以创建文件夹，改名，设权限.....

HTTP POST组内OP参数可以APPEND, CONCAT, TRUNCATE

HTTP DELETE组内OP参数可以删除文件和快照.....

```
{
  "FileStatuses": [
    {
      "FileStatus": {
        "accessTime": 0, "blockSize": 0, "childrenNum": 6, "fileId": 16389, "group": "hadoop", "length": 0, "modificationTime": 1434635241699, "owner": "yarn", "pathSuffix": "app-logs", "permission": "777", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 3, "fileId": 16403, "group": "hdfs", "length": 0, "modificationTime": 1430373809715, "owner": "hdfs", "pathSuffix": "apps", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 2, "fileId": 145301, "group": "hdfs", "length": 0, "modificationTime": 1435719592612, "owner": "hdfs", "pathSuffix": "data", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 0, "fileId": 437109, "group": "hdfs", "length": 0, "modificationTime": 1434351238766, "owner": "flume", "pathSuffix": "flume", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 1, "fileId": 16395, "group": "hdfs", "length": 0, "modificationTime": 1430245449189, "owner": "hdfs", "pathSuffix": "hdp", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 1, "fileId": 2268282, "group": "hdfs", "length": 0, "modificationTime": 1447818808124, "owner": "hdfs", "pathSuffix": "home", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 1, "fileId": 16390, "group": "hdfs", "length": 0, "modificationTime": 1430245405207, "owner": "mapred", "pathSuffix": "mapred", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 2, "fileId": 16392, "group": "hdfs", "length": 0, "modificationTime": 1430245405307, "owner": "hdfs", "pathSuffix": "mr-history", "permission": "777", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 1, "fileId": 16400, "group": "hdfs", "length": 0, "modificationTime": 1437481112419, "owner": "hdfs", "pathSuffix": "system", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 1, "fileId": 2282779, "group": "hdfs", "length": 0, "modificationTime": 1447926770271, "owner": "hdfs", "pathSuffix": "tmpdata1", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 8, "fileId": 16386, "group": "hdfs", "length": 0, "modificationTime": 1448362128914, "owner": "hdfs", "pathSuffix": "tmp", "permission": "777", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"},
        "accessTime": 0, "blockSize": 0, "childrenNum": 13, "fileId": 16387, "group": "hdfs", "length": 0, "modificationTime": 1436341722627, "owner": "hdfs", "pathSuffix": "user", "permission": "755", "replication": 0, "storagePolicy": 0, "type": "DIRECTORY"}
  ]
}
```

`curl -i -X PUT`

`"http://x.x.x.x:50070/webhdfs/v1/TokyoHot?op=MKDIRS&user.name=xianglei"`

Doc: <http://hadoop.apache.org/docs/r2.7.1/hadoop-project-dist/hadoop-hdfs/WebHDFS.html>

RESTful YARN/MR

通过curl访问YARN/MapReduce

```
curl -i -L "http://x.x.x.x:8088/ws/v1/cluster/info"
```

```
HTTP/1.1 200 OK
Cache-Control: no-cache
Expires: Tue, 29 Dec 2015 13:38:59 GMT
Date: Tue, 29 Dec 2015 13:38:59 GMT
Pragma: no-cache
Expires: Tue, 29 Dec 2015 13:38:59 GMT
Date: Tue, 29 Dec 2015 13:38:59 GMT
Pragma: no-cache
Content-Type: application/json
Transfer-Encoding: chunked
Server: Jetty(6.1.26.hwx)

{"clusterInfo":{"id":1451059230551,"startedOn":1451059230551,"state":"STARTED",
"haState":"ACTIVE","rmStateStoreName":"org.apache.hadoop.yarn.server.resourcemana
ger.recovery.ZKRMStateStore","resourceManagerVersion":"2.6.0.2.2.6.0-2800","reso
urceManagerBuildVersion":"2.6.0.2.2.6.0-2800 from acb70ecfae2c3c5ab46e24b0caebce
aec16fdcd0 by jenkins source checksum 65301b226e4b9e8135f2f93f9887ea63","resourc
eManagerVersionBuiltOn":"2015-05-18T20:28Z","hadoopVersion":"2.6.0.2.2.6.0-2800"
,"hadoopBuildVersion":"2.6.0.2.2.6.0-2800 from acb70ecfae2c3c5ab46e24b0caebceaec
16fdcd0 by jenkins source checksum a25c30f622eb057f47e2155f78dba5e","hadoopVersi
onBuiltOn":"2015-05-18T20:21Z"}}xianglei@xianglei-u303ub:~$
```

可以获取所有我们想要的作业相关信息，也可以提交作业，不过，首先你得把作业的jar包放进HDFS，然后再写一个json文件来提交。详情稍后

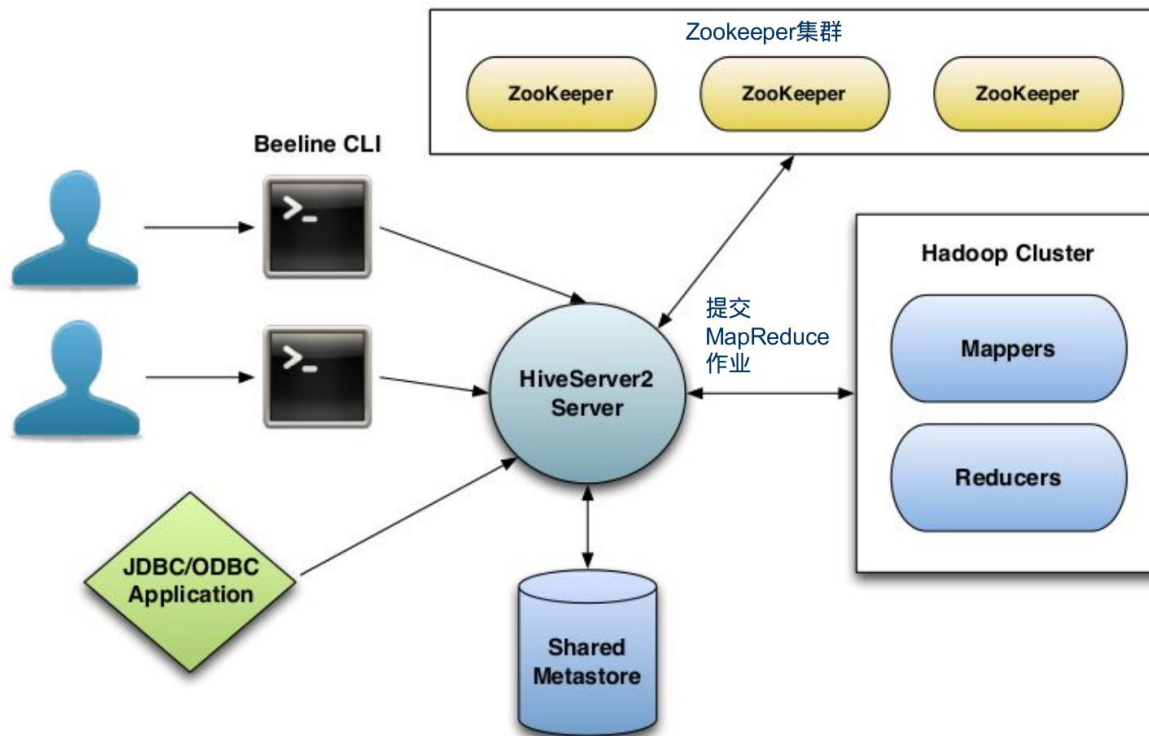
Doc: <http://hadoop.apache.org/docs/r2.7.1/hadoop-yarn/hadoop-yarn-site/ResourceManagerRest.html>

关于Hive和HBase的Thrift

Hive, HBase, SparkSQL都可以通过thrift / thrift2来访问

- 实际上SparkSQL用的就是Hive的接口
- thrift2实际上是thrift接口的一个增强版，增加了安全性(cyrus-sasl authentication)，并发性(Zookeeper Quorum)...

Hive有个HiveMetastore的9083端口的什么鬼，还有个10000端口的HiveServer / HiveServer2的什么鬼。



Hive thrift/2能做什么？

可以获取数据库，表，字段定义

可以提交SQL作业

可以创建数据库，表，字段定义

可以修改库，表，字段定义

可以创建修改分区，桶，和其他一切操作

```
public $hive_host;  
public $hive_port;  
public $socket;  
public $transport;  
public $protocol;  
public $hive;
```

```
public function __construct()  
{  
    parent::__construct();  
    $GLOBALS['THRIFT_ROOT'] = __DIR__ . "/../libs/";  
    include_once $GLOBALS['THRIFT_ROOT'] . 'packages/hive_service/ThriftHive.php';  
    include_once $GLOBALS['THRIFT_ROOT'] . 'transport/TSocket.php';  
    include_once $GLOBALS['THRIFT_ROOT'] . 'protocol/TBinaryProtocol.php';  
  
    $this->hive_host = $this->config->item('hive_host');  
    $this->hive_port = $this->config->item('hive_port');  
    $this->socket = new TSocket($this->hive_host, $this->hive_port);  
    $this->socket->setSendTimeout(30000);  
    $this->socket->setRecvTimeout(30000);  
    $this->transport = new TBufferedTransport($this->socket);  
    $this->protocol = new TBinaryProtocol($this->transport);  
    $this->hive = new ThriftHiveClient($this->protocol);  
}
```

```
public function drop_database($db_name, $del = FALSE)  
{  
    try  
    {  
        $this->transport->open();  
        $this->hive->drop_database($db_name, $del);  
        $this->transport->close();  
    }  
    catch (Exception $e)  
    {  
        echo 'Caught exception: '. $e->getMessage(). "\n";  
    }  
}
```

```
public function create_database($db_name, $db_desc = '')  
{  
    $sql = "CREATE DATABASE IF NOT EXISTS ".$db_name." COMMENT '".$db_desc.'";"  
    try  
    {  
        $this->transport->open();  
        $this->hive->execute($sql);  
        $this->transport->close();  
        return $sql;  
    }  
    catch (Exception $e)  
    {  
        echo 'Caught exception: '. $e->getMessage(). "\n";  
    }  
}
```

Hive的thrift/2应该做什么？

除了提交SQL作业之外的所有事情

之前不是说thrift可以用来提交HQL查询吗？

Why Not?

Thrift/2无法返回作业的状态和进度，一个MR可能会延续几分钟或几小时，即便是你用Spark引擎跑Hive，一个大任务也会花费很长时间，坐在那什么也不干，只等着屏幕刷新是很无趣也是不可被接受的。特别是你还不知道它多长时间才会返回结果。

另外，thrift也不支持并发查询，thrift2虽然支持并发，但需要配置Zookeeper。

How to do?

我们在开发交互查询时，可以用Thrift来管理Hive/SparkSQL的元数据，但查询可以用别的方式。

使用CLI方式提交也许会更好

https://github.com/xianglei/phpHiveAdmin/blob/master/application/models/hive_model.php
774 - 880行

原理：在网页中单独开启一个php子进程，将SQL提交到CLI接口，然后通过子进程获取stdout和stderr的输出。然后将获取内容倒排后以JSON数组形式发给网页，网页只要不断的刷新就可以了。

其实更好的办法也许是用WebSocket，但是我对前端没那么熟。

One Notice:

由于你的webserver使用的用户也许没有配置HADOOP相关的环境变量，所以你需要把命令行里加上输出HADOOP_HOME, JAVA_HOME和HIVE_HOME的三个环境变量的export。

```
$descriptorspec = array(
    0 => array("pipe", "r"), // stdin is a pipe that the child will read from
    1 => array("pipe", "w"), // stdout is a pipe that the child will write to
    2 => array("pipe", "w") // stderr is a file to write to
);

$pipes= array();

$process = proc_open($command, $descriptorspec, $pipes);

$output= "";

if (!is_resource($process))
{
    return false;
}

#close child's input imidiately
fclose($pipes[0]);

stream_set_blocking($pipes[1],0);
stream_set_blocking($pipes[2],0);

$todo= array($pipes[1],$pipes[2]);
try
{
    $fp = fopen($file_name, "w");
    #fwrite($fp,$time_stamp."\n\n");
    while( true )
    {
        $read= array();
        #if( !feof($pipes[1]) ) $read[]= $pipes[1];
        if( !feof($pipes[$type]) )
            $read[]= $pipes[$type]; // get system stderr on real time

        if (! $read)
        {
            break;
        }

        $ready= stream_select($read, $write=NULL, $sex= NULL, 2);

        if ($ready === false)
        {
            break; #should never happen - something died
        }

        foreach ($read as $r)
        {
            $s= fread($r,128);
            $output .= $s;
            fwrite($fp,$s);
        }
    }

    fclose($fp);
}
```

MR/Pig/Mahout/Spark

1. 都可以使用CLI或Restful方式来提交并运行作业，作业状态通过读取stdout和stderr或HTTP GET返回前端。

2. 把作业结果保存在文件中，可以保存在本地或保存在HDFS，在一个数据库里记录每个作业的提交，这样，你即使关闭浏览器，也可以从历史信息中找到你之前提交的作业结果。

```
public function cli_query($sql, $finger_print)
{
    $this->load->model('utilities_model', 'utils');
    $LANG = " export LANG=" . $this->config->item('lang_set') . " ";
    $JAVA_HOME = " export JAVA_HOME=" . $this->config->item('java_home') . " ";
    $HADOOP_HOME = " export HADOOP_HOME=" . $this->config->item('hadoop_home') . " ";
    $HIVE_HOME = " export HIVE_HOME=" . $this->config->item('hive_home') . " ";
```

浏览HDFS HiveQL模板 用户管理 历史查询 登出

文件名	文件内容	文件大小
<input type="checkbox"/> admin_2015-06-18-15-29-36_7ec66eba5728ae7c7b645e94cc84b2f6f7cf66fe.log	USE ruian_mac;select count(distinct user_mac) from ruian_mac.ruian_mac_20150618 where from_unixtime(cast(log_time as int),'HH') >= "14"	136 Bytes
<input type="checkbox"/> admin_2015-06-18-15-27-33_1e6f14d398c81e116a106f076464963364190b03.log	USE ruian_mac;select from_unixtime(cast(log_time as int),'HH') , count(distinct user_mac) from ruian_mac.ruian_mac_20150618 where from_unixtime(cast(log_time as int),'HH') >= "14" group by log_time	198 Bytes
<input type="checkbox"/> admin_2015-06-18-15-24-50_a650fd2ac35642e9ec00c29f59ae3b55df8dba1.log	USE ruian_mac;select from_unixtime(cast(log_time as int),'HH') as time, count(distinct user_mac) as cnt from ruian_mac.ruian_mac_20150618 where from_unixtime(cast(log_time as int),'HH') >= "14" group by time order by cnt desc	228 Bytes
<input type="checkbox"/> admin_2015-06-18-15-21-03_ee24c9130a1ac5c9f530b623028f0ba809bb69b9.log	USE ruian_mac;select count(*) from ruian_mac.ruian_mac_20150618 where from_unixtime(cast(log_time as int),'HH') >= "14"	119 Bytes
<input type="checkbox"/> admin_2015-06-18-15-18-33_b3455841e1dea589ef9c030e6dd5e34d532140e6.log	USE ruian_mac;select count(*) from ruian_mac.ruian_mac_20150618 where from_unixtime(cast(log_time as int),'HH') >= "14"	119 Bytes
<input type="checkbox"/> admin_2015-06-18-15-15-43_372652cbdbed212dd83ec6968bca1b23744f7476.log	USE ruian_mac;select count(*) from ruian_mac.ruian_mac_20150618 where from_unixtime(log_time,"HH") >= "14"	106 Bytes
<input type="checkbox"/> admin_2015-06-18-15-11-09_79991cd78459ef61c3c335649ac00d71cb82732f.log	USE ruian_mac;select count(*) from ruian_mac.ruian_mac_20150618 where from_unixtime(log_time,"HH") >= "14"	106 Bytes
<input type="checkbox"/> lxuan_2015-06-18-13-21-03_e84c517acb2d124ef0b32b12ef37d3c37d65735c.log	USE ruian_mac;select * from ruian_mac.ruian_mac_20150618 group by user_mac	77 Bytes

```
}
catch (Exception $e)
{
    echo 'Caught exception: '. $e->getMessage(). "\n";
}
}
```

我们了解了这最基本的三种接口

RESTful

- 基于HTTP/HTTPS的json/xml接口

Thrift

- 基于thrift的接口，事实上，Hive的JDBC/ODBC也不过就是对 thrift 的再次封装

CLI

- 最普通的命令行接口，简单，却非常实用。另外，可以并发查询。

把这些接口都整合在一起

下面以Python为例，说明开发步骤。

Otherside: 为何不用Java写？

1. 我太懒，Java代码太啰嗦，人生苦短，我用Python
2. 我对Java属于会看不会写，理由参看上一条
3. RESTful的JSON可直接转换Python的dict类型处理, dict可直接转换JSON

构建可视化交互查询的步骤

1. 通过Thrift获取Hive/SparkSQL的库表定义，并以JSON形式返回前端

```
class GetDatabaseRestHandler(BaseHandler):
    @tornado.web.authenticated
    def get(self):
        config = Config.get_config()
        db = self.get_argument('db')
        if config['ServerPort']:
            try:
                transport = TSocket.TSocket(config['hive_server_ip'], config['ServerPort'])
                transport = TTransport.TBufferedTransport(transport)
                protocol = TBinaryProtocol.TBinaryProtocol(transport)

                client = ThriftHive.Client(protocol)
                transport.open()

                dbi = client.get_database(db)

                transport.close()
                '''
                Database is an instance,
                dbi.name str
                dbi.parameters dict
                privileges ?
                ownerType int
                ownerName str
                locationUri str
                description str
                '''
                dbj = {}
                dbj['name'] = dbi.name
                dbj['parameters'] = dbi.parameters
                dbj['privileges'] = dbi.privileges
                dbj['ownerType'] = dbi.ownerType
                dbj['ownerName'] = dbi.ownerName
                dbj['locationUri'] = dbi.locationUri
                dbj['description'] = dbi.description
                #print dbi
                self.write(json.dumps(dbj).encode('unicode_escape'))
            except Thrift.TException, tx:
                self.write('%s' % (tx.message))
        elif config['Server2Port']:
            return
        else:
            self.write('Not set Hive thrift port')
        return
```

构建可视化交互查询的步骤

2. 在网页上搞个输入框用来输入SQL和提交查询的按钮
3. 把HQL或pig脚本写入到文本文件，既是提交查询所需，也是罪证
4. 提交查询，然后写入到stdout和stderr输出临时文件中

```
def write_hive_log(self, hql):
    JAVA_HOME=self.config['JAVA_HOME']
    HADOOP_HOME=self.config['HADOOP_HOME']
    HIVE_HOME=self.config['HIVE_HOME']
    SPARK_SQL_HOME=self.config['SPARK_SQL_HOME']
    logger = Logger()

    try:
        hql_file = open(self.hql_filename, 'a')
        hql = 'set hive.cli.print.header=true;' + 'use ' + self.db_name + ';' + hql
        hql_file.write(hql)
        hql_file.close()
        if self.config['spark_sql'] == '1':
            cmd = 'export JAVA_HOME=' + JAVA_HOME + '; export HADOOP_HOME=' + HADOOP_HOME + '; export HIVE_HOME=' + HIVE_HOME + '; ' + 'export SPARK_HOME=' + SPARK_SQL_HOME +
            '; ' + SPARK_SQL_HOME + '/bin/spark-sql -f ' + self.hql_filename + ' 2>' + self.log_filename + ' 1>' + self.tmp_filename + ' &'
        else:
            cmd = 'export JAVA_HOME=' + JAVA_HOME + '; export HADOOP_HOME=' + HADOOP_HOME + '; export HIVE_HOME=' + HIVE_HOME + '; ' + HIVE_HOME + '/bin/hive -f ' +
            self.hql_filename + ' 2>' + self.log_filename + ' 1>' + self.tmp_filename + ' &'
            #debug command below
            #cmd = '/usr/bin/hive 2>' + self.log_filename + ' 1>' + self.tmp_filename + ' &'
            print cmd
            os.system(cmd)

        logger.info('Executor : write_hive_log : ' + self.log_filename)

        return self.log_filename
    except IOError, e:
        logger.error('Executor : write_hive_log : ' + e.replace('\n', ''))
        return e
```

构建可视化交互查询的步骤

5. 读取stdout和stderr的输出临时文件，并且返回给前端。

```
def read_hive_log(self, log_filename):  
    logger = Logger()  
    try:  
        log_file = open(log_filename, 'r')  
        f = log_file.readlines()  
        log_file.close()  
  
        logger.info('Executor : read_hive_log : ' + log_filename)  
  
        return f #return is a list  
    except IOError, e:  
        logger.error('Executor : read_hive_log : IOError' + log_filename)  
        return e
```

构建可视化交互查询的步骤

6. 用正则方式获取mr进度或其他内容，返回给前端做进度条用

```
def get(self, *args, **kwargs):
    logger = Logger()
    #log_filename = self.get_argument('log_filename')
    username = self.get_secure_cookie('username')
    finger_print = self.get_argument('finger_print')
    db_name = self.get_argument('db_name')

    config = Config.get_config()
    conn = sqlite3.connect(config['database'])
    cursor = conn.cursor()
    sql = 'select date_time from hive_history_jobs where finger_print = ?'
    cursor.execute(sql, (finger_print,))
    row = cursor.fetchone()
    date = row[0]

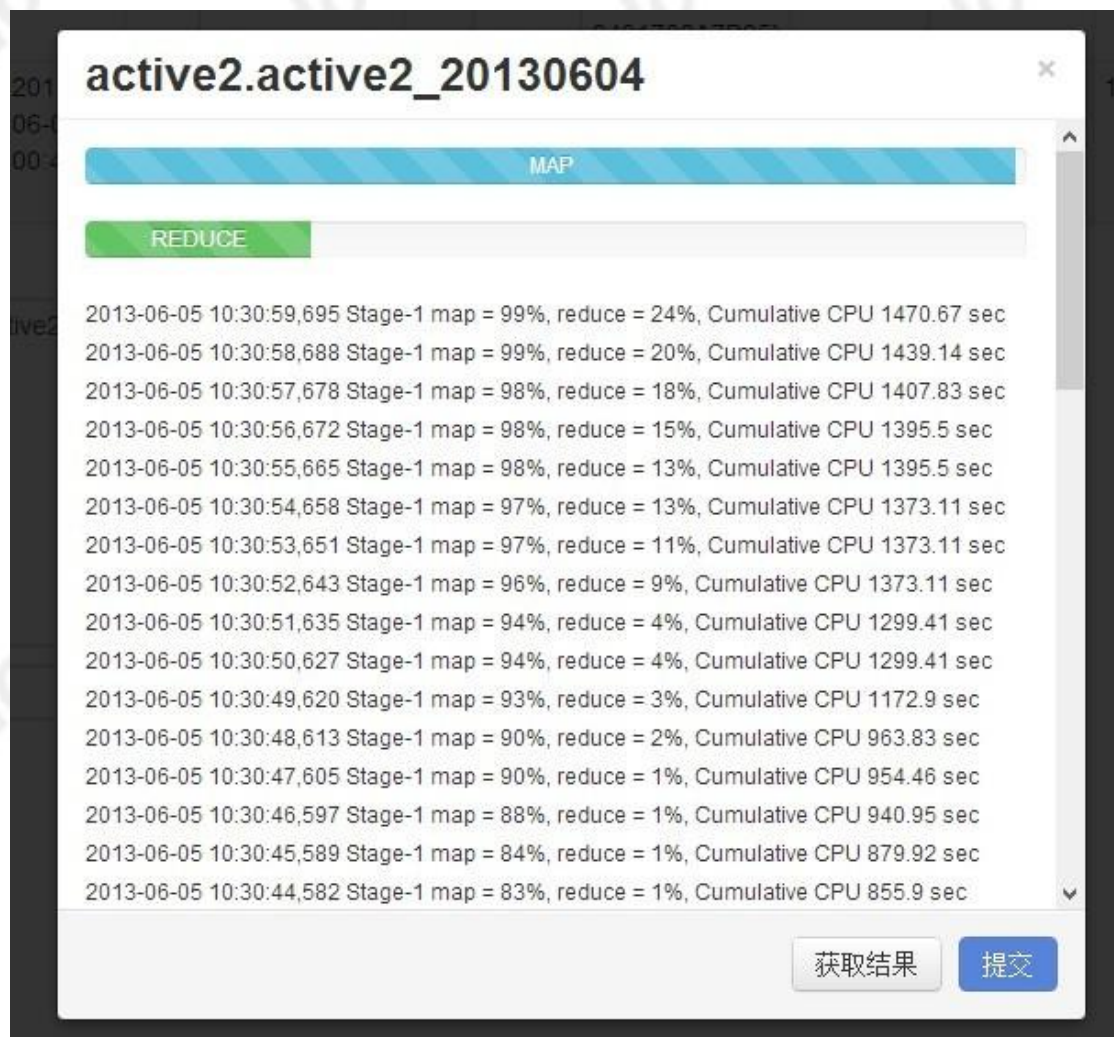
    run = Executor(username, finger_print, db_name)
    file_basename = username + '_' + date + '_' + finger_print
    log_filename = config['log_dir'] + file_basename + '.log'
    log = {}
    try:
        log_reverse = run.read_hive_log(log_filename)[::-1] # or log.reverse()
        if config['Tez'] == '0' and config['spark_sql'] == '0':
            m_start = log_reverse[0].find('map = ') + 6
            m_end = log_reverse[0].find('%')
            m_per = log_reverse[0][m_start: m_end]

            r_start = log_reverse[0].find('reduce = ') + 9
            r_end = log_reverse[0].rfind('%')
            r_per = log_reverse[0][r_start: r_end]

            if r_per.isdigit() is not True:
                r_per = 100
            if m_per.isdigit() is not True:
                m_per = 100
            log['map_per'] = m_per
            log['reduce_per'] = r_per
            log['log'] = log_reverse
        else:
            log['map_per'] = '0'
            log['reduce_per'] = '0'
            log['log'] = log_reverse
        logger.info('QueryExecutorHandler : get : ' + run.log_filename)
    except IOError, e:
        log = e.split('\n')
        logger.error('QueryExecutorHandler : get : ' + e.replace('\n', '') + run.log_filename + sql)
    self.write(json.dumps(log, ensure_ascii=False))
```

构建可视化交互查询的步骤

7. 前端JS渲染从后端来的stdout和stderr数据



MR/Pig/Mahout/Spark/HBase

整合HDFS和YARN/MR RESTful API，那么数据分析师就可以自己选择HDFS目录和数据创建自己分析所需要的表。甚至可以自己从本地或服务上传数据做分析。对于没有REST和Thrift的生态组件，我们可以采用CLI方式。

当然，再加上HCatalog，Pig也就可以使用Hive的元数据定义。

```
'''
send a hive cli execution with stderr redirect to log_filename and result to tmp_filename
return nothing
'''

def write_pig_log(self, pig):
    JAVA_HOME=self.config['JAVA_HOME']
    HADOOP_HOME=self.config['HADOOP_HOME']
    PIG_HOME = self.config['PIG_HOME']
    logger = Logger()

    try:
        pig_file = open(self.pig_filename, 'a')
        pig_file.write(pig)
        pig_file.close()
        cmd = 'export JAVA_HOME=' + JAVA_HOME + '; export HADOOP_HOME=' + HADOOP_HOME + '; export PIG_HOME=' + PIG_HOME + '; ' + PIG_HOME + '/bin/pig -f ' + self.pig_filename
+ ' 2>' + self.log_filename + ' 1>' + self.tmp_filename + ' &'
        #debug command below
        #cmd = '/usr/bin/hive 2>' + self.log_filename + ' 1>' + self.tmp_filename + ' &'
        print cmd
        os.system(cmd)

        logger.info('Executor : write_pig_log : ' + self.log_filename)

    except IOError, e:
        logger.error('Executor : write_pig_log : ' + e.replace('\n', ''))
        return e

'''
read from log_filename that redirected by write_pig_log
'''

def read_pig_log(self, log_filename):
    logger = Logger()
    try:
        log_file = open(log_filename, 'r')
        f = log_file.readlines()
        log_file.close()

        logger.info('Executor : read_pig_log : ' + log_filename)

        return f_#return is a list
    except IOError, e:
        logger.error('Executor : read_pig_log : IOError' + log_filename)
        return e
```

MR提交的特别之处

我们可以通过YARN的RESTful接口来提交作业，需要以下步骤

1. 通过HDFS RESTful把jar包传到HDFS

2. 提交一个HTTP POST获取app-id

curl -v -X POST 'http://localhost:8088/ws/v1/cluster/apps/new-application'

```
{
  application-id: application_1409421698529_0012",
  "maximum-resource-capability":{"memory":"8192", "vCores":"32"}
}
```

3. 获取ID后编写一个JSON并POST提交给RESTful接口

curl -v -X POST -d @example-submit-app.json -H "Content-type: application/json" 'http://x.x.x.x:8088/ws/v1/cluster/apps'

```
{
  "application-id":"application_1404203615263_0001",
  "application-name":"test",
  "am-container-spec":
  {
    "local-resources":
    {
      "entry":
      [
        {
          "key":"AppMaster.jar",
          "value":
          {
            "resource":"hdfs://hdfs-namenode:9000/user/testuser/DistributedShell/demo-app/AppMaster.jar",
            "type":"FILE",
            "visibility":"APPLICATION",
            "size": "43004",
            "timestamp": "1405452071209"
          }
        }
      ]
    },
    "commands":
    {
      "command":"{{JAVA_HOME}}/bin/java -Xmx10m org.apache.hadoop.yarn.applications.distributedshell.ApplicationMaster --container_memory 10 --container_vcores 1 --num_containers 1",
      "environment":
      {
        "entry":
        [
          {
            "key": "DISTRIBUTEDSHELLSCRIPTTIMESTAMP",
            "value": "1405459400754"
          },
          {
            "key": "CLASSPATH",
            "value": "{{CLASSPATH}}<CPS>./*<CPS>{{HADOOP_CONF_DIR}}<CPS>{{HADOOP_COMMON_HOME}}/share/hadoop/common/*<CPS>{{HADOOP_COMMON_HOME}}/share/hadoop/common/lib/*<CPS>{{HADOOP_COMMON_HOME}}/share/hadoop/common/lib/*<CPS>{{HADOOP_COMMON_HOME}}/share/hadoop/yarn/lib/*<CPS>{{HADOOP_YARN_HOME}}/share/hadoop/yarn/lib/*"
          },
          {
            "key": "DISTRIBUTEDSHELLSCRIPTLEN",
            "value": "6"
          },
          {
            "key": "DISTRIBUTEDSHELLSCRIPTLOCATION",
            "value": "hdfs://hdfs-namenode:9000/user/testuser/demo-app/shellCommands"
          }
        ]
      }
    }
  },
  "unmanaged-AM":"false",
  "max-app-attempts":"2",
  "resource":
  {
    "memory":"1024",
    "vCores":"1"
  },
  "application-type":"YARN",
  "keep-containers-across-application-attempts":"false"
}
```


MR提交的特别之处

然后你拿到ID号，还可以查看任务状态

curl

```
'http://localhost:8088/ws/v1/cluster/apps/application_1409421698529_0012/state'
```

或者蛋疼一下，再杀掉这个任务

```
curl -v -X PUT -d '{"state": "KILLED"}
```

```
"http://x.x.x.x:8088/ws/v1/cluster/apps/application_1409421698529_0012"
```

关于Spark作业RESTful提交与监控

监控

```
curl -i -L "http://x.x.x.x:4040/api/v1/applications/[app-id]/storage/rdd"
```

历史作业信息

```
curl -i -L "http://x.x.x.x:18080/api/v1/applications"
```

如果是 on YARN

```
curl -i -L "http://x.x.x.x:4040/api/v1/applications/[app-id]/[attempt-id]"
```

Doc: <http://spark.apache.org/docs/latest/monitoring.html>

Spark作业提交

ooyala公司有个叫spark-jobserver的项目

先sbt打包并upload 作业到spark

```
curl --data-binary @job-server-tests/target/scala-2.10/job-server-tests-$VER.jar localhost:8090/jars/test
```

然后

```
curl -d "input.string = a b c a b see"
```

```
'localhost:8090/jobs?appName=test&classPath=spark.jobserver.WordCountExample'
```

Spark返回:

```
{
  "status": "STARTED",
  "result": {
    "jobId": "5453779a-f004-45fc-a11d-a39dae0f9bf4",
    "context": "b7ea0eb5-spark.jobserver.WordCountExample"
  }
}
```

Doc: <https://github.com/spark-jobserver/spark-jobserver>

HBase

对于HBase来说界面化就更简单了，既支持restful，也支持thrift和cli，不再赘述

关于集群本身的监控

请使用HDFS/YARN/HBase/Spark的jmx接口和qry参数选择需要监控的内容。

```
5:8088/jmx
{
  "beans": [{
    "name": "java.lang:type=Memory",
    "modelerType": "sun.management.MemoryImpl",
    "HeapMemoryUsage": {
      "committed": 530055168,
      "init": 1054735680,
      "max": 1908932608,
      "used": 364548456
    },
    "NonHeapMemoryUsage": {
      "committed": 67764224,
      "init": 24576000,
      "max": 136314880,
      "used": 56575856
    },
    "ObjectPendingFinalizationCount": 0,
    "Verbose": false,
    "ObjectName": "java.lang:type=Memory"
  }, {
    "name": "java.lang:type=MemoryPool,name=PS Eden Space",
    "modelerType": "sun.management.MemoryPoolImpl",
    "CollectionUsage": {
      "committed": 131072000,
      "init": 264765440,
      "max": 673710080,
      "used": 0
    },
    "CollectionUsageThreshold": 0,
    "CollectionUsageThresholdCount": 0,
    "MemoryManagerNames": [ "PS MarkSweep", "PS Scavenge" ],
    "PeakUsage": {
      "committed": 264765440,
      "init": 264765440,
      "max": 714604544,
      "used": 264765440
    },
    "Usage": {
      "committed": 131072000,
      "init": 264765440,
      "max": 673710080,
      "used": 109995016
    }
  }
]
```

```
@Override
public void doGet(HttpServletRequest request, HttpServletResponse response) {
  try {
    if (!isInstrumentationAccessAllowed(request, response)) {
      return;
    }
    JsonGenerator jg = null;
    PrintWriter writer = null;
    try {
      writer = response.getWriter();

      response.setContentType("application/json; charset=utf8");
      response.setHeader(AccessControlAllowMethods, "GET");
      response.setHeader(AccessControlAllowOrigin, "");

      jg = jsonFactory.createJsonGenerator(writer);
      jg.disable(JsonGenerator.Feature.AUTO_CLOSE_TARGET);
      jg.useDefaultPrettyPrinter();
      jg.writeStartObject();

      // query per mbean attribute
      String getmethod = request.getParameter("get");
      if (getmethod != null) {
        String[] splitStrings = getmethod.split("\\:\\:");
        if (splitStrings.length != 2) {
          jg.writeStringField("result", "ERROR");
          jg.writeStringField("message", "query format is not as expected.");
          jg.flush();
          response.setStatus(HttpServletResponse.SC_BAD_REQUEST);
          return;
        }
        listBeans(jg, new ObjectName(splitStrings[0]), splitStrings[1],
          response);
        return;
      }

      // query per mbean
      String qry = request.getParameter("qry");
      if (qry == null) {
        qry = "*:*";
      }
      listBeans(jg, new ObjectName(qry), null, response);
    } finally {
      if (jg != null) {
        jg.close();
      }
      if (writer != null) {
        writer.close();
      }
    }
  }
}
```


Doc: <http://slaytanic.blog.51cto.com/2057708/1179108>

关于集群本身的监控

HDFS Counter

Total DFS Space: 16.29 GB
Free DFS Space: 14.41 GB
NonDFS Space: 1.88 GB
DFS Space: 28.01 KB
Free percent: 88 %
NonDFS percent: 12 %
DFS percent: 0 %

HDFS Usages



Exadoop.com

Free

NonDFS

#	Hostname	IP	HDFS Status
---	----------	----	-------------

1	node2	192.168.128.152
---	-------	-----------------

Free

Used: 28.01 KB / Total: 16.29 GB

<http://blog.csdn.net/fansy1990>

EasyHadoop 管理中心

首页

安装教程

配置教程

操作教程

节点监控

用户操作

退出

- 节点状态
- 集群状态
- CPU状态
- 内存状态
- 网络状态

Map/Reduce real time monitoring

Total Map Slots: 2
Total Reduce Slots: 2
Running Map Slots: 0
Running Reduce Slots: 0

Map/Reduce



Map

Free 2

Reduce

Free 2

#	主机名称	IP地址	Map状态	Map状态	Reduce状态	Reduce状态
1	easy-slave	102.165.1.91	Free	Used: 0 / Total: 2	Free	Used: 0 / Total: 2

开发这样一个平台需要多少人

全栈式？一个就够了

或者

一个牛逼前端加一个牛逼后端

平台开发的核心是要知道有这些接口可以使用。

平台最好做成前后端分离的方式，方便扩展新的接口。

Spark和**Hadoop**还留有一些隐藏的接口没有写入文档，需要看源码来发现或简化接口的调用。

Conclusion & Thanks a lot

It's very simple to build an AD-HOC query/management system on Hadoop & Spark.

My recently opensource project plan:

pyHUI - Hadoop User Interface based on python-tornado

Contact me @:

<https://github.com/xianglei>

or

<mailto:horseman@163.com>

C'ya

ADMaster