

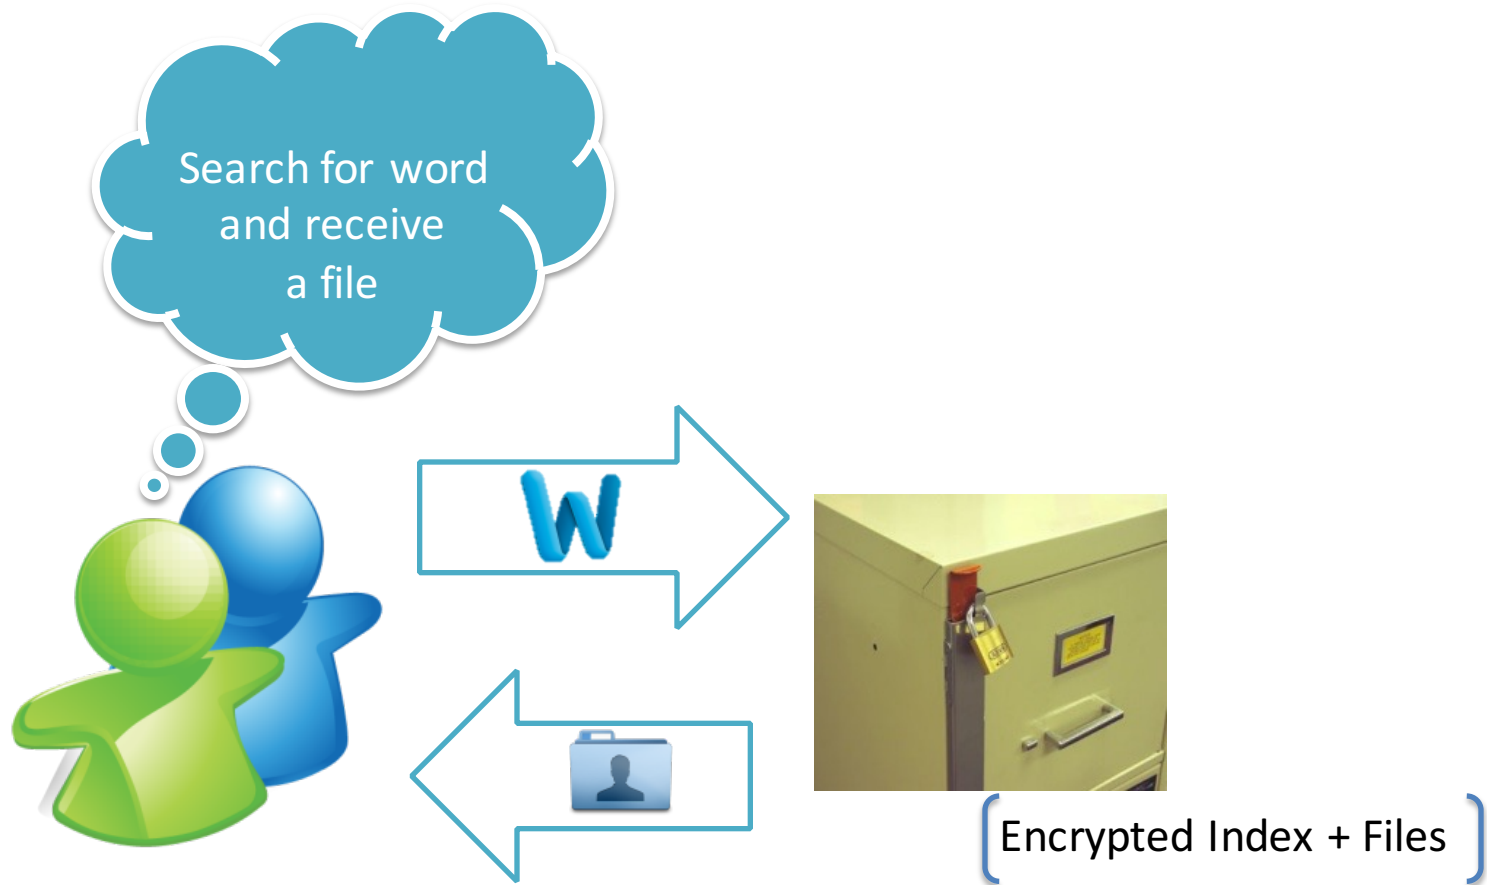
Constant-Time Dynamic Symmetric Searchable Encryption from Constrained Functional Encryption

Prof. Dr. Sebastian Gajek
NEC Research Labs and FUAS



is a game changer

Searchable Encryption (SENC)

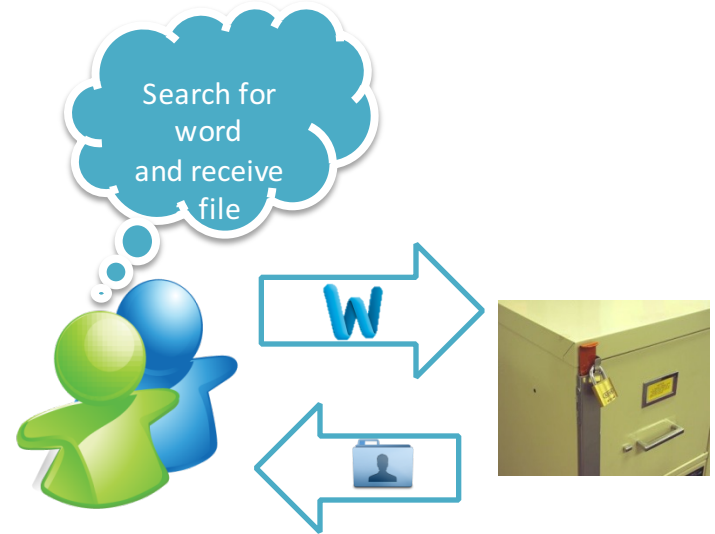


Designing a SENC system means to **juggle** with requirements



Privacy

- Pattern (Information leaked)
 - **One-Time¹ (weakest)**: Database learns after first query the search token
 - **Search**: Database does not learn that same search word is queried
 - **Access² (strongest)**: Database does not learn that same data is queried
- Attack Model (Database is our foe)
 - **Honest-but-curious (weakest)**: Honest, but tries to infer information from protocol executions (passive)
 - **Covert**: When dishonest, some odds to detect curiosity (rational)
 - **Malicious (strongest)**: Dishonest, tries everything to derive information (active)

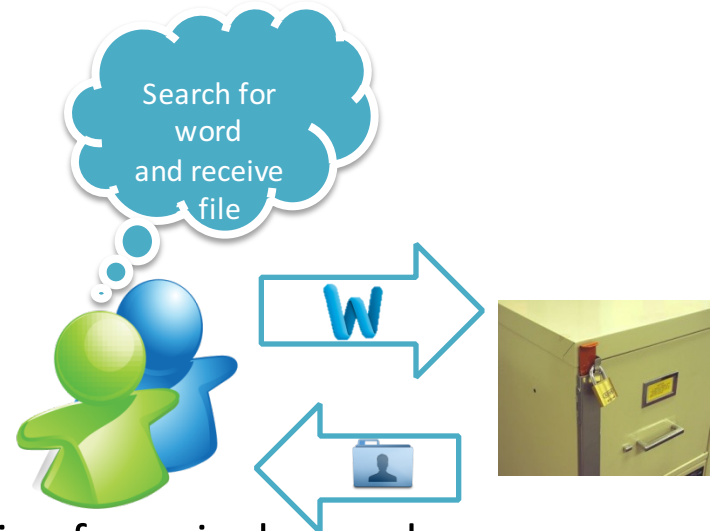


¹ Curtmola et al. Symmetric Searchable Encryption, CCS'06

² Goldreich, Ostrovsky: Software Protection and Simulation on Oblivious RAMs, STOC'95

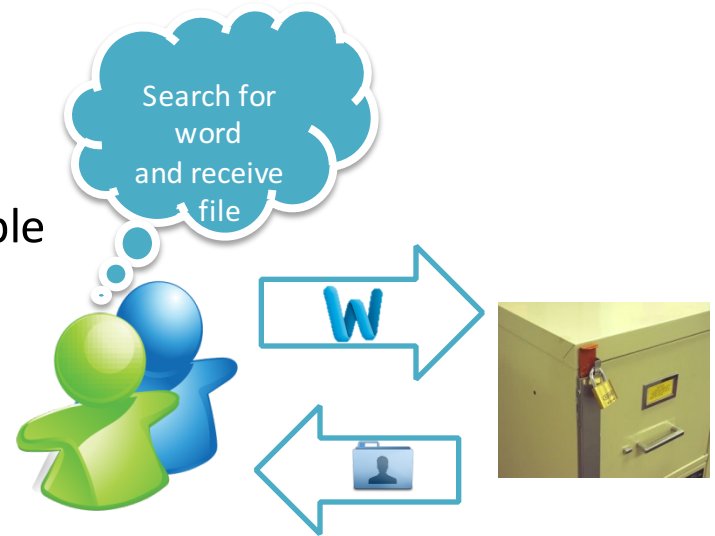
Functionality

- Database
 - **Static:** Database supports search and retrieval of (multiple) files
 - **Dynamic:** Database supports addition and deletion of (multiple) words and files
- Query language
 - **Single-word:** each token allows for searching for a single word
 - **Multi-word:** each token allows for searching multiple words (ideally, query for some CNF/DNF formula)
 - **Nearest-word:** each token allows for searching multiple words, each word w_i close to some word w_j , i.e. $|w_i - w_j| < \epsilon$ (e.g., range queries)



Scalability

- Performance
 - **Non-parallelizable:** no gain by sharing search over multiple clouds
 - **Parallelizable:** performance gain by multiple cloud
- Generality
 - **Specific:** SENC system is a mash-up of cryptographic algorithms and data structures (e.g. CryptDB)
 - **Self-contained:** SENC system is a framework of cryptographic algorithms and data structures



Our Schemes

Privacy	Honest-Curious	Covert	Malicious
<i>One-Time</i>	✓	✓	✓
<i>Search</i>	✓	✓	✓
<i>Access</i>	✗	✗	✗

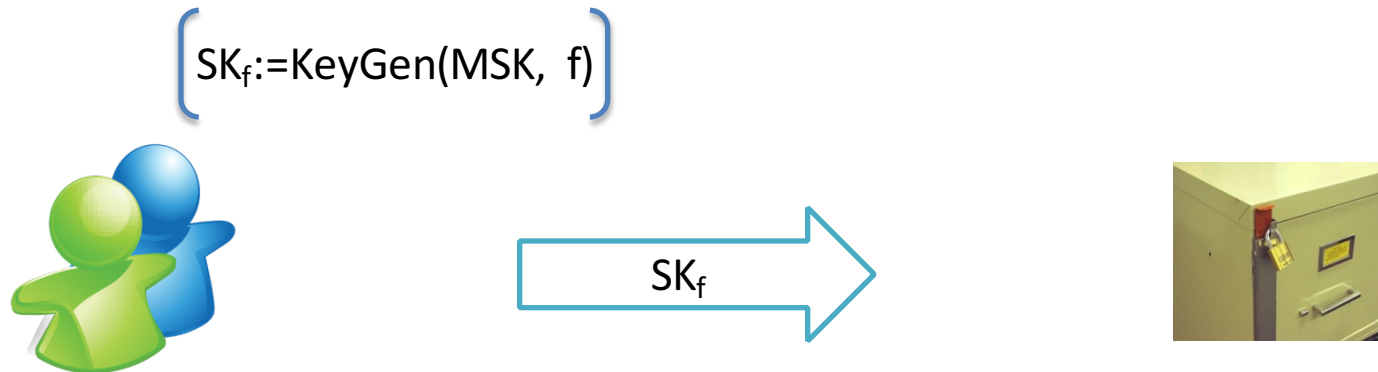

Functionality	Single-Word	Multiple-Word	Nearest-Word
<i>Search/Retrieve</i>	✓	✓	✓
<i>Add/Delete words</i>	✓	✓	✓
<i>Add/Delete files</i>	✓	✓	✓

Scalability	Parallelizable	Self-Contained	
<i>Performance</i>	✓		
<i>Generality</i>		✓	

Key Idea(s)

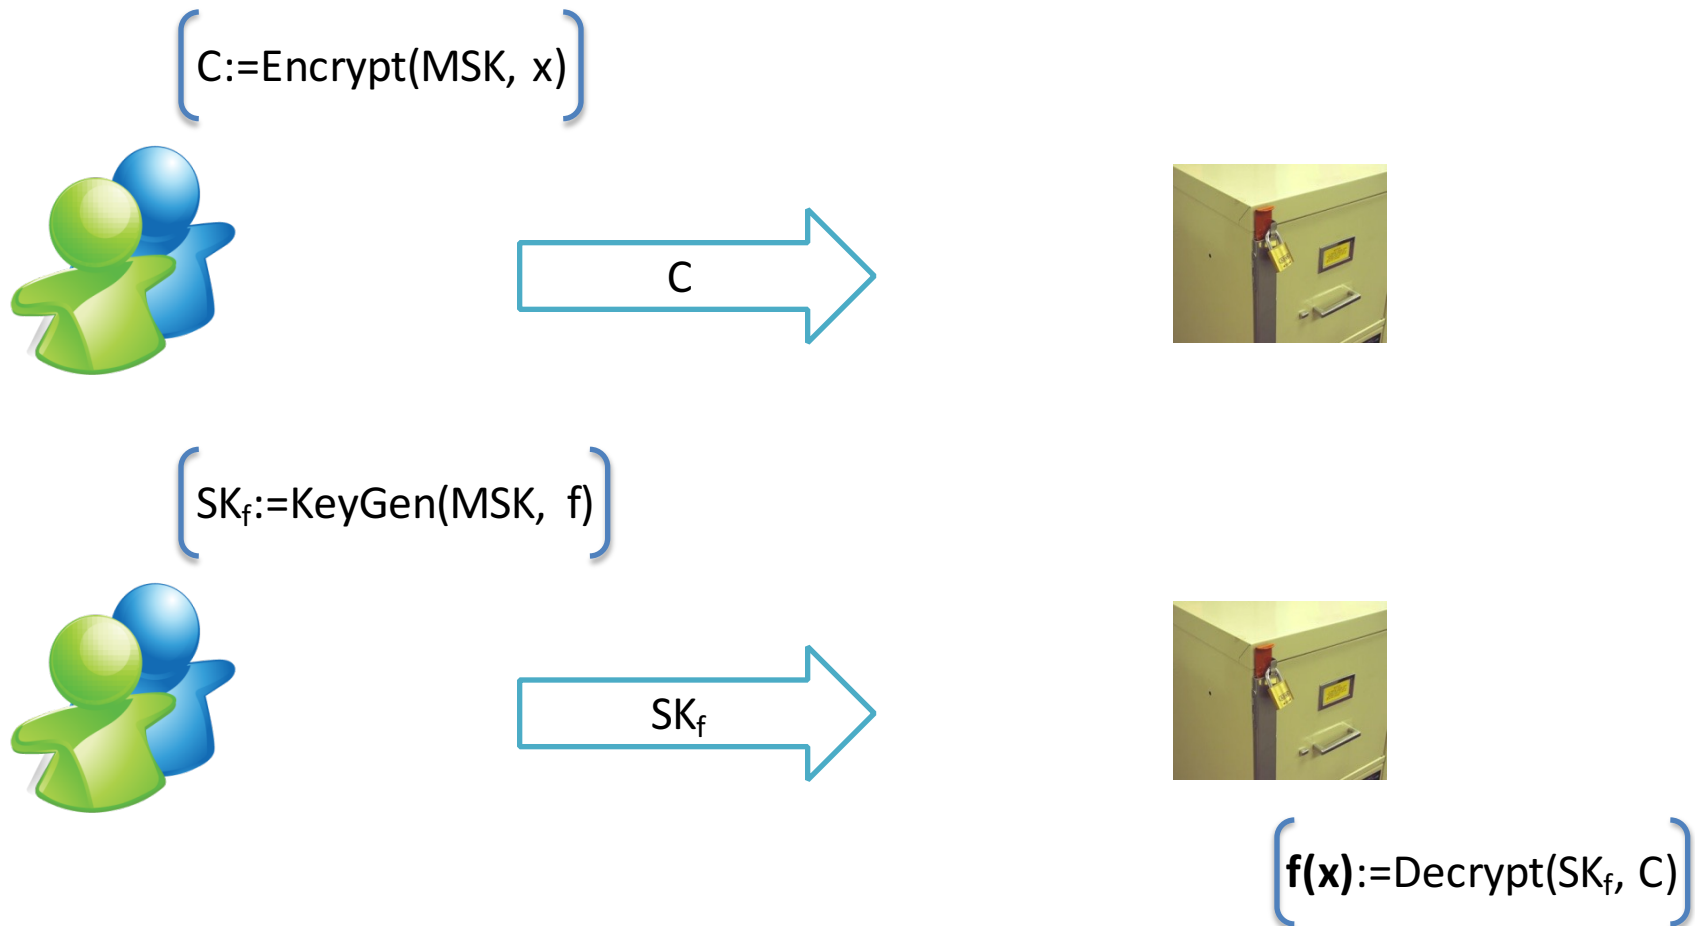
- A searchable encryption **framework**
 - Cryptographic layer
 - provides functionality and privacy
 - implemented by constrained functional encryption (for inner-product functions)
 - Data (structure) layer
 - provides functionality and scalability
 - implements search on trees, (unlinked) lists, matrices, graphs, ...

(Constrained) Functional Encryption



$$f(x) := \text{Decrypt}(\text{SK}_f, x)$$

(Constrained) Functional Encryption



SK_f **constrained** to decrypt a particular ciphertext

Our Result

Assume the subgroup membership problem holds, then there exists a secure¹ constrained functional encryption system for the class of inner product functions

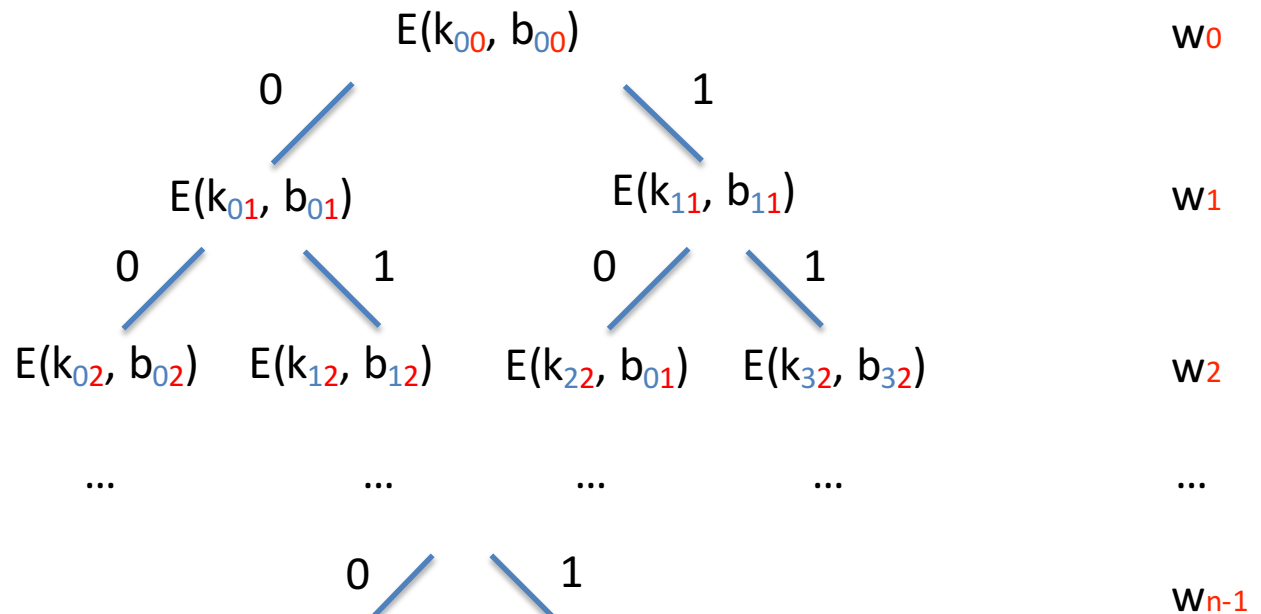
¹ Security game similar to predicate-private encryption [Shi-Waters, TCC'09]

Scheme #1


Data Structure: Binary Tree

Technique (1)

- Binary tree of depth $\log |W| = n$




bucket containing all
pointers matching
word $w=(w_0, \dots, w_{n-1})$



Our Technique (2)

- Search query q for single $w=(w_0, \dots, w_{n-1})$



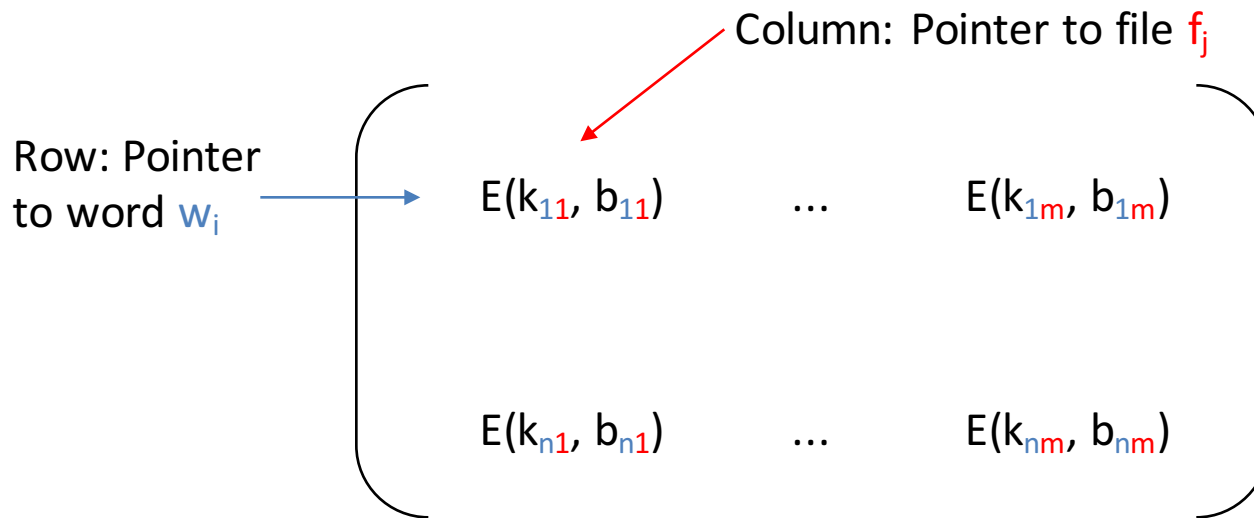
- Decryption key $sk_i = k_i * k_i$
 - Can only decrypt if correct evaluation of product performed

Scheme #2

Data Structure: (Unlinked) List

Our Technique (1)


- Database look-up Matrix $M^{n \times m}$



$i \leq n, j \leq m$: $b_{ij}=1$ if and only if f_j contains word w_i

Our Technique (2)

- Search query q for multiple words w_i



$i \leq n$: $b_i=1$ if and only if we query for word w_i

Our Technique (3)


- Search on Encrypted Data
 - Matrix multiplication $p=q^T A$

$$\begin{pmatrix} E(k_{0}, b_{0}) \\ \dots \\ E(k_{n}, b_{n}) \end{pmatrix}^T \begin{pmatrix} E(k_{11}, b_{11}) & \dots & E(k_{1m}, b_{1m}) \\ \dots & \dots & \dots \\ E(k_{n1}, b_{n1}) & \dots & E(k_{nm}, b_{nm}) \end{pmatrix}$$

- Decryption key $sk_j = \text{SUM}(k_i * k_{ij})$
 - Can only decrypt if correct evaluation of inner product performed

Performance Evaluation

- Implemented scheme based on group G
 - of composite order $N=pqr$ and symmetric pairing
 - of prime order $N=p$ and asymmetric pairing



Group	Ciphertext Size	Search word	Security
$N=pqr$	3076 bit	1.23 s	128 bit
$N=p$	2*254 bit	0.73 ms	128 bit

Conclusion

- Searchable encryption requires to find a trade-off between privacy, functionality and scalability
- Our protocol framework is tailored towards privacy and functionality
- Yet many optimisations are not explored (e.g. clustering of matrices)

RSAConference2016

San Francisco | February 29 – March 4 | Moscone Center

SESSION ID: CRYPT-T11

Private Large-Scale Databases with Distributed Searchable Symmetric Encryption



Connect **to**
Protect

Authors:

Yuval Ishai

Eyal Kushilevitz

Steve Lu

Rafail Ostrovsky

Speaker:

Steve Lu

Senior Researcher
Stealth Software Technologies, Inc.



#RSAC

A Story



#RSAC

■ In 2004...



A Story




#RSAC

- In 2004...



- Whatever happens in Vegas stays in Vegas?



Motivating Problem





Motivating Problem



#RSAC



Privacy Goals




- Protect the privacy of the Server holding the Database (from Cloud and Client)
- Protect the privacy of Client queries (from Cloud and Server)
- Server can specify private policies that enforce queries
- Modeled as a 3-party computation, with security in the semi-honest (honest-but-curious) model against any single corrupted party

Additional Goals



- Query Functionality
- Leverages the Cloud
- Scales! Goes Fast! ...









Background and Model



Data Format



#RSAC

First	Last	DOB	Photo	Fingerprint
John	Smith	12/3/45		
Jane	Doe	6/7/89		
...

Searchable Attributes
(First, Last, DOB, ...)

Rich Data (Photo, Fingerprint, ...)

Review of Crypto Primitives



#RSAC

Name	Properties	Complexity	Leakage
Private Information Retrieval(PIR) [Chor-Kushilevitz-Goldreich-Sudan95, Kushilevitz-Ostrovsky97,...]			
Oblivious RAM(ORAM) [Goldreich-Ostrovsky96,...]			
Searchable Symmetric Encryption(SSE) [Curtmola-Garay-Kamara-Ostrovsky06,...]			

Review of Crypto Primitives



#RSAC

Name	Properties	Complexity	Leakage
Private Information Retrieval(PIR) [Chor-Kushilevitz-Goldreich-Sudan95, Kushilevitz-Ostrovsky97,...]	Privately fetch index i from array	Low comm. Linear server work, mostly bitwise ops	Sizes only
Oblivious RAM(ORAM) [Goldreich-Ostrovsky96,...]			
Searchable Symmetric Encryption(SSE) [Curtmola-Garay-Kamara-Ostrovsky06,...]			

Review of Crypto Primitives



#RSAC

Name	Properties	Complexity	Leakage
Private Information Retrieval(PIR) [Chor-Kushilevitz-Goldreich-Sudan95, Kushilevitz-Ostrovsky97,...]	Privately fetch index i from array	Low comm. Linear server work, mostly bitwise ops	Sizes only
Oblivious RAM(ORAM) [Goldreich-Ostrovsky96,...]	Compiles program into one that hides data and access	Low comm. polylog work, symmetric-key ops	Sizes only
Searchable Symmetric Encryption(SSE) [Curtmola-Garay-Kamara-Ostrovsky06,...]			

Review of Crypto Primitives



#RSAC

Name	Properties	Complexity	Leakage
Private Information Retrieval(PIR) [Chor-Kushilevitz-Goldreich-Sudan95, Kushilevitz-Ostrovsky97,...]	Privately fetch index i from array	Low comm. Linear server work, mostly bitwise ops	Sizes only
Oblivious RAM(ORAM) [Goldreich-Ostrovsky96,...]	Compiles program into one that hides data and access	Low comm. polylog work, symmetric-key ops	Sizes only
Searchable Symmetric Encryption(SSE) [Curtmola-Garay-Kamara-Ostrovsky06,...]	Search on encrypted data	Very low comm/comp, symmetric-key ops	Sizes and some access "metadata"



- (2-server) PIR
 - Subpolynomial [Dvir-Gopi15]
 - Distributed Point Functions [Gilboa-Ishai14, Boyle-Gilboa-Ishai15]
- ORAM
 - Asymptotic [Kushilevitz-Lu-Ostrovsky12]
 - Circuit complexity [Wang-Chan-Shi15]
 - Distributed [Lu-Ostrovsky13A]
 - Non-interactive (Garbled RAM) [Lu-Ostrovsky13B, Gentry-Halevi-Lu-Ostrovsky-Raykova-Wichs14, Garg-Lu-Ostrovsky-Scafuro15, Garg-Lu-Ostrovsky15]
- Searchable Symmetric Encryption
 - Large Scale [Cash-Jarecki-Jutla-Krawczyk-Rosu-Steiner13, Cash-Jaeger-Jarecki-Jutla-Krawczyk-Rosu-Steiner14, Pappas-Krell-Vo-Kolesnikov-Malkin-Choi-George-Keromytis-Bellovin14, Faber-Jarecki-Karwczyk-Nguyen-Rosu-Steiner15, Fisc-Vo-Krell-Kumarasubramanian-Klesnikov-Malkin-Bellovin15]
 - Dynamic [Kamara-Papamanthou-Roeder12, Kamara-Papamanthou13, Naveed-Prabhakaran-Gunter14]



- (2-server) PIR
 - Subpolynomial [Dvir-Gopi15]
 - Distributed Point Function [Garg-Lu-Ostrovsky15]
- ORAM
 - Asymptotic [Kushilevitz-Ostrovsky15]
 - Circuit complexity [Wang15]
 - Distributed [Lu-Ostrovsky15]
 - Non-interactive (Garbled RAM) [Garbled RAM15, Scafuro15, Garg-Lu-Ostrovsky15]
- Searchable Symmetric Encryption
 - Large Scale [Cash-Jarecki-Jutla-Krawczyk-Rosu-Steiner13, Cash-Jaeger-Jarecki-Jutla-Krawczyk-Rosu-Steiner14, Pappas-Krell-Vo-Kolesnikov-Malkin-Choi-George-Keromytis-Bellovin14, Faber-Jarecki-Karwczyk-Nguyen-Rosu-Steiner15, Fisc-Vo-Krell-Kumarasubramanian-Klesnikov-Malkin-Bellovin15]
 - Dynamic [Kamara-Papamanthou-Roeder12, Kamara-Papamanthou13, Naveed-Prabhakaran-Gunter14]

Can we leverage
these advances
and get the best
of all worlds?



- We create a private SSE scheme that only leaks sizes (and query types), assuming semi-honest (honest-but-curious) model with no collusion.
- Privacy is modeled via the ideal/real paradigm from secure computation literature
- Supports range, substring, Boolean,... queries
- Supports simple deny policies
- Supports updates



Our Construction



Overview

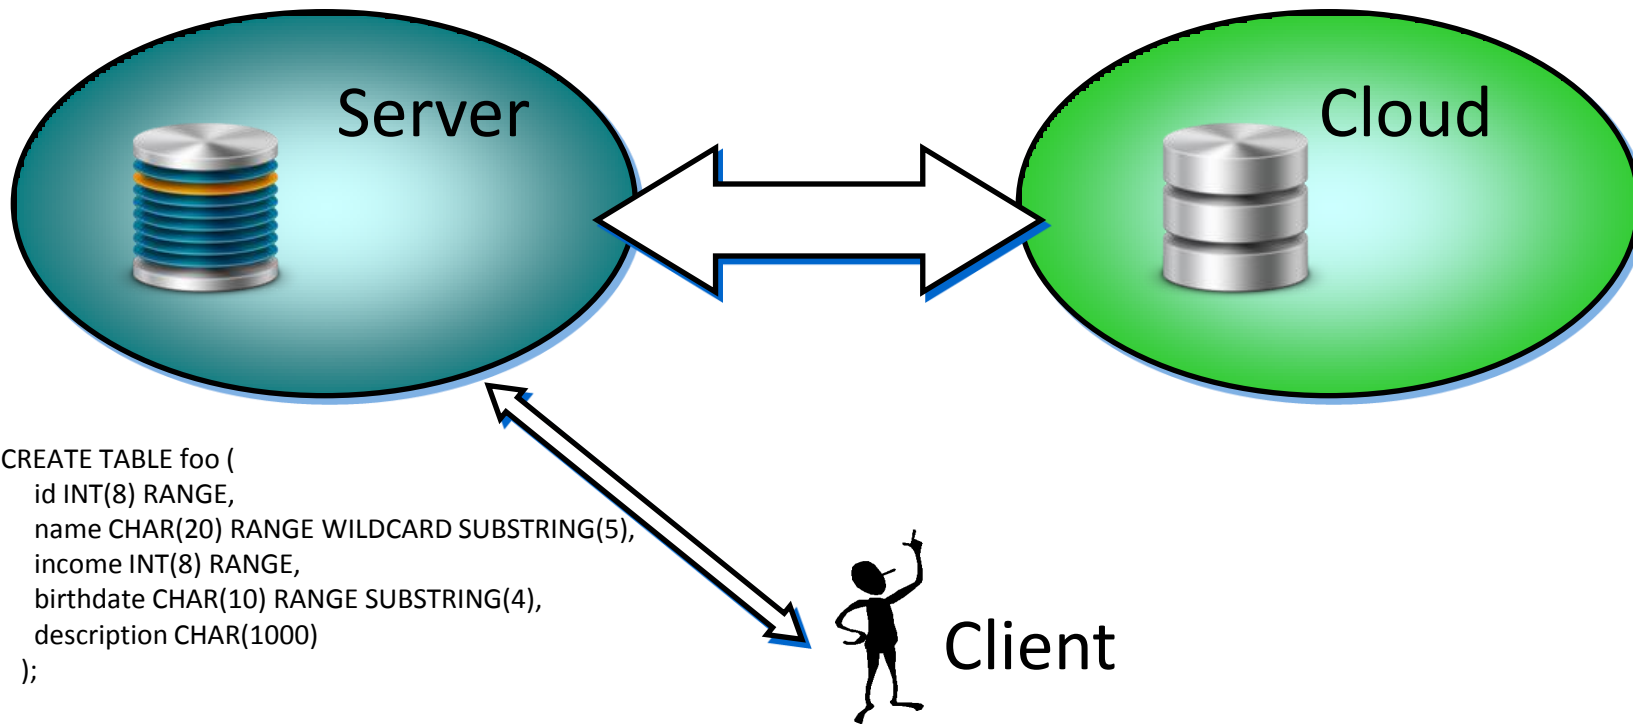



- Setup
- Query
- Updates
- Policy

Setup



#RSAC






Update



#RSAC




```
INSERT INTO table VALUES  
(4, 'Carol', 123456, '1970-05-17', NULL);
```

```
DELETE FROM table WHERE id = 4;
```



Client



Deny all range queries

Deny wildcard queries on 'birthdate'

Deny all queries on 'gender'



Client

High Level Idea

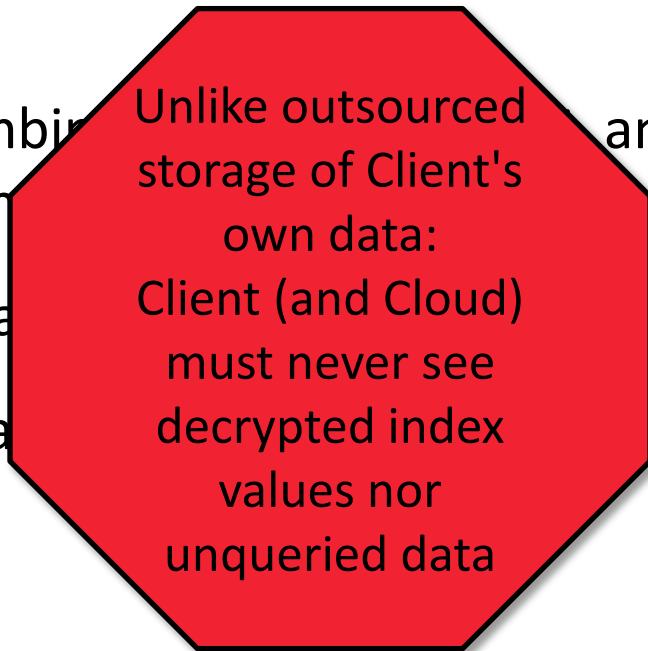


#RSAC

- Setup: Server stores encrypted database and encrypted indices (B-tree) in Cloud
- Query: Use a combination of PIR, ORAM, and secure computation techniques to traverse tree privately
- Updates: Create a mini-database and intermittently merge
- Policies: Interweave policy enforcement with query mechanism



- Setup: Server stores encrypted database and encrypted indices (B-tree) in Cloud
- Query: Use a combination of secure computation techniques and secure storage of Client's own data: privately
- Updates: Create a new encrypted database and incrementally merge
- Policies: Interweave encrypted data with query mechanism






- Setup: Server stores encrypted database and encrypted indices (B-tree) in Cloud
- Query: Use a combination of PIR, ORAM, and secure computation techniques to traverse tree privately
- Updates: Create a mini-database and intermittently merge
- Policies: Interweave policy enforcement with query mechanism

Structure of Setup




#RSAC




Binary Tree Example For Range Query



#RSAC




Binary Tree Example



Binary Tree Example



#RSAC



Main Traversal Technique



#RSAC

PIR to fetch

Secure "Millionaires" to
step down

.....

ORAM to fetch

Must be secret-shared

Reminder: Secret Sharing/One-Time Pad




#RSAC

- $M \oplus R = C$
 - One-time pad, R is random key, used only once, C perfectly hides M
- $M = C \oplus R$
 - C and R form a secret sharing of M, each perfectly hiding M

Main Traversal Technique (Details)



#RSAC

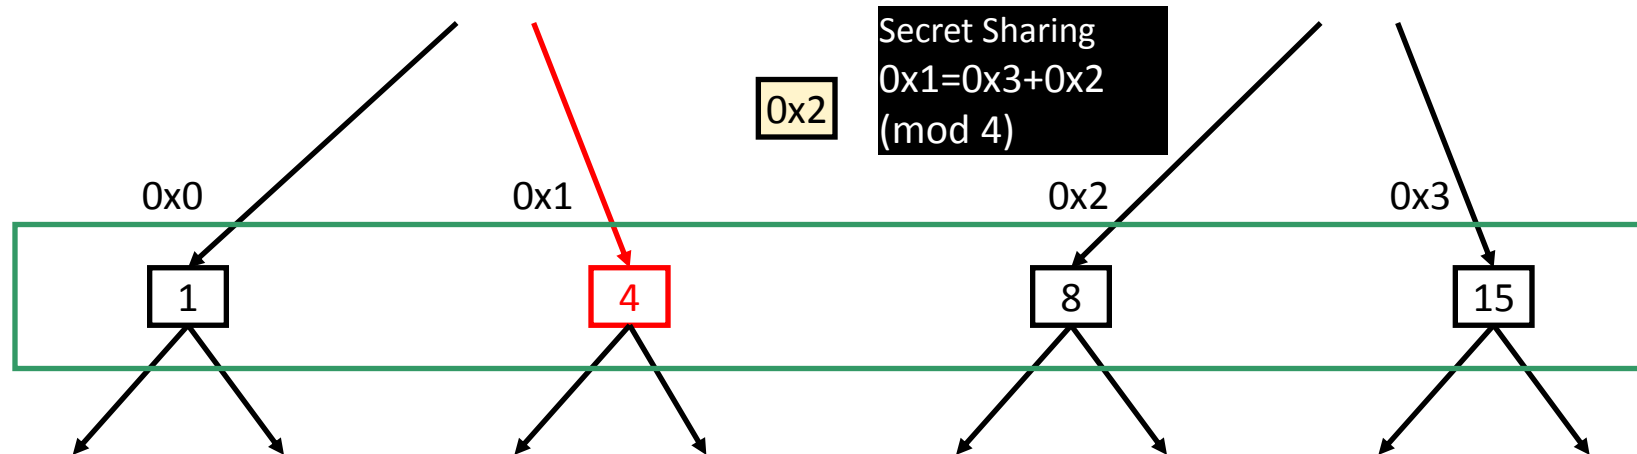


Query:5

Recursive Assumption



#RSAC



0x3

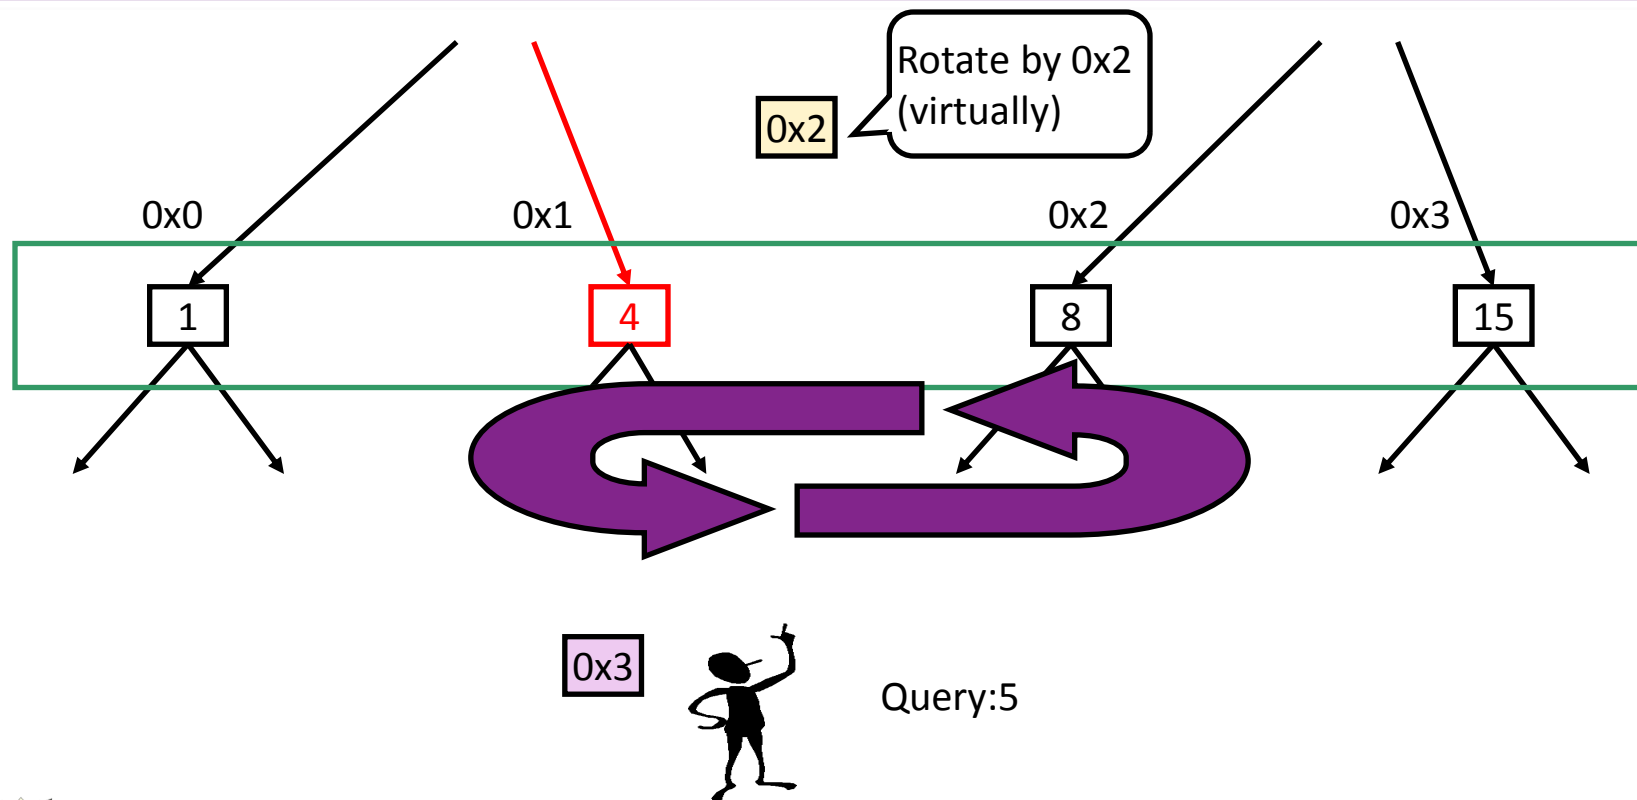


Query:5

Virtual Rotation



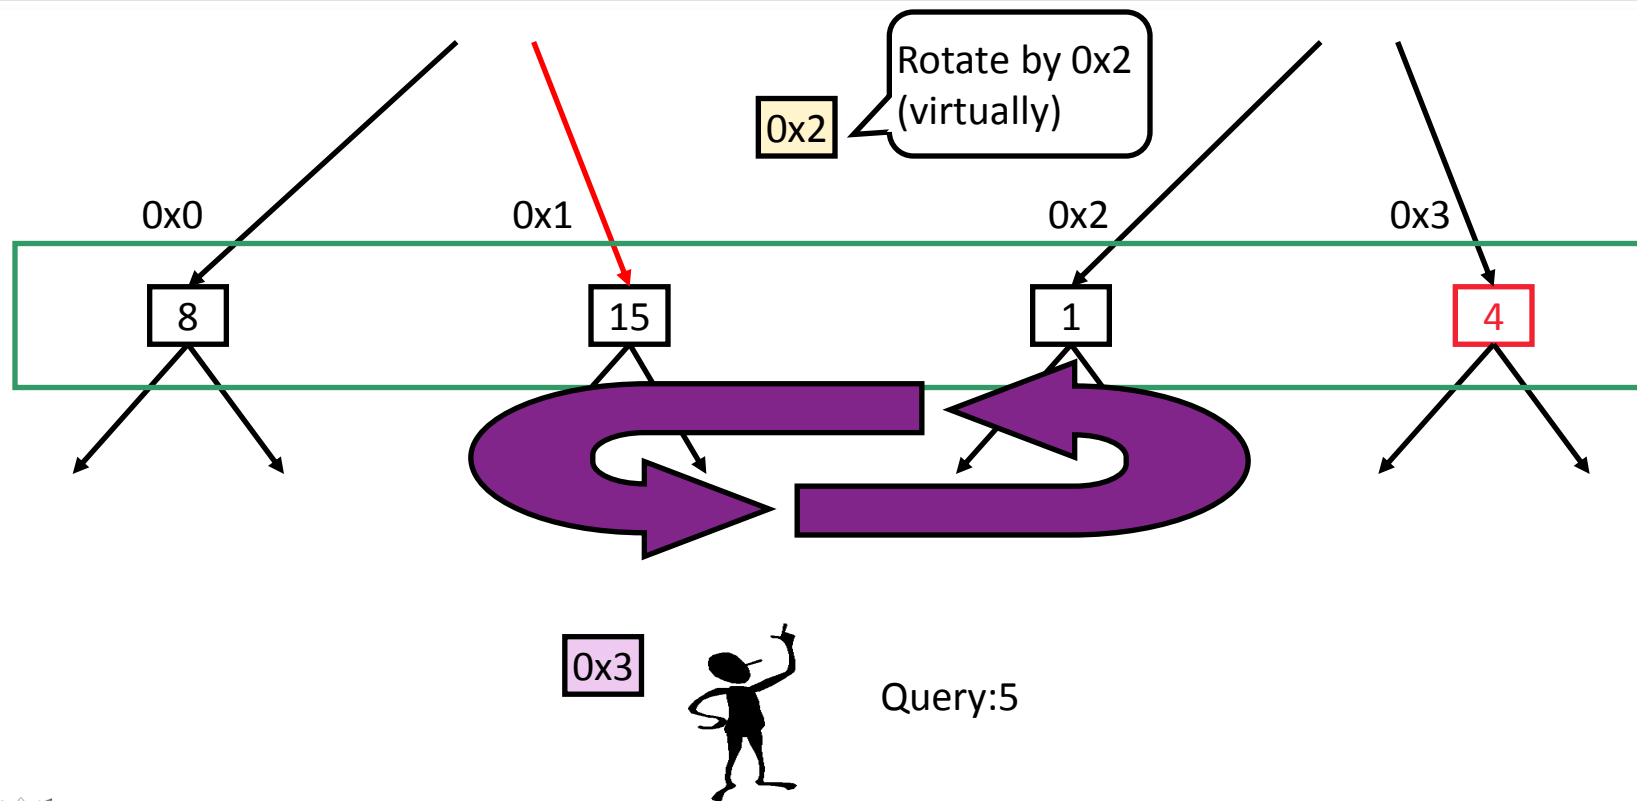
#RSAC



Virtual Rotation



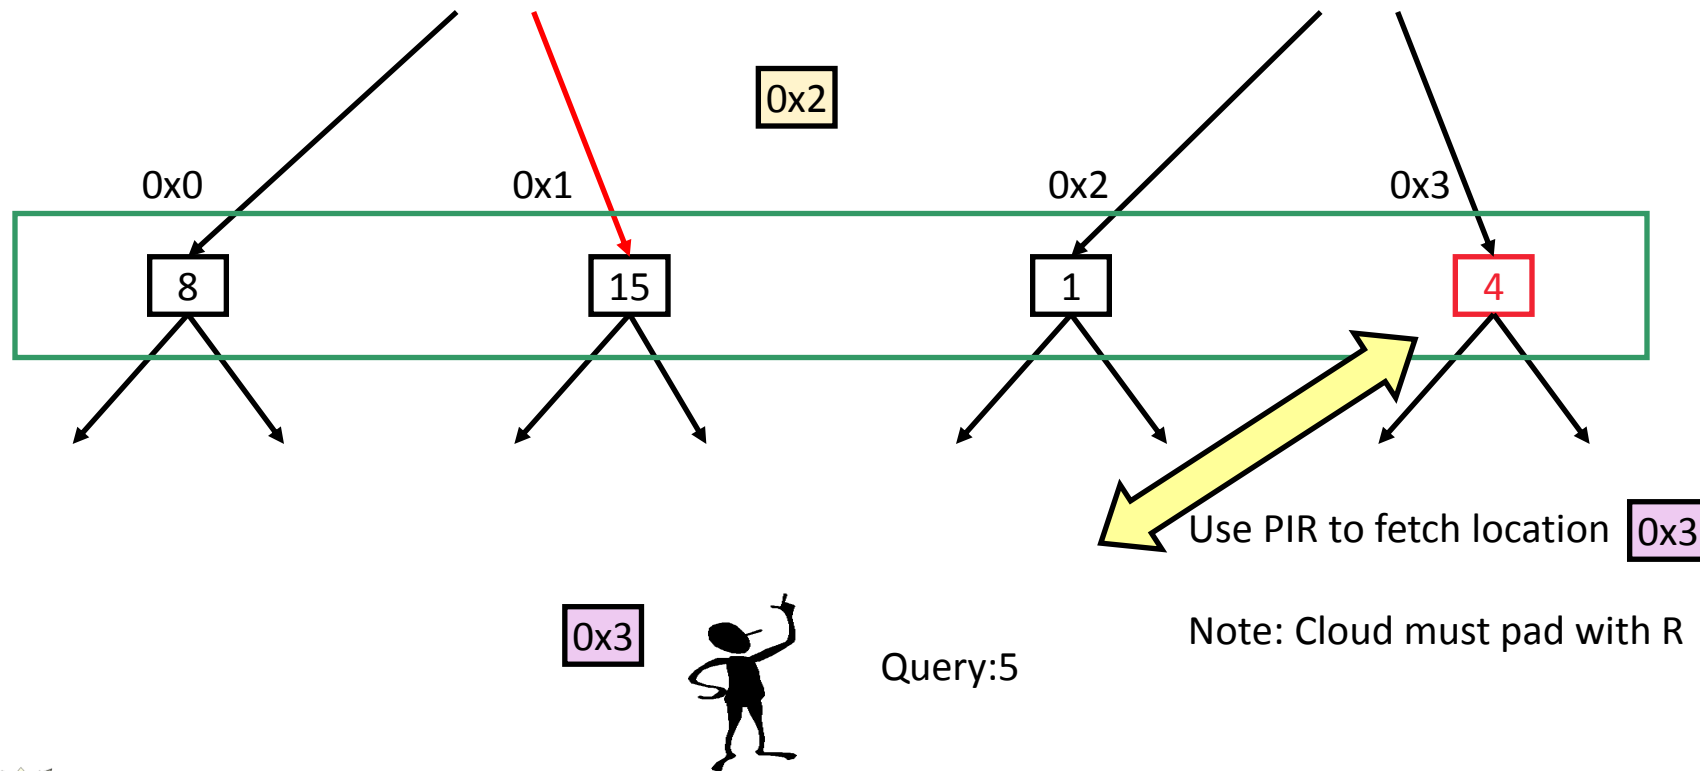
#RSAC



Use PIR to fetch



#RSAC



Special Purpose Secure Computation




- PIR returns a block B containing (value=4,Lptr,Rptr)
- Client holds: Secret key k , Query=5, Result of PIR $X=E_k(B) \oplus R$
- Cloud holds: R
- Need to securely compute $F(k, \text{Query}, X; R)$:
 - Set $B=D_k(X \oplus R)$ (Custom protocol for this)
 - Return $(q>B.\text{value}) ? B.Rptr : B.Lptr$ (Millionaires problem!)
- Final caveat: must return secret shared output

Completing the Recursion




#RSAC



Completing the Recursion



#RSAC

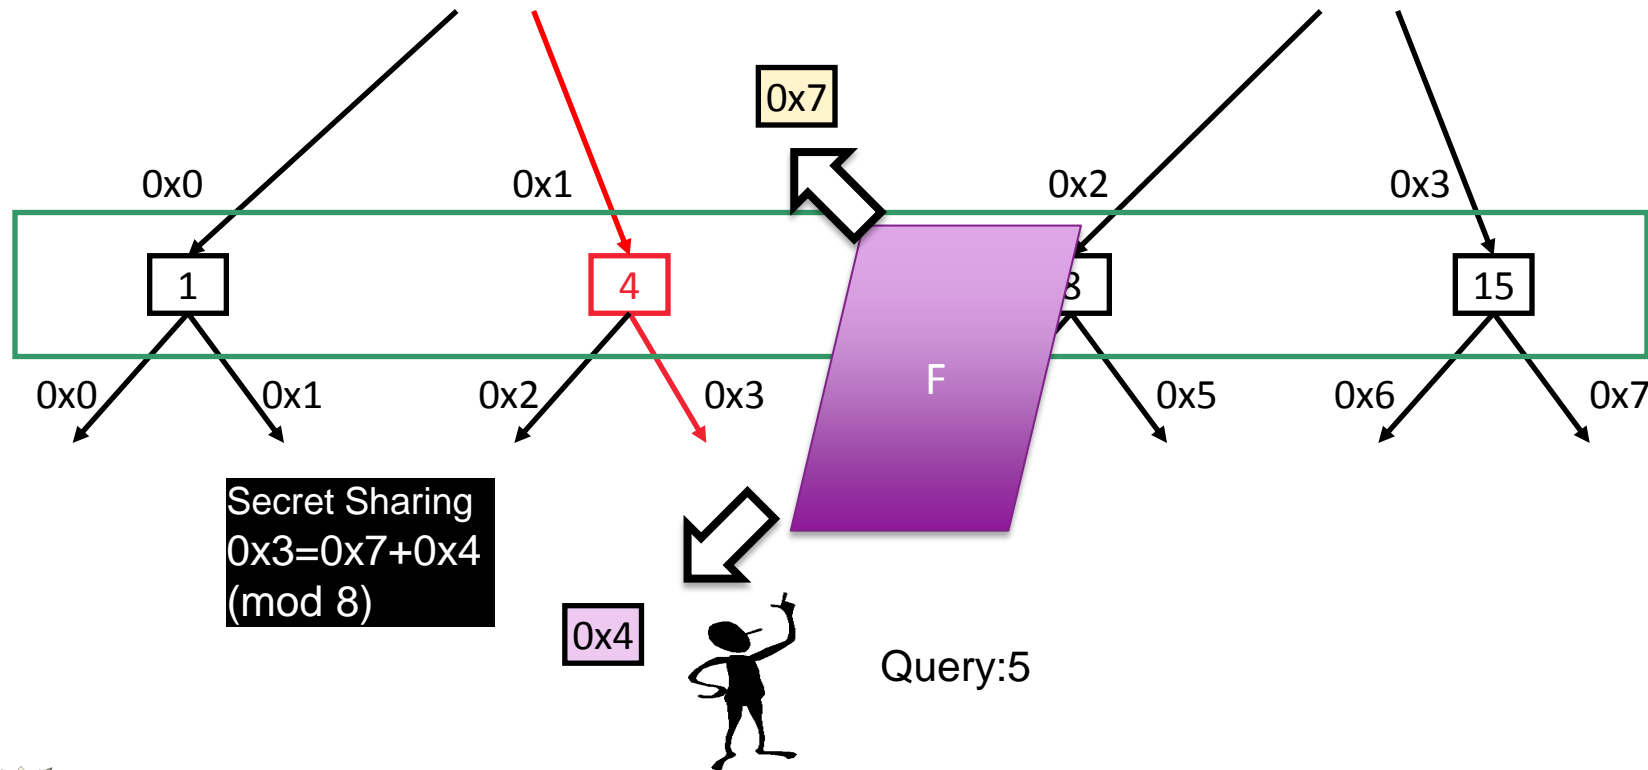


Query:5

Completing the Recursion



#RSAC



Summary and Conclusion



- We constructed a new SSE scheme that supports a wide variety of queries and can enforce policies and support updates
 - Combination of several techniques including PIR, ORAM, and secure computation
- We implemented the solution, and for large queries we are only 5x slower than MySQL (smaller queries have overhead of up to 100x, but actual time is under 1 second)

Apply What You Have Learned Today



- Within 1 month you should:
 - Identify scenario where this setting occurs
 - Further research our paper and others
- Within 3 months you should:
 - Understand and identify acceptable and unacceptable leakage amongst secure database solutions
- Within 6 months you should:
 - Have a broader understanding of different solutions
 - Discuss applying this technology to suit your needs



THANK YOU!

**Full version available on ePrint (2015/1190):
eprint.iacr.org/2015/1190**