



# Exciting, To-Be-Announced Platform Session

We can't tell you about it now, but trust us - it's awesome.

Philipp Drieger | Staff Machine Learning Architect

October 2018





# PHILIPP DRIEGER

Staff Machine Learning Architect

[philipp@splunk.com](mailto:philipp@splunk.com)




# Forward-Looking Statements

During the course of this presentation, we may make forward-looking statements regarding future events or the expected performance of the company. We caution you that such statements reflect our current expectations and estimates based on factors currently known to us and that actual events or results could differ materially. For important factors that may cause actual results to differ from those contained in our forward-looking statements, please review our filings with the SEC.

The forward-looking statements made in this presentation are being made as of the time and date of its live presentation. If reviewed after its live presentation, this presentation may not contain current or accurate information. We do not assume any obligation to update any forward-looking statements we may make. In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. Splunk undertakes no obligation either to develop the features or functionality described or to include any such feature or functionality in a future release.

Splunk, Splunk>, Listen to Your Data, The Engine for Machine Data, Splunk Cloud, Splunk Light and SPL are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names, or trademarks belong to their respective owners. © 2018 Splunk Inc. All rights reserved.





# AI | ML | Deep Learning

---

# Let's start with AI

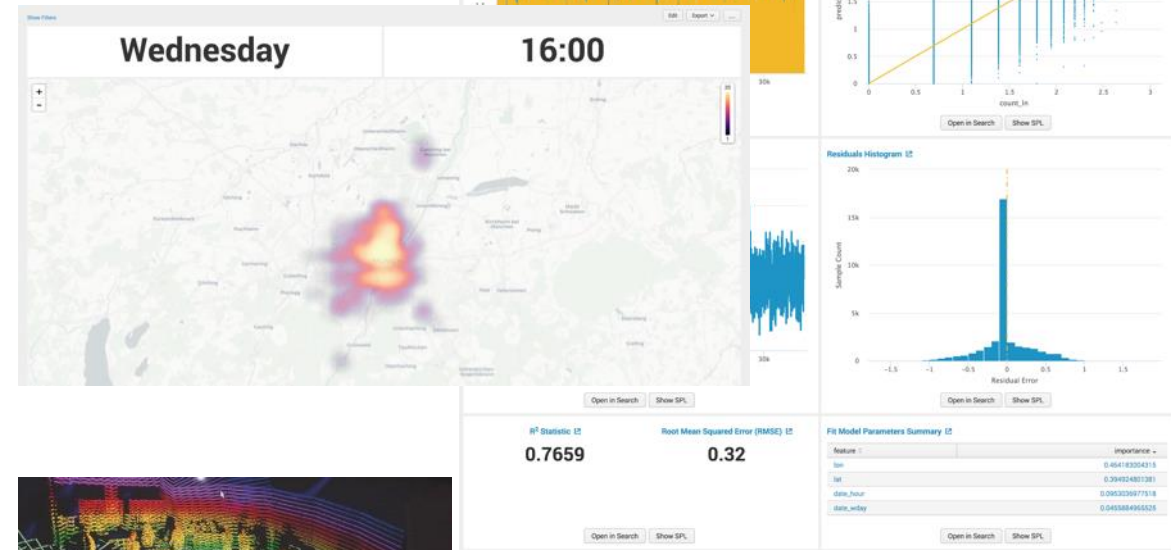
## Artificial Intelligence

- ▶ **Artificial Intelligence (AI)**  
The self driving car




## Machine Learning

- ▶ **Machine Learning (ML)**  
Predicting car demand based on past history
- ▶ Example given by BMW



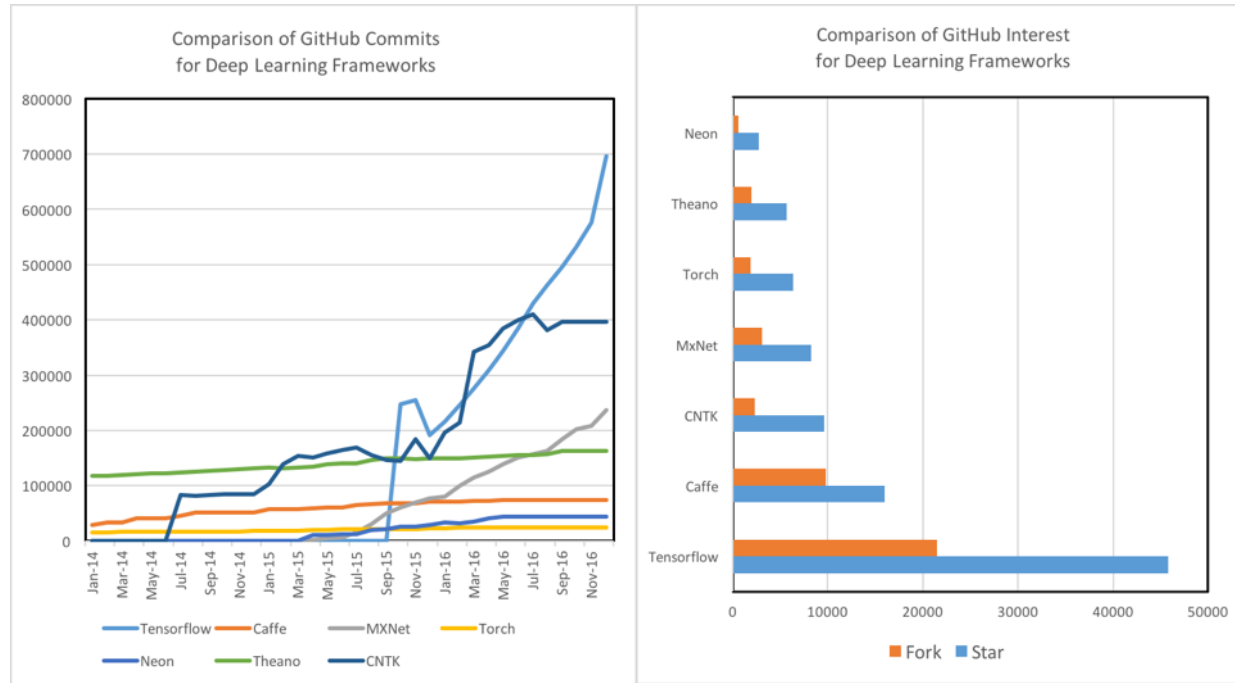
## Deep Learning

- ▶ **Deep Learning (DL)**  
~~Not in our product today...~~ but we still do NO image processing






# Popular ML/DL Frameworks in the Python Landscape



<https://www.kdnuggets.com/2017/03/getting-started-deep-learning.html>

Top 20 Python AI and Machine Learning projects on Github



<https://www.kdnuggets.com/2018/02/top-20-python-ai-machine-learning-open-source-projects.html>


# MLTK Container for TensorFlow™

# MLTK Container for TensorFlow™ – Why?

Because our customers ask. We listen. Simple as that!



Popular deep learning frameworks help to extend MLTK for specific use cases.




Freedom for Data Scientists and Developers to bring in custom code and models



Flexibility to run compute intense model trainings on GPU accelerated hardware

130.60.4 - - [07/Jun 18:10:57:153] "GET /category.screen?category\_id=GIFTS&JSESSIONID=5D1SLAFF10ADFF10 HTTP 1.1" 404 720 "http://buttercup-shopping.com/cart.do?action=view&itemId=EST-6&product\_id=FI-SW-01" "Opera/9.20 (Win  
128.241.220.82 - - [07/Jun 18:10:57:123] "GET /product.screen?product\_id=FL-DSH-01&JSESSIONID=5D5SL7FF6ADFF9 HTTP 1.1" 404 3322 "http://buttercup-shopping.com/cart.do?action=purchase&itemId=EST-26&product\_id=K9-CW-01" "Mozilla/4.0 (Compaq i486 Win  
ows NT 5.1; SV1; .NET CLR 1.1.4322)" 468 125.17.14.189 "GET /category.screen?category\_id=FLOWERS&JSESSIONID=5D5SL7FF6ADFF9 HTTP 1.1" 200 2423 "http://buttercup-shopping.com/cart.do?action=changequantity&itemId=EST-1B&product\_id=AV-CB-01&JSESSIONID=5D1SL8FF2ADFF9" "Mozilla/4.0 (Compaq i486 Win  
itemId=EST-16&product\_id=RP-LI-02" "0  
buttercup-shopping.com/cart.do?action=purchase&itemId=EST-26&product\_id=K9-CW-01" "Mozilla/4.0 (Compaq i486 Win  
buttercup-shopping.com/cart.do?action=purchase&itemId=EST-26&product\_id=K9-CW-01" "Mozilla/4.0 (Compaq i486 Win

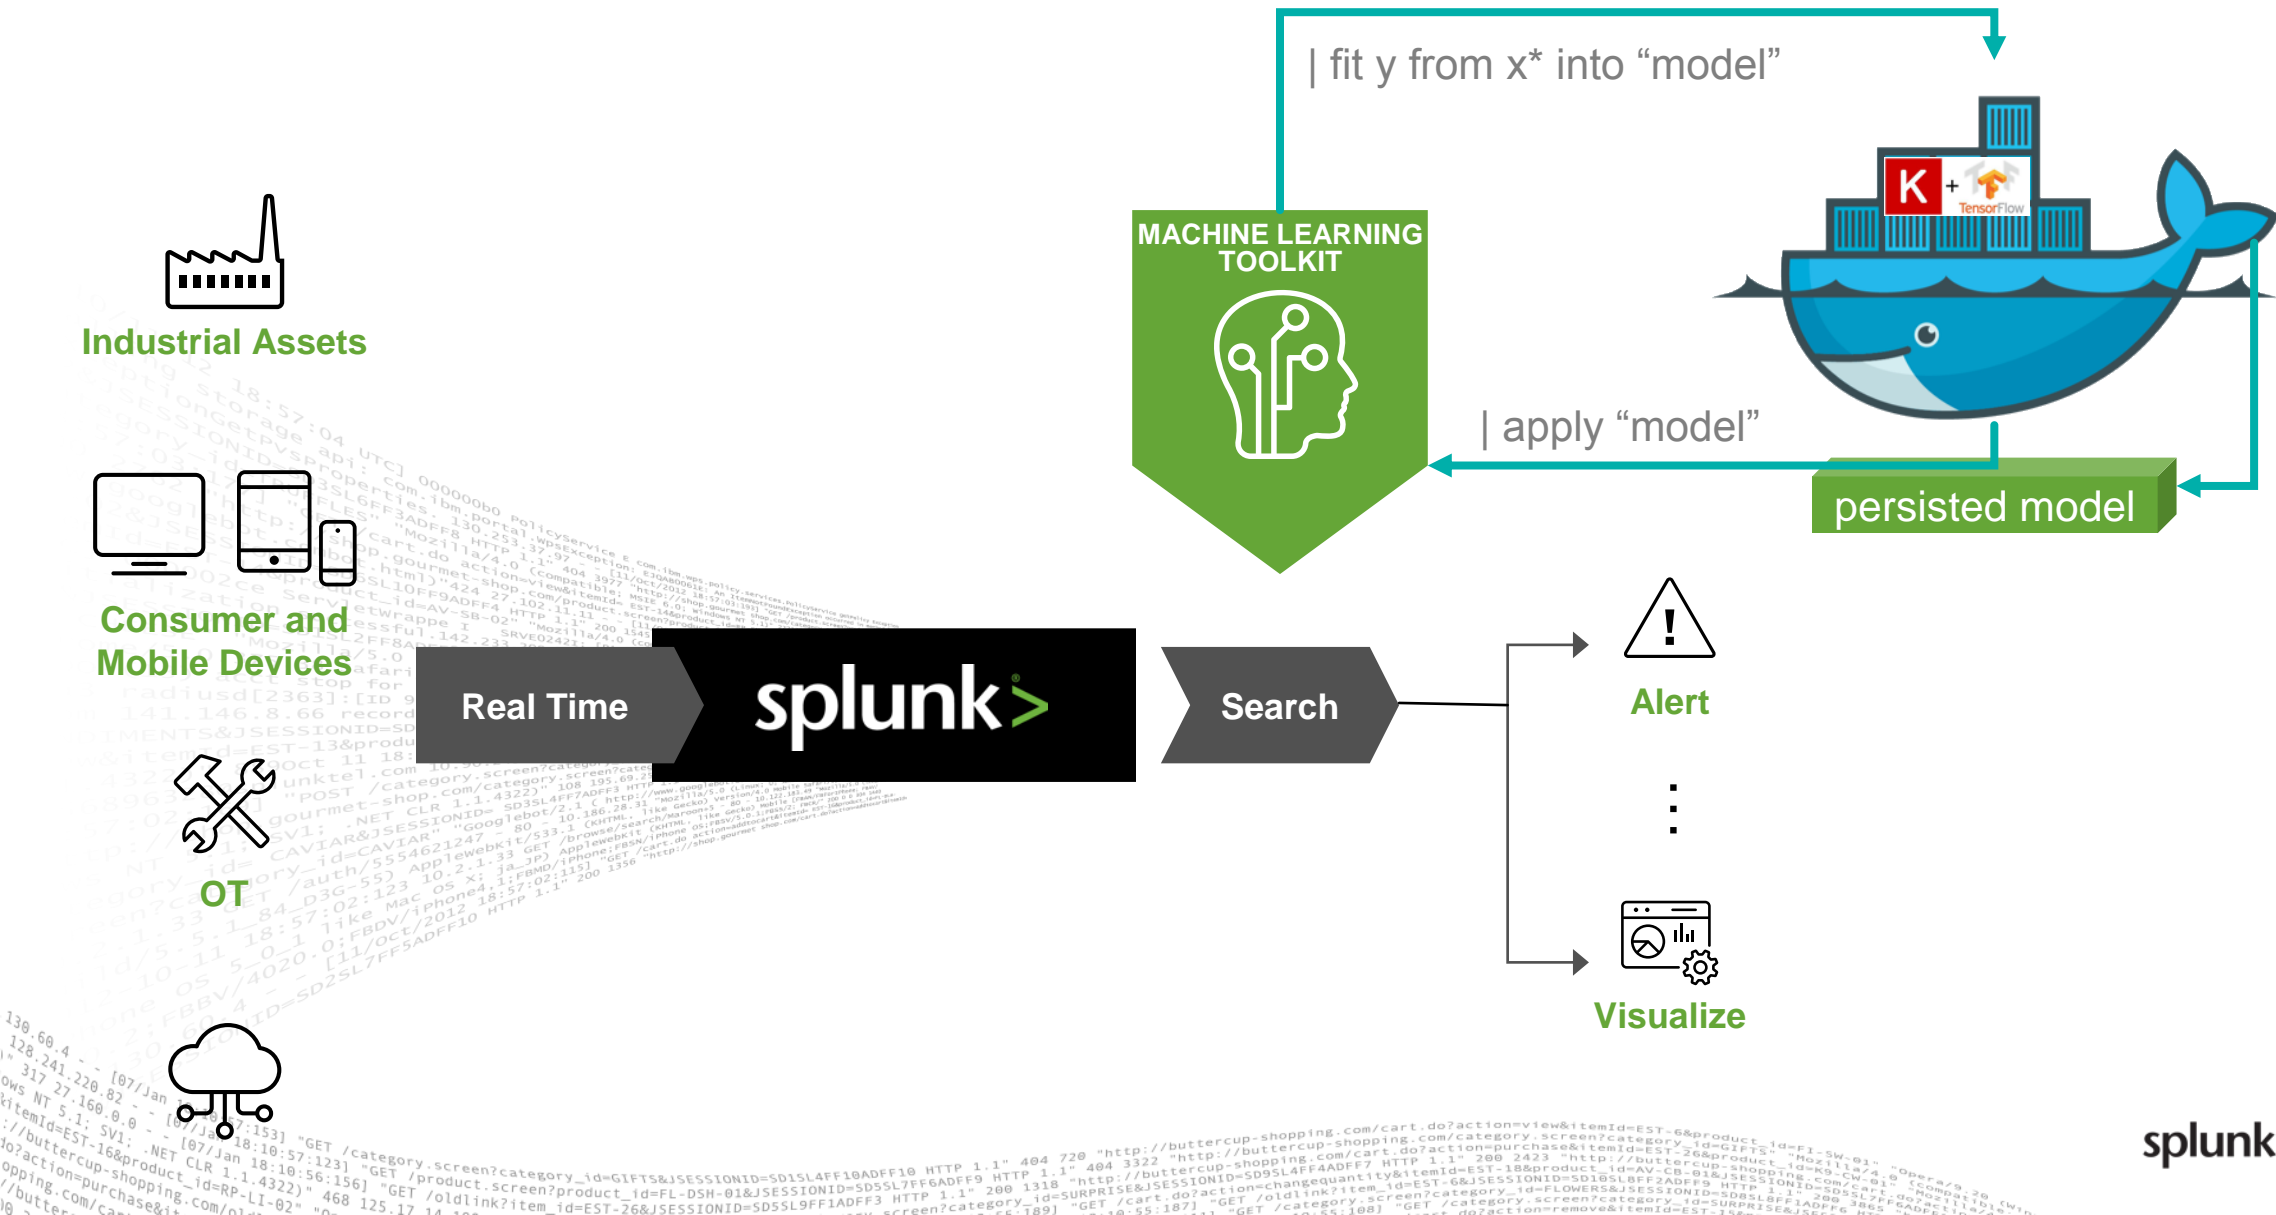




# Architecture Overview

---

# Splunk > MLTK > Dockerized Deep Learning





[illegible]

# Splunk App for the MLTK Containers



# MLTK Container for TensorFlow™

- Overview
- Classifier
- Regressor

## Overview




### Setup Instructions

Learn about how to setup the MLTK Container step by step



### Container Management

Controls to start and stop the MLTK Container and check its status



### MLTK Container Overview

Description of the architecture

## Classifier

serum_insulin										skin_thickness									
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
110	32																		
3	4	5	6	7	8	9	10	next >											
Predicted 0										Predicted 1									
79 (77.5%)										23 (22.5%)									
11 (20.8%)										42 (79.2%)									

### Neural Network Classifier Example

This example shows how to use a binary neural network classifier build on keras and TensorFlow™

serum_insulin										skin_thickness									
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
110	32																		
3	4	5	6	7	8	9	10	next >											
Predicted 0										Predicted 1									
79 (77.5%)										23 (22.5%)									
11 (20.8%)										42 (79.2%)									

### Linear Classifier Example


This example shows a linear classifier using the TensorFlow™ estimator class

serum_insulin										skin_thickness									
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
110	32																		
3	4	5	6	7	8	9	10	next >											
Predicted 0										Predicted 1									
79 (77.5%)										23 (22.5%)									
11 (20.8%)										42 (79.2%)									

### LSTM Example


This example shows the results of a LSTM to classify DGA domains

## Regressor



### Linear Regressor Example

This example shows the results of a simple linear regression using the TensorFlow™ estimator class



### Random Forest Regressor Example

This example shows the results of a Random Forest Regressor using TensorFlow™

# Container Management

Edit

Export ▾

...

Container Name

Select... ▾

Container endpoint

http://localhost:5000

Submit

Hide Filters

Container Controls

RUN

RUN with logging\*

STOP


\* logging currently only works for mac osx

Containers Running (update delay 5s)

1

Information

This dashboard provides simple container controls to start and stop a container and retrieve status information



## Container Status (update delay 5s)

action ⇅	
1	07545ddbff1 mltk-container-tensorflow "./bootstrap.sh" 18 seconds ago Up 17 seconds 0.0.0.0:5000->5000/tcp mltk-container-tensorflow

## Container Logs (update delay 5s)

_time ⇅	line ⇅
2018-09-04 12:27:27.506	WARNING: Do not use the development server in a production environment.
2018-09-04 12:27:27.506	* Environment: production
2018-09-04 12:27:27.506	* Serving Flask app "/srv/container/index.py"
2018-09-04 12:27:27.170	Docker bootstrap entry point set to index.py
2018-09-04 12:26:40.650	172.17.0.1 - - [04/Sep/2018 10:26:40] "POST /fit HTTP/1.1" 200 -
◀ prev 1 2 3 4 5 6 7 8 9 10 next ▶	

## MLTK Logs

_time ⇅	_raw ⇅
2018-09-04 12:26:40.630	1536056800.630000 PID 90477 2018-09-04 12:26:40,630 DEBUG [mlspl.MLTKClassifier] [endpoint] POST endpoint [http://localhost:5000/fit] returned with payload (16484 bytes) with status 200
2018-09-04 12:25:43.343	1536056743.343556 PID 90477 2018-09-04 12:25:43,343 DEBUG [mlspl.MLTKClassifier] [endpoint] POST endpoint [http://localhost:5000/fit] called with payload (44042 bytes)



# MLTK Container Classifier Example

Edit

Export

...

Model Name

diabetes\_test

×

Container endpoint

http://localhost:5000

Submit

Hide Filters

This example shows the interaction with Splunk> Machine Learning Toolkit and a custom container that runs a (multi layer fully connected) neural network classifier. Make sure to check the [setup page](#) and perform all steps needed to run this dashboard successfully.

response	response_prediction	response_prediction_raw	BMI	age	blood_pressure	diabetes_pedigree	glucose_concentration	number_pregnant	serum_insulin	skin_thickness
1	1	0.9126277566	33.6	50	72	0.627	148	6	0	35
0	0	0.0446447767	26.6	31	66	0.351	85	1	0	29
1	0	0.4475494921	23.3	32	64	0.672	183	8	0	0
0	0	0.0420317426	28.1	21	66	0.167	89	1	94	23
1	1	0.5078072548	43.1	33	40	2.288	137	0	168	35

« prev

1

2

3

4

5

6




7

8


9

10

next »

confusion matrix			Classification statistics				
	Predicted actual 	Predicted 0 	Predicted 1 	accuracy	precision	recall	f1
	0	468	32	0.77	0.78	0.77	0.77
	1	143	125				

Information about existing model (click to inspect the model)					Model summary for diabetes_test						
app	name	owner	sharing	type	acc	batch_size	epochs	input_shape	loss	model_name	summary
SA-MLTK-Container-Tensorflow	diabetes_test	admin	user	MLTKClassifier	0.772135416667	None	50	X(768, 8) Y(768, 1)	0.455660077433	diabetes_test	diabetes_test





# Tests and Benchmarks

---



# Fit on 100k DGA training dataset

splunk> App: DGA App for Splunk ▾ Administrator ▾ 2 Messages ▾ Settings ▾ Activity ▾ Help ▾ Find

DGA App for Splunk Dashboards ▾ Search More ▾ DGA Analysis

New Search Save As ▾ New Table Close

```

1 | inputlookup dga_domains_features
2 | fields class ut* PC*
3 | eval dga=if('class'="legit",0,1)
4 | fit SPLDLClassifier dga from ut* PC* into "dga100k20" epochs=20 endpoint_url="http://localhost:5000"
5 | rename "predicted(dga)" as "probability(dga)"
6 | eval "predicted(dga)"=if('probability(dga)'>0.5,1,0)
7 | `confusionmatrix("dga","predicted(dga)")`

```

✓ 2 results (1/26/18 6:04:00.000 PM to 1/26/18 7:04:16.000 PM) No Event Sampling ▾

Events Patterns Statistics (2) Visualization

20 Per Page ▾ Format Preview ▾

Predicted actual ▾ ↕	Predicted 0 ▾ ↕	Predicted 1 ▾ ↕
0	48210	1790
1	1211	48789

SPLDL: Model fitted successfully with epochs=20, batch\_size=None and evaluated with loss=0.0816066845539 accuracy = 0.96999

Edit Job Settings...

Send Job to Background

Inspect Job

Delete Job

splunk> .conf18

# Apply on test dataset

## 1.7M events (ca. 11K EPS throughput)

splunk> App: DGA App for Splunk

Administrator Messages Settings Activity Help Find

DGA App for Splunk Dashboards Search More

DGA Analysis

New Search

1 | inputlookup dga\_test\_features  
2 | apply "dga100k50"

1,728,112 results (1/25/18 8:32:00.000 AM to 1/25/18 9:32:16.000 AM)

Events Patterns Statistics (1,728,112) Visualize

20 Per Page Format Preview

PC_1	PC_2	PC_3	class
0.139327626858	0.651454874929	0.312195717152	dga
0.752265006343	-0.431404065393	0.133834700086	dga
0.137354024813	0.665703828981	0.318043684523	dga
0.0458532959218	0.107512530751	-0.116416633238	dga
0.140218961551	0.649686200457	0.309892080456	dga
0.094706824356	0.26342031993	-0.886041020165	dga
0.752265006343	-0.431404065393	0.133834700086	dga
-0.0470368386525	0.0220646908229	-0.0297032852795	dga
0.0465212038804	0.118249702976	-0.125595502852	dga
0.0877973615283	0.239381212319	-0.799703192605	dga
0.094706824356	0.26342031993	-0.886041020165	dga
0.0507778983211	0.120690549458	-0.129967555063	dga
0.0597513291773	0.109118908816	-0.108723550064	dga
0.162973196327	0.724572195014	0.349486521074	dga
-0.0582370731805	0.0112318507722	-0.026693956543	dga

Search job inspector

This search has completed and has returned 1,728,112 results by scanning 0 events in 154.149 seconds  
(SID: 1516869136.2901) [search.log](#)

Execution costs

Duration (seconds)	Component	Invocations	Input count	Output count
117.50	command.apply	1	1,728,112	1,728,112
15.52	command.inputlookup	1	-	1,750,000
0.00	dispatch.check_disk_usage	1	-	-
0.00	dispatch.createdSearchResultInfrastructure	1	-	-
0.77	dispatch.evaluate	1	-	-
0.77	dispatch.evaluate.apply	1	-	-
0.00	dispatch.evaluate.inputlookup	1	-	-
0.00	dispatch.evaluate.noop	1	-	-
0.77	dispatch.optimize.FinalEval	1	-	-
0.87	dispatch.optimize.matchReportAcceleration	1	-	-
0.03	dispatch.optimize.optimization	1	-	-
0.00	dispatch.optimize.reparse	1	-	-
0.80	dispatch.optimize.toJson	1	-	-


srtpv957j66a0op.org 1 0.9883615375 chinad 0.5 0.3  
46yey4sxzqde8cpi.cn 0 0.9966819882 chinad 0.631579 0.210526315789

length ut\_meaning\_ratio ut\_shannon ut\_vowel\_ratio

20 0.1 3.88418371978 0.15  
20 0.2 4.22192809489 0.1  
20 0.2 4.02192809489 0.1  
21 0.238095238095 3.97541801791 0.142857142857  
20 0.3 3.92192809489 0.2  
20 0.4 3.98418371978 0.25  
20 0.15 3.58418371978 0.05  
19 0.0 3.93213803976 0.0526315789474  
21 0.190476190476 3.82088885135 0.190476190476  
20 0.1 3.92192809489 0.15  
20 0.1 3.92192809489 0.2  
21 0.0952380952381 4.01136504183 0.0952380952381  
19 0.0 3.82687488186 0.105263157895  
20 0.3 3.78418371978 0.2  
19 0.0526315789474 3.82687488186 0.157894736842

splunk> .conf18



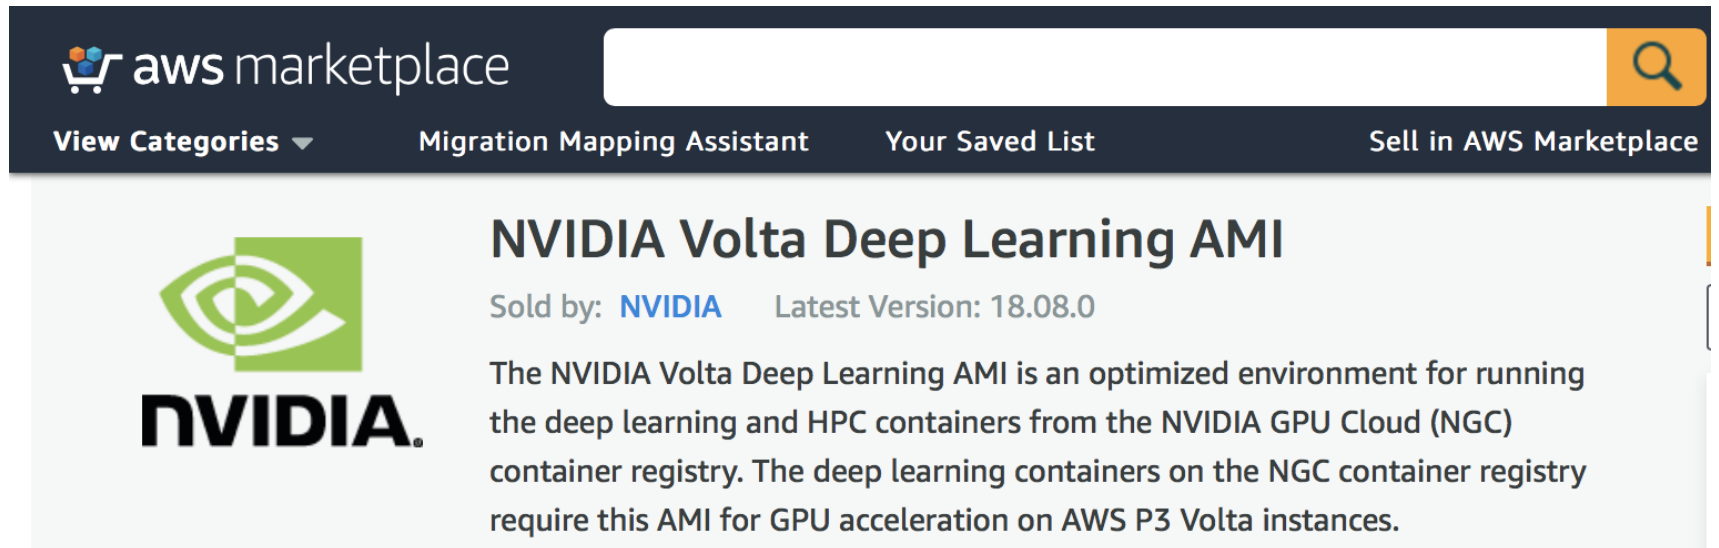


# CPU vs GPU Benchmark

---

# Benchmarking Model Training

- ▶ AWS Instance: p3.2xlarge (64GB, 8vCPU, NVIDIA V100 GPU 16GB)
- ▶ DGA Dataset: 100K events with 100 dimensions + 1 target dimension
- ▶ Neural Network: 10 layer deep neural network with 886K trainable parameters,
- ▶ 100 layer deep with 9M trainable parameters



The screenshot shows the AWS Marketplace interface. At the top, there's a navigation bar with the AWS Marketplace logo, a search bar, and links for 'View Categories', 'Migration Mapping Assistant', 'Your Saved List', and 'Sell in AWS Marketplace'. The main content area features the NVIDIA logo on the left and the title 'NVIDIA Volta Deep Learning AMI' on the right. Below the title, it says 'Sold by: NVIDIA' and 'Latest Version: 18.08.0'. A paragraph of text describes the AMI as an optimized environment for running deep learning and HPC containers from the NVIDIA GPU Cloud (NGC) container registry, noting that these containers require this AMI for GPU acceleration on AWS P3 Volta instances.

**aws marketplace**

View Categories ▾ Migration Mapping Assistant Your Saved List Sell in AWS Marketplace

**NVIDIA**

## NVIDIA Volta Deep Learning AMI

Sold by: **NVIDIA** Latest Version: 18.08.0

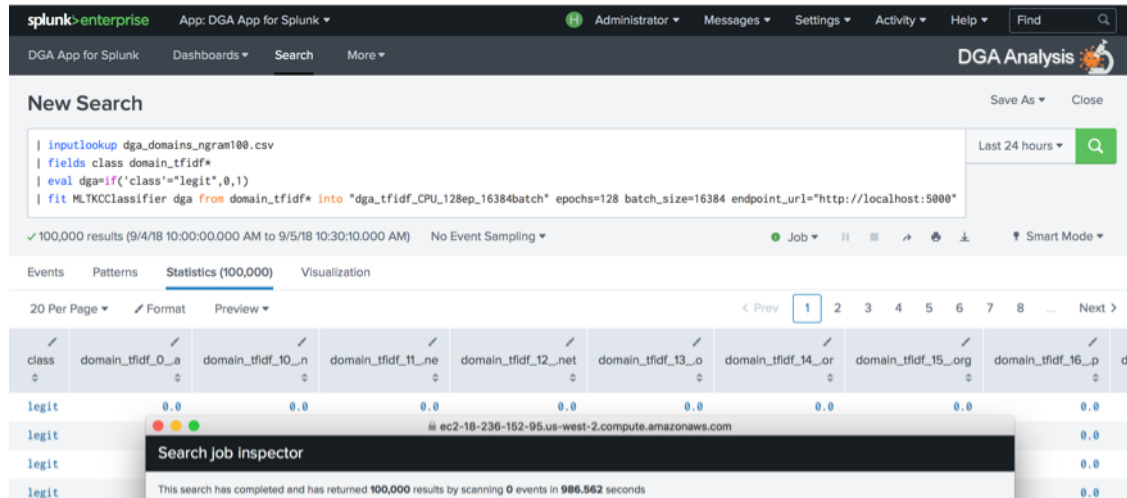
The NVIDIA Volta Deep Learning AMI is an optimized environment for running the deep learning and HPC containers from the NVIDIA GPU Cloud (NGC) container registry. The deep learning containers on the NGC container registry require this AMI for GPU acceleration on AWS P3 Volta instances.




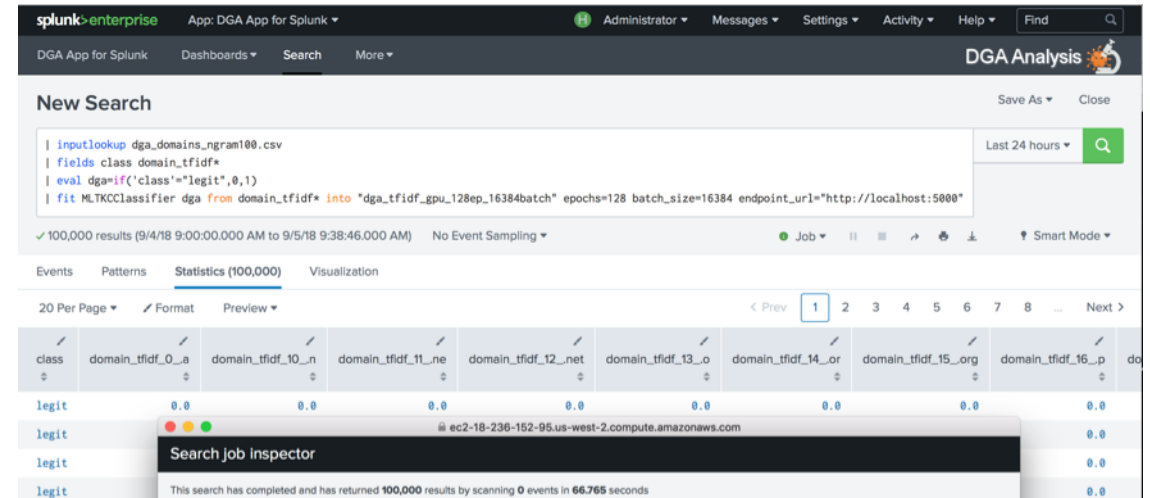
# CPU vs GPU > 15x speedup on search runtime

100K dataset | 100 dimensions | 10 layer NN

**CPU: 986 seconds (00:16:26)**



**GPU: 66 seconds (00:00:66)**



# CPU vs GPU > 25x speedup on model fitting

100K dataset | 100 dimensions | 10 layer NN

splunk>enterprise App: SA-MLTK-Co... Administrator Messages Settings Activity Help Find

MLTK Container for TensorFlow™ Dashboards Examples Search App SA-MLTK-Container

## Benchmark CPU vs GPU : 25x speedup

Save Save As View Close

```
| summary dga_tfidf_CPU_128ep_16384batch
| append
  [| summary dga_tfidf_GPU_128ep_16384batch]
| fields model_name time_fit_duration
| transpose header_field=model_name
```

✓ 1 result (before 9/5/18 10:03:51.000 AM) No Event Sampling Job

Events Patterns **Statistics (1)** Visualization

20 Per Page Format Preview

column	dga_tfidf_CPU_128ep_16384batch	dga_tfidf_GPU_128ep_16384batch
time_fit_duration	0:15:52.876940	0:00:36.857751



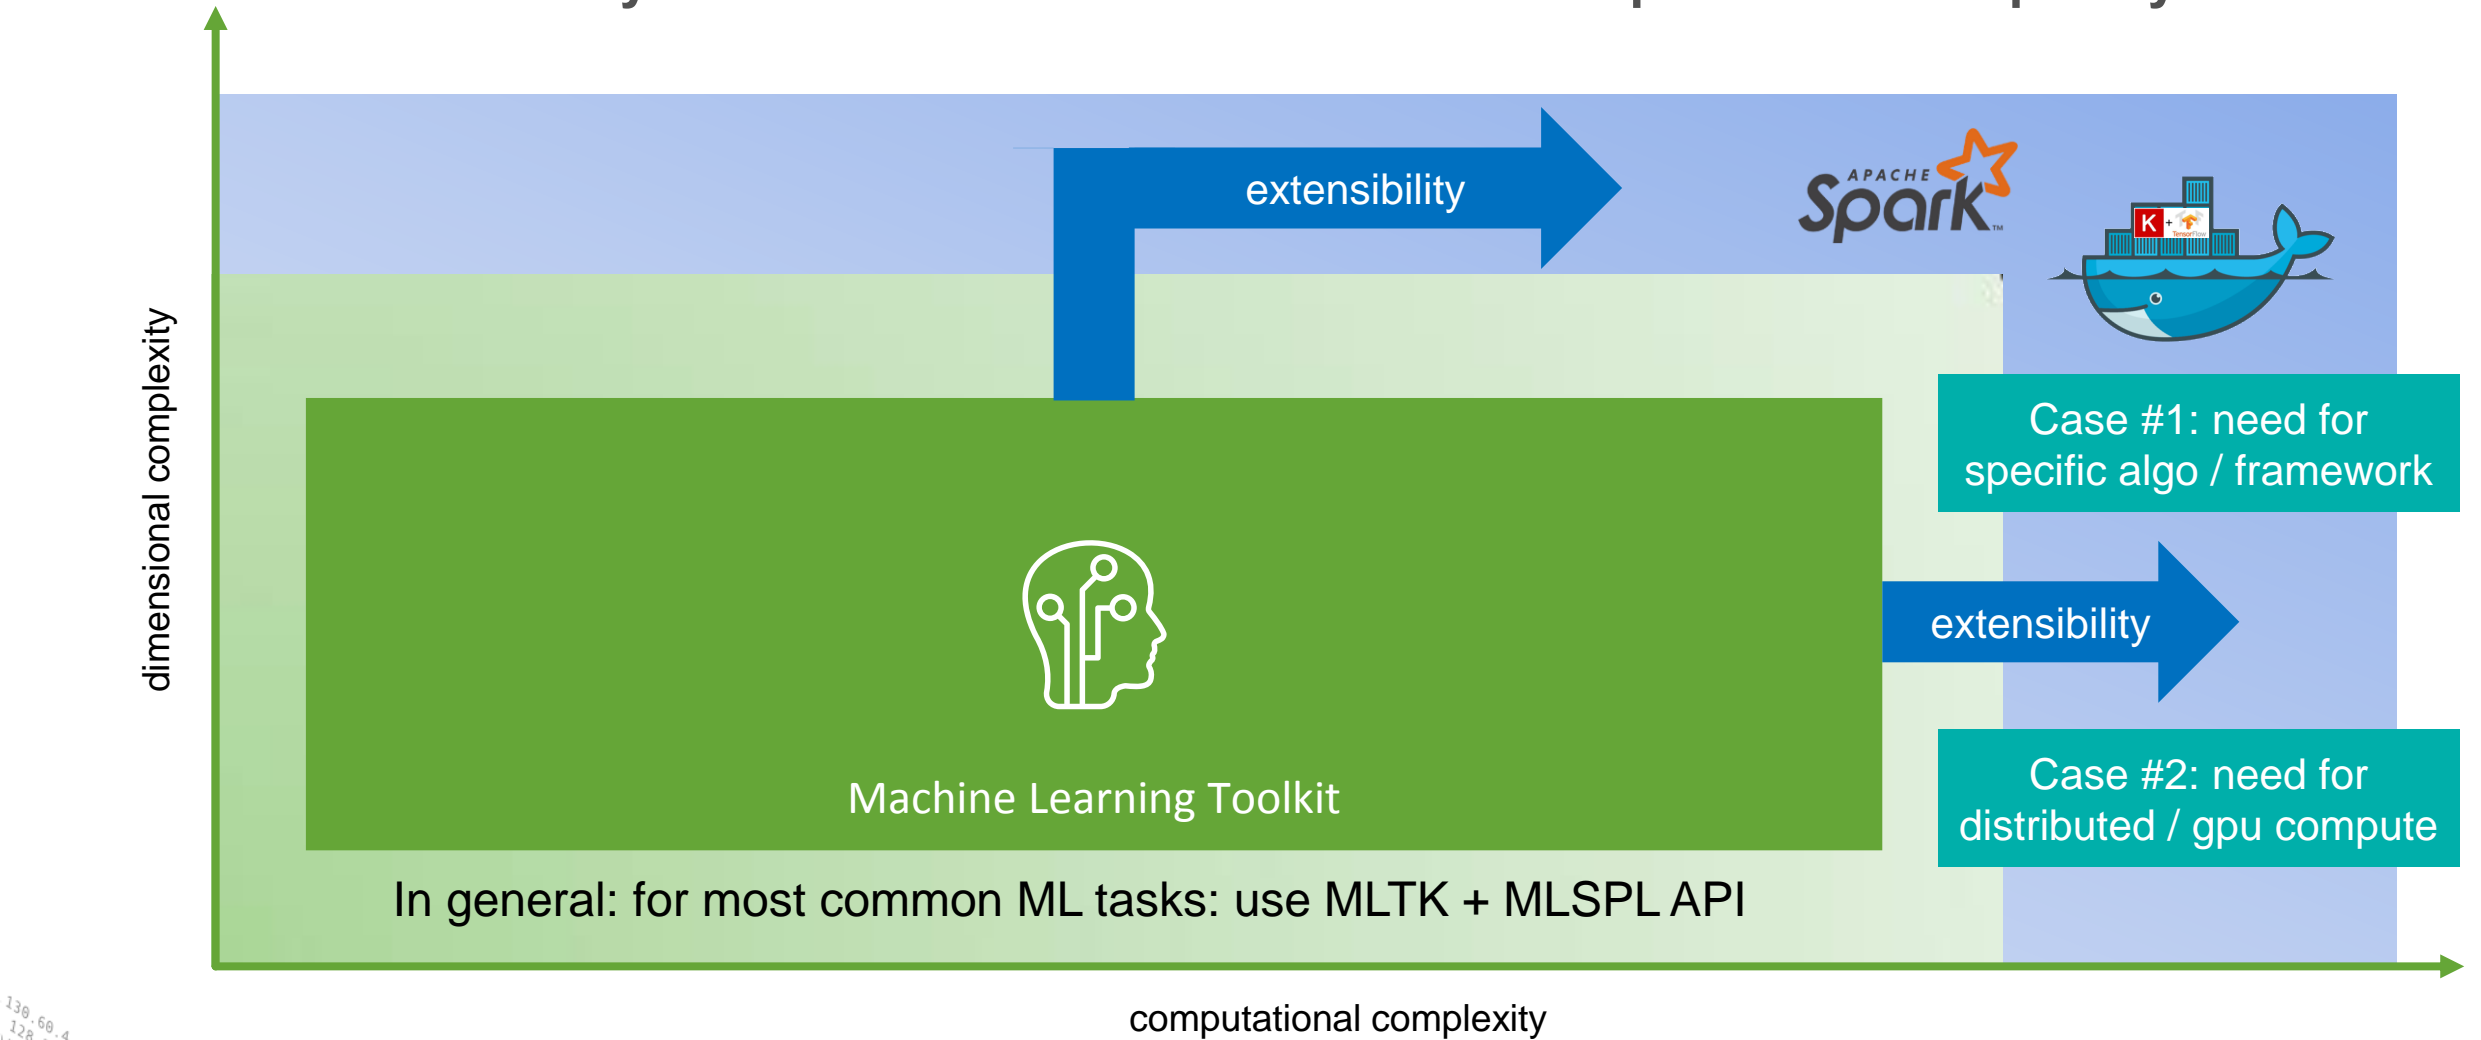


# Wrap up

---

# Recommendation Matrix

consider your ML dataset's dimensional and computational complexity






# Key Take Aways

What are the benefits of MLTK Container for TensorFlow™



Easy install and prebuilt examples in the MLTK Container for TensorFlow™ App



Customize containerized code for specific use cases using any ML/DL frameworks of choice



Flexibility to run compute intense model trainings on GPU accelerated hardware

“Available as Splunk Professional Services  
(Whiteglove Program) that you can engage  
as of today – let us know!”

MLTK Container for TensorFlow™





# Thanks | Q&A

[philipp@splunk.com](mailto:philipp@splunk.com)



# Thank You

Don't forget to **rate this session**  
in the **.conf18** mobile app

**.conf18**

**splunk>**



# Thanks to ...

**... so many amazing colleagues supporting this idea and helping in getting it real**

- ▶ Customers and Partners
- ▶ ML PMs: Manish, Andrew, Harsh, ...
- ▶ ML Eng: Xander, Lin, Chang, ...
- ▶ Other: James, Duncan, ...