



Splunking the Transaction Waterfall

Scott Garcia | Principal Architect - Enterprise Solution Delivery, T-Mobile

Andrew Koo | Manager – Enterprise Solution Delivery, T-Mobile

Gary Burgett | Staff Sales Engineer, Splunk

October 1-4 2018 | Orlando, FL

Forward-Looking Statements

During the course of this presentation, we may make forward-looking statements regarding future events or the expected performance of the company. We caution you that such statements reflect our current expectations and estimates based on factors currently known to us and that actual events or results could differ materially. For important factors that may cause actual results to differ from those contained in our forward-looking statements, please review our filings with the SEC.

The forward-looking statements made in this presentation are being made as of the time and date of its live presentation. If reviewed after its live presentation, this presentation may not contain current or accurate information. We do not assume any obligation to update any forward-looking statements we may make. In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. Splunk undertakes no obligation either to develop the features or functionality described or to include any such feature or functionality in a future release.

Splunk, Splunk>, Listen to Your Data, The Engine for Machine Data, Splunk Cloud, Splunk Light and SPL are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names, or trademarks belong to their respective owners. © 2018 Splunk Inc. All rights reserved.



Introduction

Agenda

- ▶ T-Mobile Splunk @ a Glance
 - Splunk Evolution
 - Fun Facts and Environment
- ▶ Journey to Splunk Unlimited
 - IT Transformation
- ▶ Waterfall Use Case Walk-Through
- ▶ Data Maturity
- ▶ Utilizing the Data
 - Searching and Visualizing the Data

T-Mobile @ a Glance

- ▶ 1 million+ net customer adds – 19 quarters in a row
- ▶ 10s of thousands of employees
- ▶ 100s of business and engineering applications
- ▶ Splunk customer for 7 years, strategic partner for 5 years
- ▶ 100s of teams using Splunk



Evolution of Splunk @ T-Mobile

GEN 1: Initial

- 2 TB license
- Siloed Search Heads

GEN 2: Growth

- 5.5 TB license
- Increased adoption
- Larger, more complex use cases: CDR, Retail

GEN 3: Commoditize Data

- 11.5 TB license
- Bare-metal right-sizing
- Geo redundancy added
- Mediation gen 1 + 2
- Scaling of syslog
- More structured self-serve at scale
- Broad and deep adoption
- Beginning of tiered environments (HW)

GEN 4: Advanced Self-Serve


- 'Unlimited' license
- User Engagement
- Advanced Self-Serve
- Launch new analytics capabilities (Viz & ML)
- Streamlined HW provisioning & tiered options
- Improved Availability & Performance
- Workflow Automation
- Show-back Reporting



Fun Facts & Environments

- ❑ 45 TB per day **INGESTION** (+125% Q1-Q2)
- ❑ 14,000,000 served **SEARCHES** (Q1-Q2)
- ❑ 112,000 distinct **DEVICES** supported
- ❑ 15,000 supported **FORWARDERS**
- ❑ 1,900 supported **DASHBOARDS**
- ❑ 1,500 supported **SOURCETYPES**
- ❑ 4000 distinct **USERS** (+18% Q1-Q2)
- ❑ 200 supported **APPS**
- ❑ ~ 5 per day **ONBOARDINGS**

- ❑ **PCI/Security**
- ❑ **SOX**
- ❑ **Enterprise**
 - Class A1
 - Class A2
 - Class B
 - Class X



Journey to Splunk Unlimited



IT Transformation

- ▶ 3 years ago
 - Utilized an In-house logging application
 - Organization couldn't take full advantage of the data
 - Solutions were developed with a narrow focus
 - Was not scalable and therefore unsustainable
 - Investment front loaded on the network at the expense of IT
 - Disparity in investment can lead to a delay in maturity
 - IT Transformation
 - A multi-year phased introduction of a new Enterprise IT solution
 - Consolidation of 100's of applications through new development and technology



Transition to **Splunk** Unlimited

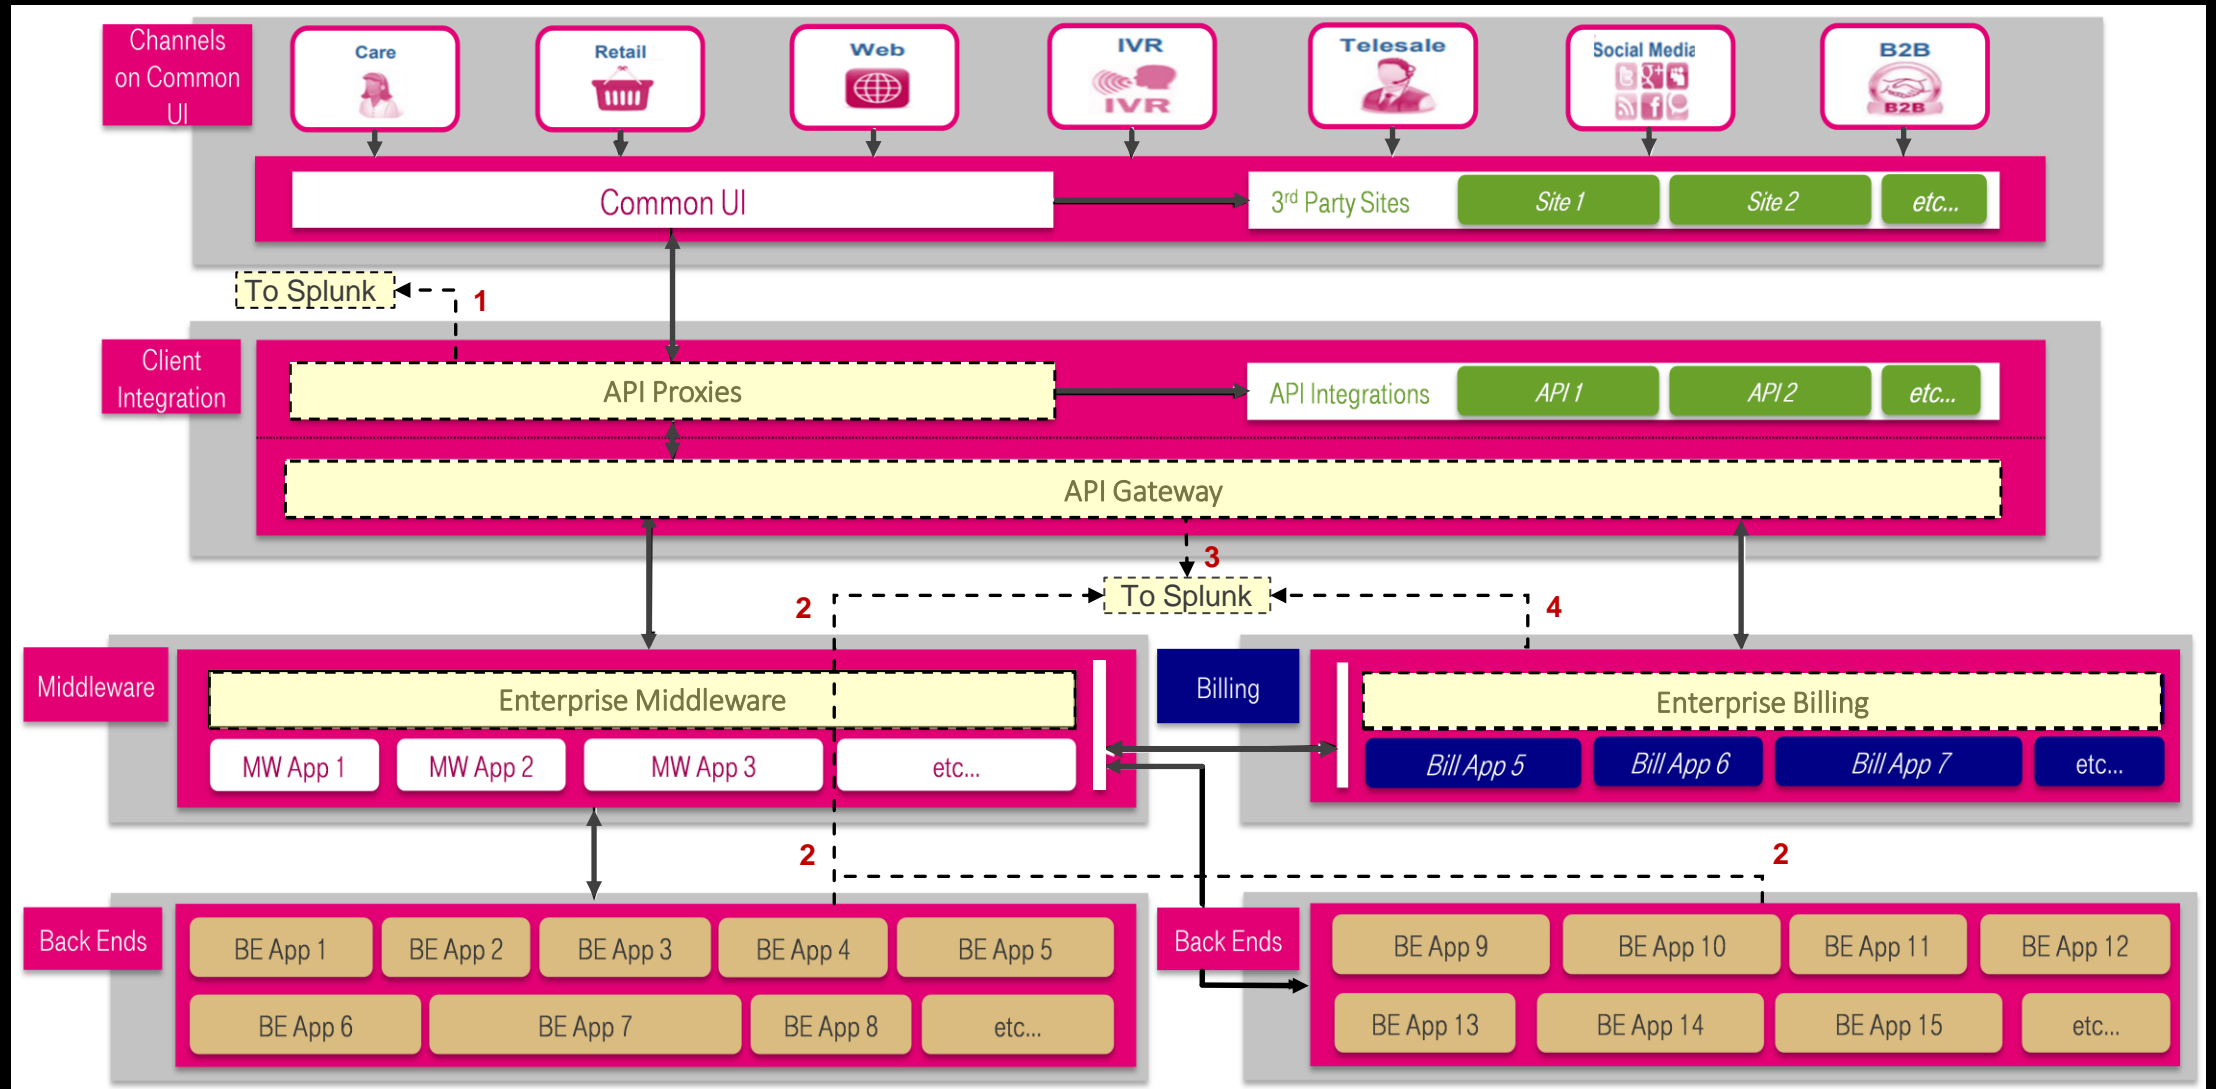
- ▶ Next Gen Centralized Logging and Monitoring
 - Value proposition and timing
 - Evaluation of solutions

***Splunk** named an enterprise solution*

- ▶ With success came challenges
 - Support of an enterprise solution vs. a targeted solution
 - Splunk knowledge transfer for general use and dashboard development.
 - Infrastructure expansion was needed to support the trend of increasing usage
 - The acceleration of Splunk adoption places pressure on license capacity

*Transitioned to an **unlimited licensing model** – 12/31/17*

IT Transformation

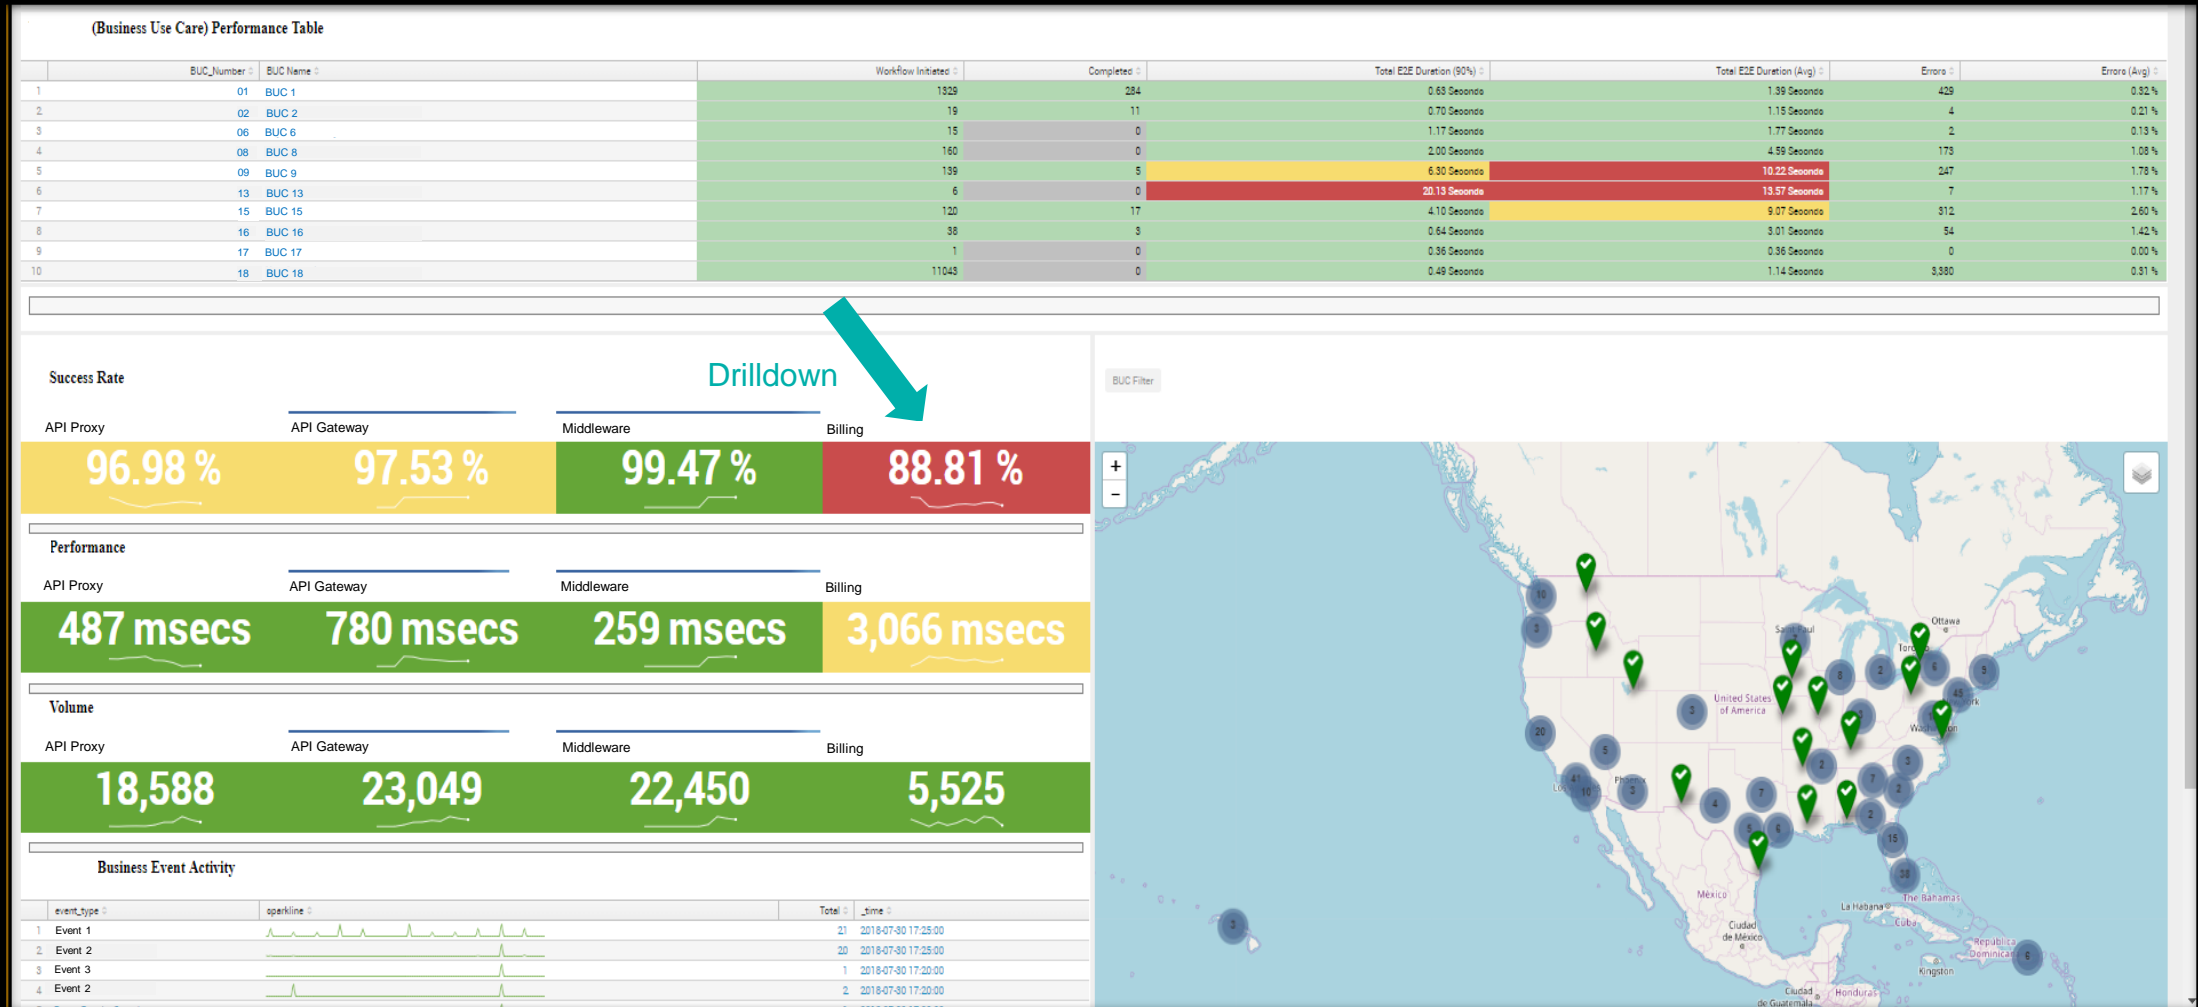


Waterfall Use Case

End to End Transactional Flows

Dashboard

Top Level Performance and Availability Dashboard



Dashboard

1st Level Drilldown – Billing Success Rate

E// Core Apps - Error Messages Resolved as Negative Responses (True) or Not (False)

Negative Response ▾	Target App ▾	PIER Queue ▾	Error Message ▾	# of Errors ▾
FALSE	App 1	Queue 1	Activity timed out	11478
FALSE	App 1	Queue 1	Username or password wrong	52
FALSE	App 1	Queue 1	JED020: JavaScript error	37
FALSE	App 1	Queue 1	Error From null: Order is NOT allowed. There is already an in-progress Customer level order.	27
FALSE	App 1	Queue 1	; ERROR_MESSAGE_EXPLANATION: UNKNOWN	23
FALSE	App 1	Queue 1	Field lm_crm_channel_code is mandatory	23
FALSE	App 1	Queue 1	ON-OC-ED0310[validateBusinessLogic]	14
FALSE	App 1	Queue 1	Channelid is not valid	14

« prev 1 2 3 4 5 6 next »

Success Rate Statistics

Number of Outlier(s)

7

Number of Hourly Intervals Below min Threshold

Errors - (15 minutes)

736

Number of Errors - p/15 min

Success Rate- Avg (15 Minute)

95.68

Avg Success Rate - p/15 min
MIN MAX

Predicted Success Rate vs Actual


Sensitivity (max 1)

.97


Confidence Interval (max 100)

97.5

Include MTC Window:

☒ Include MTC Window☐ Non-MTC Window Only

Weekly Historical View








Dashboard

2nd Level Drilldown – Activity Timeout

Time Setting: Custom time Target Application: APP-1 Service: Service 1 Operation: Operation 1 Error Message: All [Submit](#) [Hide Filters](#)

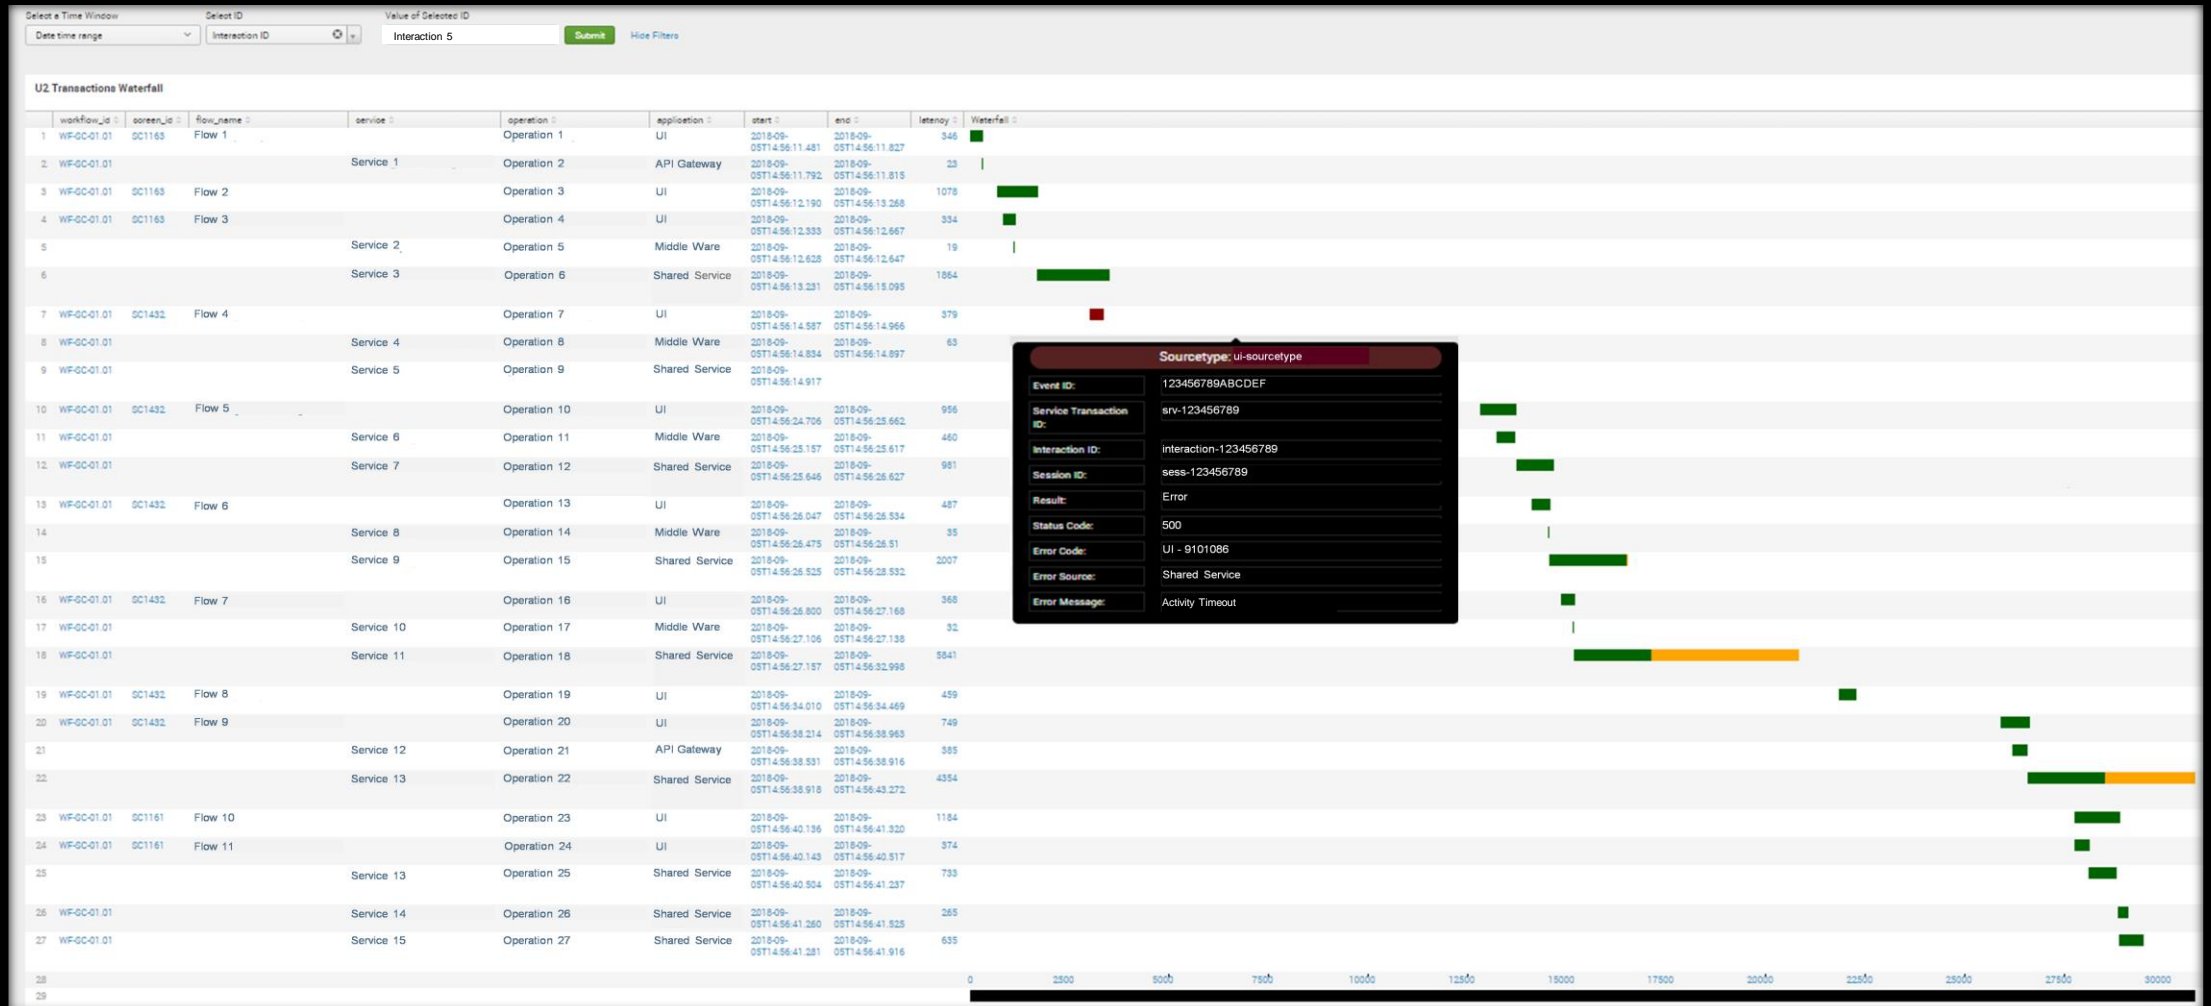
Error Result Table

	_time	zone	Target	service	operation	result	negative_response	error_message	error_code	rtt	Interaction		end	start
1	2018-09-05 15:45:14.398	Zone 1	APP-1	Service 1	Operation 1	0	FALSE	Activity Timeout	9101086	29	Interaction 1		1536187544.398	1536187484.398
2	2018-09-05 16:15:16.772	Zone 1	APP-1	Service 1	Operation 1	0	FALSE	Activity Timeout	9101086	58	Interaction 2		1536189346.772	1536189286.772
3	2018-09-05 16:15:18.403	Zone 1	APP-1	Service 1	Operation 1	0	FALSE	Activity Timeout	9101086	125	Interaction 3		1536189348.403	1536189288.403
4	2018-09-05 16:27:42.976	Zone 1	APP-1	Service 1	Operation 1	0	FALSE	Activity Timeout	9101086	65	Interaction 4		1536190092.976	1536190032.976
5	2018-09-05 16:06:38.893	Zone 1	APP-1	Service 1	Operation 1	0	FALSE	Activity Timeout	9101086	37	Interaction 5		1536188828.893	1536188768.893

[Drilldown](#) [Q](#) [↓](#) [i](#) [↺](#) 1m ago

Dashboard

3rd Level Drilldown – Waterfall Dashboard





Waterfall | Value Observed

- ▶ Tier 1 and service desk organizations with little knowledge of application processing can easily reference a transactional flow.
- ▶ Tier's 2 & 3 are provided with the ability to visually analyze the sequence and contents of events
- ▶ The waterfall view is reusable allowing for its integration within many dashboards across T-Mobile's monitoring solutions
- ▶ Interactions between Interfaced applications can be observed
- ▶ Performance and validation teams are capturing performance related metrics
- ▶ Both production and non production teams have access to the waterfall view



Data Maturity

Data Maturity

Logging best practices necessary to obtain transactional flows

Data Content Standards

- ▶ Cross-platform transaction IDs
- ▶ Field naming conventions
- ▶ Distribute ownership

Log/Event Structural Standards

- ▶ Transport (input) method standards
- ▶ Drive data providers to preferred input methods
- ▶ Handling complex payloads w/ Complex Event Processing (CEP)



Utilizing the Data

Searching the Data

Now that we have the data, what can we do with it

Different Platforms, Different formats

- ▶ Multiple middleware technologies where instrumentation takes place, each with their own format
 - XML
 - JSON
 - Raw text/KV pairs
- ▶ Schema-on-the-fly saves the day
- ▶ If the information is there, we can normalize it

Rich fields for filtering/grouping

- ▶ Sessions
- ▶ Activities
- ▶ Interactions
- ▶ Applications
- ▶ Channels
- ▶ Customers
- ▶ Operations
- ▶ Latencies
- ▶ Return Codes
- ▶ Error Messages

Searching the Data

Now that we have the data, what can we do with it

Consistent Fields = Simple Searches

index=transactions (sourcetype=mw OR sourcetype=core_gw OR sourcetype=core_proxy)
) cust_id=\$cust_id\$

Not so simple in practice

Myriad Use Cases

- ▶ Support can filter on a particular Customer ID to troubleshoot individual user issues
- ▶ Latency and Error Rates can be measured and baselined by Operation, Environment, Channel, etc.
- ▶ See individual transactions end-to-end

Visualizing the Data

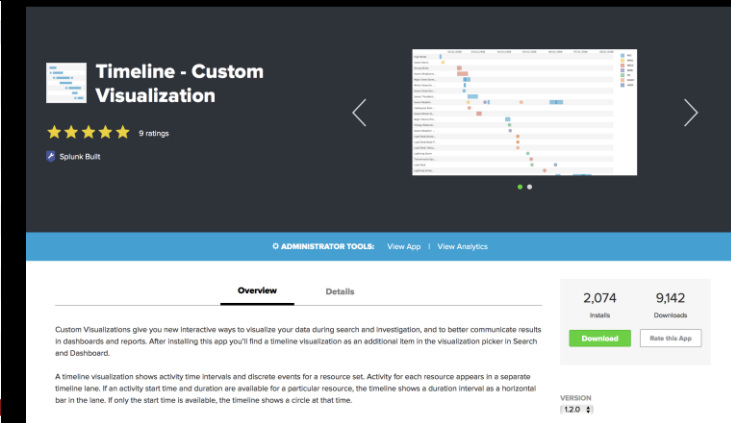
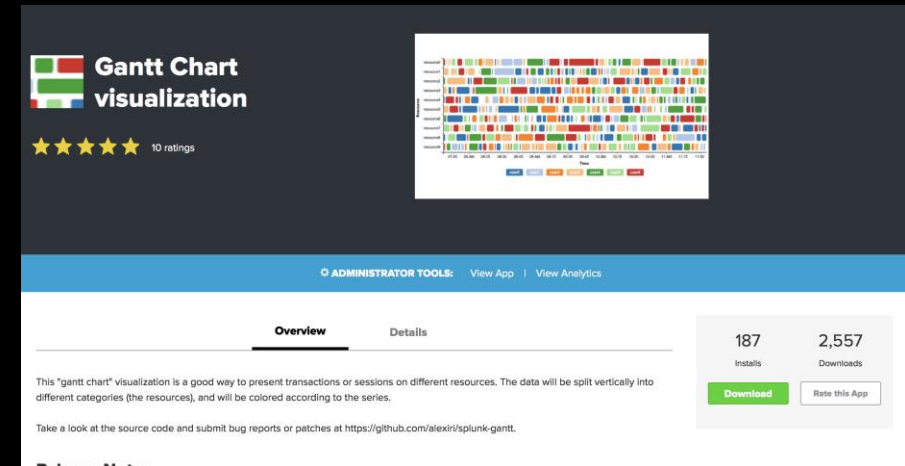
How do I visualize transaction durations in Splunk?

Options for nonstandard viz

- ▶ Splunkbase
- ▶ Custom Visualization Framework
- ▶ Extend Current Visualizations with JS Extensions

Settled on Extension of Table View

- ▶ Simple, minimal code
- ▶ Leverage existing functionality
- ▶ Data Density

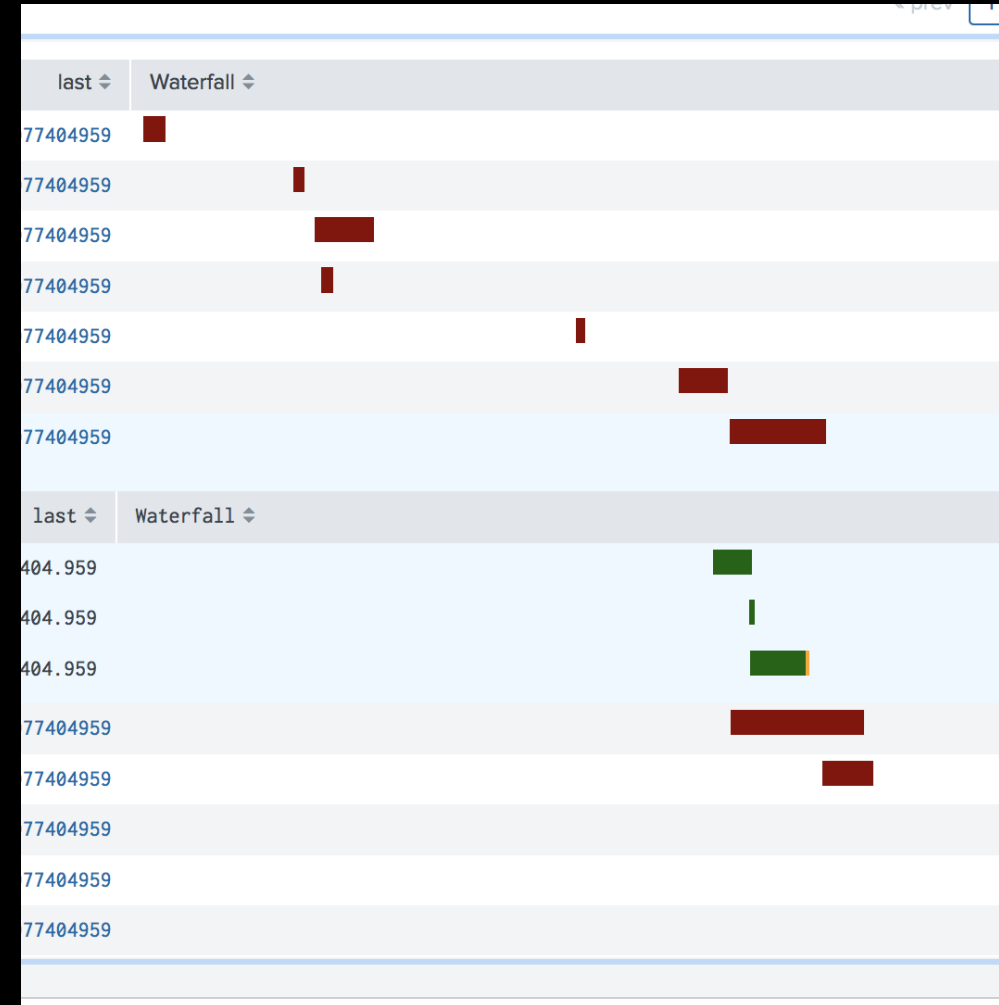


Visualizing the Data

How do I visualize transaction durations in Splunk?

Searching to build the Waterfall

- ▶ eval and rex
 - Your best friends for data manipulation
- ▶ eventstats
 - Perform aggregate stats while maintaining the original data
- ▶ appendpipe
 - Perform aggregate stats and append them to the table



Bring it together

```
index=soa (sourcetype=api_gateway OR
sourcetype=messaging_gateway OR
sourcetype=backend_gateway) transaction_id=e3q29ir1u29zoi
```

```
| eventstats first(_time) AS first last(_time) AS last
```

```
| eval duration=first-last
```

```
| eval buffer=100*(_time-last)/duration
```

```
| eval diff=100*latency/(1000*duration)
```

```
| reverse
```

```
| eval Waterfall=tostring(buffer)+"," +tostring(diff)
```

```
| eval spent=latency+" ms"
```

```
| appendpipe
```

```
    [stats max(eval("total," +tostring(duration))) AS Waterfall
max(eval(tostring(duration*1000)+" ms")) AS spent
```

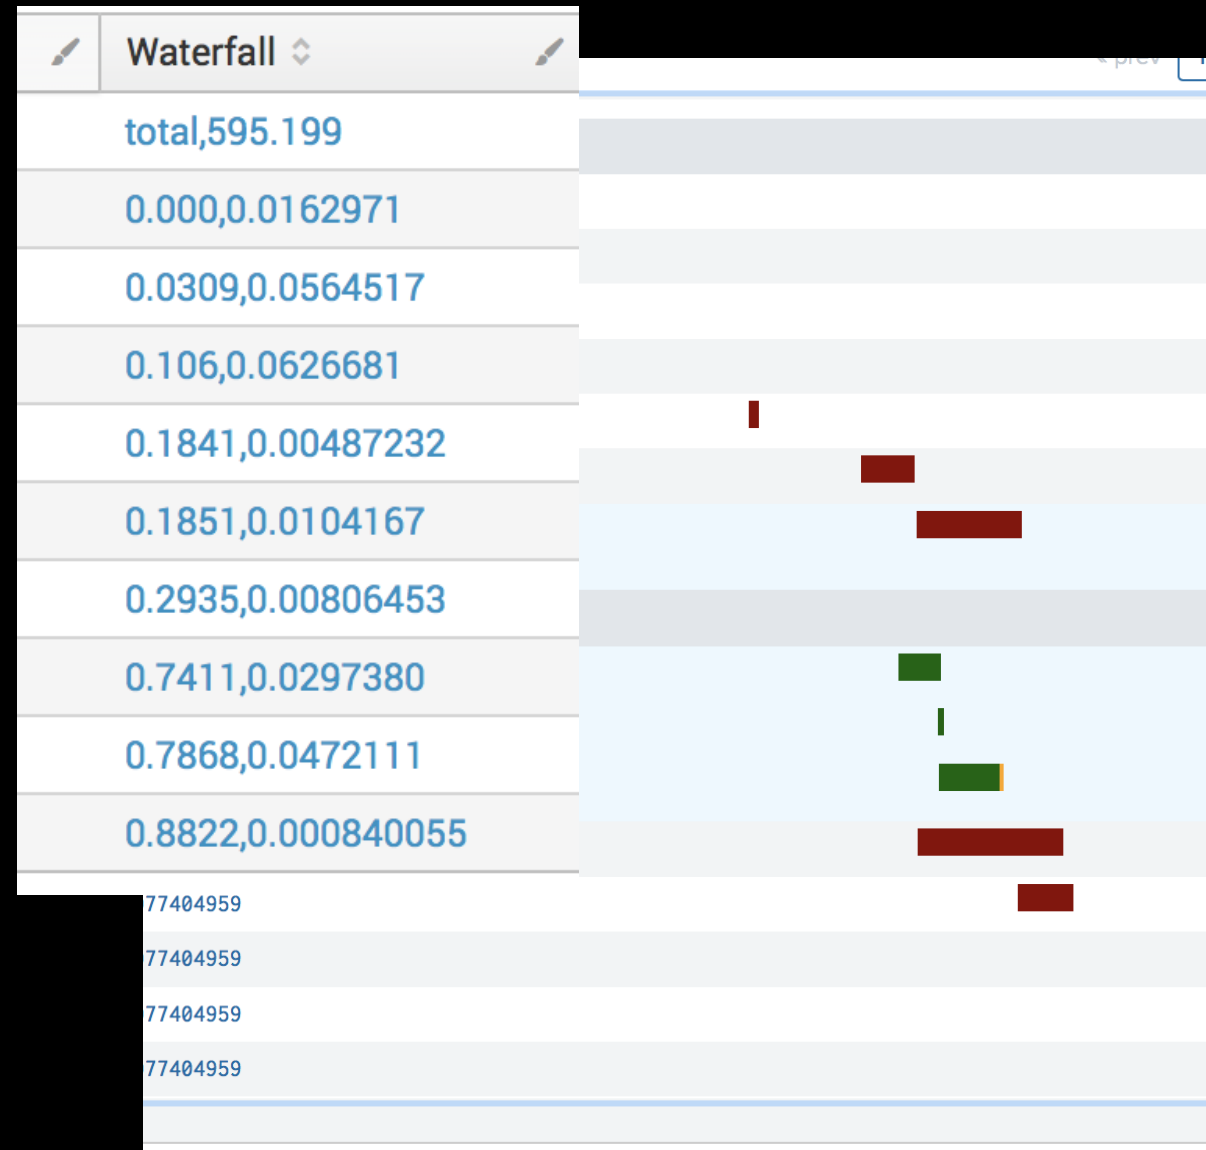
```
    | eval buffer=tonumber(-1)]
```

```
| appendpipe
```

```
    [stats min(eval("markers," +tostring(duration))) AS Waterfall
```

```
    | eval buffer=tonumber(101)]
```

```
| sort buffer | fields - buffer, diff
```



Key Takeaways

Good Logging Practices

1. Employ Good Logging Practices

2. Related Splunk .conf2018 Sessions

- **IT1256 - Zipkin & Splunk: Tracing Transactions Across Your Ecosystem**


Tuesday, Oct 02, 4:45 p.m. - 5:30 p.m.

Tom Martin, Staff Practitioner, Splunk

- **IT1847 - Splunking Application Performance: Traditional APM and Beyond**

Wednesday, Oct 03, 12:45 p.m. - 1:30 p.m.

Gary Burgett, Staff Sales Engineer, Splunk



Q&A

Scott Garcia
Andrew Koo
Gary Burgett



Thank You

Don't forget to **rate this session**
in the **.conf18** mobile app

