Using Splunk ML for Threat Hunting

Joe Partlow, ReliaQuest

>whoami

Joe Partlow – CTO, ReliaQuest

Joe has been in the IT and informaxtion Security industry for 20+ years and most recently been working with simulated attack & defense networks, security analytics, building big data platforms and machine learning. Reliaquest partners with the worlds largest enterprise splunk customers performing analyst, engineering/architecture and content development functions.

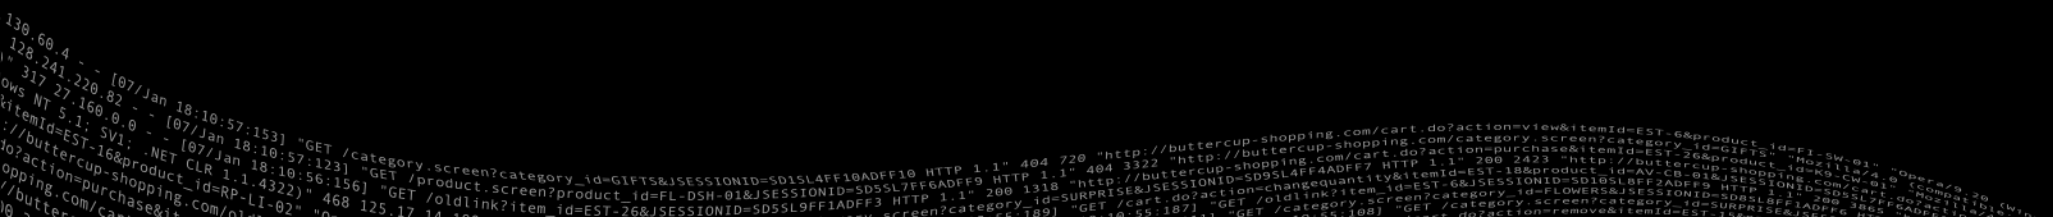

splunk> .conf18

Inspired by gibson

What ml/ai is not

- Algos will not replace good analysts, just another resource
  - But when can I have wintermute in my soc???
- Many times, successful hunt campaigns achievable with just effective searches
- Many good products/models fail because of poor and incomplete data
- Haven’t even scratched the surface of basic ML, let alone “ai” …. yet
- Gone too soon:

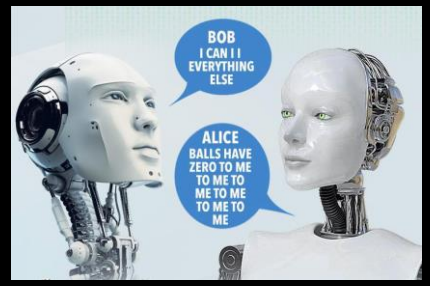

Ensuring success

- Clean data will make or break your training models
  - Include filtering, black/white lists to remove large datasets that could skew results
- Data normalization across the various sources
- Know your data! (supervised and unsupervised learning)
- Numeric data shouldn’t always be treated as such (ie. port numbers)
- Incorporate red teamers with your data science team for better domain knowledge
- Continuously retrain your models as the environment changes

(most) Relevant algorithms

- YMMV but below are some algorithms well-suited for common hunt campaigns:
  - Detect Numeric Outliers – Useful for determining weird status codes or abnormal event IDs
  - Detect Categorical Outliers – Useful for finding weird DNS queries or abnormal user-agents/page requests
  - Cluster Numeric Events – Helpful for finding outlying host login counts or application usage counts
  - Prediction and Forecasting algorithms might be better suited until the environment is stable and baselined

Web attack use case

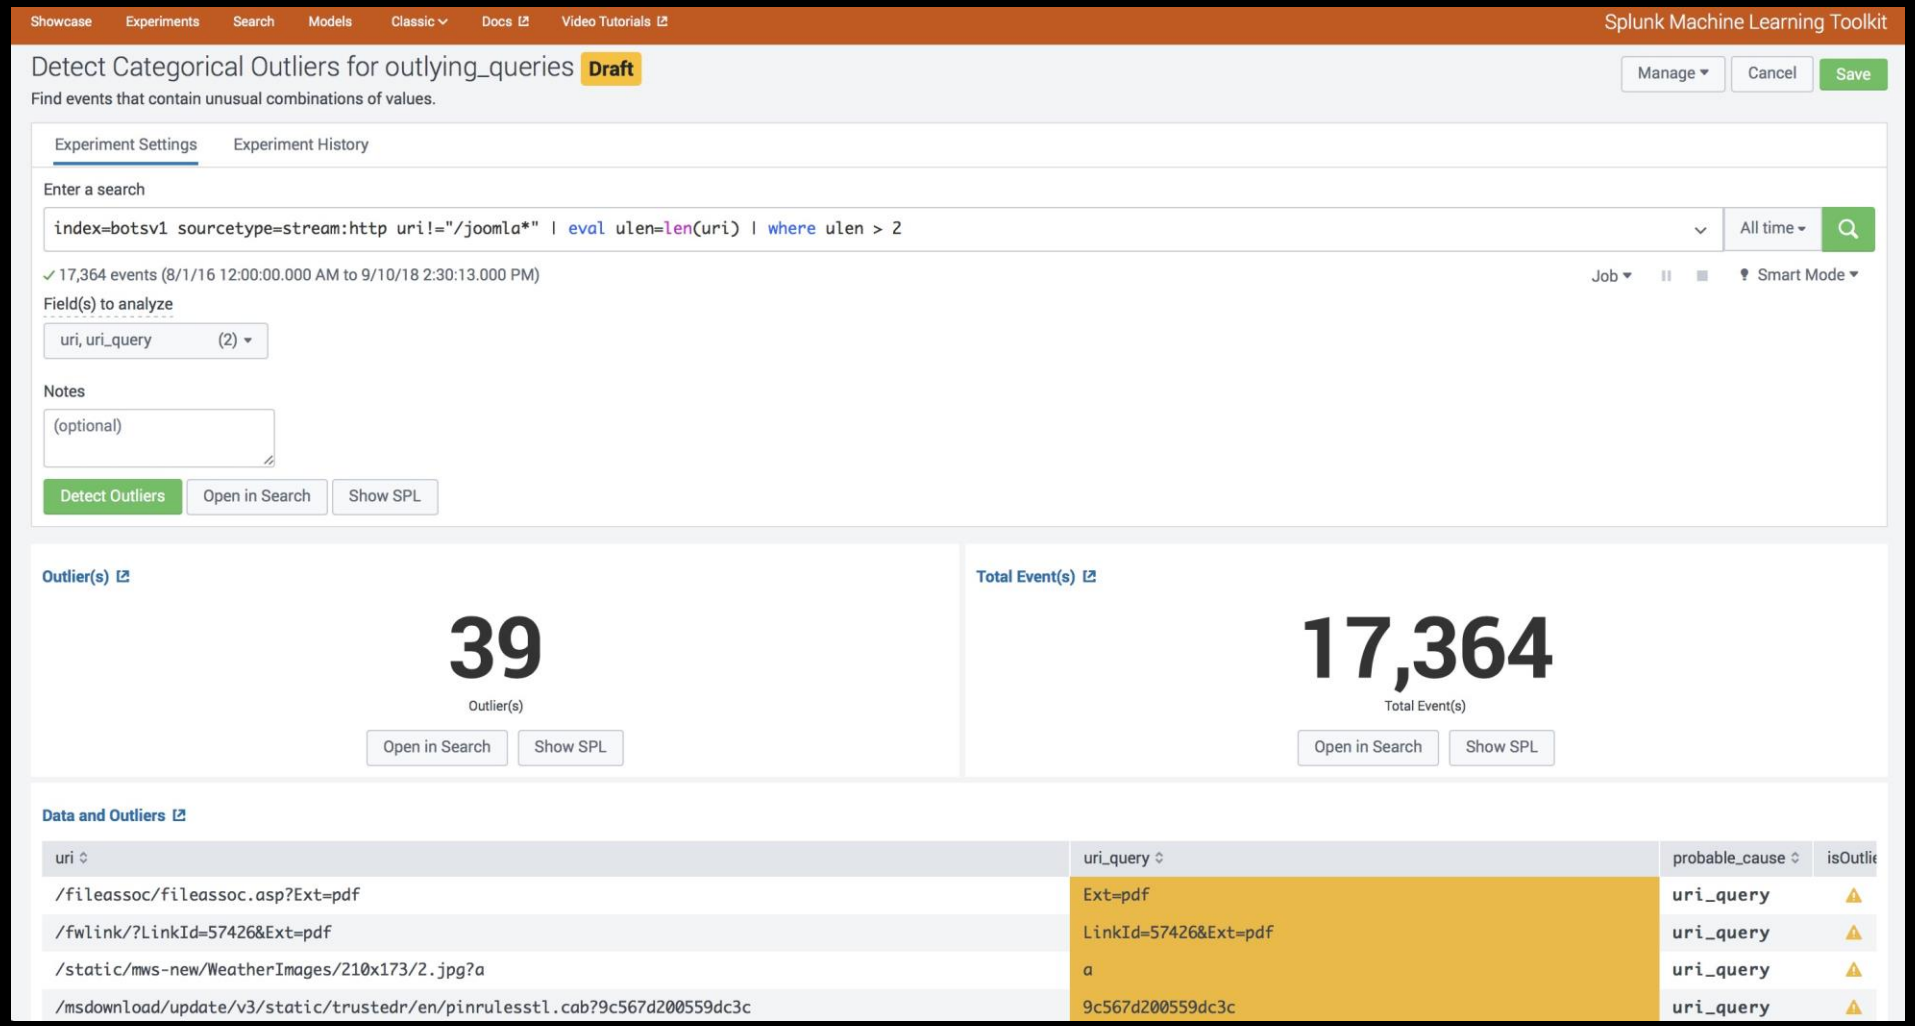

splunk> .conf18

Web attack use case

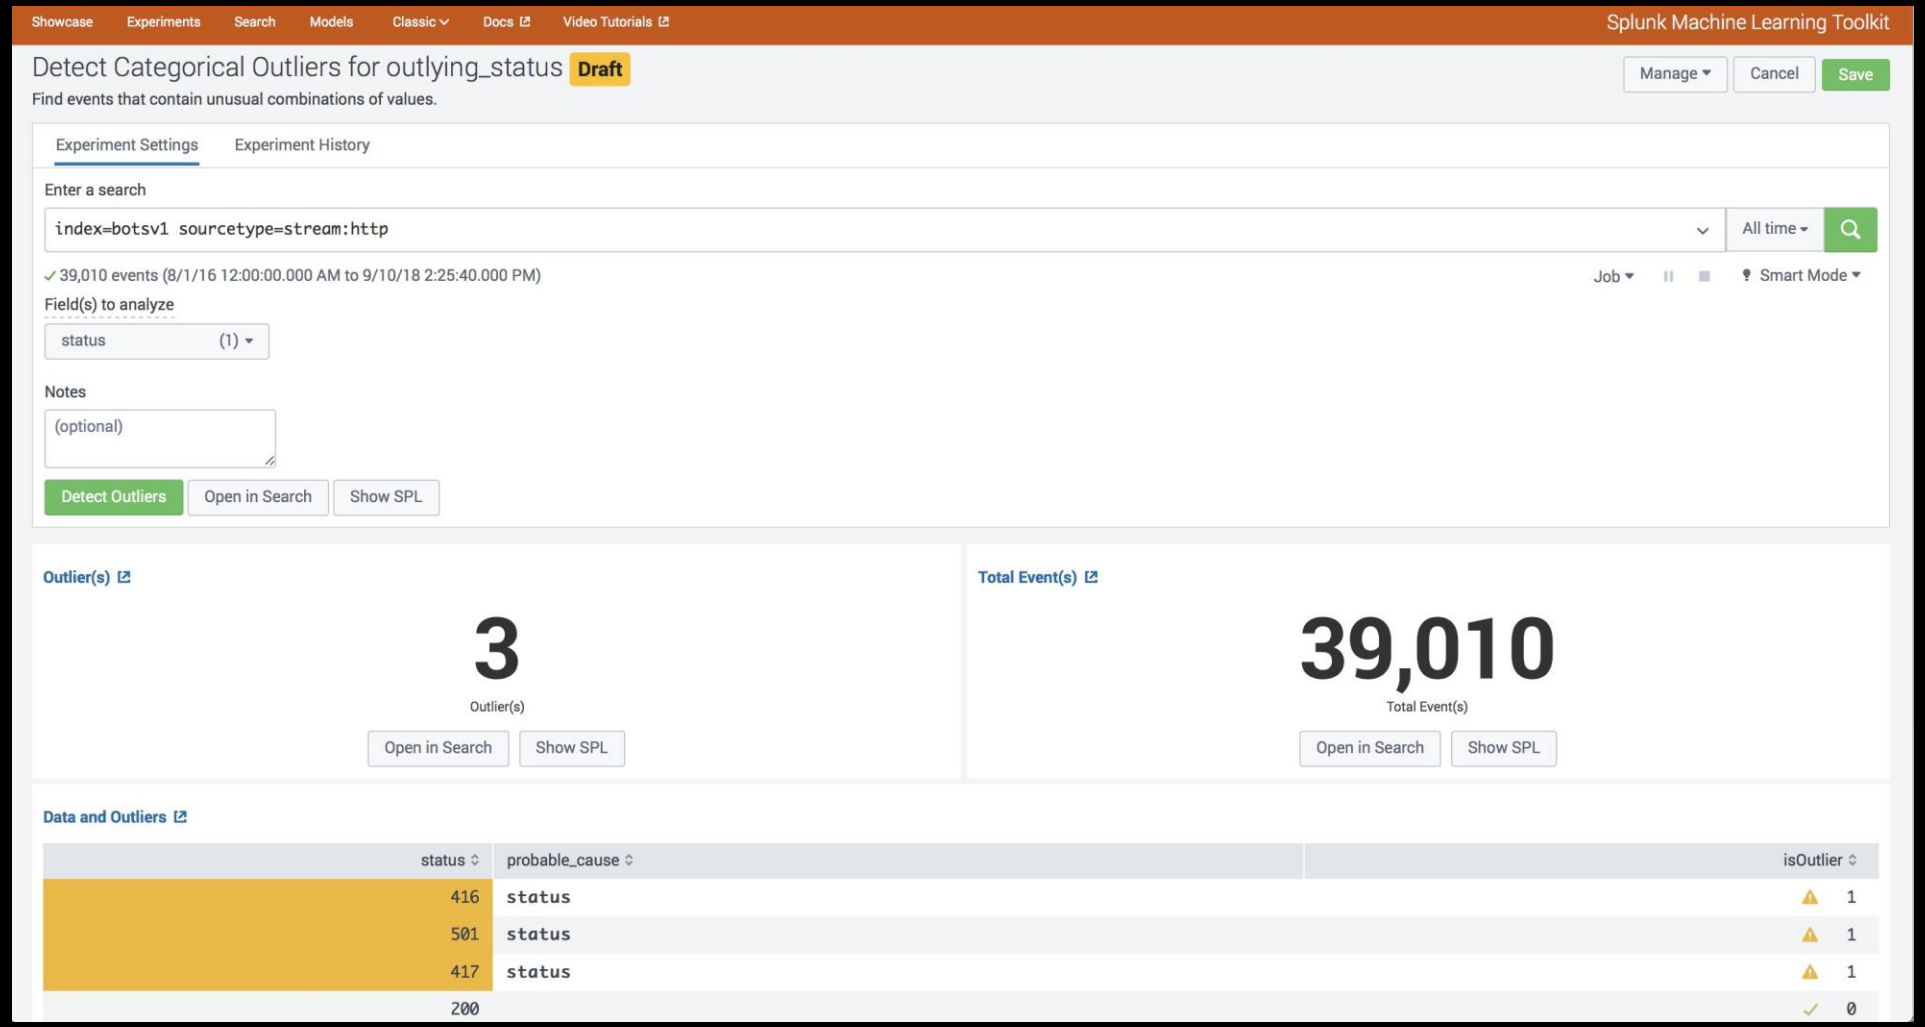

splunk> .conf18

Web attack use case

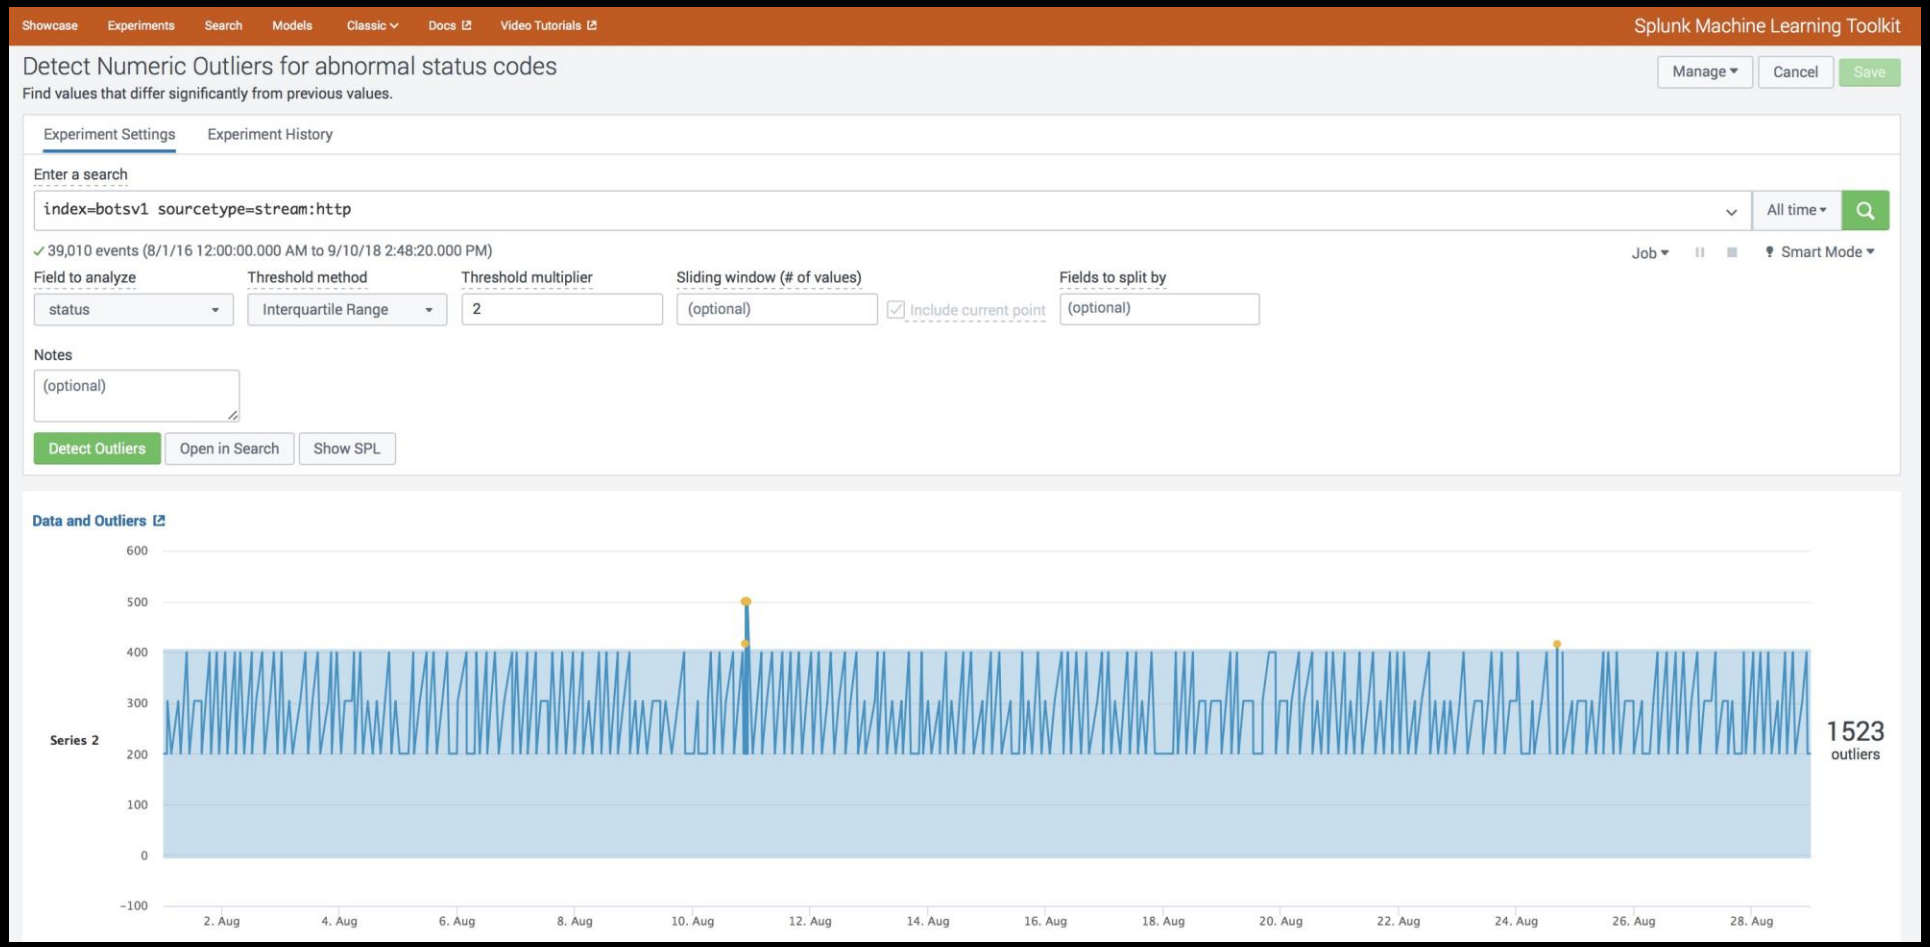

splunk> .conf18

Web attack use case

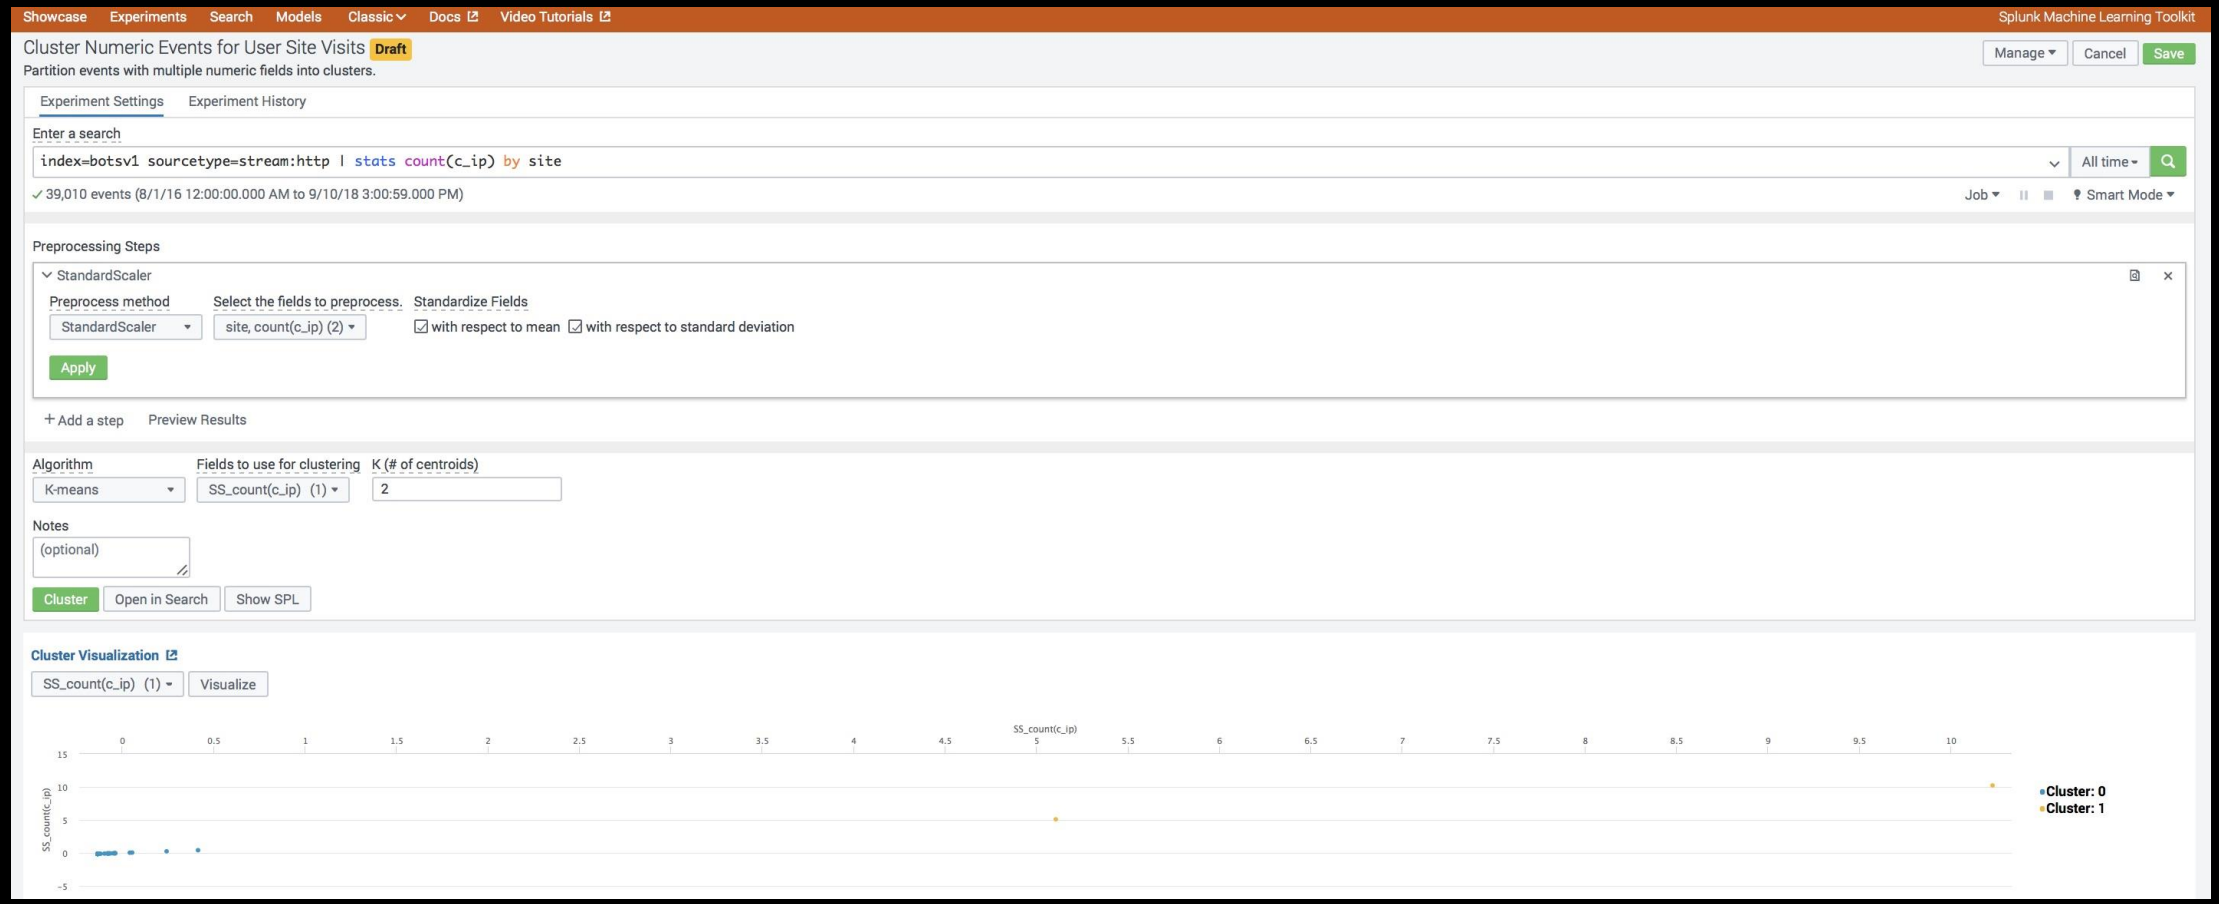

splunk> .conf18

Breaking the models

- Already proven for image classification deep learning*
- Similar to SIEM issues, overload the data ingestion with enough noise that the “abnormal becomes normal”
- Attackers are already good at blending in (living off the land, pivoting, etc.)

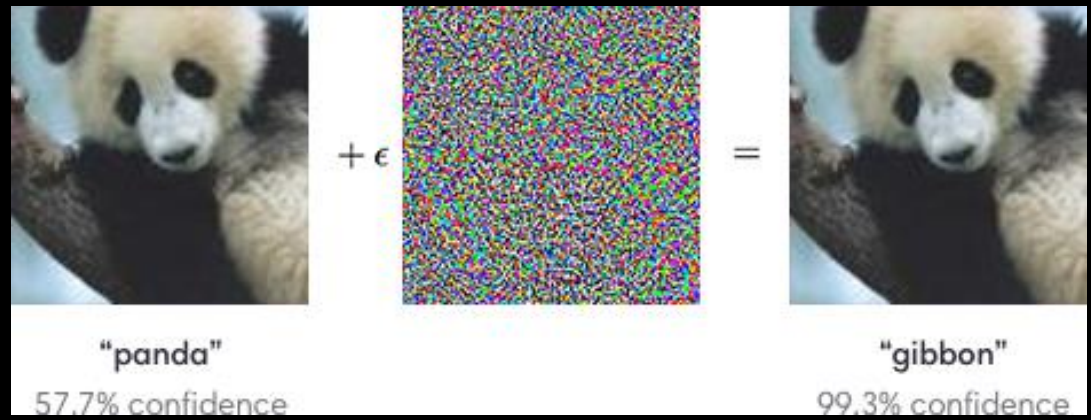

** https://blog.openai.com/adversarial-example-research/*

Future enhancements

- Field is progressing amazingly fast. Just because something isn't possible now, give it 6 months!
- Move towards stacked/ensemble learning to avoid "jack of all trades, master of none" algos
- Build up and better utilize belief networks
  - Attempts to increase accuracy by adding conditional dependencies
  - Excellent blackhat talk by raffael marty - https://www.slideshare.net/zrlram

THANK YOU

Questions? jpartlow@reliaquest.com

Don't forget to rate this session in the .conf18 mobile app