



# ElasticIntel



Scalable Threat Intel Aggregation in  
AWS



# Obligatory Who I Am slide..

Presenter: Matt Jane

- Builder/Automator
- I put things in clouds
- Open Source Advocate
- <Insert credential alphabet soup here if it makes you feel better>



**The ideas and opinions expressed  
in this talk....**

Are not those of my employer, my imaginary  
friend, or anyone else

---

Warning: There will be Memes

# The slide before the other slides...

---

Backstory and why I  
started this project

Goals

---

Current status of the  
project

Future plans

# Why Build Elastic Intel?

- Researched products and services that could provide “Threat Intelligence”
  - The findings we really bleak
  - REALLY expensive
  - The automation possibilities for most solutions were BAD
-

# Value of threat intelligence (Stolen from Scott Roberts)

---

- |                       |                                |
|-----------------------|--------------------------------|
| 1. Your own incidents | 4. Peers/Sharing Communities   |
| 2. Vendor Reports     | 5. 3rd Party Paid Intelligence |
| 3. Honeypots          |                                |

Most vendor's "Threat Intelligence" was just IOCs

Most IOCs were just recycled open-source feeds

Roughly 5-10% of IOCs produced by vendors were original material



A meme featuring Woody and Buzz Lightyear from the movie Toy Story. Woody is on the left, looking concerned. Buzz is on the right, holding a green and purple light gun and making a hand gesture with his fingers spread. The background is a blurred indoor setting.

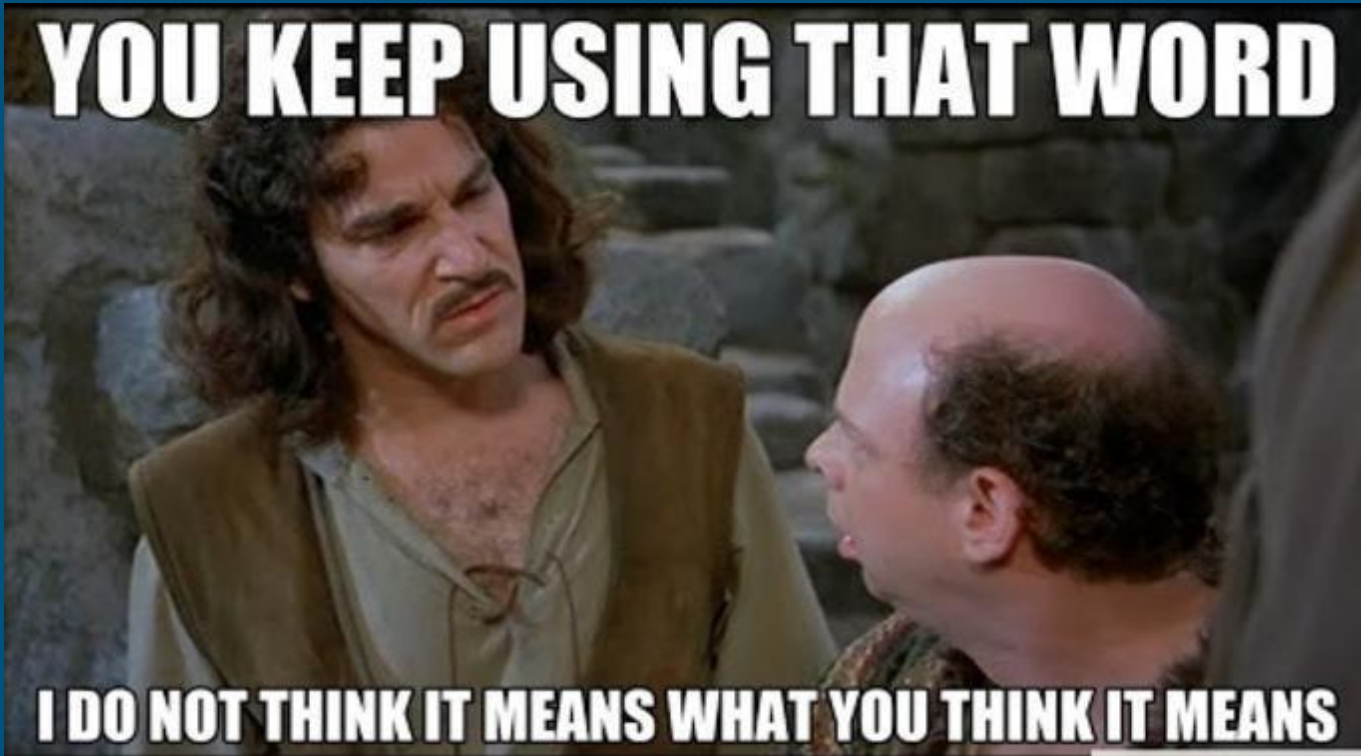
**IOCS**

**IOCS EVERYWHERE**

# “Intelligence”

- IOCs. JUST IOCs
- Occasionally we get WHOIS info
- Sometimes even LOSE what little context may have been available

# “Intelligence”



# Goodbye, context

- BotScout SSH brute force → “Malicious Activity”
- Phishing URL targeting 0365 → “Credential Harvesting”

# Targeted Intelligence!

Turns out to be a string search for some keywords

- Limited to 5 or 10 keywords
  - Searches only across certain types of IOCs
  - Google search results....
  - Marketing campaigns...
-

What if we charged for  
**EVERYTHING!**

## Entry cost

---

Ranged between  
\$100k - \$400k

## Gets you...


---

A (Sometimes) pretty  
web GUI  
Ability to make basic  
queries through web UI  
Limited # of users  
(usually 1-5)

# Pay To Play Model

---

- Want API access? \$\$
- Want more users? \$\$
- Want to add a feed? \$\$
- Want more contextual searches? \$\$\$
- Want BULK api access? That's your first born + \$\$\$
- Want a chat bot? LOLOLOLOLOL





# Automation & Integrations

We have an API but....

- You can only query it for single indicators
  - You can only query it once a minute
  - You can only query it 20 times a day (seriously)
  - If you want more than 1 result in a list, you need to query it for every result
  - If you want more than 50 queries, you can buy packs of 50/day for \$10k
-

Well...



# There's Got to be a Better Way!



# ElasticIntel

- Primary component is Elasticsearch
  - Aggregates Threat Feeds (IOC feeds)
  - Provides API for automating searches
  - Slack bot for quick search
  - Low Barrier to Entry (stand up with one command)
-

# GOALS

- Gather roughly the same data (minus the 5-10% of custom data)
- Make it cost effective (I run a personal version)
- Make it performant (1k queries/min or better).



# Continued...

---

- Make it usable
- Make it valuable regardless of company size
- Zero Maintenance (Or as close as possible)
- API as a first class citizen

# “Serverless”

---



- 
- Don't worry about patching
  - Do worry about credentials and sensitive information



# AWS Services

- Elasticsearch service
  - Lambda
  - SNS
  - API Gateway
  - S3
  - Cloudwatch
  - KMS
  - IAM
-

# Elasticsearch service

---

Managed Elasticsearch Cluster

# Lambda

---

Serverless compute: You give it code, it runs it for you

# SNS

---

Simple Notification Service – Pass Message  
and signals between resources

# API Gateway

---

Acts as a RESTful proxy to Lambda

# KMS

- Key Management Service - Store access keys securely

# IAM

Identity and Access Management -  
Permissions

S3

Storage

Cloudwatch

Logging



# Code!

EVERYTHING is written in  
Python3

Please stop writing things  
in Python2


No really, stop it.

---



# Infrastructure as Code

- Infrastructure is code too
  - Terraform controls everything
  - Python wrappers around Terraform abstract all actions
-



# Feed Scheduler Lambda

- Runs once an hour
  - Checks the schedule of all defined feeds
  - If a feed is scheduled to be retrieved, it publishes a message to the SNS topic
-

# Ingest Lambda Function

- Subscribed to SNS Topic
  - Triggers on SNS Topic message publish
  - Retrieves the data from the specified feed URL
  - Parses the feed based on its type (txt, csv, etc.)
  - Uploads to Elasticsearch Service
-

# Lambda Performance

- Lambda's are time bound (5m max)
- You only pay for used time
- Therefore....

# Lambda optimization

- Whatever you're trying to do, it better finish within 5 minutes
- The less time it takes your lambda to run, the cheaper it will be

# Lambdas and Threading

- Threading is still a thing
- Its just...slightly harder
- Process Pools and Thread Pools not available
- Named pipes still work just fine though!

Sauce: <https://aws.amazon.com/blogs/compute/parallel-processing-in-python-with-aws-lambda/>

# IO bound tasks can still be Parallelized in Lambda

- You just have to try a little harder
- (PS, WHOIS lookups are pretty good candidates for this 😊 )



Let's talk performance...

- What's the difference between:
- Making 100 queries
- Making 10 queries for 10 documents each

## TCP/IP overhead is a real thing

- It doesn't matter much when making a few requests
- When you make 5-10k in a few seconds...

# Whois Lambda

- Runs every 3 minutes, across 15 different regions
- Grabs all IP addresses without WHOIS information
- Retrieves the whois information for them
- Updates the IOC in Elasticsearch with the whois data

# Feeds and Feeds.d/

- All feeds are described as json blobs
  - The json files representing feeds are all located in the feeds.d/ directory of the project
  - Adding a feed is as simple as adding to an existing json file, or creating one in the directory
-

# Example Feed

```
1 {
2   "feed_name": "Bambenek C2 IP Master High Confidence",
3   "feed_type": "csv",
4   "indicator_type": "ipaddress",
5   "feed_url": "http://osint.bambenekconsulting.com/feeds/c2-ipmasterlist-high.txt",
6   "field_mapping": {
7     "separator": ",",
8     "source": ["ipaddress", "description", "created", "source_url"],
9     "destination": ["ipaddress", "description", "created", "source_url"],
10    "has_headers": false
11  },
12  "check_interval": [8]
13 }
```

# End Result

- 37 feeds = 250k IOCs/day
  - TCO = (roughly) \$1,600/yr.
  - (Personal = \$40-\$50/month)
  - Unlimited API access
  - 10k+ queries/min
  - Full-text search
-

# Contributing (in order of difficulty)

## Get your feet wet

---

- Add more feeds
- Build some Dashboards in Kibana
- Submit Feature Requests for ChatBot

## Go for a brisk swim

---

- Work on API-based feed handlers
- Write unit tests
- Add a new integration like Shodan or Pastebin!

# Contributing Continued...

---

- Don't see anything up there that appeals to you or that you feel comfortable with? No problem!
- Feature requests are very welcome. Just tag with [feature request] in the issue title on GitHub.
- Questions? Hit me up on twitter or drop me an email



# Links

---

Repo: <https://github.com/securityclippy/elasticintel>

Blog: <https://blog.securelyinsecure.com/>

Twitter: @PansyMcCoward

Email: [SecurityClippy@securelyinsecure.com](mailto:SecurityClippy@securelyinsecure.com)



# THANKS