

**It's assembler,
Jim, *but not as
we know it!***

Morgan Gangwere

DEF CON 26



Dedication

Never forget the
shoulders you stand on.

Thanks, Dad

whoami

Hoopy Frood

Been fiddling with Linux SOCs since I
fiddled with an old TS-7200
EmbeddedARM board

I've used ARM for a lot of services:
IRC, web hosting, etc.

I've built CyanogenMod/LineageOS,
custom ARM images, etc.



A word of note

There are few concrete examples in this talk. I'm sorry.

This sort of work is

- One part science

- One part estimation

- Dash of bitter feelings towards others

- Hint of "What the *fuck* was that EE thinking?"

A lot comes from experience. I can point the way, but I cannot tell the future.

There's a lot of seemingly random things. Trust me, It'll make sense.

ARMed to the teeth

From the BBC to your home.

Short history of ARM

Originally the Acorn RISC machine
Built for the BBC Micro!

Acorn changed hands and became
ARM Holdings

Acorn/ARM has never cut silicon!

Fun fact: Intel has produced ARM-
based chips (StrongARM and
XScale) and still sometimes does!

The ISA hasn't changed *all that much*.







Embedded Linux 101

Anatomy of an Embedded Linux device.

Fundamentally 3 parts

Storage

SoC/Processor

RAM

Everything else? Bonus.

PHYs on everything from I2C, USB to SDIO

Cameras and Screens are via MIPI-defined protocols, CSI and DSI respectively

At one point, they all *look* mostly the same.

What is an SoC?

Several major vendors:

- Allwinner, Rockchip (China)
- Atheros, TI, Apple (US)
- Samsung (Korea)

80-100% of the peripherals and possibly storage is right there on die

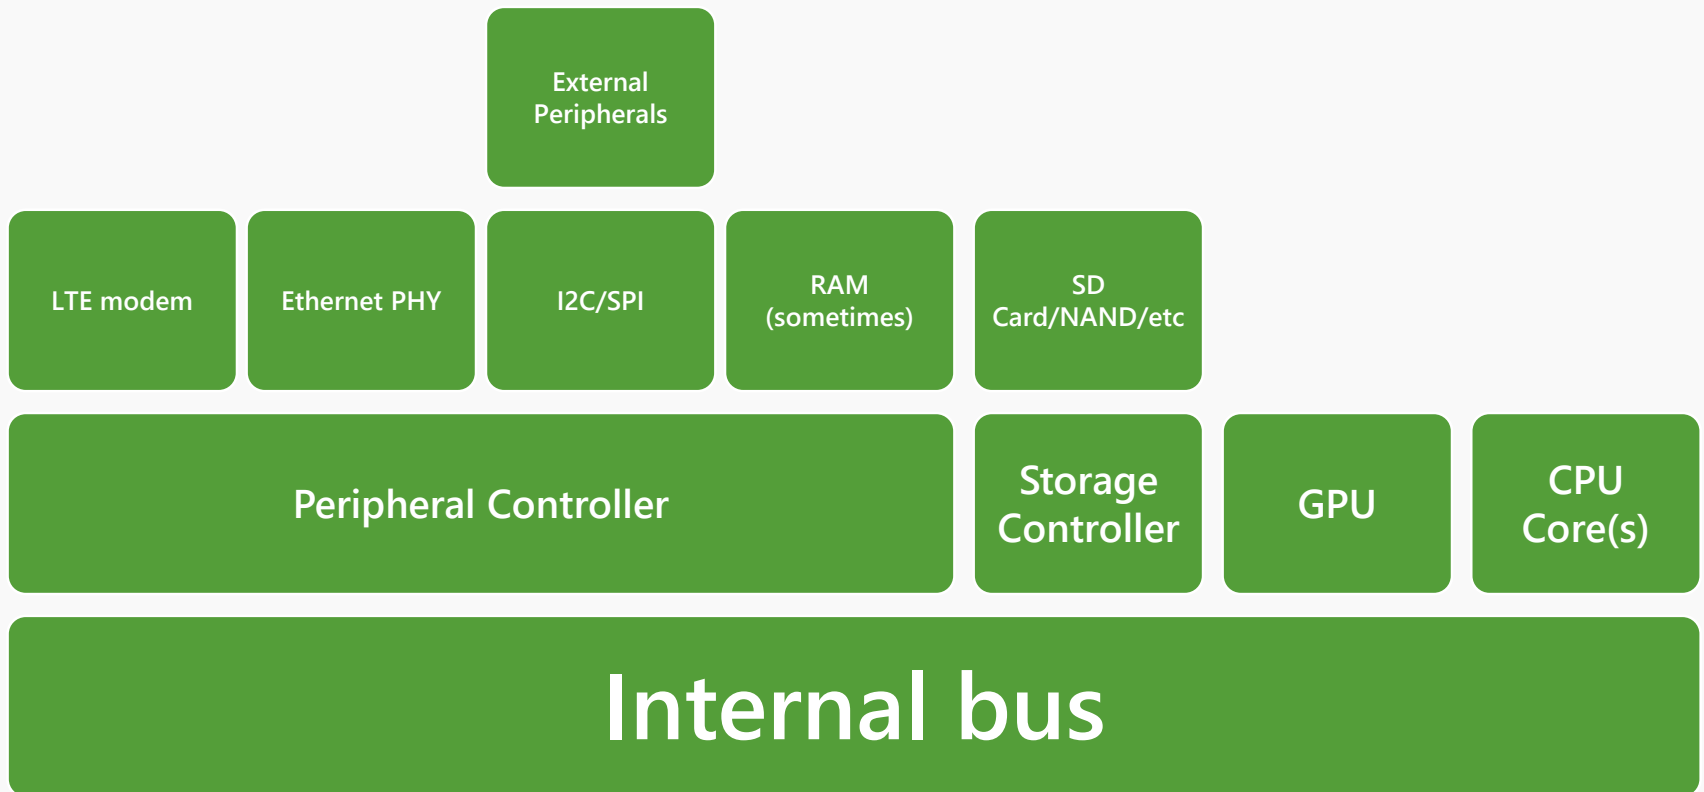
Becomes a “just add peripherals” design

Some vendors include SoCs as a part of other devices, such as TI's line of DSPs with an ARM SoC used for video production hardware and the like.

In some devices, there may be multiple SoCs: A whole line of Cisco-owned Linux-based teleconferencing hardware has big banks of SoCs from TI doing video processing on the fly alongside a DSP.



What is an SoC?



3. BLOCK DIAGRAM

Figure 3-1 shows the block diagram of the R8.

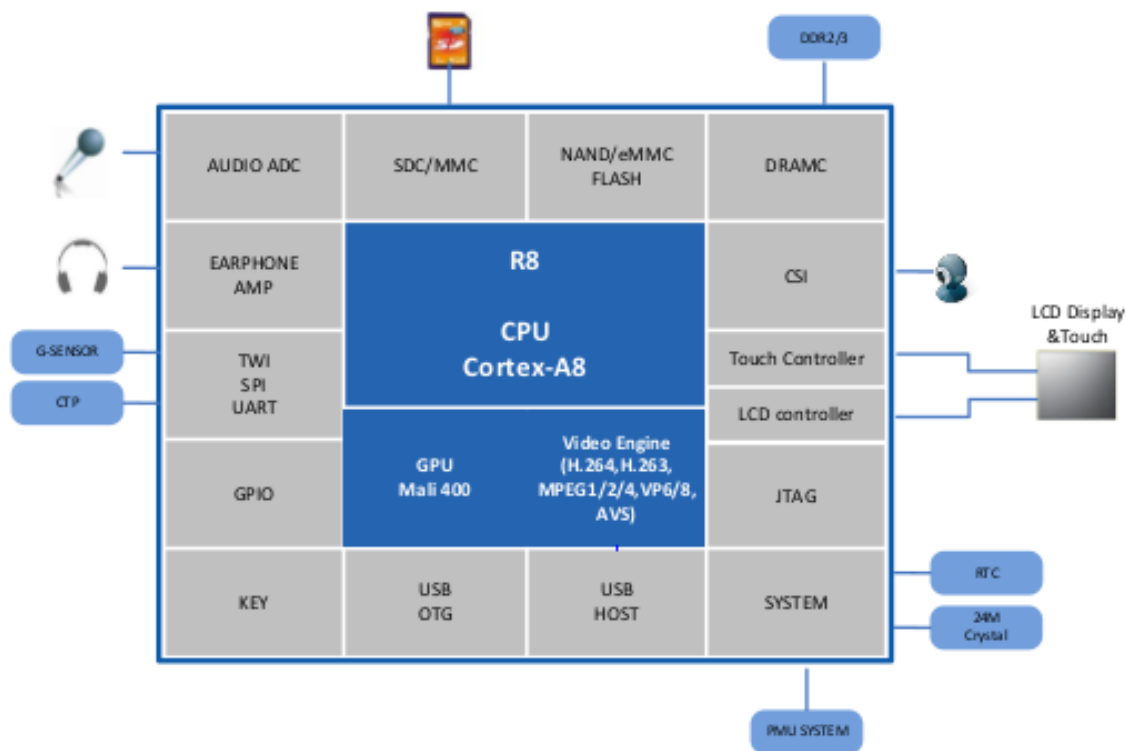
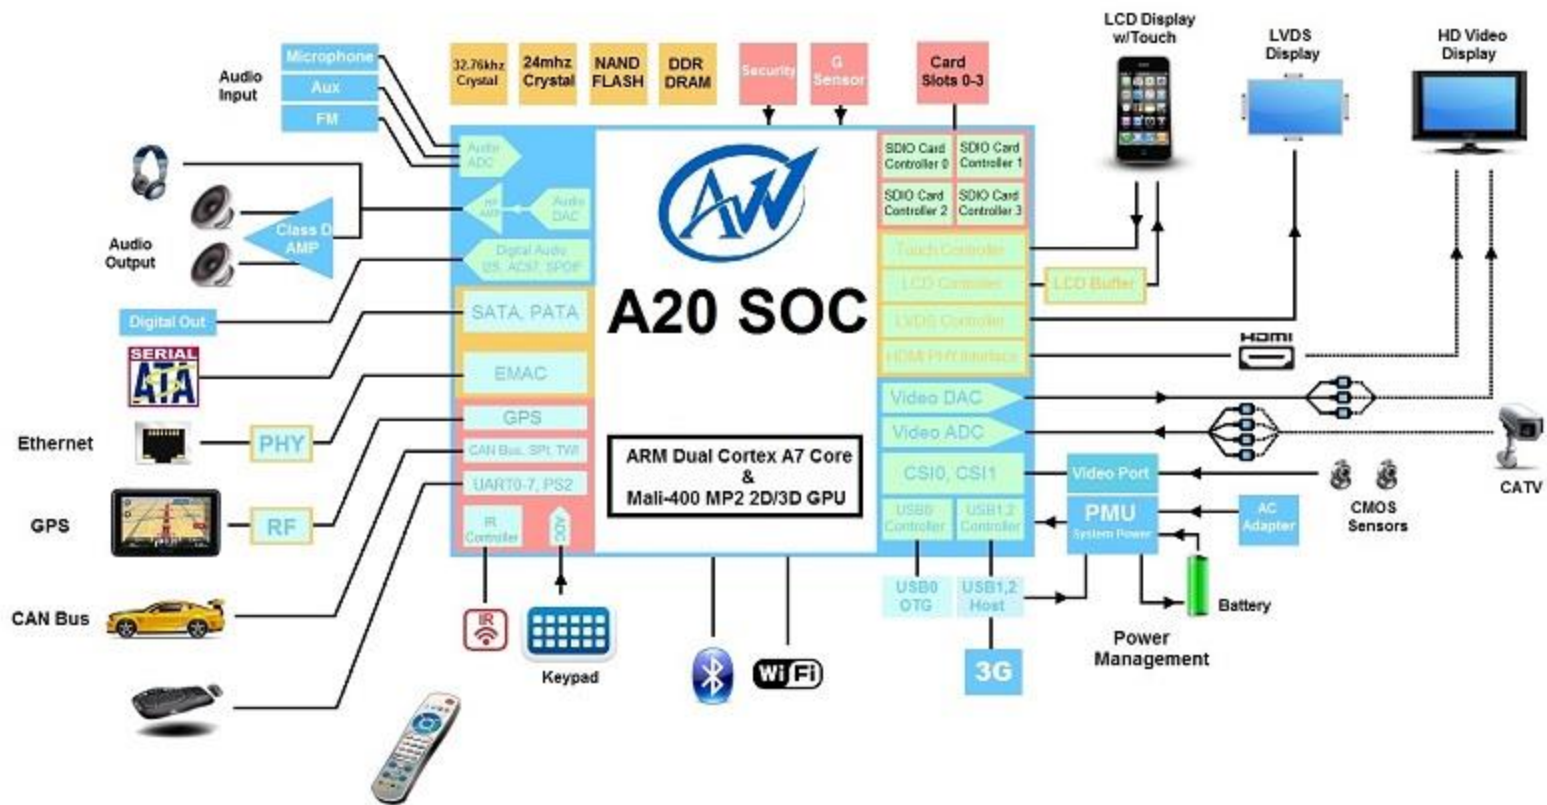
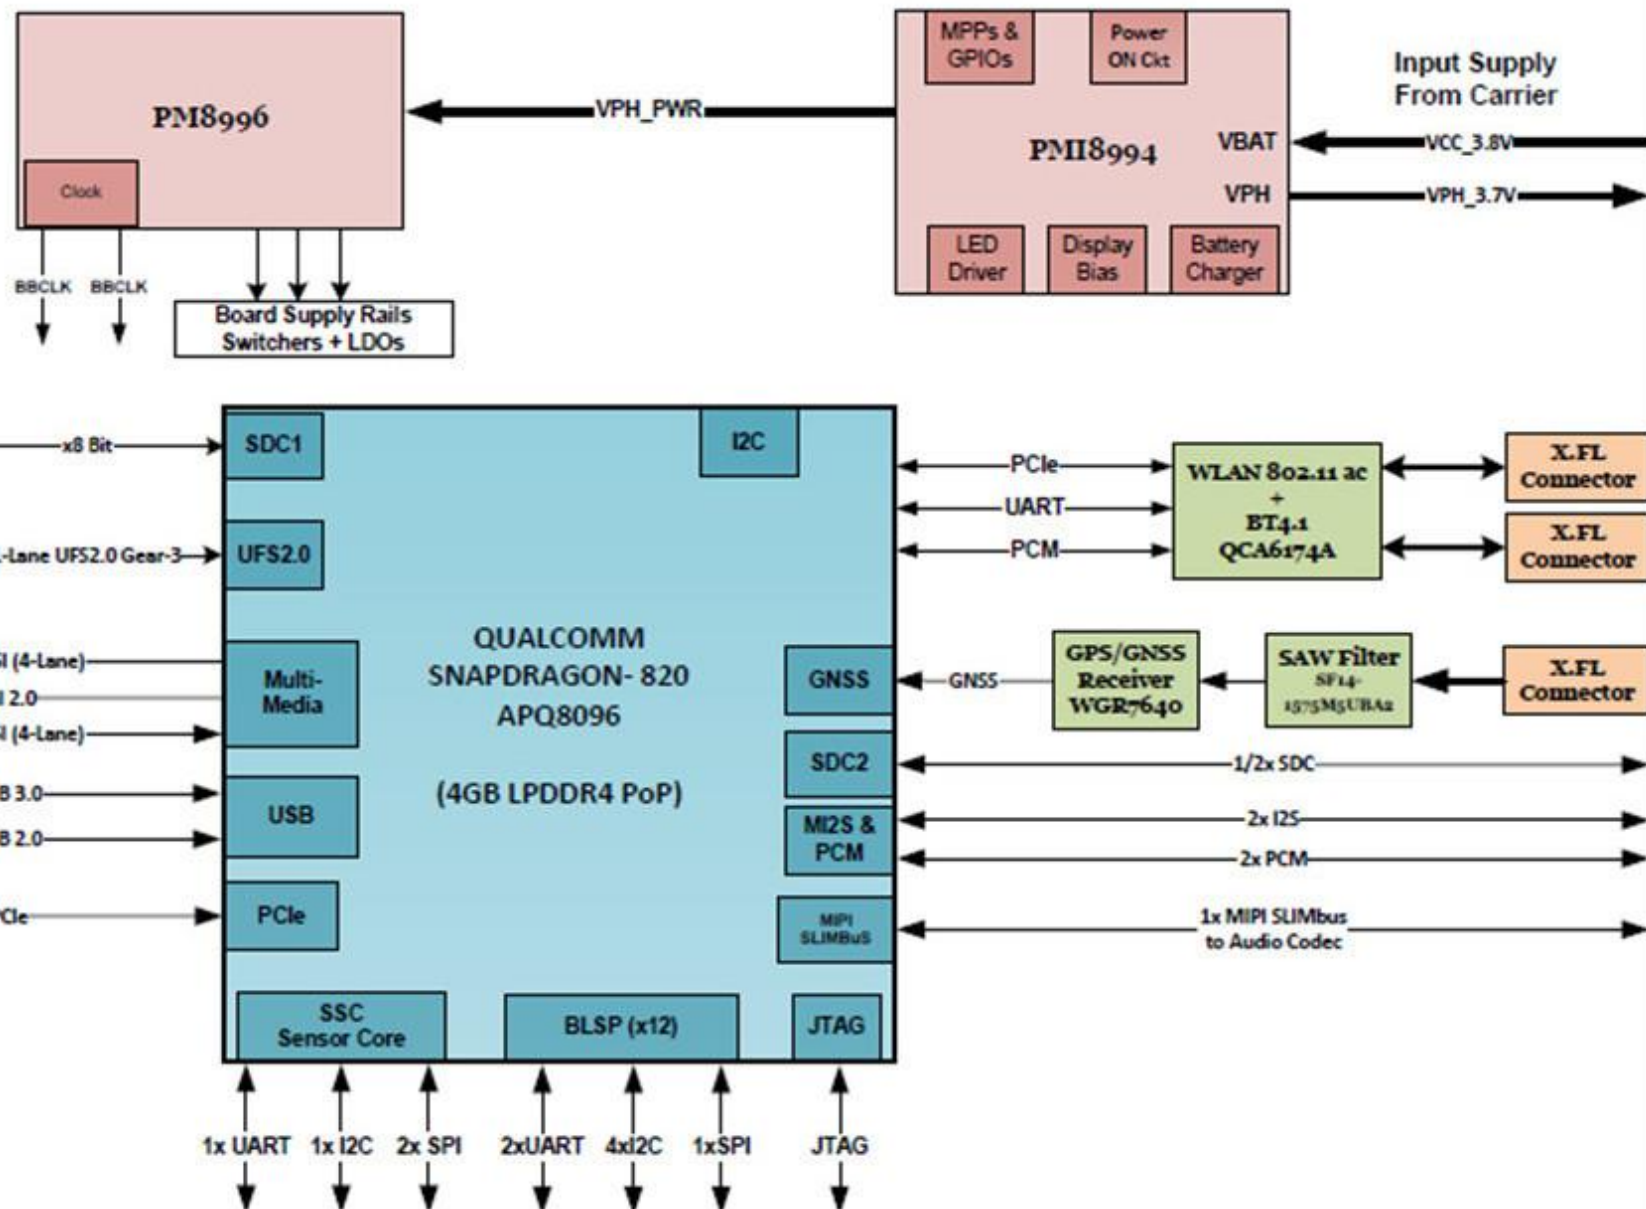


Figure 3-1. R8 Block Diagram

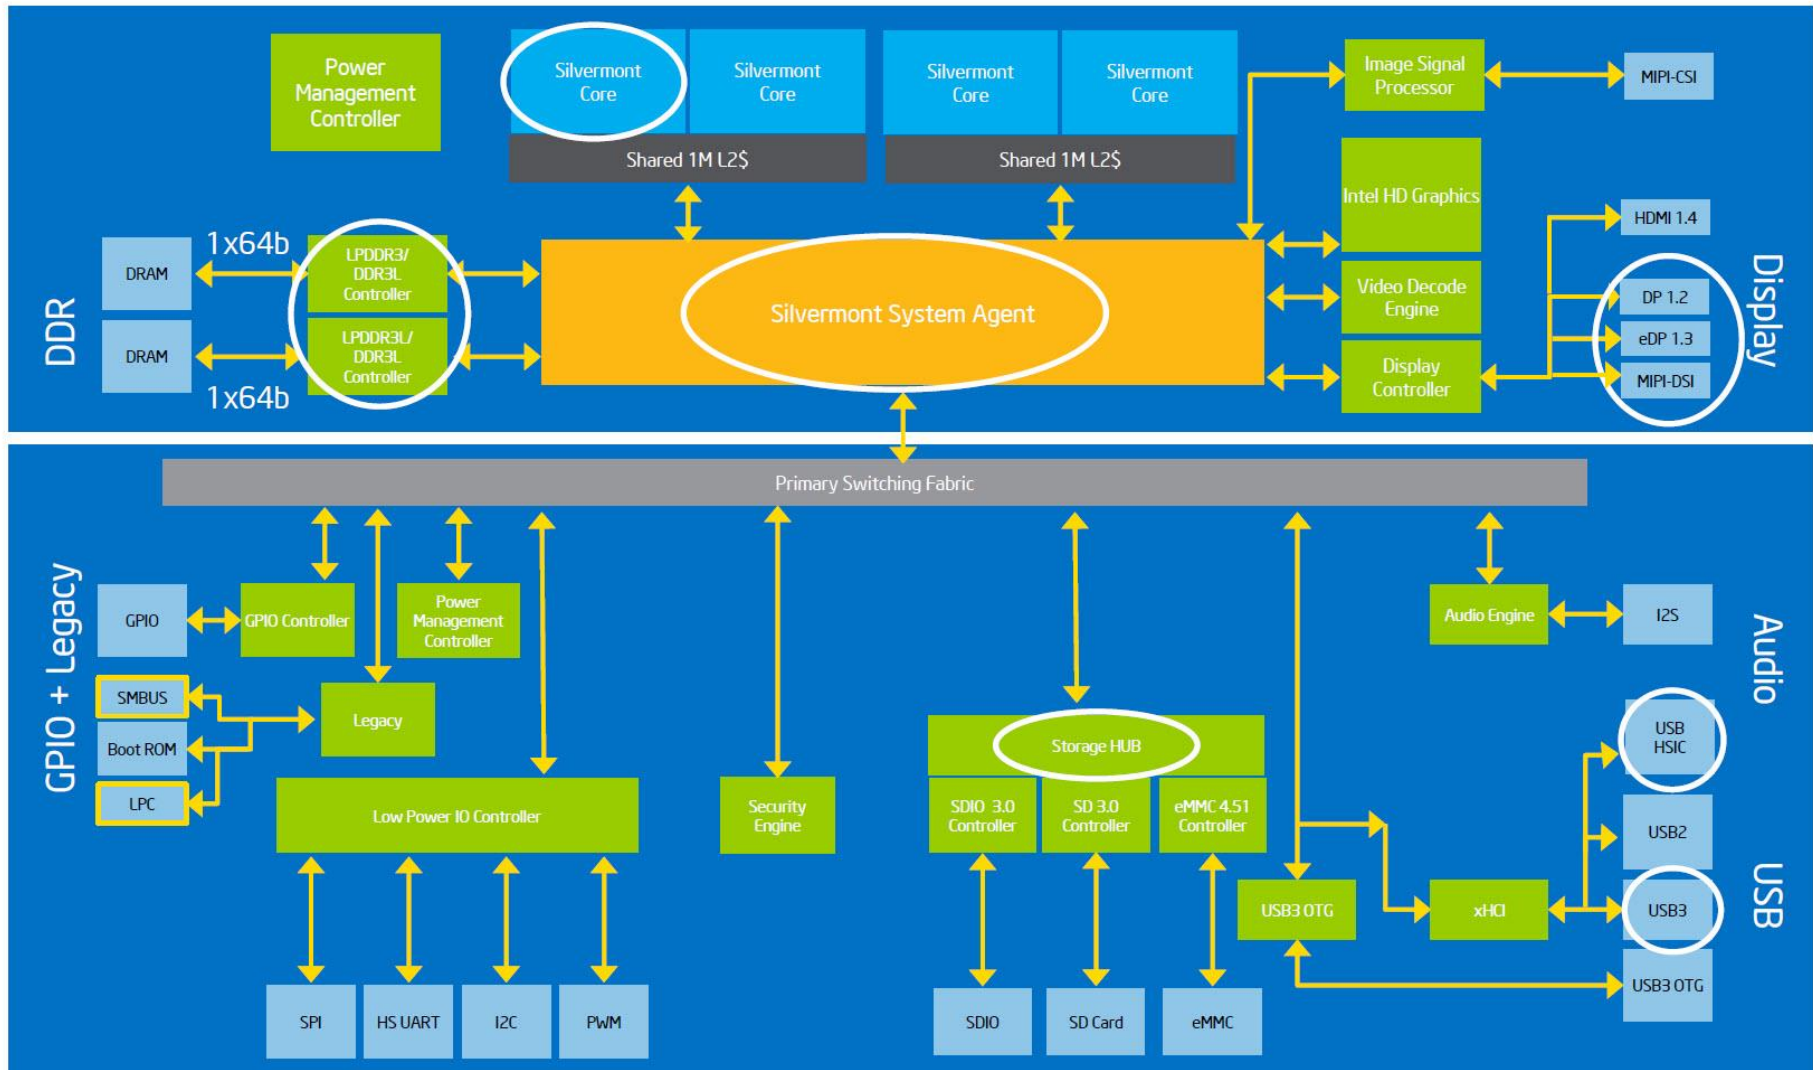


Snapdragon 820 - APQ8096 SOM Block Diagram

Rev0.7



Bay Trail SOC Block Diagram



Storage

Two/Three common flavors

MTD (Memory Technology Device): Abstraction of flash pages to partitions




Storage

Two/Three common flavors

MTD (Memory Technology Device): Abstraction of flash pages to partitions

eMMC: Embedded MultiMedia Card



Storage

Two/Three common flavors

MTD (Memory Technology Device): Abstraction of flash pages to partitions

eMMC: Embedded MultiMedia Card, SPI SD card
SD cards




Storage

Two/Three common flavors

MTD (Memory Technology Device): Abstraction of flash pages to partitions

eMMC: Embedded MultiMedia Card, SPI SD card
SD cards



Storage

Two/Three common flavors

MTD (Memory Technology Device): Abstraction of flash pages to partitions

eMMC: Embedded MultiMedia Card, SPI SD card
SD cards

Then there's UFS

Introduced in 2011 for phones, cameras: High Bandwidth storage devices

Uses a SCSI model, not eMMC's linear model



Variations

Some devices have a small amount of onboard Flash for the bootloader

Commonly seen on phones, for the purposes of bootstrapping everything else

Every vendor has different tools for pushing bits to a device and *they all suck*.

Samsung has at least three for Android

Allwinner based devices can be placed into FEL boot mode

Fastboot on Android devices

RAM

The art of cramming a lot in a small place

Vendors are seriously tight-assed

Can you cram everything in 8MB? Some routers do.

The WRT54G had 8M of RAM, later 4M

Modern SoCs tend towards 1GB, phones 4-6G

In pure flash storage, ramfs might be used to expand on-demand files (http content)

My RAM: *Exists*
Chrome:



Peripherals

Depends on what the hardware has: SPI, I2C, I2S, etc are common sights.

Gonna see some weird shit

- SDIO wireless cards

- "sound cards" over I2S

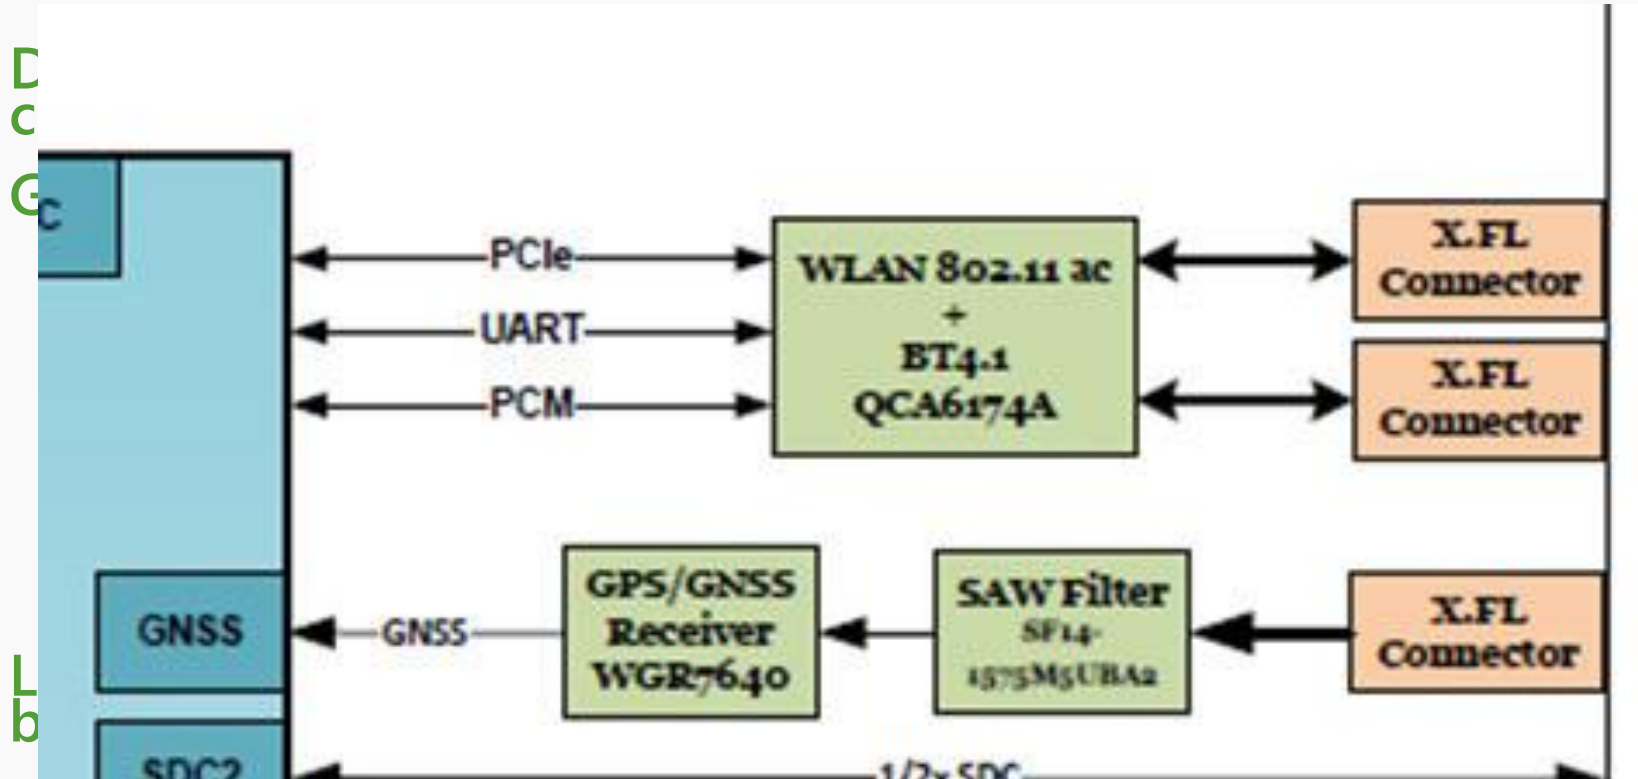
- GSM modems are really just pretending to be Hayes AT modems.

- Power management, LED management, cameras, etc.

- "We need an Ethernet PHY" becomes "We hooked an Ethernet PHY up over USB"

Linux doesn't care if they're on-die or not, it's all the same bus.

Peripherals



Peripherals

Depends on what the hardware has: SPI, I2C, I2S, etc are common sights.

Gonna see some weird shit

- SDIO wireless cards

- "sound cards" over I2S

- GSM modems are really just pretending to be Hayes AT modems.

- Power management, LED management, cameras, etc.

- "We need an Ethernet PHY" becomes "We hooked an Ethernet PHY up over USB"

Linux doesn't care if they're on-die or not, it's all the same bus.

Perip

Depends on
common si

Gonna see

SDIO wire

"sound c

GSM mod
modems.

Power m

"We need

Ethernet I

Linux does
bus.



etc are

ays AT

, etc.

ed an

re same

Bootloader

One Bootloader To Rule Them All: Das U-Boot

- Uses a simple scripting language

- Can pull from TFTP, HTTP, etc.

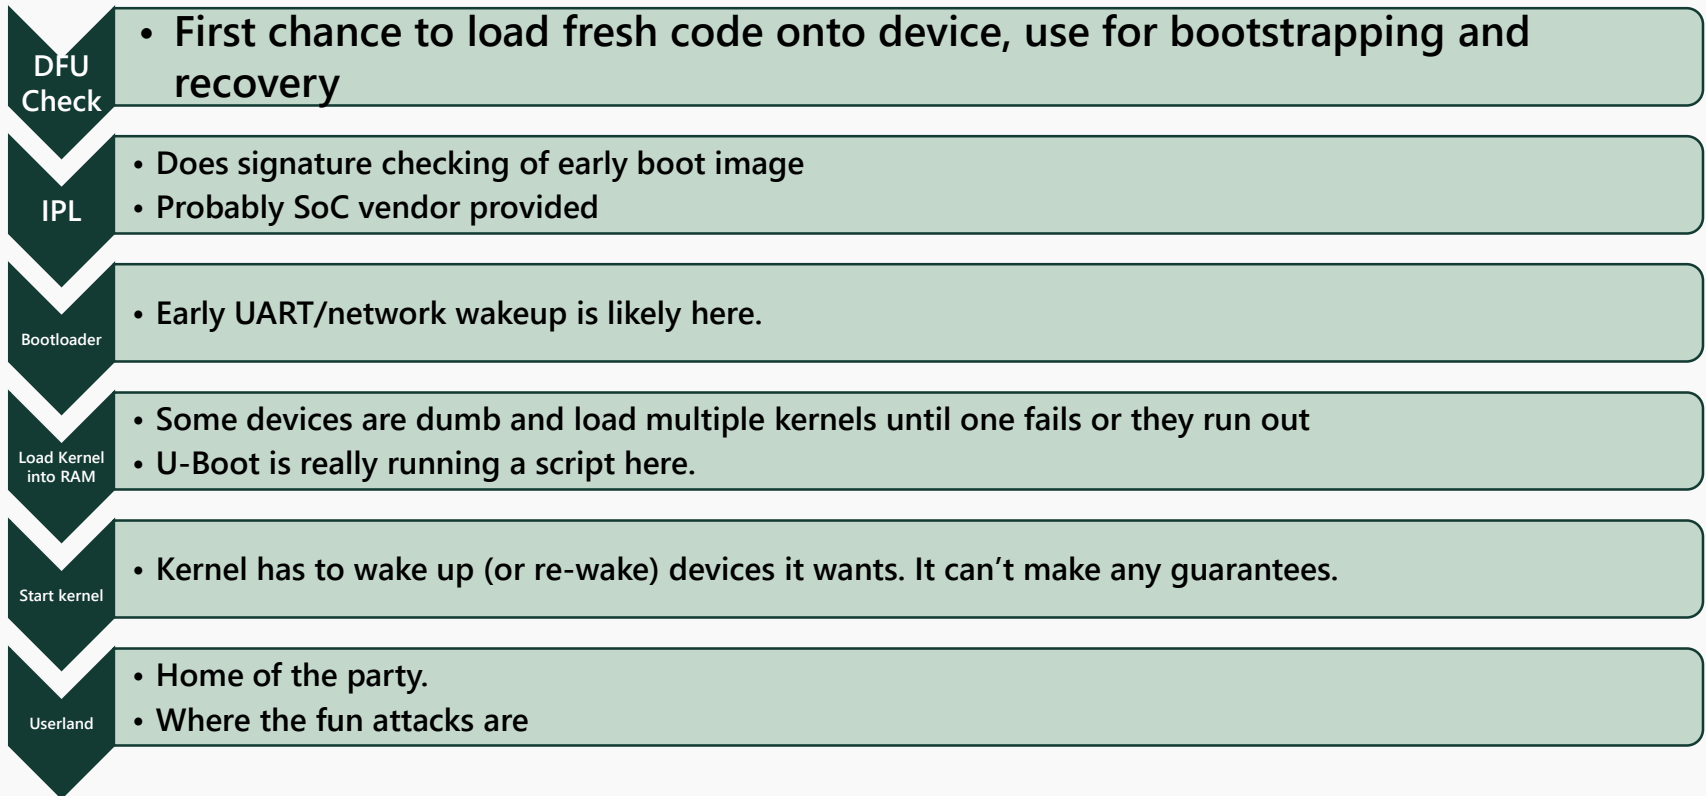
- Might be over Telnet, Serial, BT UART, etc.

Some don't use U-Boot, use Fastboot or other loaders

- Android devices are a clusterfuck of options

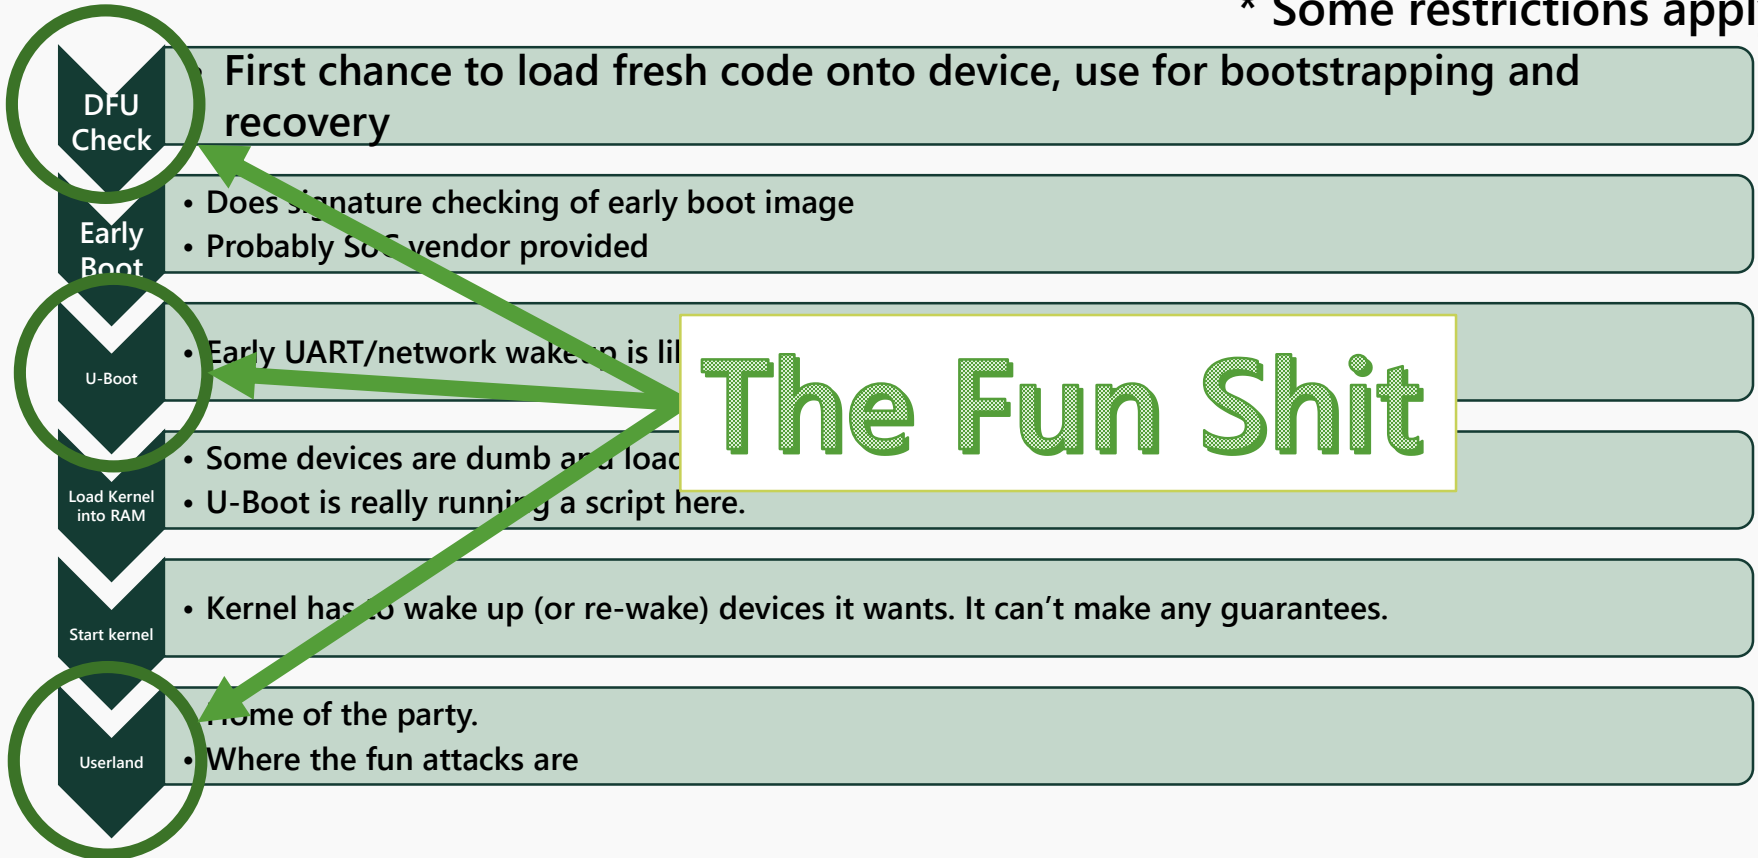
Life and death of an SoC*

* Some restrictions apply



Life and death of an SoC*

* Some restrictions apply



Root Filesystem: Home to All

A root filesystem contains the bare minimum to boot Linux: Any shared object libraries, binaries, or other content that is necessary for what that device is going to do

Fluid content that needs to be changed or which is going to be fetched regularly is often stored on a Ramdisk; this might be loaded during init's early startup from a tarball.

This is a super common thing to miss because it's a tmpfs outside of /tmp

this is a super common way of keeping / "small"

This often leads to rootfs extractions via tar that seem "too big"

There are sometimes multiple root filesystems overlaid upon each other

Android uses this to some extent: /system is where many things really are

Might be from NFS, might try NFS first, etc.

**Attacking these
devices**

Step 0: Scope out your device

Get to know what makes the device tick

- Version of Linux

- Rough software stack

- Known vulnerabilities

- Debug shells, backdoors, admin shells, etc.

ARM executables are fairly generic

- Kobo updates are very, VERY generic and the Kobo userland is *very aware of this*.

Hardware vendors are lazy: many devices likely similar to kin-devices

- Possibly able to find update for similar device by same OEM



Don't Reinvent The Wheel

Since so many embedded linux devices are similar, or run similarly outdated software, you may well have some of your work cut out for you

OWASP has a whole task set devoted to IoT security:

[https://www.owasp.org/index.php/OWASP Internet of Things Project](https://www.owasp.org/index.php/OWASP_Internet_of_Things_Project)

Tools like Firmwalker (<https://github.com/craigz28/firmwalker>) and Routersploit (<https://github.com/threat9/routersploit>) are already built and ready. Sometimes, thinking like a skid can save *time and energy* for other things, like beer!

Firmware Security blog is a great place to look, including a roundup of stuff (<https://firmwaresecurity.com/2018/06/03/list-of-iot-embedded-os-firmware-tools/>)

Option 1: It's a UNIX system, I know this

If you can get a shell,
sometimes just beating
against your target can be fun

Limited to only what is on the
target (or what you can get to
the target)

Can feel a bit like going into
the wild with a bowie knife
and a jar full of your own piss

Debugger? Fuzzer? *Compiler?*
What are those?



Option 2: Black-Box it

Lots of fun once you're used to it or live service attacks.

Safe: Never directly exposes you to "secrets" (IP)

You don't have the bowie knife, just two jars of piss.



These options suck

Can Really

Option 3: Reverse it


Pull out IDA/Radare

Grab a beer

Learn you a new ISA

The way of reversing IoT
things that don't run Linux!

*... but how the fuck do you
get the binaries?*



**Yeah but I'm fucking
lazy, asshole.**

**I don't want to learn
IDA. I want to fuzz.**


Option 4: Emulate It

You have every tool at your disposal

Hot damn is that a debugger?

Oh shit waddup it's fuzzy boiii

Once again, *how the fuck do you get your binary of choice?*



Getting root(fs)

Easy Mode: Update Packages

Updates for devices are the easiest way to extract a root filesystem

Sometimes little more than a filesystem/partition layout that gets dd'd right to disk

Android updates are ZIPs containing some executables, a script, and some filesystems

Newer android updates (small ones) are very regularly "delta" updates: These touch a known filesystem directly, and are very small but don't contain a full filesystem.

Sometimes, rarely, they're an **actual executable** that gets run on the device

Probably isn't signed

Probably fetched over HTTP

Downside: They're occasionally *very* hard to find or are incremental, incomplete patches. Sometimes they're *encrypted*.

Medium: In-Vivo extraction

You need a shell

Can you hijack an administrative interface?

Some ping functions can be hijacked into shells

Sometimes it's literally "telnet to the thing"

Refer to step 0 for more

You need some kind of packer (cpio, tar, etc)

Find is a builtin for most busybox implementations.

You need some way to put it somewhere (netcat, curl, etc)

You might have an HTTPD to fall back on

Need to do reconnaissance on your device

Might need some creativity

Wireshark, Ettercap, Fiddler, etc

Demo: Router firmware extraction (Actiontec Router)



technicolor C1100T
Modem Configuration



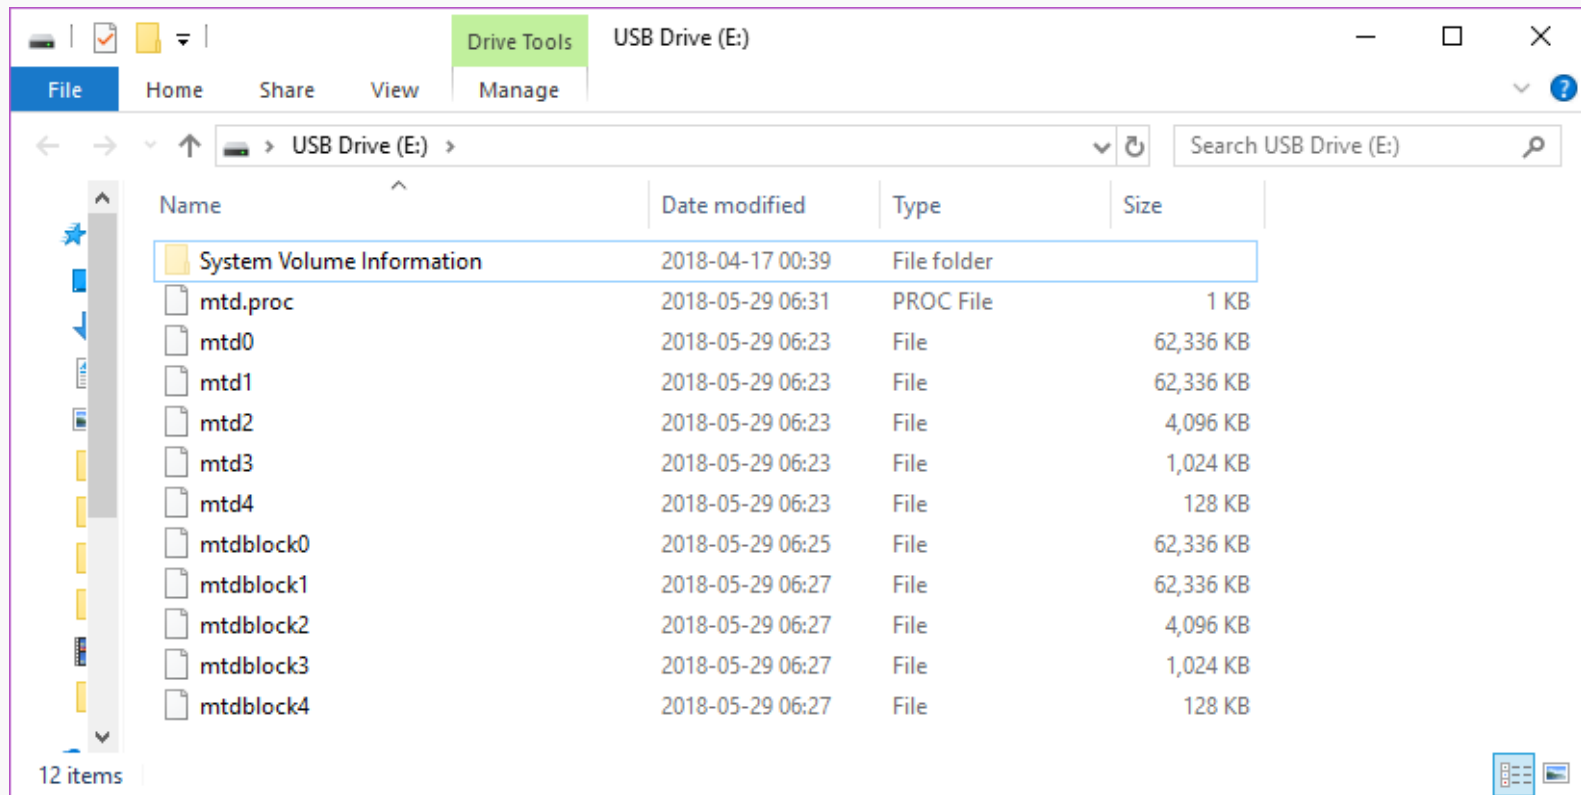
Help



```
192.168.0.1 - PuTTY
BCM963268 Broadband Router
Login: admin
Password:
>
```



What did we get?




Surprise Mode: Direct Extraction

Could be as simple as “remove SD card, image it”

Surprise Media Direct Extra

Could be



it"

Surprise Mode: Direct



CLOSED

RetroEngine Sigma - Mini Video Game Console

Retrogaming Simplified. The Greatest Retro-Station known to Man!

PROJECT OWNER



Doyodo Team
Santa Monica, United States
1 Campaign | [More](#)

I, image it"



By Rogue Foundry
First created

Sever: The Anti-Villain Box (Canceled)

Get ready to no longer exist online.


OVERVIEW



RetroEngine Sigma is as simple as plugging in a game card.



Sever



Anonymous

t, MA

Surprise Mode: Direct Extraction

Could be as simple as “remove SD card, image it”

eMMC is harder though, since you need to get to the data lines, but it can be done!

You will need to understand how the disk is laid out

Binwalk can help later, as can “standard” DOS partition tables.

Having some in-vivo information is helpful

Surprise Mode: Direct

Ex

Co


eM
da


Yo

l

p

Ha





remo
n, sing
der




partition tar
Having some



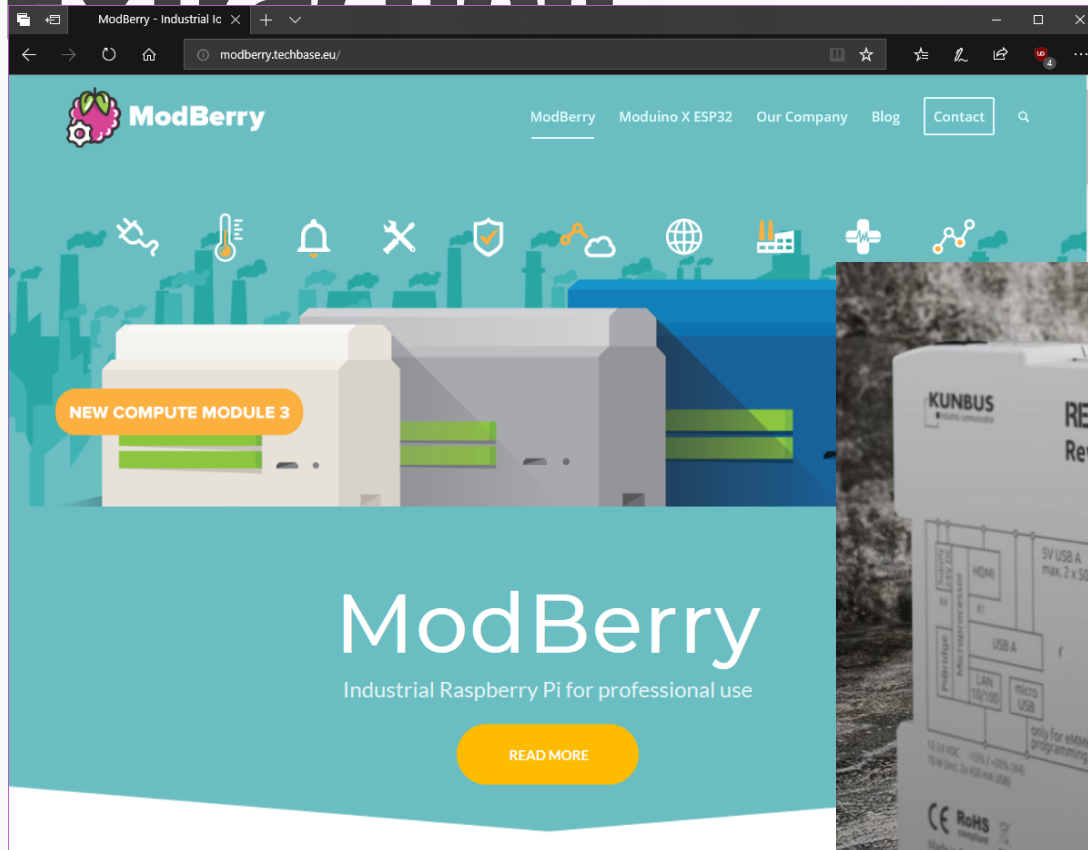
lpful

Surprise Mode: Direct



<http://blog.oshpark.com/2017/02/23/retro-cpc-dongle/>

Surprise Mode: Direct Extraction



rd, image it”
ed to get to the



Surprise Mode: Direct Extraction

Could be as simple as “remove SD card, image it”

eMMC is harder though, since you need to get to the data lines, but it can be done!

You will need to understand how the disk is laid out

Binwalk can help later, as can “standard” DOS partition tables.

Having some in-vivo information is helpful

All Else Fails: Solder to the rescue

Might need to desolder some storage, or otherwise physically attack the hardware

MTD devices are weird. Prepare to get your hands dirty

Interested in more? HHV and friends are the place to start looking.

JTAG, etc. might be the hard way out.




Logic Analyzers

Saleae makes a good one


Cheap

Basic

Runs over USB



http://dangerousprototypes.com/docs/Bus_Pirate



<https://www.crowdsupply.com/excamera/spidriver>

Hardware interfaces

Now that we have that, what do?

Try mounting it/extracting it/etc. `file` might give you a good idea of what it thinks it might be, as will `strings` and the like.

eMMCs sometimes have real partition tables

SD cards often do

Look at the reconnaissance you did

- Boot logs: lots of good information about partitions

- Fstab, /proc/mounts

Let automation do your work

Binwalk!

- Takes a rough guess at what might be in a place
- Makes educated guesses about filesystems
- High false positive rate

Photorec might be helpful

Get creative

Losetup and friends are capable of more than you give them credit for.

There are a lot of filesystems that are read-only or create-and-read, like cramfs and such. These are often spotted by binwalk but are even sometimes seen as lzma or other compressed or high-entropy data

If you're only looking to play in IDA/Radare/etc, the bulk extraction from binwalk might help.



QEMU: the Quick EMUlator

QEMU is a fast processor emulator for a variety of targets.

Targets you've never heard of?

Mainframes

ARM, MIPS

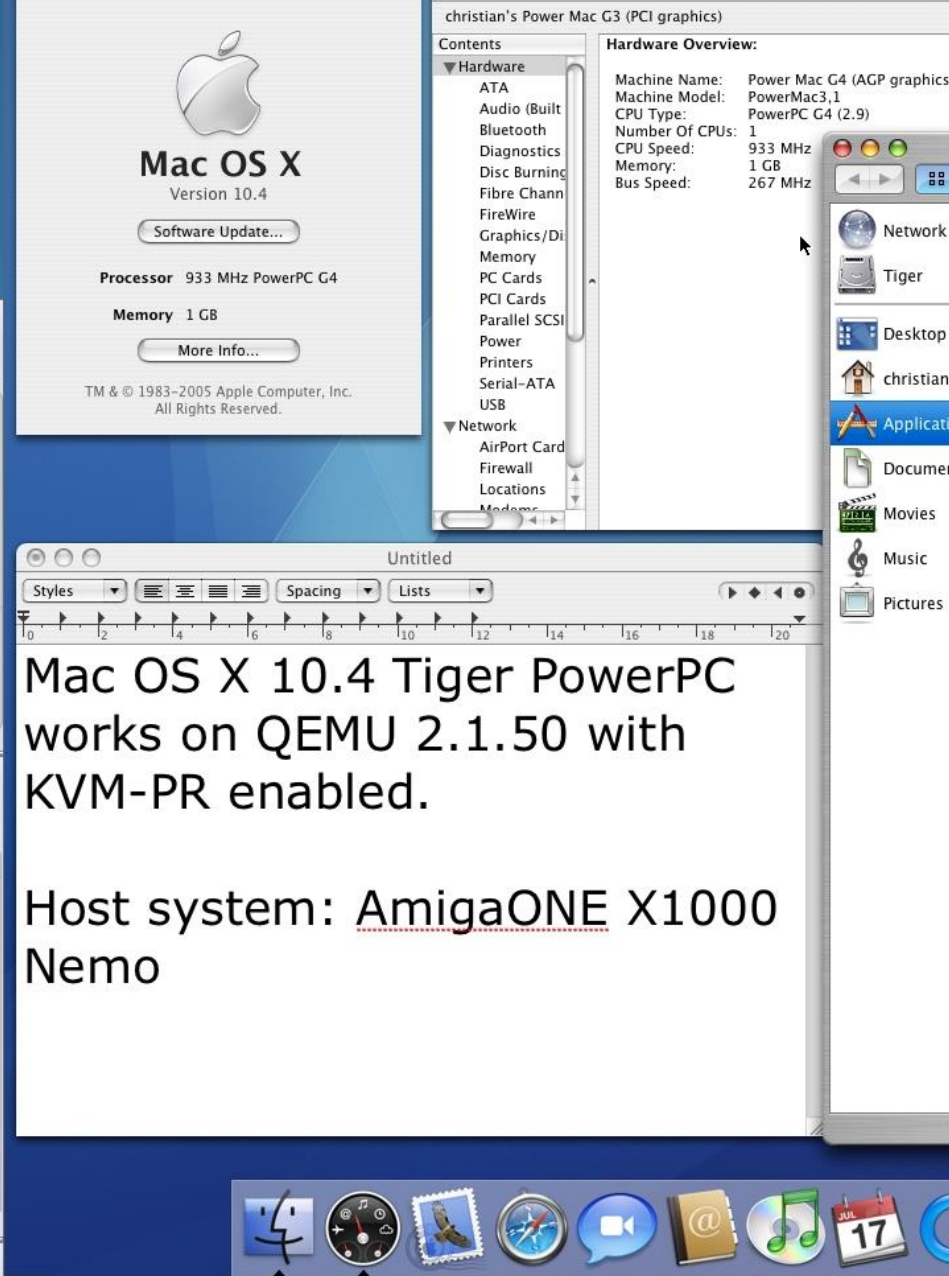
PowerPC

OpenRISC

DEC Alpha

Lots of different ports and targets have been ported.

```
r2
eX1000: ~$ una
n Reiter Hilfe
aoneX1000: ~$ qem
-version
version 2.1.50,
2003-2008 Fabric
aoneX1000: ~$
: ~$
supported
z
0102)
supported
z
0102)
cient
```



Mac OS X 10.4 Tiger PowerPC
works on QEMU 2.1.50 with
KVM-PR enabled.

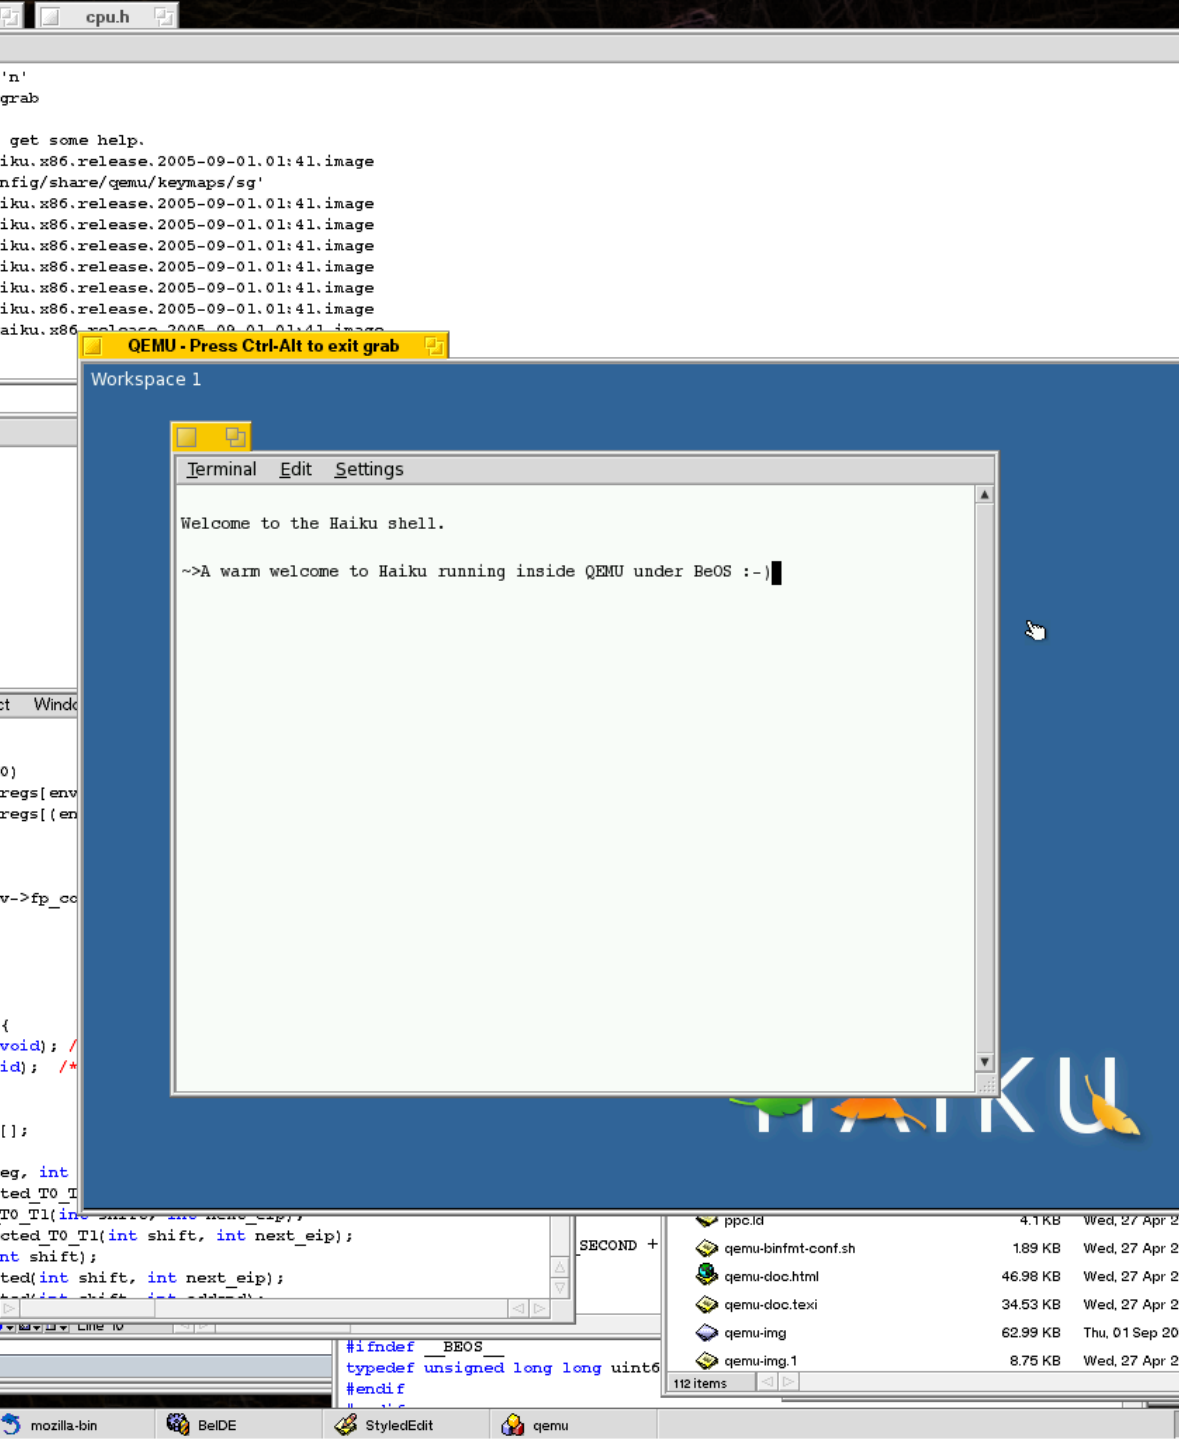
Host system: AmigaONE X1000
Nemo

**Did
someone
say "OSX on
Amiga
hardware"?**

```
1000:/home/christian/virtual_machines# qemu-system-ppc -M mac99 -cpu G4 -enable-kvm -m 1024 -hda /home/christian/virtual_machines/tiger.img -g 1
rom-env 'boot-args=-v'
```

root@Amiga... [Installing Q... QEMU christian@A... christian@A... [Paint 3.3...

13:54



Or Haiku on BeOS?

Two ways to run QEMU

AS A FULL FAT VM

You preserve full control over the whole process

You've got access to things like gdb at a kernel level

Requires zero trust in the safety of the binary

But you probably want a special kernel and board setup, though there are generic setups to get you started

Any tools are going to need to be compiled for the target environment

I hate cross-compilers.

AS A TRANSLATING LOADER

You have access to all your existing x86-64 tools (or whatever your native tools are)


They're not only native, but they're running full-speed.

You can run AFL like it's meant to

Runs nicely in containers

You don't even need a container!

Full-Fat VM: 9 tracks of DOS



Binfmt: Linux' way of loading executables

Long ago, Linux added loadable loaders

Originally for the purposes of running JARs from the command line like God and Sun Microsystems intended.

Turns out this is a great place to put emulators.

Debian ships with support for this in its *binfmt-misc* package.

QEMU can be shipped as a “static userspace” environment (think WINE)

Uses “magic numbers” – signatures from a database – to determine which ones are supposed to load what.

Fairly simple to add new kinds of binaries. You could actually execute JPEGs if you really wanted to?

QEMU as loader

WITHOUT A CONTAINER

Dumb simple to set up:

`qemu-whatever-static <binary>`

With binfmt, just call your binary.

Must trust that the executable is not malicious

Might depend on your local environment looking like its target environment

This works best for static, monolithic executables

WITH A CONTAINER

Bring that whole root filesystem along!

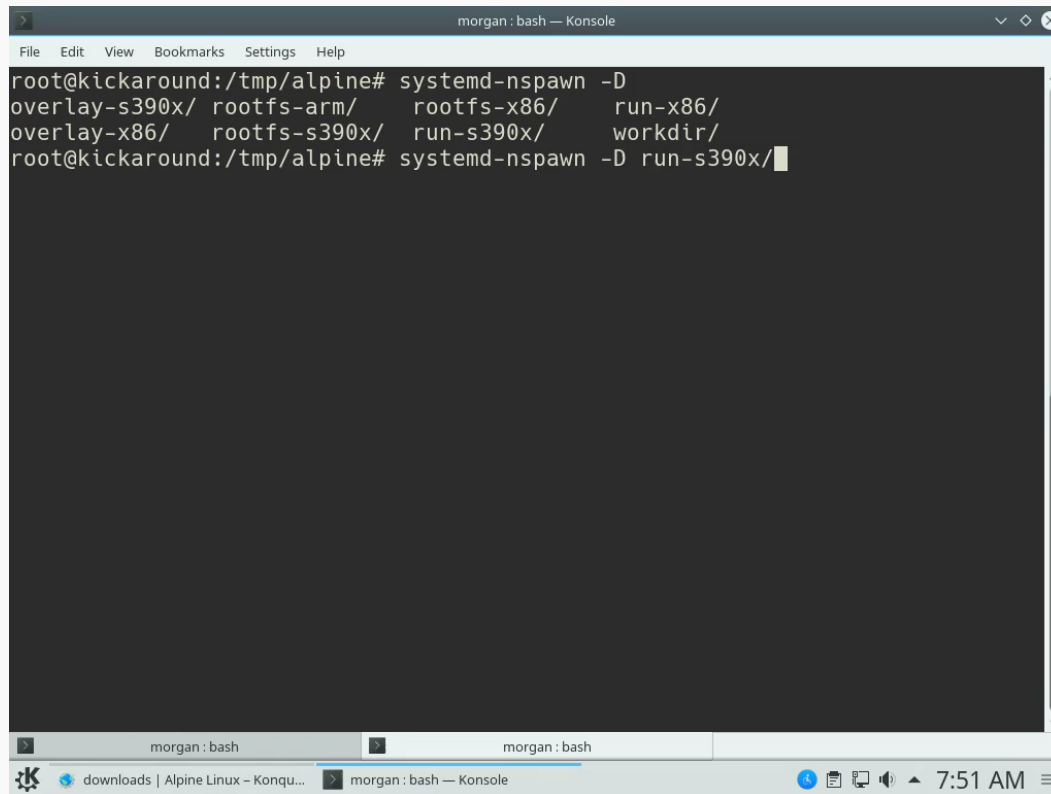
Run it in the confines of a jail, Docker instance, even something like Systemd containers

Might need root depending on your container (systemd)

Great for when your binary loads its own special versions of libraries that have weird things added to them



QEMU user Demo



The screenshot displays a QEMU user interface window titled "morgan : bash — Konsole". The window contains a terminal session with the following commands and output:

```
root@kickaround:/tmp/alpine# systemd-nspawn -D  
overlay-s390x/ rootfs-arm/   rootfs-x86/   run-x86/  
overlay-x86/   rootfs-s390x/  run-s390x/   workdir/  
root@kickaround:/tmp/alpine# systemd-nspawn -D run-s390x/
```

The terminal window has a menu bar with "File", "Edit", "View", "Bookmarks", "Settings", and "Help". The bottom of the window shows a taskbar with a "downloads" icon, a "Alpine Linux - Konqu..." icon, and a "morgan : bash — Konsole" icon. The system tray at the bottom right includes icons for network, volume, and power, along with the time "7:51 AM".

AFL setup

Oh boy. Let's talk about AFL.

AFL needs to be compiled with QEMU support

Magic sauce: `CPU_TARGET=whatever`
`./build_qemu_support.sh`


AFL needs to bring along the host's libraries

Easiest bound with `systemd-nspawn`

Don't do this in a VM

It hurts

AFL Demo



Wrapping up

What did we learn today?

- Hardware vendors are lazy

- Attacking hardware means getting creative

- QEMU is pretty neat

- AFL runs really slow when you're emulating an X86 emulating an IBM mainframe.

Going forward:

- I hope I've given you some idea of the landscape of tools

- Always remember rule 0

More Resources

Non-Linux targets:

RECON 2010 with Igor Skochinsky: Reverse Engineering for PC Reversers:

<http://hexblog.com/files/recon%202010%20Skochinsky.pdf>

JTAG Explained: <https://blog.senr.io/blog/jtag-explained>

<https://www.blackhat.com/presentations/bh-europe-04/bh-eu-04-dehaas/bh-eu-04-dehaas.pdf>

https://beckus.github.io/qemu_stm32/ (among others)

Linux targets:

eLinux.org – *Fucking Gigantic* wiki about embedded Linux.

linux-mips.org – Linux on MIPS

Thank you

Keep on hacking.