BOSTONFUSION

Anomaly Detection in Cyber Networks using Graph-node Role-dynamics and NetFlow Bayesian Normalcy Modeling

Anthony Palladino, PhD, Senior Research Scientist
Christopher Thissen, PhD, Research Scientist
Andrew Spisak, PhD, Senior Research Scientist

Boston Fusion Corp.
70 Westview Street, Suite 100
Lexington, MA 02421
www.bostonfusion.com

Presented at FloCon2018

9 January 2018

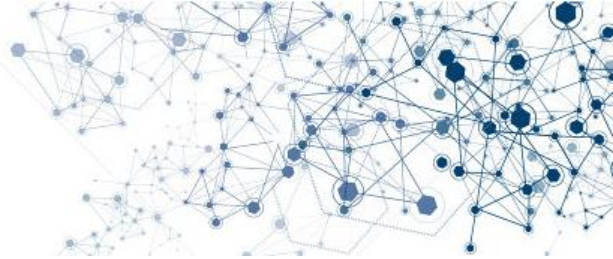

Agenda

- Introduction
- Advanced Persistent Threats
- Graph-node Role-dynamics
- Bayesian Normalcy Modeling
- Summary

Introduction

- **Context Aware INference for Advanced Persistent Threat (CAIN for APT)**
  - DARPA Phase II SBIR
- **Challenge**
  - Stealthy cyber attacks slip past state-of-the-art defenses, dealing crippling blows to critical US military and civilian infrastructure
- **Goal**
  - Rapid, automated, and accurate prioritization of cyber alerts provides timely and comprehensive cyber situational awareness (SA)
- **Technical Approach**
  - Novel graph-analytics makes sense of noisy IDS sensors
  - Novel Bayesian Dynamic Flow Model flags odd network traffic
  - Tests and evaluations with APT simulations

Agenda

- Introduction
- **Advanced Persistent Threats**
- Graph-node Role-dynamics
- Bayesian Normalcy Modeling
- Summary

Advanced Persistent Threats

- **Often associated with nation-state espionage**
- **Targets include private organizations & nation-states**
- **Low and Slow: Attack campaigns may last months**
- **Notoriously difficult to detect**

(Preprint: A. Lemay, et al. 2018)

Image: https://www.secureworks.com/blog/advanced-persistent-threats-apt-a

Simulated APT Scenarios

- **Simulation attributes**
  - Approx. 1 month of data per scenario
  - Servers, laptops, switches
  - Linux & Windows machines
  - Normal & attacked behavior
  - Generates IDS alerts and NetFlow traffic
  - Detailed attack timeline
- **Hurricane Panda simulation**
  - Attack distributed over 3 days
  - Database injection to gain credentials
  - Lateral movement and firewall deactivation
- **Energetic Bear (Crouching Yeti) simulation**
  - Attack distributed over 3 hours
  - Email phishing to redirect user to malicious website
  - Lateral movement through network using a remote-desktop exploit
  - Attacker attempted to clean-up logs and other traces

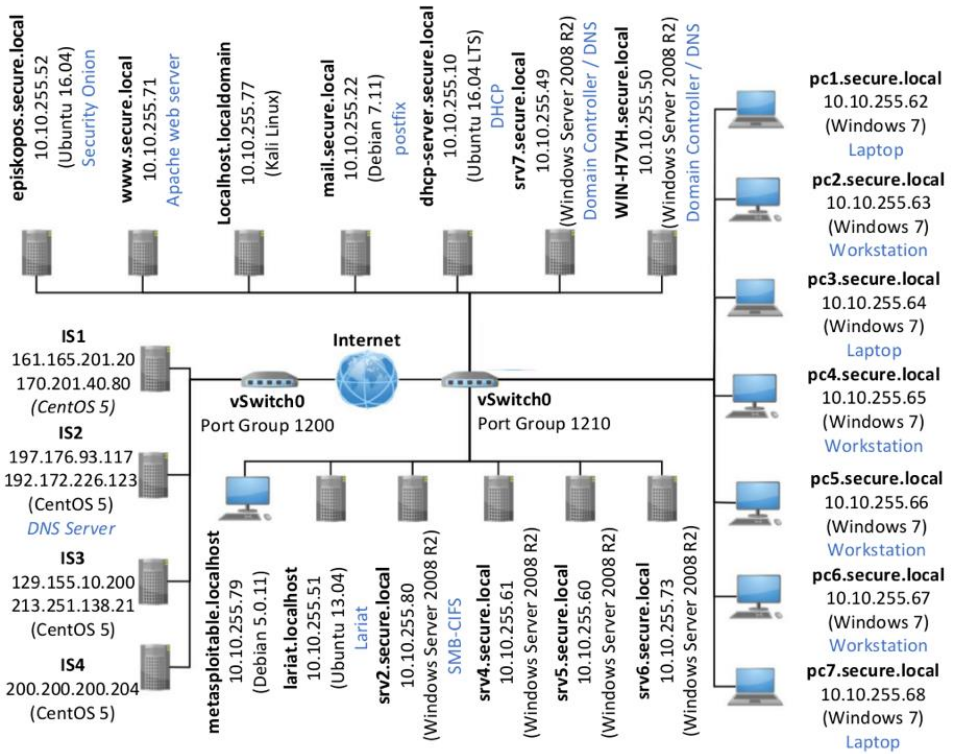

Network topology for simulations

Agenda

- Introduction
- Advanced Persistent Threats
- **Graph-node Role-dynamics**
- Bayesian Normalcy Modeling
- Summary

Graph-based Approach

- Fuses disparate IDSs
- Captures alert interdependencies
- Efficiently represents many alerts
- Robust to circumvention
- Unsupervised
- Facilitates causal analysis
- Optimal parameters determined automatically

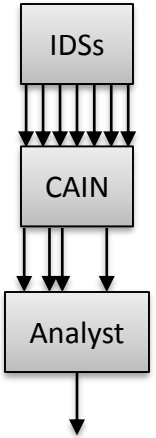

Making Sense of Noisy IDS Sensors with Graph Analytics

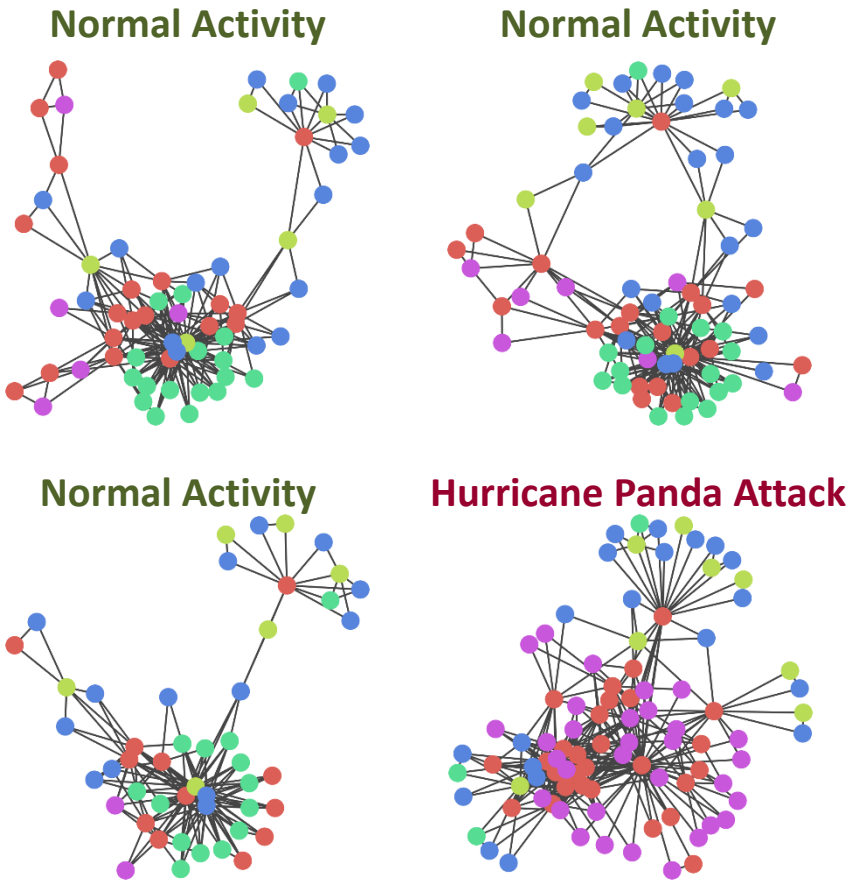

Alert Graphs from Hurricane Panda Simulation

- **Novel, graph-based analysis of IDS alerts**
  - Load IDS alerts into alert graph
  - Detect graph anomalies
- **Advantages of graph-based approach:**
  - Captures alert interdependencies
  - Fuses disparate IDSs
  - Efficiently represents alerts
  - Robust to circumvention

Akoglu et al. 2014

Alert Graphs

OSSEC Alert (Host IDS)

```
** Alert 1480536972.16316356: syslog, vsftpd,
connection_attempt 2016 Nov 30 20:16:12 (host)
10.10.255.79 -> /var/log/vsftpd.log Rule: 11401
(level 3) -> 'FTP session opened.' Src IP:
10.10.255.77 Wed Nov 30 15:17:25 2016 [pid 14562]
```

Snort Alert (Network IDS)

```
11/30-15:32:15.407340  [**] [1:2010939:2] ET
POLICY Suspicious inbound to PostgreSQL port 5432
[**] [Potentially Bad Traffic] [Priority: 2] {TCP}
10.10.255.77:38989 -> 10.10.255.50:5432
```

Alert Graph

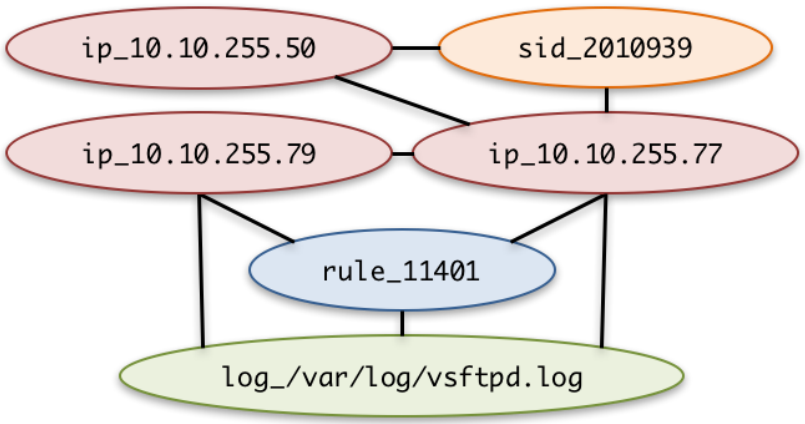

- **Graph of alerts (Not network topology)**

Alert Graphs

OSSEC Alert (Host IDS)

```
** Alert 1480536972.16316356: syslog, vsftpd,
connection_attempt 2016 Nov 30 20:16:12 (host)
10.10.255.79 -> /var/log/vsftpd.log Rule: 11401
(level 3) -> 'FTP session opened.' Src IP:
10.10.255.77 Wed Nov 30 15:17:25 2016 [pid 14562]
```

Snort Alert (Network IDS)

```
11/30-15:32:15.407340  [**] [1:2010939:2] ET
POLICY Suspicious inbound to PostgreSQL port 5432
[**] [Potentially Bad Traffic] [Priority: 2] {TCP}
10.10.255.77:38989 -> 10.10.255.50:5432
```

Alert Graph

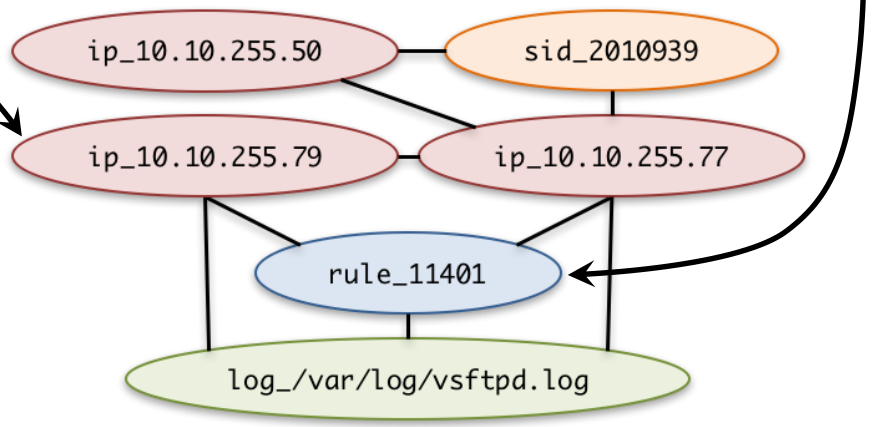

- Graph of alerts (Not network topology)
- Alert properties become nodes
- Node colors indicate property type

Alert Graphs

OSSEC Alert (Host IDS)

```
** Alert 1480536972.16316356: syslog, vsftpd,
connection_attempt 2016 Nov 30 20:16:12 (host)
10.10.255.79 -> /var/log/vsftpd.log Rule: 11401
(level 3) -> 'FTP session opened.' Src IP:
10.10.255.77 Wed Nov 30 15:17:25 2016 [pid 14562]
```

Snort Alert (Network IDS)

```
11/30-15:32:15.407340  [**] [1:2010939:2] ET
POLICY Suspicious inbound to PostgreSQL port 5432
[**] [Potentially Bad Traffic] [Priority: 2] {TCP}
10.10.255.77:38989 -> 10.10.255.50:5432
```

Alert Graph

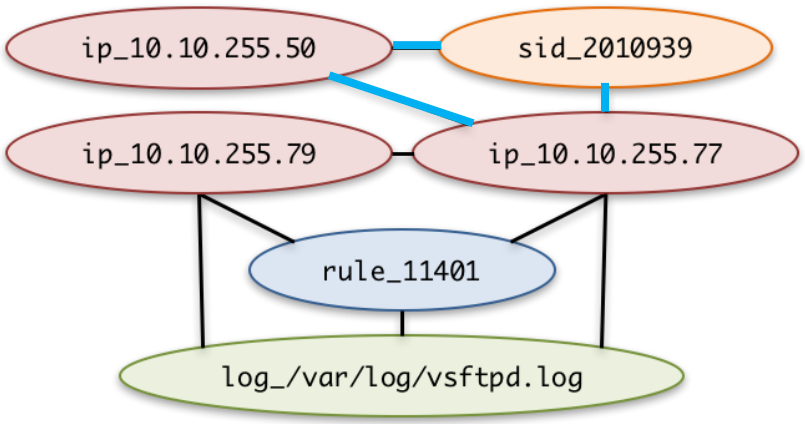

- Graph of alerts (Not network topology)
- Alert properties become nodes
- Node colors indicate property type
- Edges connect nodes that co-occur in alerts
- Edges weighted by frequency of co-occurrence

Alert Graphs

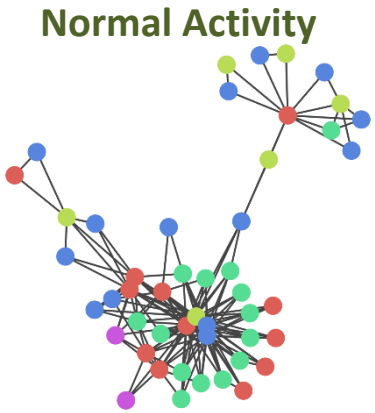

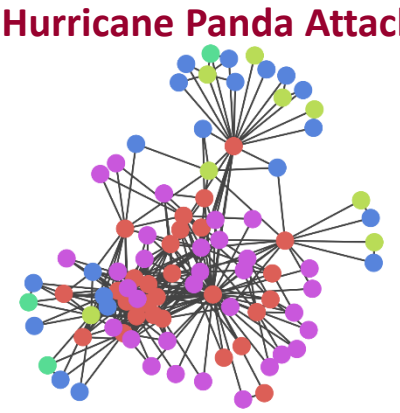

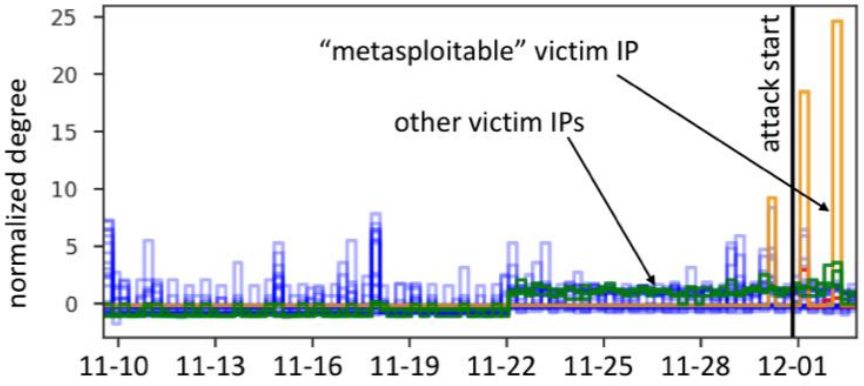

- **Cyber attacks change IDS alert logs**
- **Intuition**
  - Changes in alert logs modify alert graph
  - Anomalies in the graph features (properties) may indicate cyber attacks
- **Quick test**
  - Degree of IP nodes shows marked changes during simulated attack
  - But a single feature is likely insufficient
  - What features should we track?
  - Should we model all features for anomalies?

Role Dynamics

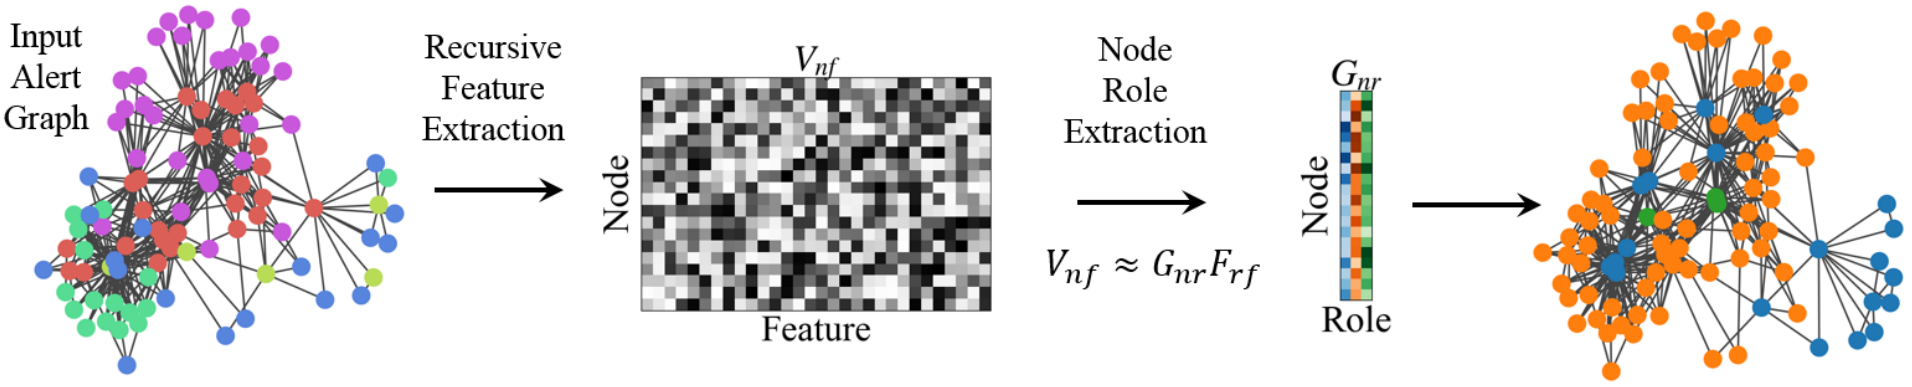

- **Infeasible to model every feature of every node**
- **Instead, use graph-based anomaly detection algorithms**
- **Role dynamics (Rossi et al., 2012)**
  - Collect features and factorize as roles
  - Roles provide a succinct, integrated summary across a large number of features
  - Output is probability of membership in each role, for each node
  - Application to IDS alerts is novel

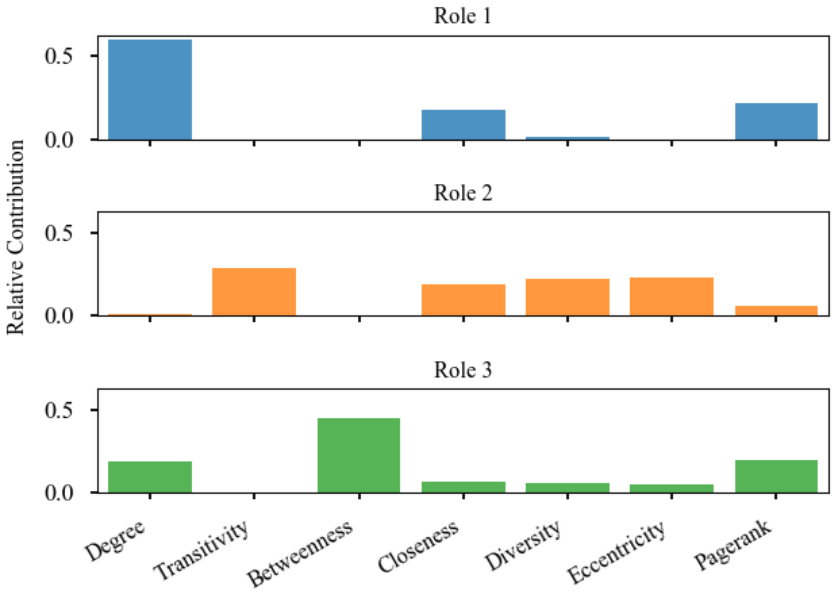

Role Dynamics

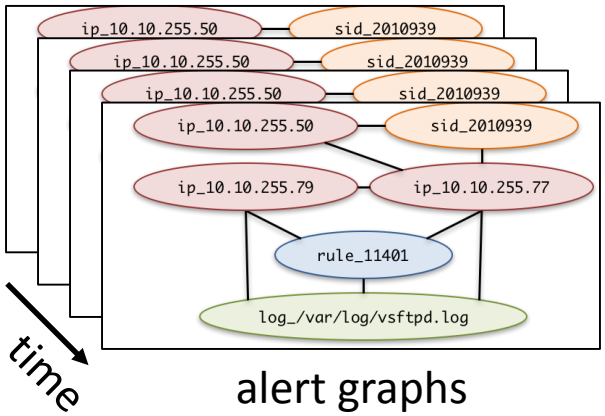

alert graphs

- **Infeasible to model every feature of every node**
- **Instead, use graph-based anomaly detection algorithms**
- **Role dynamics (Rossi et al., 2012)**
  - Collect features and factorize as roles
  - Roles provide a succinct, integrated summary across a large number of features
  - Output is probability of membership in each role, for each node
  - Application to IDS alerts is novel
  - **Track role memberships over time**

Roles for nodes in Hurricane Panda simulation

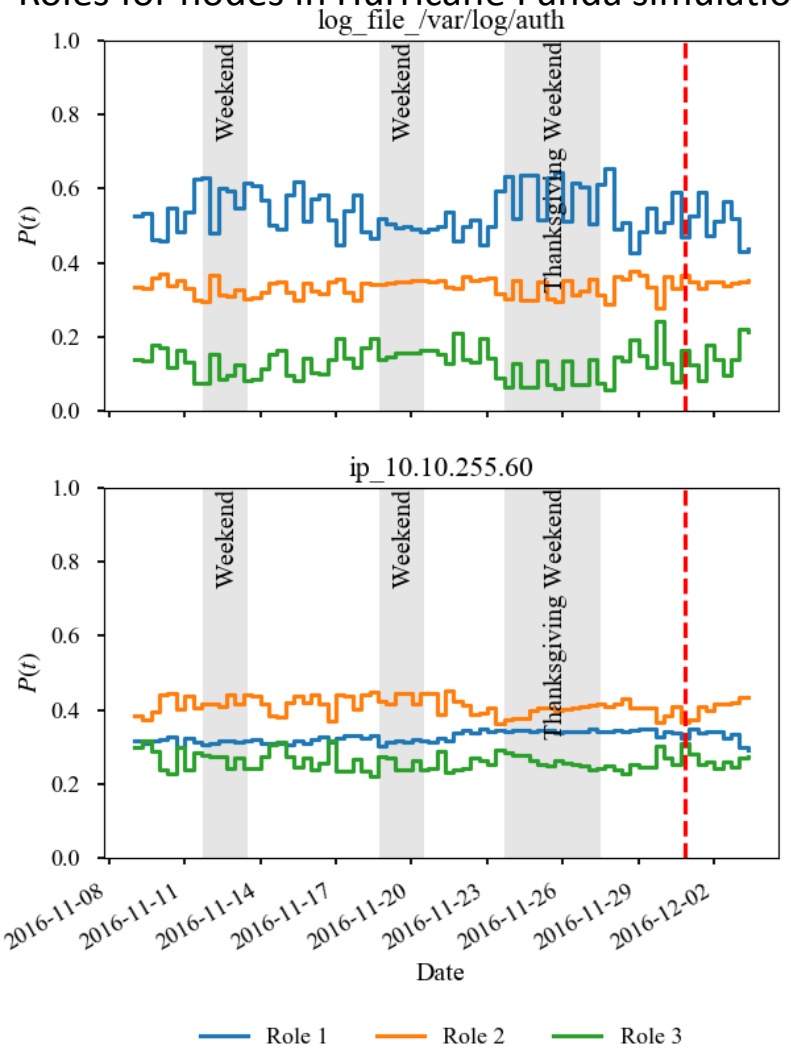

Role Dynamics

- **Why role dynamics?**
  - Linear
  - Weighted
  - Dynamic
  - Attributed
  - Unsupervised
  - Explainable
  - Extensible
  - Automated parameter selection
  - Available

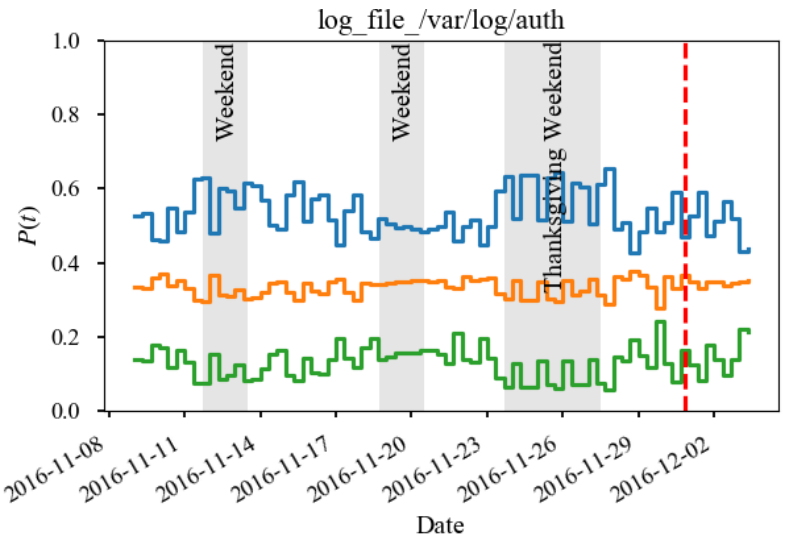

- **Explainable**
  - Identifies anomalous nodes
  - Helps with causal analysis
- **Automated parameter selection**
  - Recursive features
  - Optimal number of roles
  - Set during a training period

Finding Role Anomalies

- **Role anomalies**
  - Now we have roles over time for all nodes in graph
  - How to identify anomalies in the roles?
- **Aggregate changes into a few useful metrics**
  - For example, average magnitude of the rate of change in role membership: $\sum_{n=1}^{N} |P_n(t) - P_n(t-1)|/N$
  - Monitor metrics for anomalies

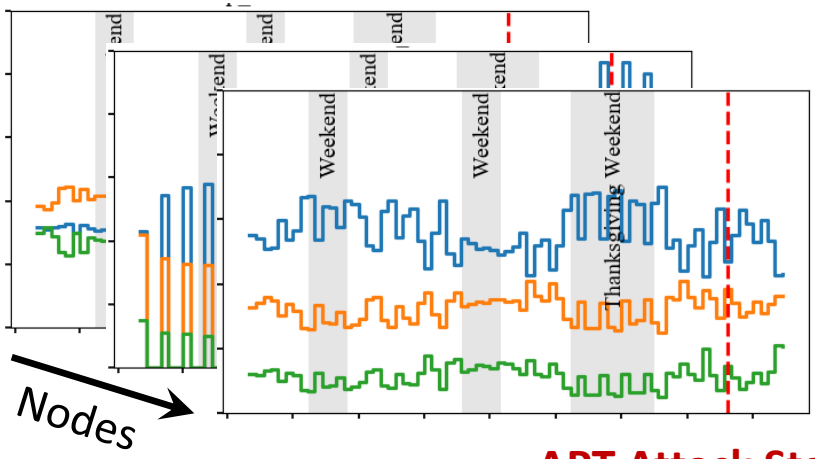

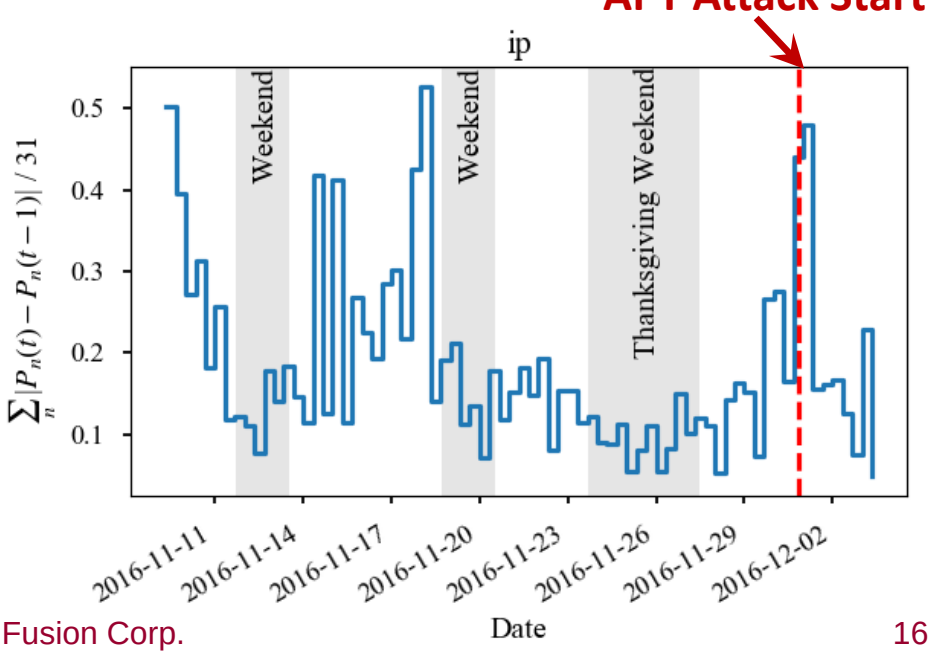

Results: APT Scenario 1

- **Hurricane Panda scenario**
  - Virtual network of servers, laptops, switches, etc.
  - Linux & Windows machines
  - 9 Nov 2016 – 3 Dec 2016
  - Attack distributed 30 Nov – 2 Dec
  - Snort (NIDS) & OSSEC (HIDS)
  - Database injection to gain credentials
  - Lateral movement and firewall deactivation
- **Results**
  - Using threshold at 0.3, CAIN identified 4 anomalies
  - Second two anomalies relate to machines coming online for the first time
  - Last anomaly corresponds with the start of Hurricane Panda's attack

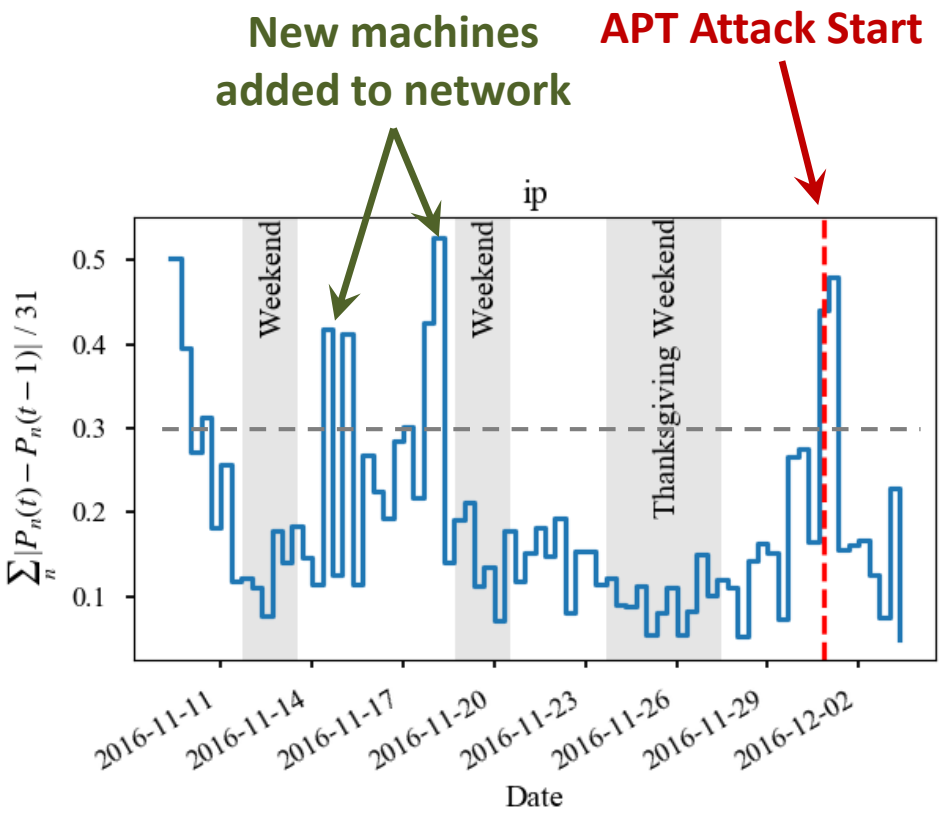

Results: APT Scenario 2

- **Energetic Bear scenario**
  - Same network as Hurricane Panda
  - 1 Jan 2017 – 4 Feb 2016
  - Attack on Jan 31, 2017
  - 644,067 OSSEC (HIDS) alerts
  - Email phishing to redirect user to malicious website
  - Lateral movement through networl using a remote-desktop exploit
  - Attacker attempted to clean-up logs and other traces
- **Results**
  - Using threshold at 0.3, CAIN identified 2 anomalies
  - Third anomaly corresponds with the start of the Energetic Bear attack

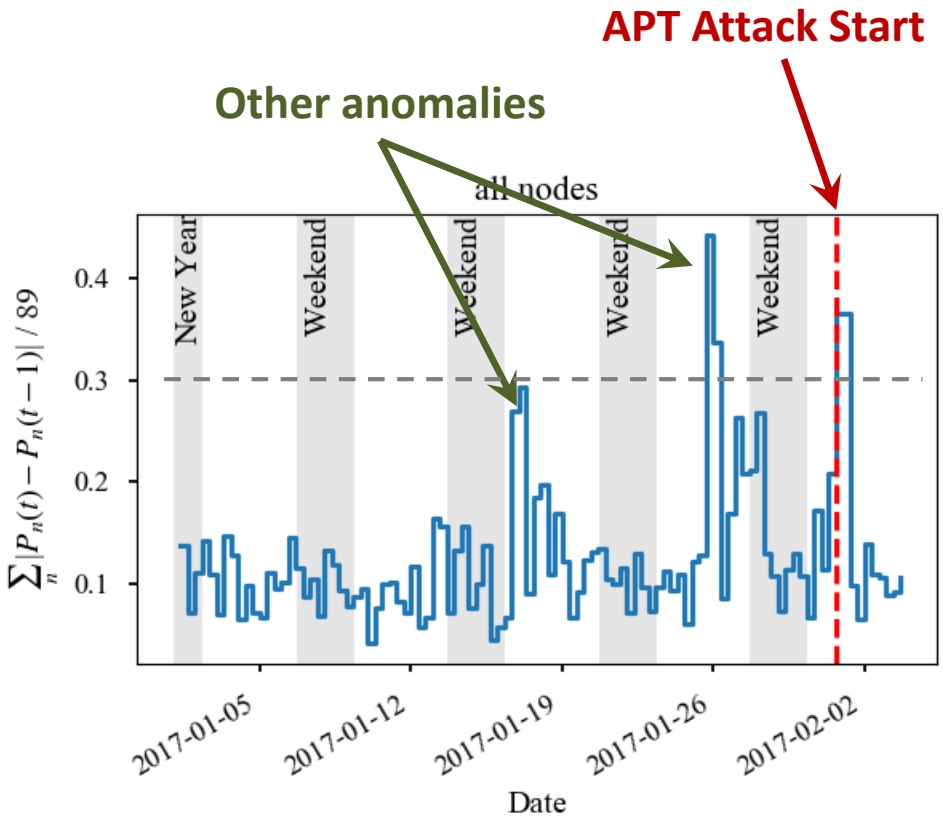

Conclusions: Making Sense of Noisy IDS Sensors with Graph Analytics

- **Graph-based Role-dynamics:**
  - Fuses IDS sensor alerts
  - Reduces >750k alerts to a handful of anomalies
  - Identifies anomalies in IDS alerts during APT attacks
- **Success in 2 APT scenarios demonstrates:**
  - Robust to different types of APTs and attack vectors
  - Insensitive to IDS systems

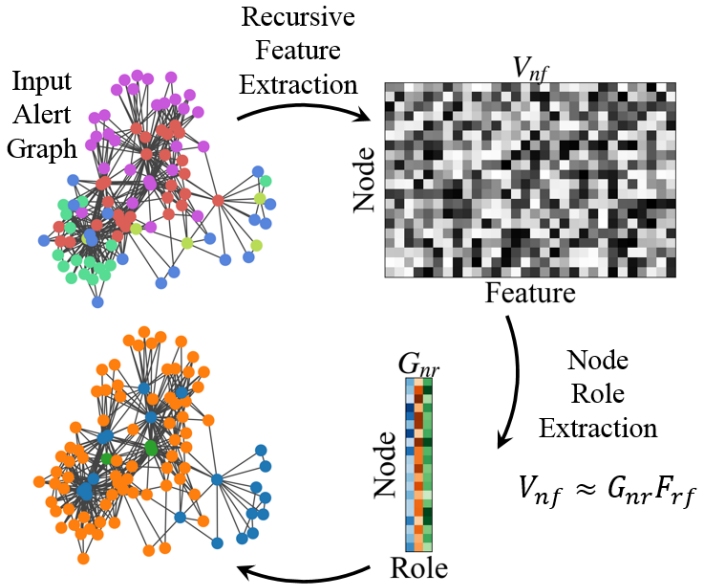

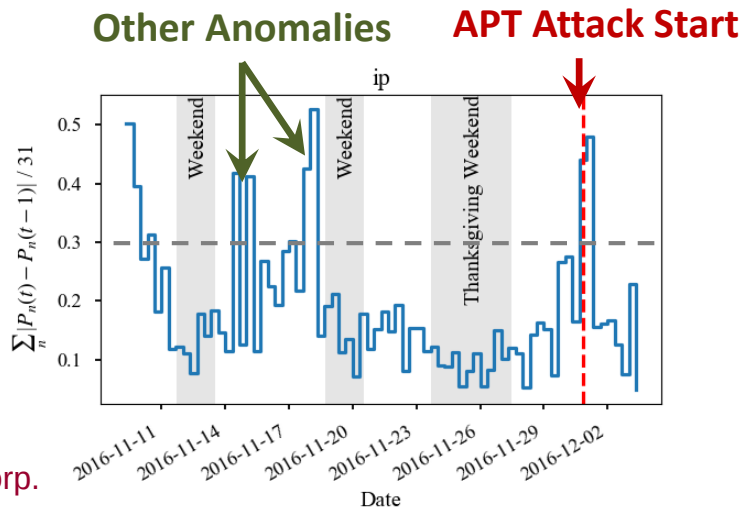

Agenda

- Introduction
- Advanced Persistent Threats
- Graph-node Role-dynamics
- **Bayesian Normalcy Modeling**
- Summary

Bayesian Dynamic Flow Model

- **Unsupervised model of NetFlow traffic dynamics**
- **Assume data follows Poisson distribution**

$$x_t \sim \text{Poisson}(\phi_t)$$

- **Model temporal evolution as Gamma-Beta discount model**
  - Prior: $x_t \sim P(\phi_t | x_{0:(t-1)}) = \Gamma(\delta_t r_{t-1}, \delta_t c_{t-1})$
  - Posterior: $x_t \sim P(\phi_t | x_{0:t}) = \Gamma(\delta_t r_t, \delta_t c_t)$

(X. Chen, et al. 2016)

Results
Bayesian Dynamic Flow Model

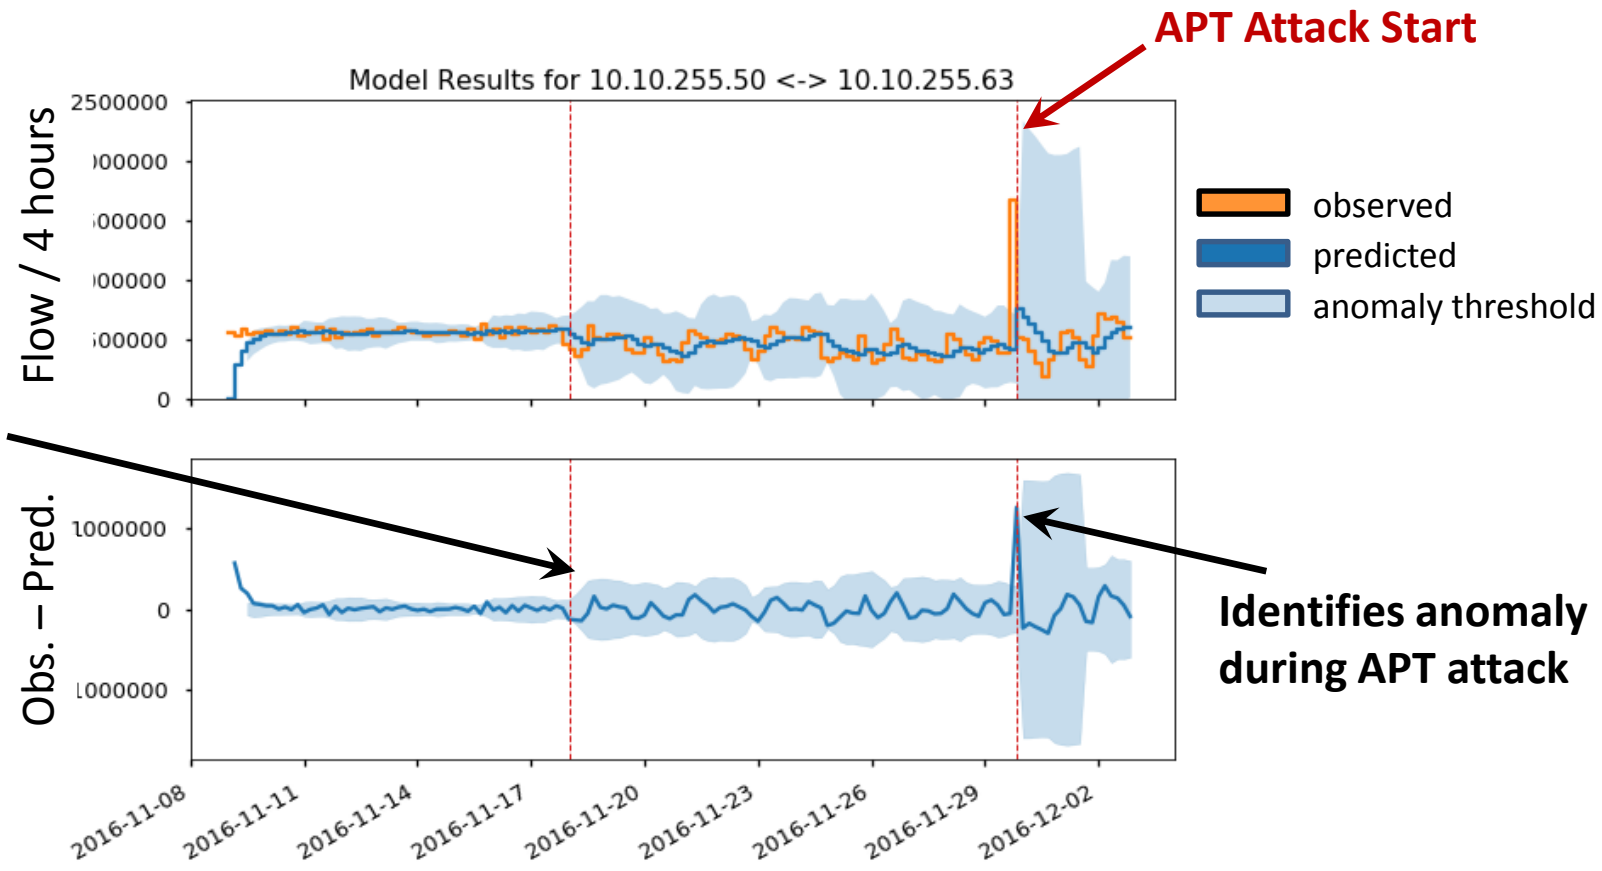

Identifies anomalous change in packet flow volatility

Identifies anomaly during APT attack

- Complementary to graph-based role-dynamics
- Multiple methods corroborate detection

Agenda

- Introduction
- Advanced Persistent Threats
- Graph-node Role-dynamics
- Bayesian Normalcy Modeling
- **Summary**

Summary

- **Developed two complementary anomaly detection techniques**
  - IDS: Graph-based Role Dynamics
  - NetFlow: Bayesian Dynamic Flow Model
- **Tested on two APT scenarios**
  - Hurricane Panda
  - Energetic Bear (a.k.a. Crouching Yeti)
- **Successful anomaly detection in two APT scenarios suggests:**
  - Robust to different types of APTs and attack vectors
  - Insensitive to IDS systems

References

- A. Lemay, et al., "*Survey of publicly available reports on advanced persistent threat actors*," Computers & Security, 72 (2018) 26–59
- L. Akoglu, H. Tong, and D. Koutra, "*Graph-based Anomaly Detection and Description: A Survey*," (2014)
- K. Henderson, et al., "*It's Who You Know: Graph Mining Using Recursive Structural Features*," Proc. 17th ACM SIGKDD Int. Conf. Knowl. Discov. Data Min., (2011) 663–671.
- K. Henderson, et al., "*RolX : Structural Role Extraction & Mining in Large Graphs*," kdd, (2012) 1231–1239
- R. Rossi, B. Gallagher, J. Neville, and K. Henderson, "Role-dynamics: Fast mining of large dynamic networks," *arXiv Prepr.*, (2012)
- X. Chen, et al., "*Scalable Bayesian Modeling, Monitoring and Analysis of Dynamic Network Flow Data*" arXiv:1607.02655 [stat.ME] (2016)

Statements / Disclaimers

- **Copyright 2018 Boston Fusion Corp.**
- **This research was developed with funding from the Defense Advanced Research Projects Agency (DARPA)**
- **The views, opinions and/or findings expressed are those of the author and should not be interpreted as representing the official views or policies of the Department of Defense or the U.S. Government**
- **Distribution Statement "A" (Approved for Public Release, Distribution Unlimited)**