



A Container Adventure: Scaling and Monitoring Kubernetes Logging Infrastructure

Gimi Liang | Sr. Software Engineer
Matthew Modestino | IT Practitioner
David Baldwin | Product Manager

October 2018

Forward-Looking Statements

During the course of this presentation, we may make forward-looking statements regarding future events or the expected performance of the company. We caution you that such statements reflect our current expectations and estimates based on factors currently known to us and that actual events or results could differ materially. For important factors that may cause actual results to differ from those contained in our forward-looking statements, please review our filings with the SEC.

The forward-looking statements made in this presentation are being made as of the time and date of its live presentation. If reviewed after its live presentation, this presentation may not contain current or accurate information. We do not assume any obligation to update any forward-looking statements we may make. In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. Splunk undertakes no obligation either to develop the features or functionality described or to include any such feature or functionality in a future release.

Splunk, Splunk>, Listen to Your Data, The Engine for Machine Data, Splunk Cloud, Splunk Light and SPL are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names, or trademarks belong to their respective owners. © 2018 Splunk Inc. All rights reserved.

Agenda

- ▶ The container landscape
- ▶ Choose your Adventure
 - Splunk Connect for Docker
 - Splunk Universal Forwarder
 - Splunk Connect for Kubernetes
 - Amazon EKS
 - Openshift
- ▶ Wrap-up
- ▶ Q&A



Container Landscape

Database and Data Warehouse



Streaming




Source Code Management



Application Definition and Image Build



Continuous Integration / Continuous Delivery (CI/CD)



App Definition and Development

Orchestration & Management

Runtime

Provisioning

Cloud

Scheduling & Orchestration



Coordination & Service Discovery



Service Management



Cloud-Native Storage




Container Runtime



Cloud-Native Network



Host Management / Tooling



Infrastructure Automation



Container Registries



Secure Images



Key Management



Platforms


Certified Kubernetes - Distribution



Certified Kubernetes - Hosted



Certified Kubernetes - Installer



Non-Certified Kubernetes



PaaS/Container Service



Observability & Analysis

Monitoring



Logging



Tracing



Serverless



This landscape is intended as a map through the previously uncharted terrain of cloud native technologies. There are many routes to deploying a cloud native application, with CNCF Projects representing a particularly well-traveled path.

CLOUD NATIVE Landscape
CLOUD NATIVE COMPUTING FOUNDATION
Redpoint Amplify

Special



Kubernetes Certified Service Provider

Kubernetes Training Partner

Select the tools to deploy and maintain your apps...



Select the tools to deploy and maintain your container cluster...



Host on the public, private, or hybrid cloud...



splunk>



Splunk Enterprise Security™



Splunk IT Service Intelligence™



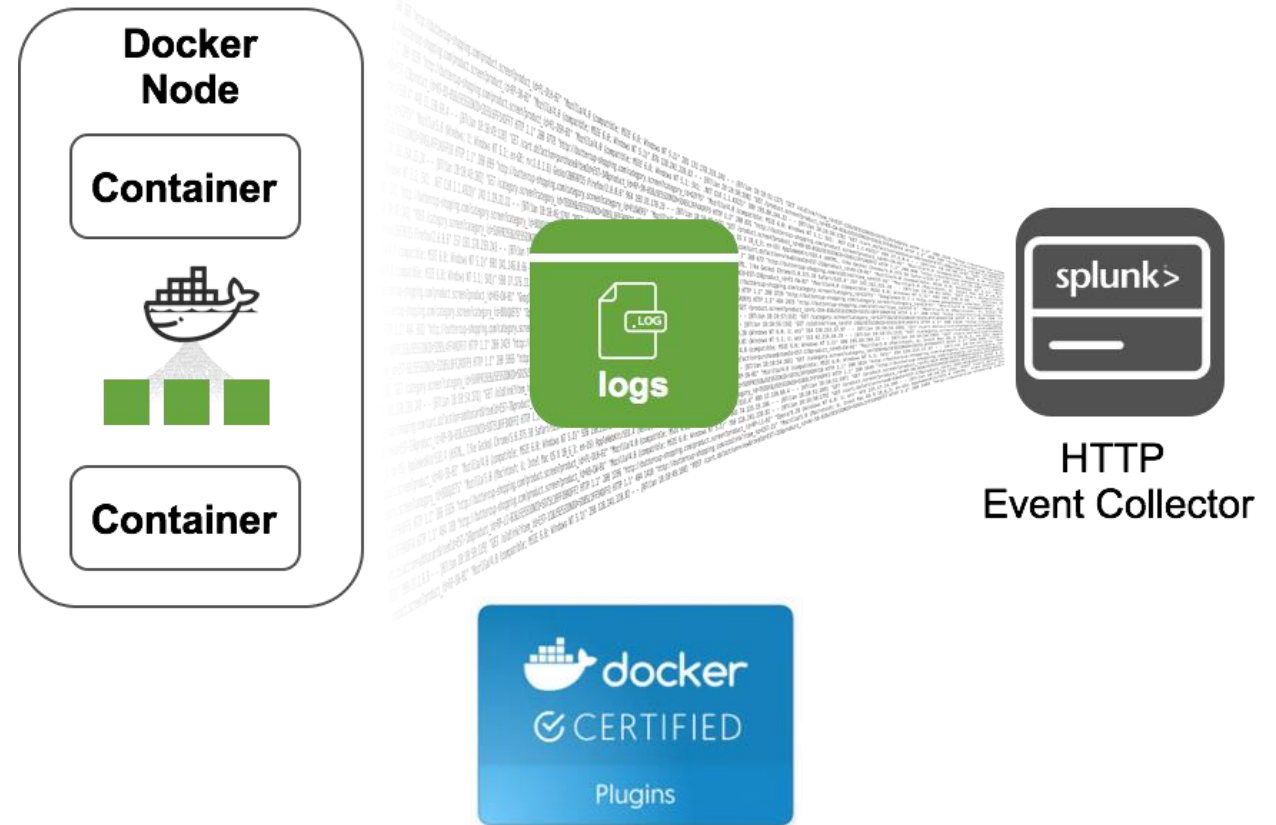
Splunk User Behavior Analytics™

“Gartner predicts that, by 2020, more than 50% of global organizations will be running containerized applications in production, up from less than 20% today”.

**Market Guide for Public Cloud Container
Services August 2017 ID: G00317096**

Splunk Connect for Docker

- ▶ 2015:
 - Introduced Docker Logging Driver
- ▶ 2018:
 - Certified Docker Logging Plug-in
 - Replacement for Docker Logging Driver
 - Supported Open Source*
 - Optimize consumption into Splunk through HEC



Search | Splunk 7.2.0

Splunk Connect for Docker - x

splunk/docker-logging-plugin: x

David

Secure

https://store.docker.com/plugins/splunk-connect-for-docker

☆

★ Bookmarks

Splunk Inc.

Cont

Git-Fingals

DA Kanban

FDSE Kanban


DA Current Progress

TA 3rd Party


DA Priorities

DA Landing

» Other Bookmarks

 docker store

Explore Publish Feedback Log In



Splunk Connect for Docker

By [Splunk](#)

Splunk Connect for Docker is a Splunk supported Docker Logging Plug-in


Plugin


Docker Certified

Linux

x86-64

Logging



Free Plan 

\$0.00

Splunk Connect for Docker is an open source plug-in available for users of Splunk Enterprise, Splunk Light, and Splunk Cloud.

[Terms of Service](#)

Setup Instructions

DESCRIPTION

REVIEWS

RESOURCES

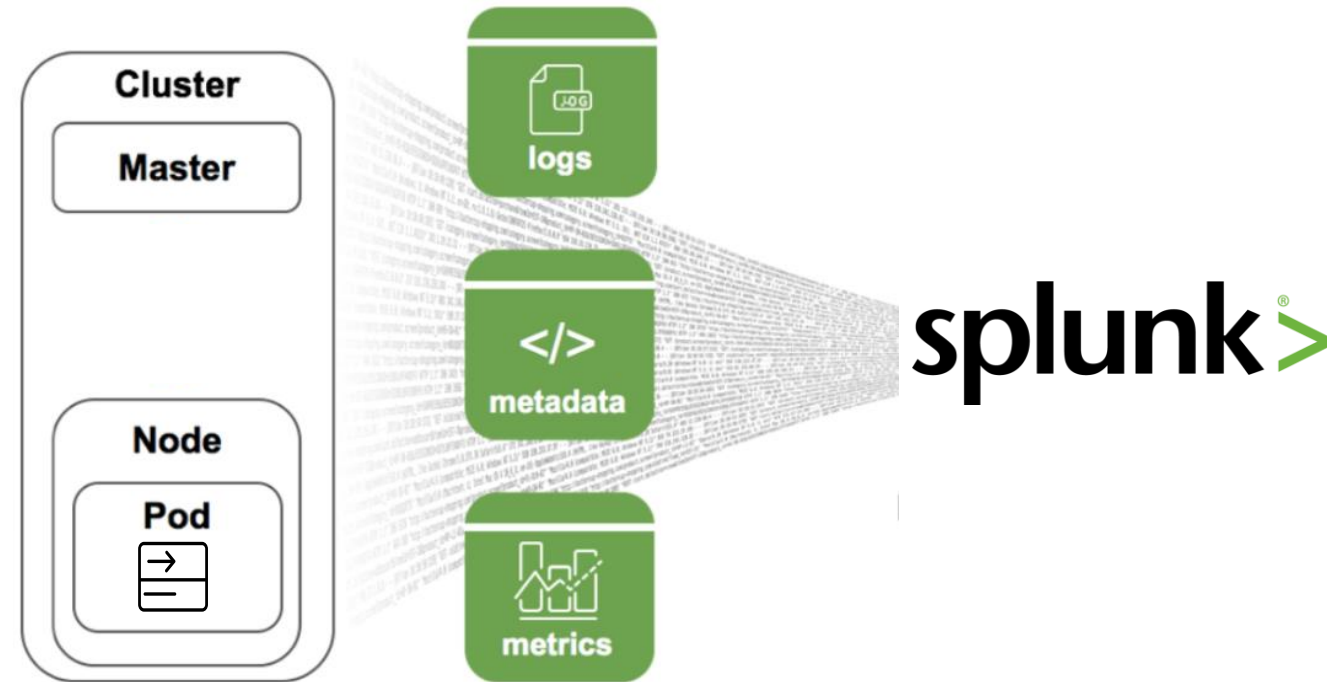
Online Chart Makerhtm

Show All

X

Universal Forwarder

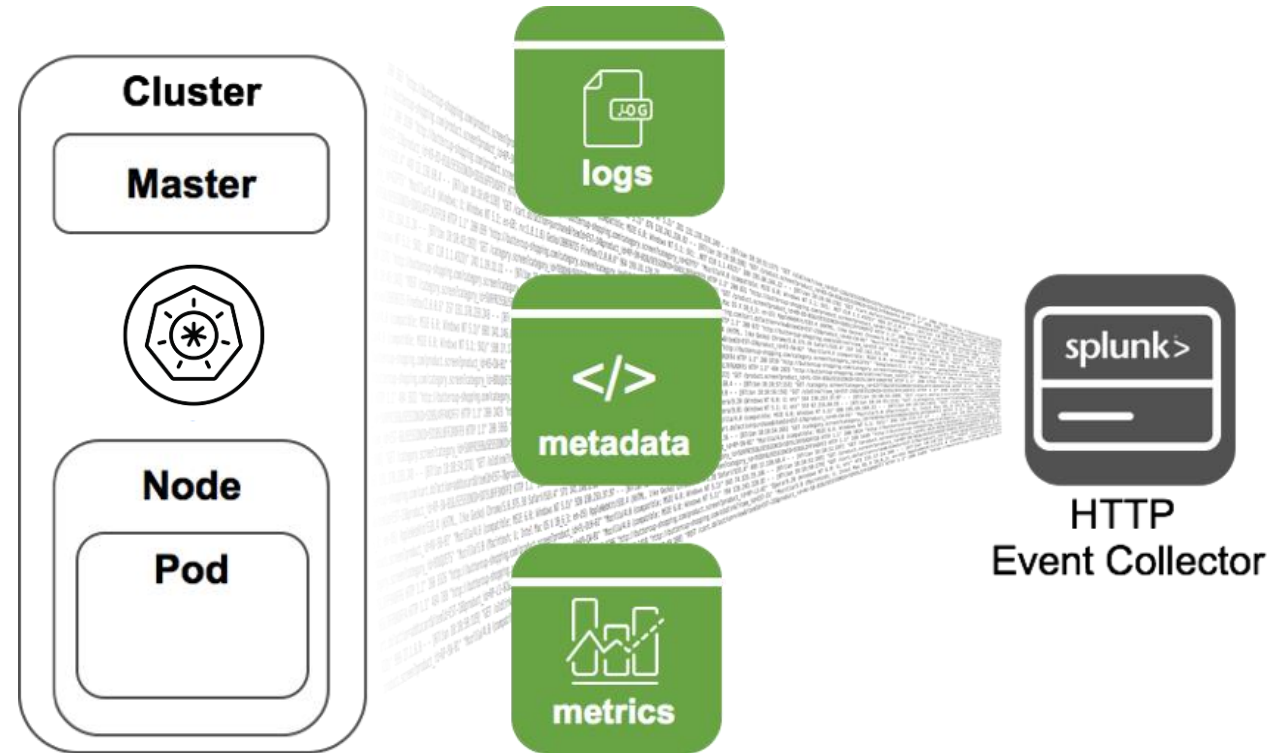
- ▶ Run directly on Kubernetes nodes or as a daemonset using docker image
- ▶ Easy way for Splunk teams to get started with container logs
- ▶ Can move a lot of data, reliably and securely using existing Splunk process
- ▶ Needs modifications to make integration easier with json driver and journald




```
mmodestino-mbp:kubernetes mmodestino$  
mmodestino-mbp:kubernetes mmodestino$  
mmodestino-mbp:kubernetes mmodestino$  
mmodestino-mbp:kubernetes mmodestino$  
mmodestino-mbp:kubernetes mmodestino$
```

Splunk Connect for Kubernetes

- ▶ Secure, Simple, Scale, Configurable
- ▶ 3 Components
 - Logging
 - Metadata/Objects
 - Metrics
- ▶ Leveraged CNCF sponsored projects
- ▶ Managed through Helm
- ▶ Supported Open Source*
- ▶ Optimize consumption into Splunk through HEC



REFERENCE DEPLOYMENT

Red Hat OpenShift on AWS

Container application platform with Kubernetes orchestration on the AWS Cloud

[View deployment guide](#)

This Quick Start sets up a cloud architecture and deploys Red Hat OpenShift Container Platform on AWS.

Red Hat OpenShift Container Platform is based on Docker-formatted Linux containers, Google Kubernetes orchestration, and Red Hat Enterprise Linux (RHEL).

The Quick Start includes AWS CloudFormation templates that build the AWS infrastructure using AWS best practices, and then pass that environment to Ansible playbooks to build out the OpenShift environment. The deployment provisions OpenShift master instances, etcd instances, and node instances in a highly available configuration.

This deployment also includes AWS Service Broker, which provides direct access to AWS services on the Red Hat OpenShift Container Platform.



This Quick Start was developed by
AWS solutions architects.

[Request AWS credits for this deployment](#)



What you'll build



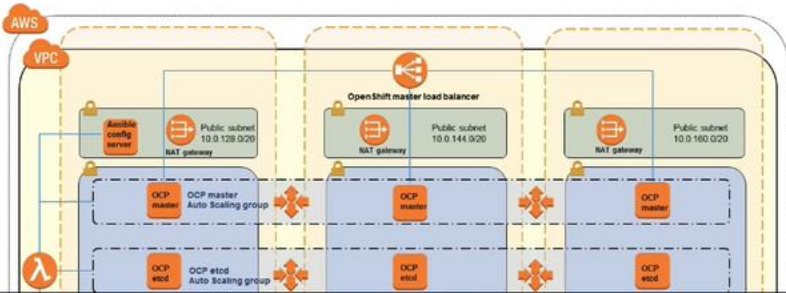
How to deploy



Cost and licenses

Use this Quick Start to automatically set up the following Red Hat OpenShift environment on AWS:

- A virtual private cloud (VPC) that spans three Availability Zones, with one private and one public subnet in each Availability Zone.*
- An internet gateway to provide internet access to each subnet.*
- In one of the public subnets, an Ansible config server instance.
- In the private subnets:
 - Three OpenShift master instances in an Auto Scaling group





Testing

Environment and Setup


- ▶ 3 Node Kubernetes Cluster
 - Master - m4.large
 - Nodes - m4.16xlarge
- ▶ Splunk Deployment
 - Focus is on a single node indexer performance
- ▶ Tested with different message sizes
 - 256 Byte
 - 1 KiB
- ▶ 30 containers generating 1000 messages/sec each
 - 30K messages/sec/node

Benchmarking Results


256 byte Message Size

- ▶ **Consistent ~7.5 Mbps and ~30K EPS through indexer**
- ▶ Results repeatable with different buffer size
- ▶ Executed 4X with 60 sec pause in between

Splunk Indexing Rate



Event Indexed per Second




Benchmarking Results


1KiB Message Size

- ▶ **Consistent ~20 Mbps and ~20K EPS through indexer**
- ▶ Results repeatable with different buffer size
- ▶ Executed 4X with 60 sec pause in between

Splunk Indexing Rate



Event Indexed per Second



Results in a less controlled environment

- 

Testing in the Wild

Results in a less controlled environment

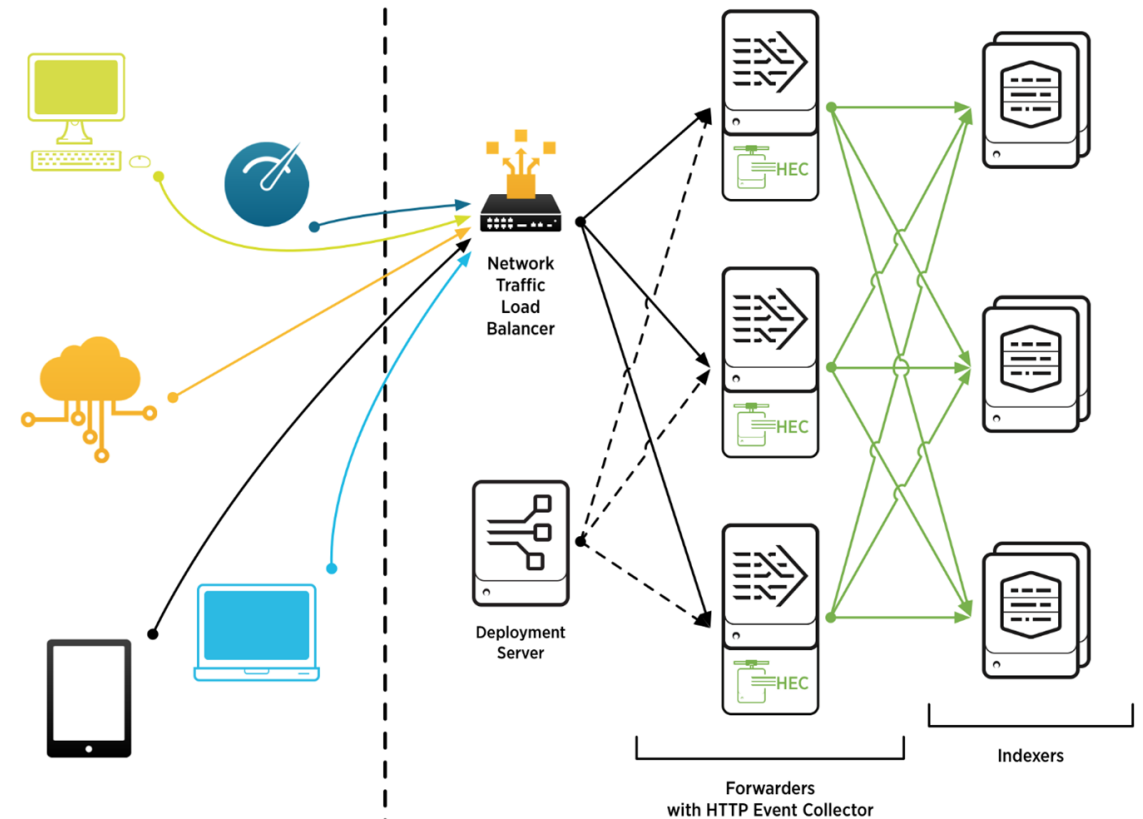
- ▶ Indexing throughput
 - Maxed ~3MB/s - Not bad for public internet!
- ▶ Indexing Latency
 - Saw up to ~7 mins indexing delay, but no message loss



Benchmarking Guidelines

Load Balancing

- ▶ HEC Architecture
- ▶ Single Connection (HF)
 - Each logging pod has 1 connection to HEC
 - HF or LB can be used to spread logs across indexer tier
 - HEC Deploy
 - <http://dev.splunk.com/view/event-collector/SP-CAAAE73>
 - NGINX
 - <http://dev.splunk.com/view/event-collector/SP-CAAAE9Q>



Key Takeaways

- ▶ Container orchestration is the foundation for current and future applications
- ▶ Splunk is the go to solution for container insights for every container deployment
- ▶ Extremely scalable and flexible

Q&A

Gimi Liang | Engineering

Matthew Modestino | Practitioner & Evangelist

David Baldwin | Product Management

Looking for More?

Check out these sessions here at .conf18!

All Skill Levels

IT Operations

⊖ IT1501 - Deep Dive into Boss of the NOC 2018 with a Splunk IT Specialist

SCHEDULE

Thursday, Oct 04, 11:00 a.m. - 11:45 a.m.

Advanced

Foundations/Platform

⊖ FN1089 - Dockerizing Splunk at Scale 2: The Container Strikes Back

SCHEDULE

Thursday, Oct 04, 2:45 p.m. - 3:30 p.m.

Join us in the Innovation Lab for more on Splunk and containers!

Resources Discussed During Session

- Repo: <https://github.com/splunk/splunk-connect-for-kubernetes>
- Gems:
 - <https://rubygems.org/gems/fluent-plugin-kubernetes-objects>
 - <https://rubygems.org/gems/fluent-plugin-splunk-hec>
- Docker Hub:
 - <https://hub.docker.com/r/splunk/fluentd-hec/>
 - <https://hub.docker.com/r/splunk/kube-objects/>
- ▶ **Splunk Connect for Docker**
 - Repo: <https://github.com/splunk/docker-logging-plugin>
 - Docker Store: <https://store.docker.com/plugins/splunk-connect-for-docker>
- ▶ **Universal Forwarder:**
 - <https://store.docker.com/community/images/splunk/universalforwarder>



Resources Discussed During Session

Thank You

Don't forget to **rate this session**
in the **.conf18** mobile app

.conf18

splunk>