

RSA®Conference2020

San Francisco | February 24 – 28 | Moscone Center

HUMAN
ELEMENT

SESSION ID: ACB-F03

Scaling Key Management: Thousands of Clients, Trillions of Operations







Yaron Sheffer

Director, Security Technologies Product Development
Intuit

@yaronsheffer

#RSAC

Intuit

- A global financial platform company
- Maker of  **turbotax**  **quickbooks**  **mint**
- More than **50M customers** trust us with their financial data
- Steadfast commitment to security
- Strategic move to the AWS cloud
- Acquired Israeli security startup Porticor in Feb. 2015 

IDPS, Intuit Data Protection Service

- Secret and Key Management service
 - Storage for secrets and encryption keys
 - Cryptographic operations




- For highly sensitive data we require both transparent disk-level and **application-level data encryption**

RSA®Conference2020

Core Features


Where to Encrypt

- Fetch the key into the client, then encrypt
 - Alternatively, send the data to be encrypted on the KM service
-
- Instead: RDLE – **Remote Derive, Local Encrypt**
-
- Fine-grain keys: implications on threat model
 - Spatial granularity
 - Also on performance, scalability



Temporal Granularity: Key Versions

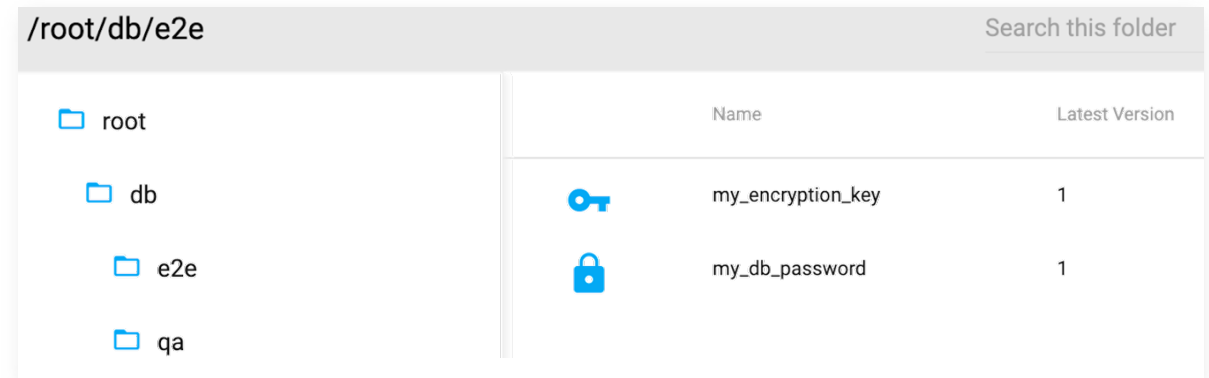
- New key version every 90 days
- Key versioning vs. re-encryption (vs. “key rotation”)
- **Versioning is automated**, saved as part of a ciphertext header
 - No industry standard, sigh





- Versioning works well with probabilistic encryption
- More challenging with deterministic encryption

Managing Keys

- Each project receives a **strictly segregated namespace**
- A namespace is a hierarchy of folders and keys
 - Keys consist of multiple key versions
- Keys can be created, listed etc. with a UI and a REST API
- Role-based access control at the folder level
- Custom permissions
e.g. read public-key only



The screenshot shows a web interface for managing keys. The breadcrumb path is `/root/db/e2e`. On the left, a folder tree shows `root` containing `db`, which contains `e2e` and `qa`. The main area displays a table of keys under the heading "Search this folder".

	Name	Latest Version
	my_encryption_key	1
	my_db_password	1

Client Authentication

- Initially, had an OAuth1-like message signing
 - Assert long-term credentials, obtain a short-term bearer token
- Found out this is not workable: turtles all the way down
- Built a separate **policy authentication** service
 - Policy: AWS role plus additional properties
 - Good for: EC2 instances, Lambda functions, Docker containers, more
- Now being reused for other internal services



The Customer View

- IDPS only used by *internal* customers
-
- SDK: Java, Python, Go, JavaScript
 - Using the REST API directly is hard
 - UI and a CLI tool for secret management
-
- No manual handling of keys
 - Enterprise-wide governance

Why Have Your Own KMS?

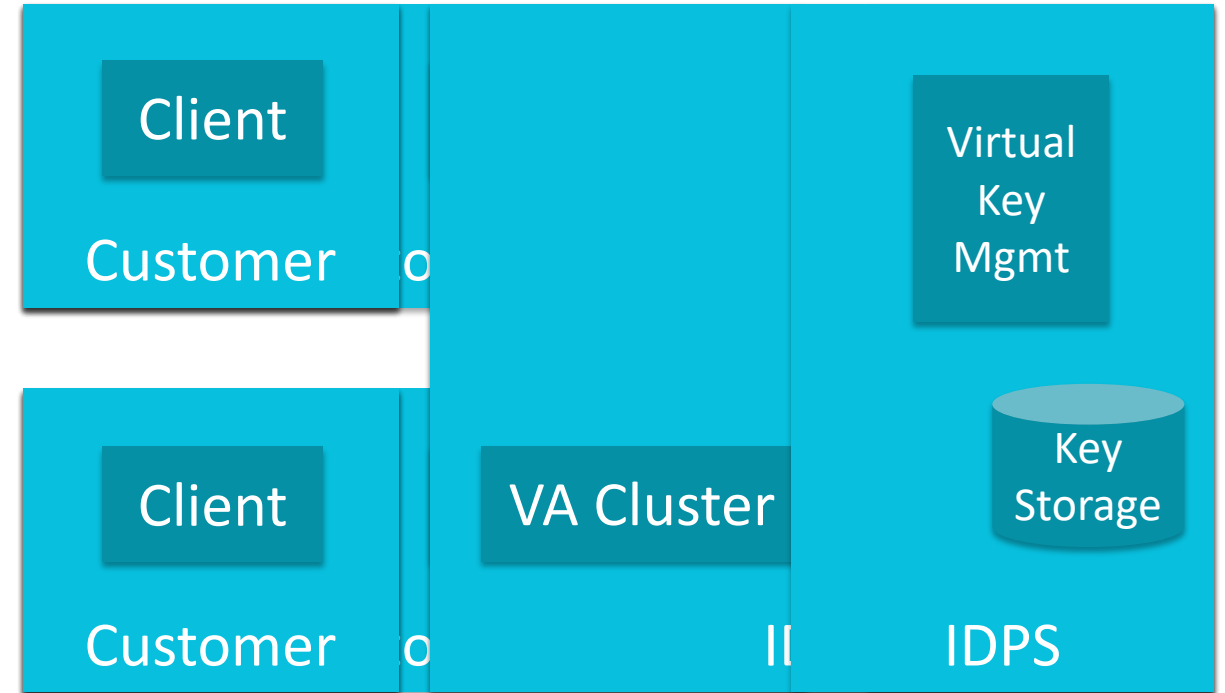


RSA[®]Conference2020

Service Architecture


The Architectural Journey

1. Virtual Appliance (VA)
owned by customer,
master key
2. Fully managed
service, including
VA clusters
3. End-to-end secrecy
4. Multitenant
key derivation



Evolving Key Derivation

- Need a deterministic function of a secret (“root key”) and a non-secret input, returning a secret derived key
 - With several security guarantees
- HKDF, RFC 5869
- How to do it in a shared environment?
Oblivious PRF!
- The service uses a blinded version of the root key to derive a subkey
 - Which is then unblinded by the client



IDPS Technology

- EC2 instances
 - Formerly also hosted in a traditional data center
- RHEL, moving to Amazon Linux 2
- Go
 - Including most crypto code
- On the client side, primarily Java
 - Bouncy Castle for the crypto code
 - Apache Commons Crypto for performance
- Moved from CloudFormation to Terraform



RSA®Conference2020

Developed for Security

Security Principles

- IDPS is primarily about securing applications, not users
- Keys are only in the clear when used
- Cryptography is bog standard
 - When in doubt, call a cryptographer

Security Principles

- Humans are fallible
 - Dev/ops separation
 - No access to encryption keys, by anybody



Security Infrastructure: Intuit

- Large scale AWS account management
 - Including centralized SSO
- Automated policy enforcement
 - Both via tickets and real-time technical intervention
- Internal red team
- Strong compliance org/culture
- Cost tracking tools

Security Infrastructure: Unique to IDPS

- Zoning and containment: multiple separate AWS accounts
- SELinux
 - A strict custom policy reduces risk from code vulns and Golang zero days
- Dedicated instances
 - They come at a cost
- The Meltdown/Spectre family of issues validated these decisions



Security Infrastructure: Unique to IDPS (Cont.)

- Very conservative: no Lambda functions, no containers
 - But we support any *customer* deployment model
- Short term key backup allows rapid recovery of customer keys
- Custom disaster recovery aimed primarily at **malicious service interruption**



RSA®Conference2020

Wrap Up

Apply What You Have Learned Today: Recommendations

- In the next **two weeks** verify:
 - That you have a valid key management strategy, one that applies to modern automated key use
 - That you can identify its true security value, beyond “compliance”

- In the next **two months** make sure:
 - That this strategy remains relevant when you transition to the cloud
 - That security-critical infrastructure in the cloud is built with cloud ab/use cases in mind

Summary

- **Security-sensitive services can be run in the cloud**
- Consider, in depth, your particular cloud threat model
- But remember it's a moving target
- Security-optimized architecture, without hindering developer productivity

**Security software in the cloud:
Not easy but definitely fun!**

