

RSA®Conference2020

San Francisco | February 24 – 28 | Moscone Center

HUMAN
ELEMENT

SESSION ID:XXX-XXXX

Secure Sandboxing in a Post-Spectre World



Tyler McMullen

CTO
Fastly
@tbmcmullen

Jonathan Foote


Principal Security Architect
Fastly
@footePGH


#RSAC

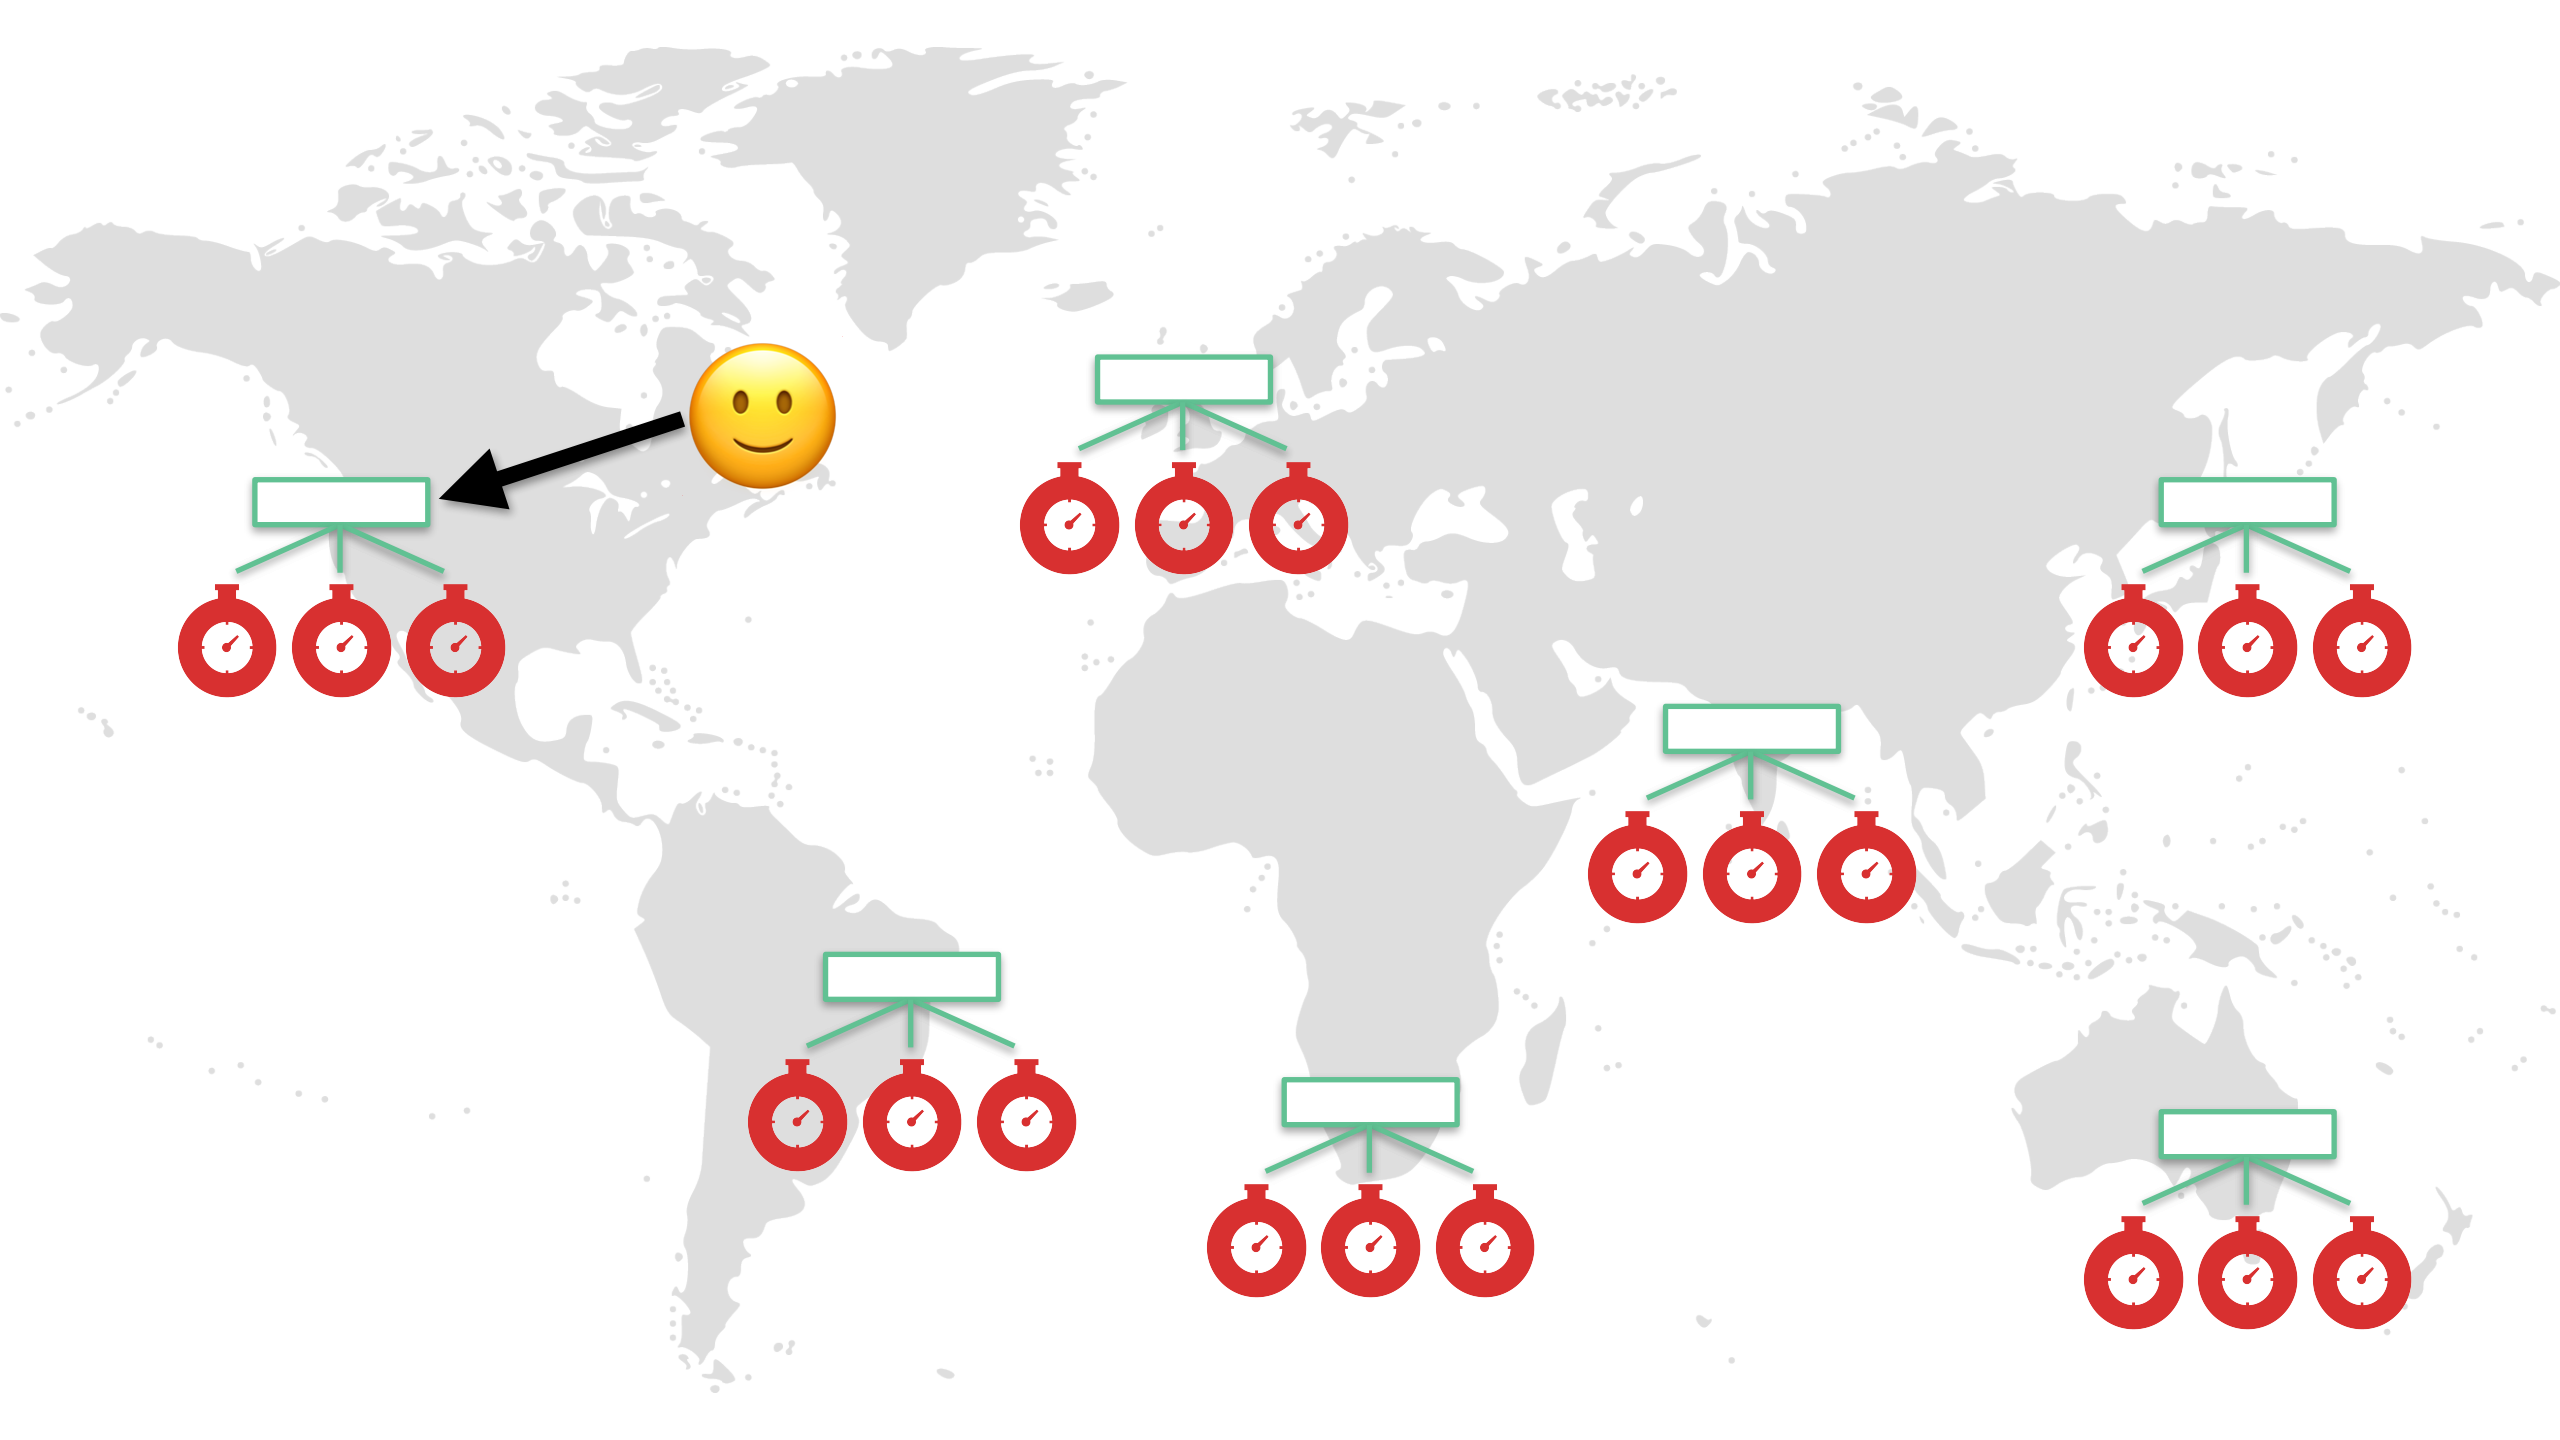


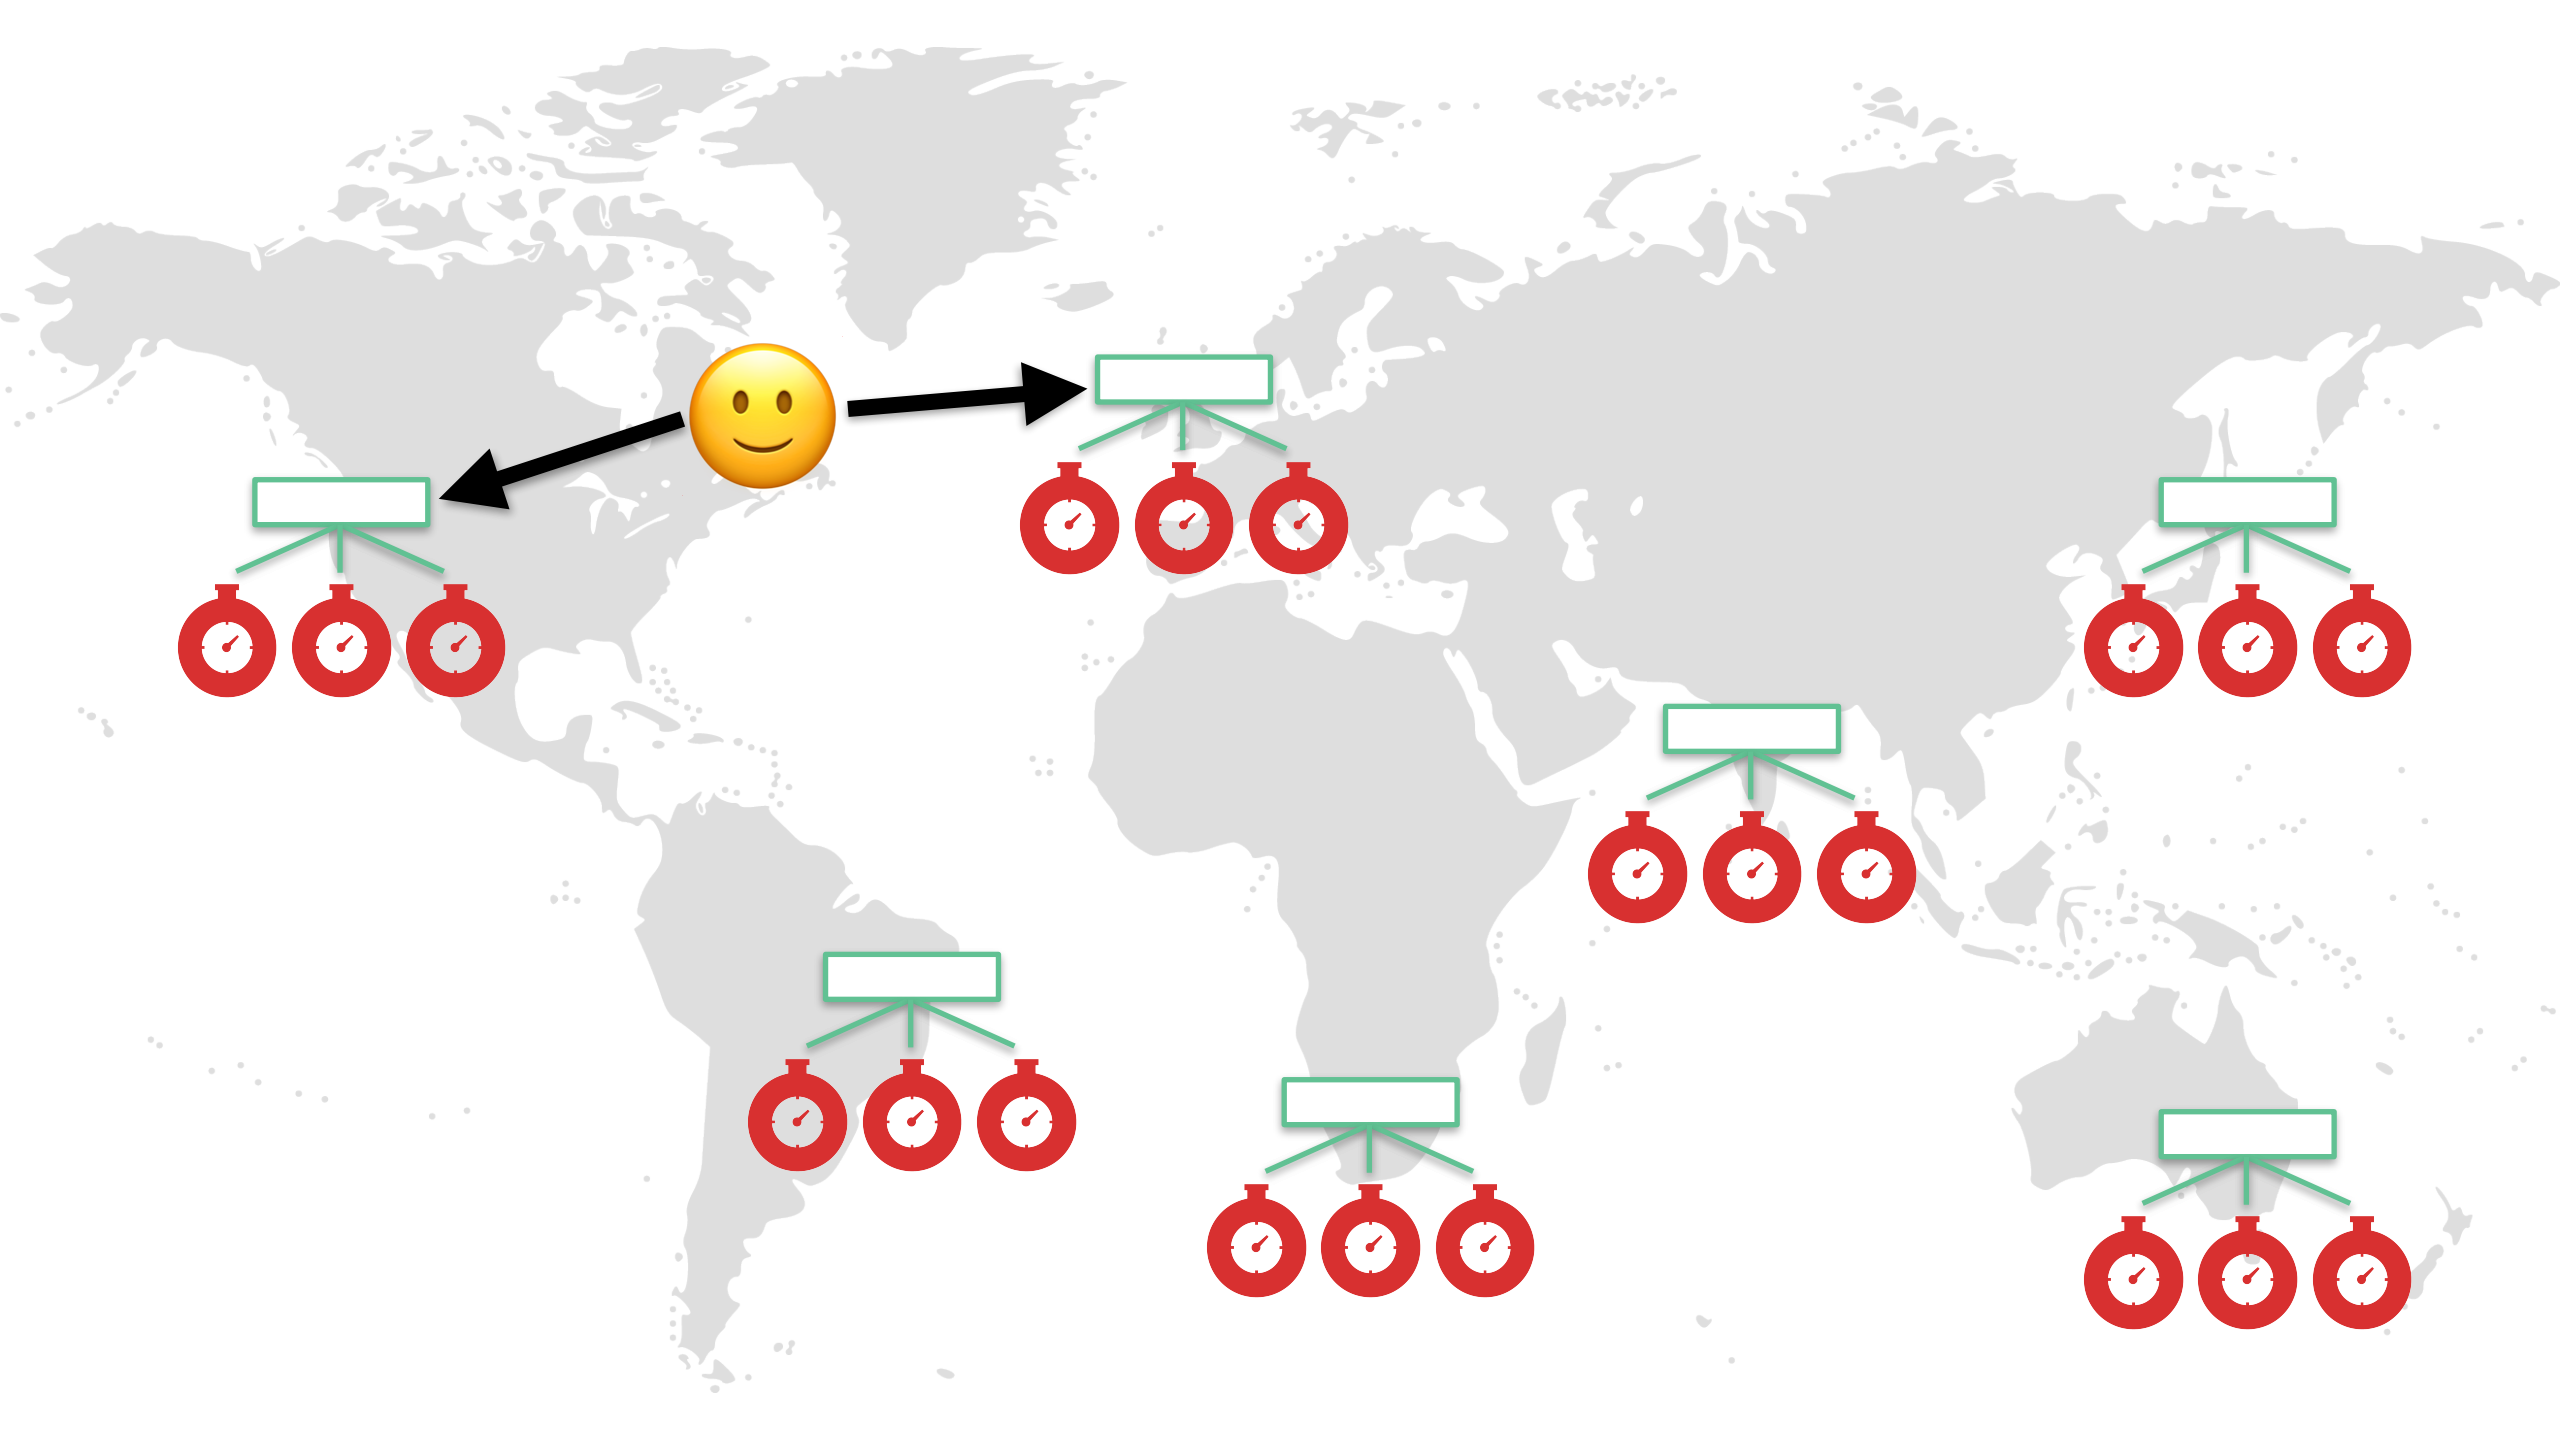
RSA®Conference2020

Edge compute










FaaS



Ruby VM





Virtualization

Process

Sandbox

Sandbox

Sandbox


Sandbox

Sandbox

Sandbox

Sandbox

Sandbox




Granular Sandboxing




bytecodealliance/lucet: Lucet, × +

← → ↺

github.com/bytecodealliance/lucet

☆ !  ⋮





Search or jump to... /


Pull requests

Issues

Marketplace

Explore

 + ↓  ↓

 bytecodealliance / lucet

Unwatch ▼

94

★ Star

2.9k

Fork

97

<> Code

! Issues 44

🔗 Pull requests 8

▶ Actions

📁 Projects 0

📖 Wiki

🛡 Security

📊 Insights

⚙ Settings

Lucet, the Sandboxing WebAssembly Compiler.

Edit

webassembly

wasi

assemblyscript

rust

wasm

Manage topics

🕒 893 commits

🌿 25 branches

📦 0 packages

🏷 9 releases

👤 19 contributors

📄 Apache-2.0

Branch: master ▼


New pull request

Create new file


Upload files

Find file

Clone or download ▼


 data-pup Merge pull request #386 from data-pup/module-bindings-tests ...

✓ Latest commit 4b59161 3 days ago

 .cargo


Replace -W,-export-dynamic with -rdynamic

9 months ago

 .github


add `make test-full` target; make `test` more focused

7 days ago

 assemblyscript

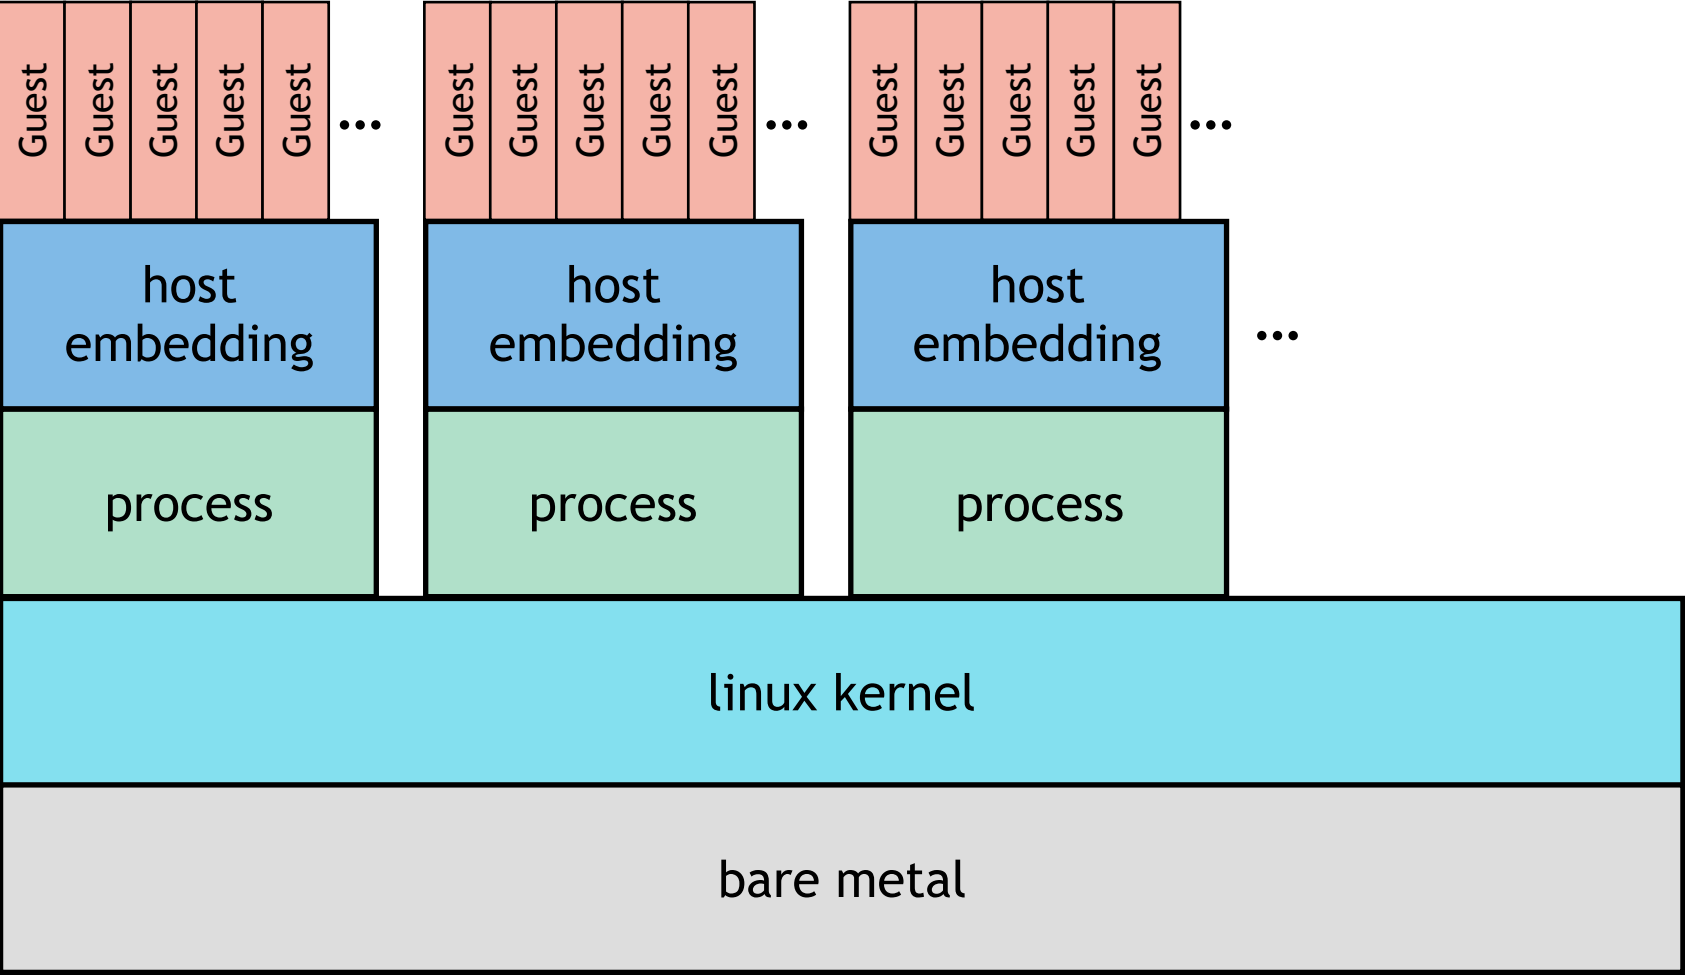
Update AssemblyScript, WASA and the AssemblyScript example and instru...

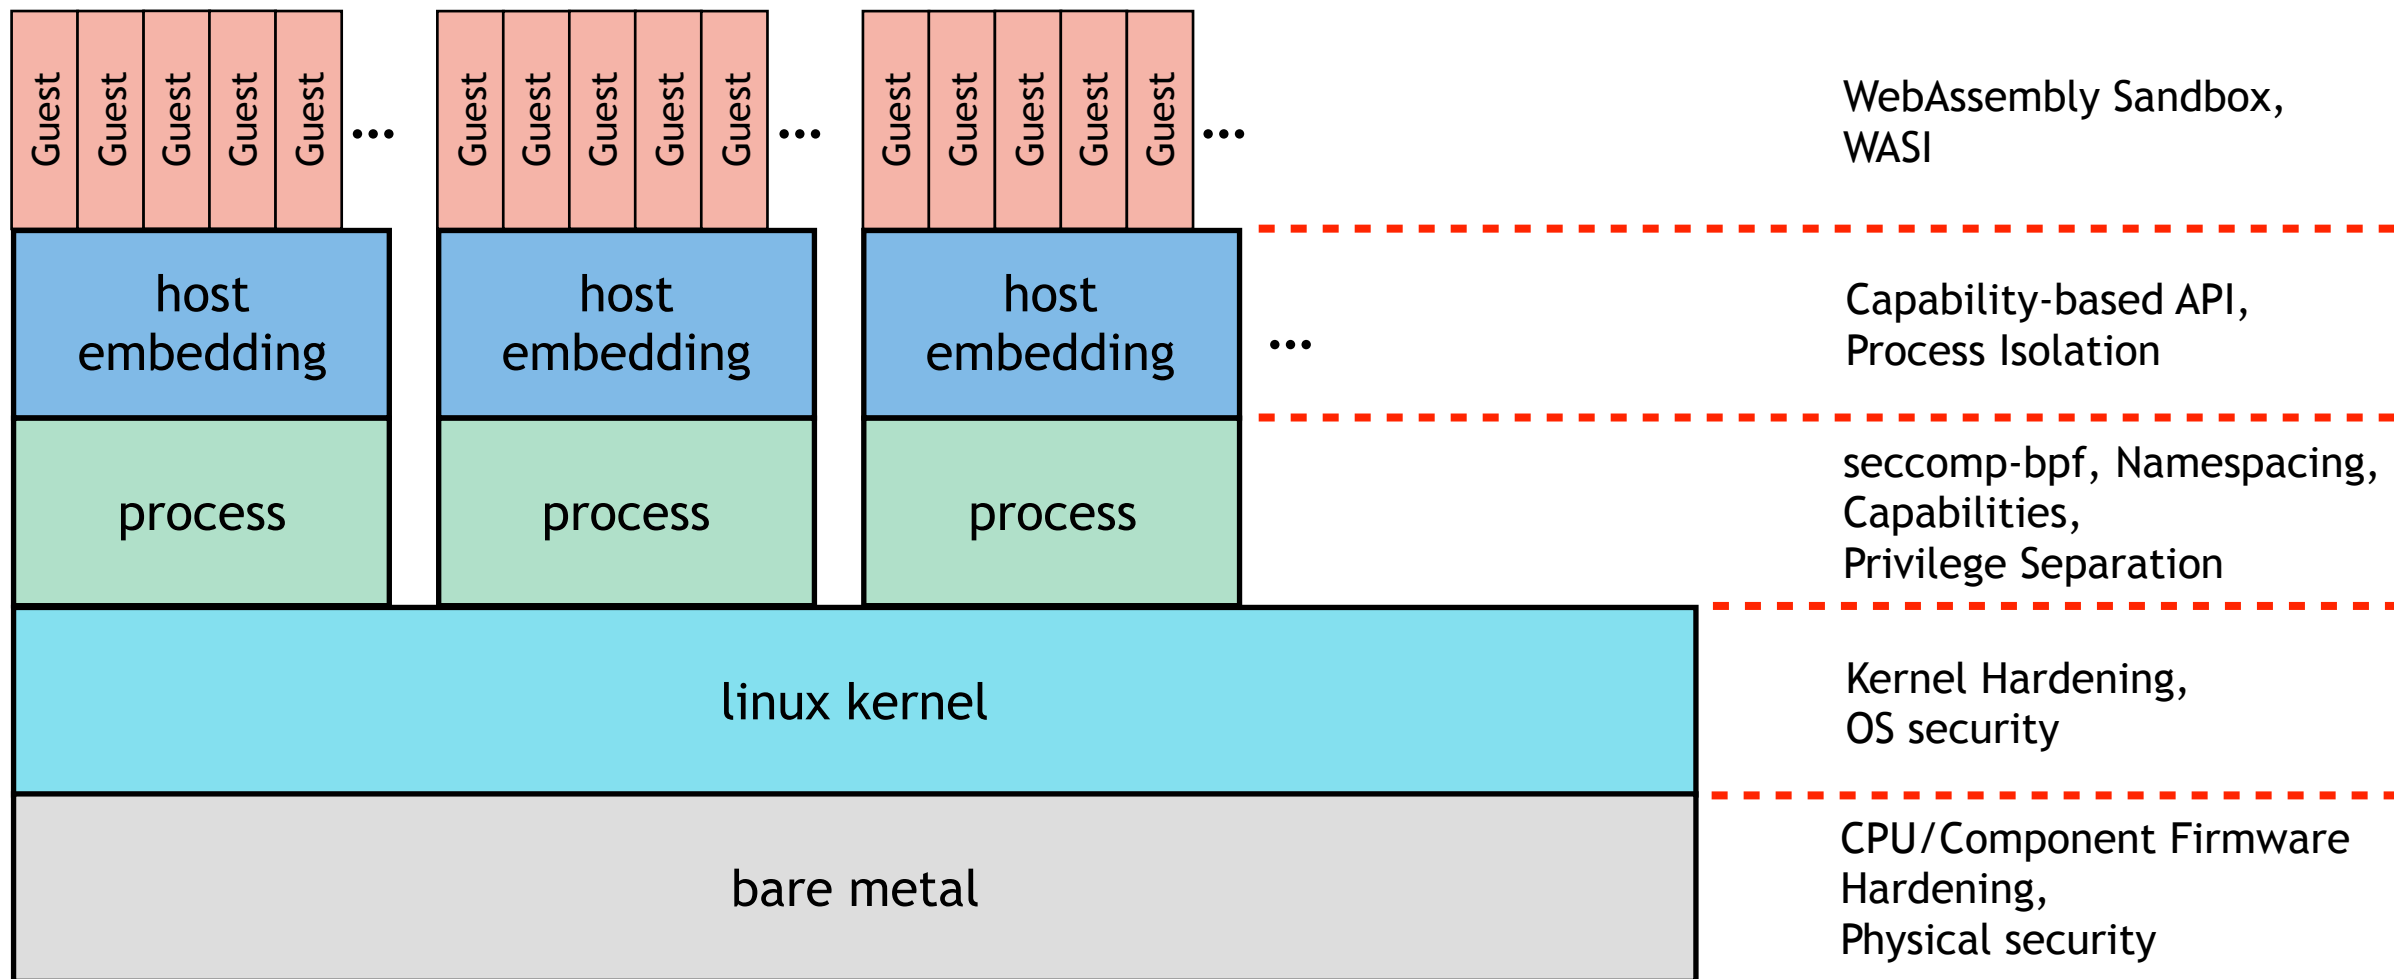
6 months ago

 assets

Don't consider assets/*.png files as text

9 months ago





RSA®Conference2020


Transient execution security

RSA®Conference2020

Research and development evolution

Attacks and defenses






2018



2019



2020

<https://cpu.fail/>

<https://meltdownattack.com/>

<https://foreshadowattack.eu/>

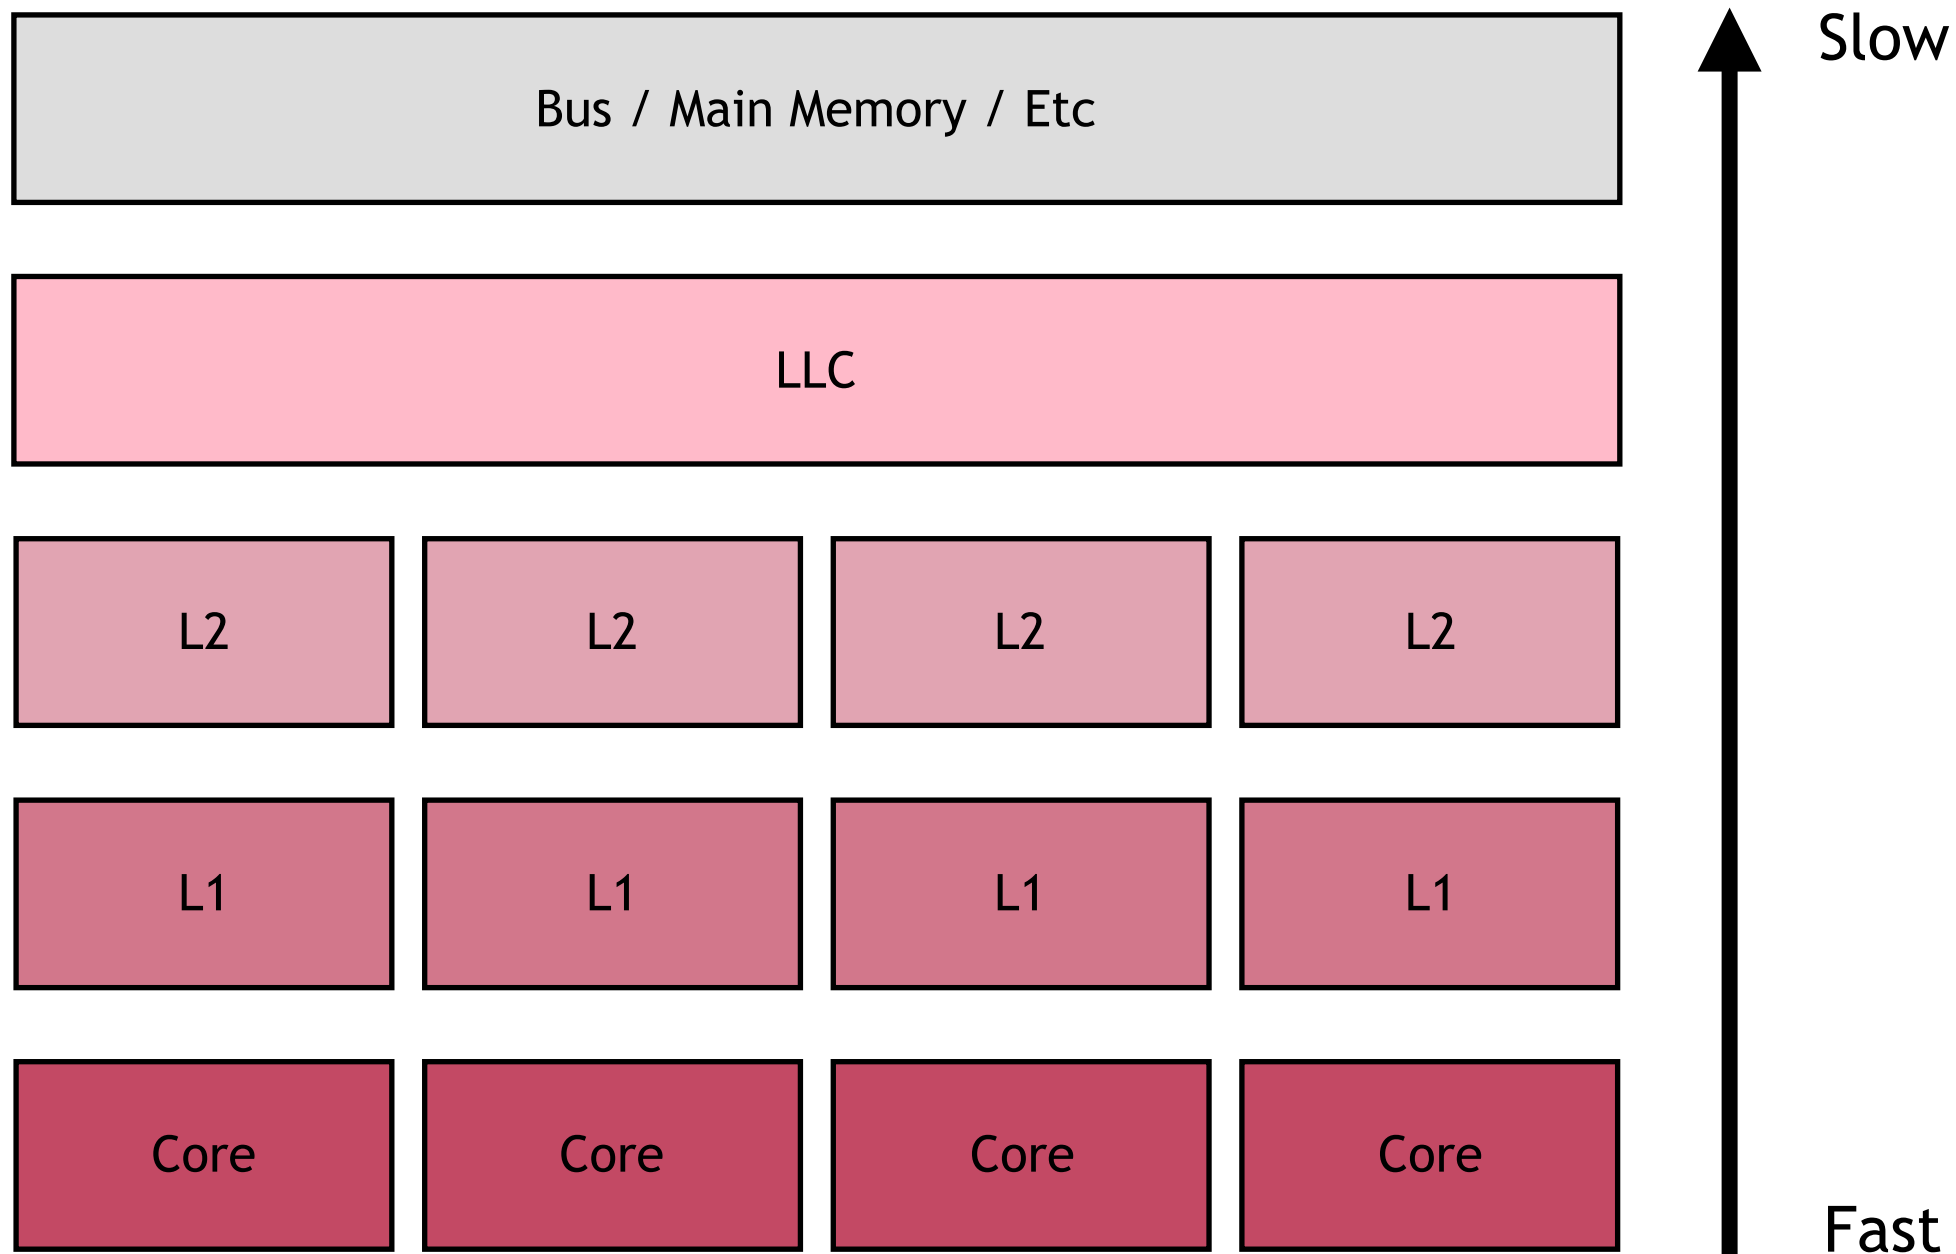
Expectations

- Advances in offensive security research
 - Unexplored CPU components, new techniques
- Significant improvements in HW-supplied defenses
 - Intel, AMD
- Continued exposure of Fastly network
 - New hardware
 - New execution models
- Malicious exploitation

RSA®Conference2020

Transient execution attacks

Concepts



arch



μ arch

RSA®Conference2020

Transient execution attacks and defenses

Analysis by example

arch

μarch

arch

μ arch





arch

μarch

```
→ secret = "secret"; // or set affinity, etc.  
→ oracle = malloc(4096 * 256);
```

```
→ for cache_line in oracle {  
    clflush(cache_line);  
}
```

L1D



arch

```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```

```
for cache_line in oracle {
    clflush(cache_line);
}
```



→ gadget:

```
// ...
ret2spec_gadget(i) {
    // ...


    *(oracle + secret[j]);
}
```

μarch

L1D



RSB



Priming the system: defenses

- Disallow clflush and proxies
- Disallow large array allocations
- Process isolation
- Secrets management
- Detect suspicious code patterns
- Detect suspicious L*D and allocation activity
- ...

arch

```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```



```
for cache_line in oracle {
    clflush(cache_line);
}
```

```
gadget:
    // ...
    ret2spec_gadget(i) {
        // ...


        *(oracle + secret[j]);
    }
```

μarch

L1D



RSB



arch

```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```

```
for cache_line in oracle {
    clflush(cache_line);
}
```

→ `ret2spec_recurse();`



gadget:

```
// ...
ret2spec_gadget(i) {
    // ...


    *(oracle + secret[j]);
}
```

μarch

L1D



RSB



Problematic transient execution: defenses

- Apply microcode patches
- Limit reach of user-supplied code
- Flush shared resources by security domain
- Co-schedule by security domain
- Detect suspicious code patterns, runtime activity
- ...

CPU component	Trans exec technique	ECP prevention controls
BPU::BTB	Bounds check bypass / Spectre v1	32-bit addressing for guest programs Array index masking or lfence/serialization for hostcalls <i>Process isolation by security domain</i>
	Branch target injection / Spectre v 2	Retpoline for indirect calls within Wasm guests Retpoline for host or IBPB on hostcall invocation <i>IBPB on hostcall return</i> <i>IBRS on guest context switch</i> Enable STIBP globally <i>Process isolation by security domain</i>
...
Memory disambiguator	Speculative Store Bypass (SSB) / Spectre v 4	32-bit addressing for guest programs Array index masking or lfence/serialization for hostcalls <i>SSBD</i>
Line Fill Buffer (LFB), Load port, Store buffer	Microarchitectural data sampling (MDS) / Zombieload (variants 1-5) / RIDL	Co-schedule physical hardware by security domain <i>Process isolation by security domain</i> VERW (or equivalent) on host+guest context switches <i>Disable TSX</i> <i>Enable KPTI</i> <i>Enable SMAP</i> <i>Disable SGX</i>
...

arch

```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```

```
for cache_line in oracle {
    clflush(cache_line);
}
```

```
ret2spec_recurse();
```

```
gadget:
```

```
// ...
```

```
ret2spec_gadget(i) {
```



```
// ...
```

```
*(oracle + secret[j]);
```


```
}
```

μarch

L1D



RSB



arch

μarch

```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```

```
for cache_line in oracle {
    clflush(cache_line);
}
```

```
ret2spec_recurse();
```

→



```
for block in oracle {
    mfence; lfence; rdtsc;
    // ...
    update_scores();
}
```

```
gadget:
```


```
// ...
ret2spec_gadget(i) {
    // ...

    *(oracle + secret[j]);
}
```

L1D



RSB



Measuring side effects: defenses

- Disallow asm
- Disallow fine-grained timers/primitives; add jitter
- Deterministic ordering of all events
- Array preloading, non-deterministic array access, buffer ASLR, etc.
- ...

arch

```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```

```
for cache_line in oracle {
    clflush(cache_line);
}
```

```
ret2spec_recurse();
```

```
for block in oracle {
    mfence; lfence; rdtsc;
    // ...
    update_scores();
}
```



```
gadget:
```

```
// ...
ret2spec_gadget(i) {
    // ...


    *(oracle + secret[j]);
}
```

μarch

L1D



RSB



```
secret = "secret"; // or set affinity, etc.
oracle = malloc(4096 * 256);
```

```
→ for k in (len(secret) * LRG_NUMBER) {
    for cache_line in oracle {
        clflush(cache_line);
    }


    ret2spec_recurse();

    for block in oracle {
        mfence; lfence; rdtsc;
        // ...
        update_scores();
    }
}

gadget:
    // ...
    ret2spec_gadget(i) {
        // ...

        *(oracle + secret[j]);
    }
```


L1D



CPU



RSB



Orchestrating the attack: defenses

- Disallow longer, continuous runtimes
- Isolate workloads by security domain
 - “Quarantine” low-trust workloads
- Detect suspicious code patterns, runtime activity
- ...

RSA®Conference2020

Transient execution attacks and defenses

Approach for Fastly Compute@Edge



flush+reload



evict+time

prime+probe

flush+flush

prime+abort

...




Prime the system

Invoke problematic
transient execution

Read side-channel

Orchestrate the attack



<https://cpu.fail/>

<https://meltdownattack.com/>

<https://foreshadowattack.eu/>

Prime the system



Disallow direct, granular flushing of cache lines
Process isolation by security domain
...

Invoke problematic
transient execution



Apply microcode patches
Limit reach of user-supplied code
...

Read side-channel



Apply coarse-grained, jittered timers
Non-deterministic array access
...

Orchestrate the attack

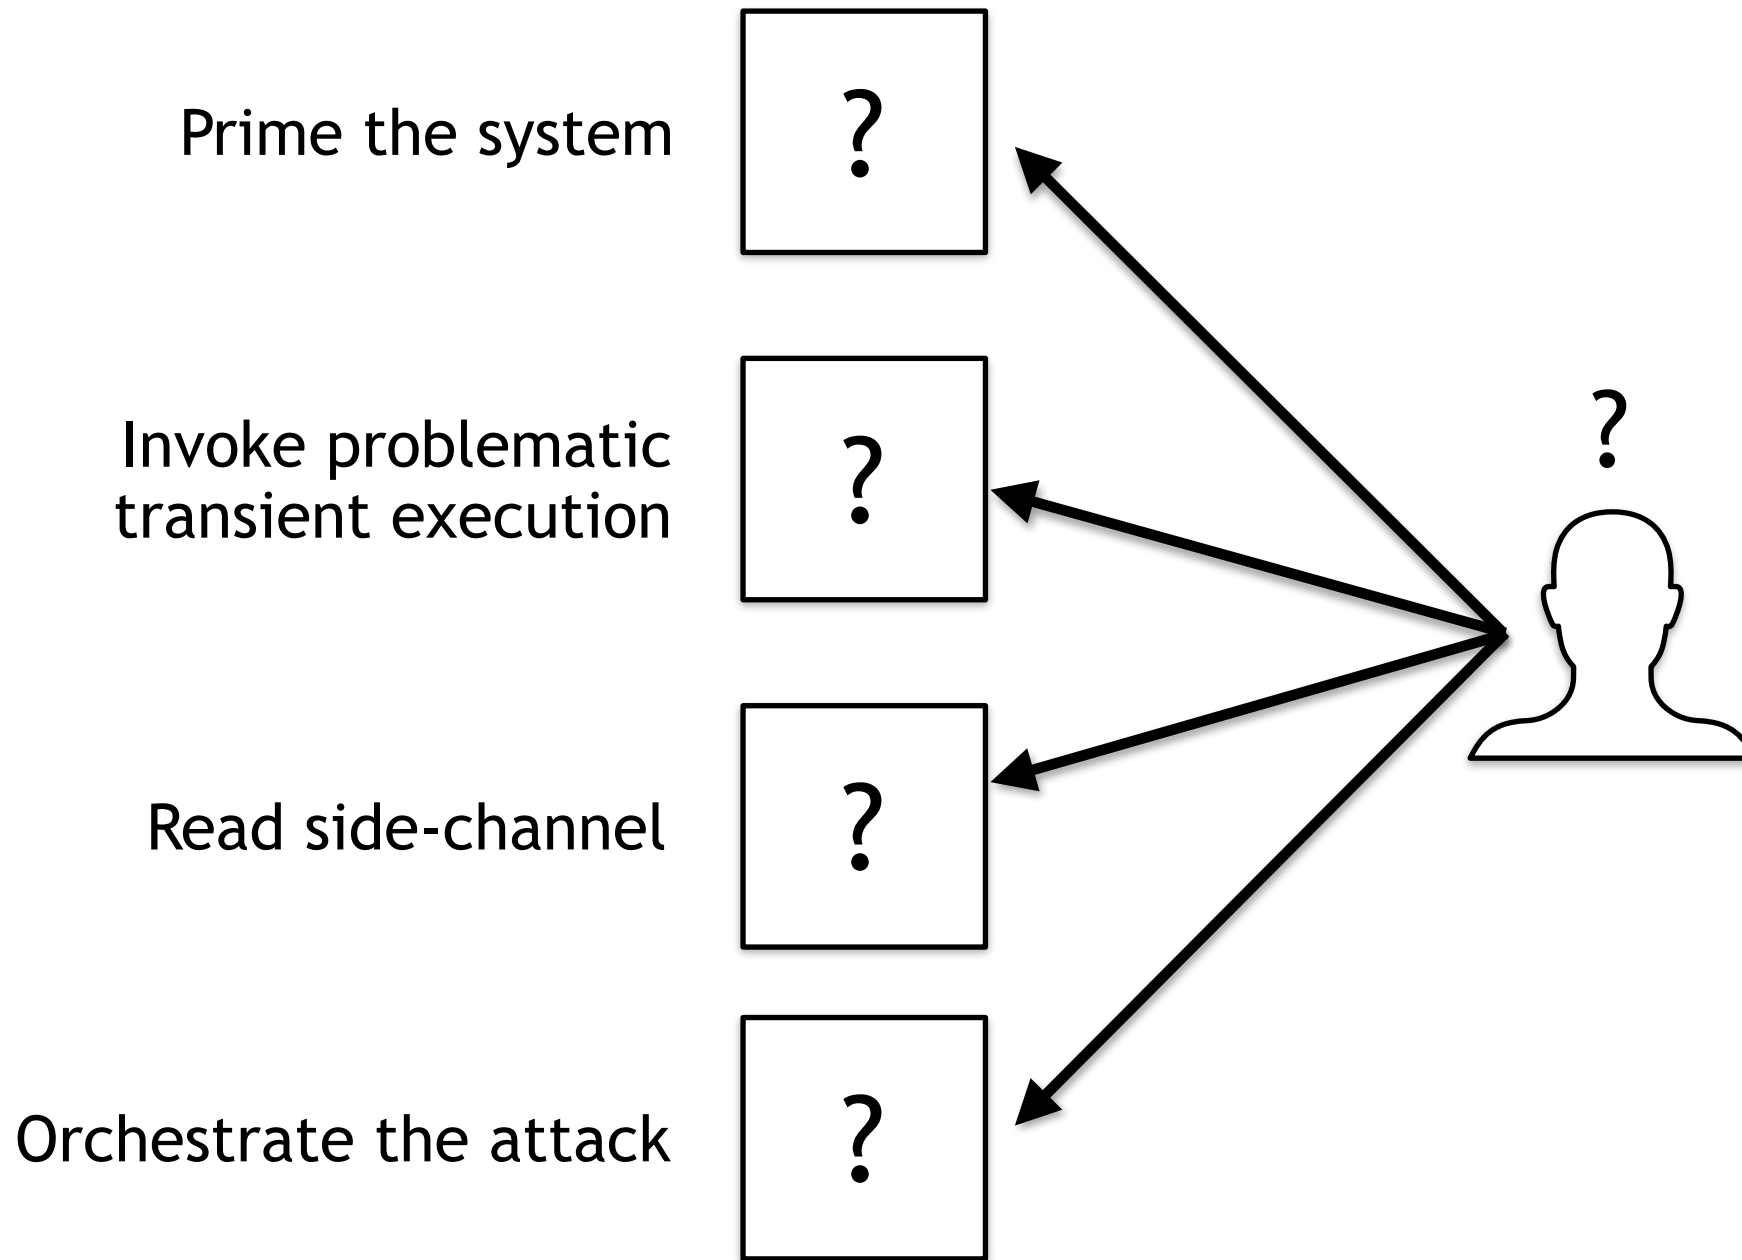


Time-limit, safely schedule workloads
Detect and respond to suspicious workloads
...

<https://cpu.fail/>

<https://meltdownattack.com/>




<https://foreshadowattack.eu/>



Prime the system



Invoke problematic
transient execution



Read side-channel



Orchestrate the attack



Expectations (review)

- Advances in offensive security research
- Continued exposure of Fastly network
- Significant improvements in HW-supplied defenses
- Malicious exploitation

Coverage of approach

- Malicious exploitation
 - → Known attacks are evaluated and prevented
- Significant improvements in HW-supplied defenses
 - → Fit into defined categories of defense
- Continued exposure of Fastly network
 - → Framework can be applied across CPU architectures and execution models
- Advances in offensive security research
 - → Mitigate risk of practical attacks while defenses are rolled out

RSA[®]Conference2020

Conclusion

Defense is best approached holistically

Validate with adversarial analysis

Transient execution attacks are here to stay

RSA[®]Conference2020

Thanks