

Attacking Client Side JIT Compilers

BlackHat 2011

Introduction

Chris Rohlf - Principal Security Consultant

@chrisrohlf
chris@matasano.com

Yan Ivnitskiy - Security Consultant

@yan
yan@matasano.com

<http://www.matasano.com/research>



Overview

- Introduction
- Firefox JIT(s)
- LLVM JIT
- JIT Code Emission Bugs
- JIT Exploitation Primitives
- JIT Hardening
- JIT Engine Comparison
- Our Tools and Techniques

Introduction to JITs

- Interpreters and JIT Engines
 - Parse high level languages
 - Generate bytecode
 - Optimize and compile bytecode to native code
- They are everywhere
 - Browsers
 - Language runtimes (Java, Ruby, C#)

Introduction to JITs

```
10 PRINT "HELLO WORLD"  
20 GOTO 10
```

Developer

"Compiler"

		User
pushq	%rbp	
movq	%rsp,%rbp	
leaq	0x0041(%rip),%rdi	
movl	\$0x0000,%eax	
callq	0x10f36	



Introduction to JITs

```
10 PRINT "HELLO WORLD"  
20 GOTO 10
```

Developer

User

"Compiler"

```
pushq   %rbp  
movq    %rsp,%rbp  
leaq    0x0041(%rip),%rdi  
movl    $0x0000,%eax  
callq   0x10f36
```



Introduction to JITs

```
10 PRINT "HELLO WORLD"  
20 GOTO 10
```

Developer

"Compiler"

JIT

```
pushq    %rbp  
movq     %rsp,%rbp  
leaq     0x0041(%rip),%rdi  
movl     $0x0000,%eax  
callq    0x10f36
```



Introduction to JITs

`a = new Array();`

JSOP_NEWARRAY

```
mov    $0x8963778,%edx
lea    0x50(%ebx),%ecx
mov    %ecx,0x14(%esp)
mov    %esp,%ecx
mov    %ebx,0x1c(%esp)
movl   $0x8962ec5,0x18(%esp)
call   0x8265670
```


Introduction to JITs

- Bytecode / Bitcode / Intermediate Representation (IR)
 - Both trusted and untrusted
 - Expressive and bloated (slower)
 - Simple and slim (faster)
 - Potentially usable to an attacker
 - Overwrite bytecode
 - Trigger interpreter

Introduction to JITs

- Untrusted bytecode
 - Can be delivered from untrusted sources
 - Flash, CLR, LLVM
 - Completely external to the compiler

Introduction to JITs

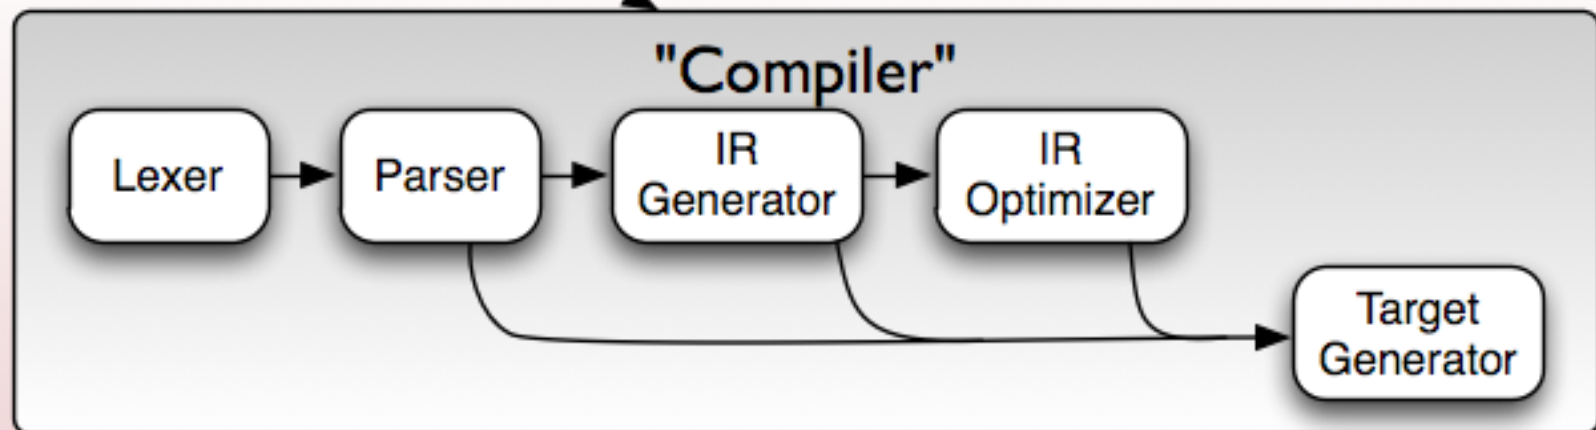
- Trusted bytecode
 - Produced internally by a trusted front end
 - Spidermonkey
 - Still potentially usable to an attacker with control of the process

Introduction to JITs

```
10 PRINT "HELLO WORLD"  
20 GOTO 10
```

Developer

JIT



```
pushq    %rbp  
movq     %rsp,%rbp  
leaq     0x0041(%rip),%rdi  
movl     $0x0000,%eax  
callq    0x10f36
```



Introduction to JITs

- JITs come in a few different designs
 - Method
 - Firefox
 - LLVM
 - V8
 - Tracing
 - Firefox
 - Rubinius

Introduction to JITs

- Tracing
 - Only JITs CPU-intensive code
 - Enables optimizations
 - Types are generally known from tracing

TraceMonkey



- Introduced in Firefox 3.5
- Tracing JIT
- Uses NanoJIT as an assembler
- SpiderMonkey Bytecode → LIR Bytecode
- NanoJIT LIR → Native code



TraceMonkey



- TraceMonkey JITs hot code blocks
 - The recorder traces execution of SpiderMonkey IR
 - Produces trace trees
 - Emits optimized code
- Doesn't handle type changes well

TraceMonkey




- CodeAlloc Class
 - Handles allocating JIT pages that will hold code
- CodeList Class
 - Inline meta-data for tracking the location of JIT pages

Introduction to JITs

- Method
 - JITs entire functions/blocks
 - Usually generates unoptimized code
 - Slow type lookups are usually required

JaegerMonkey



- Introduced in Firefox 4.0
- Method JIT
- Uses the Nitro assembler backend from WebKit
- SpiderMonkey bytecode Native Code
- Uses an Inline Cache for  handling new types

JaegerMonkey



- Inline Caching

- JavaScript is dynamically typed
- It can JIT a generic function that handles multiple types

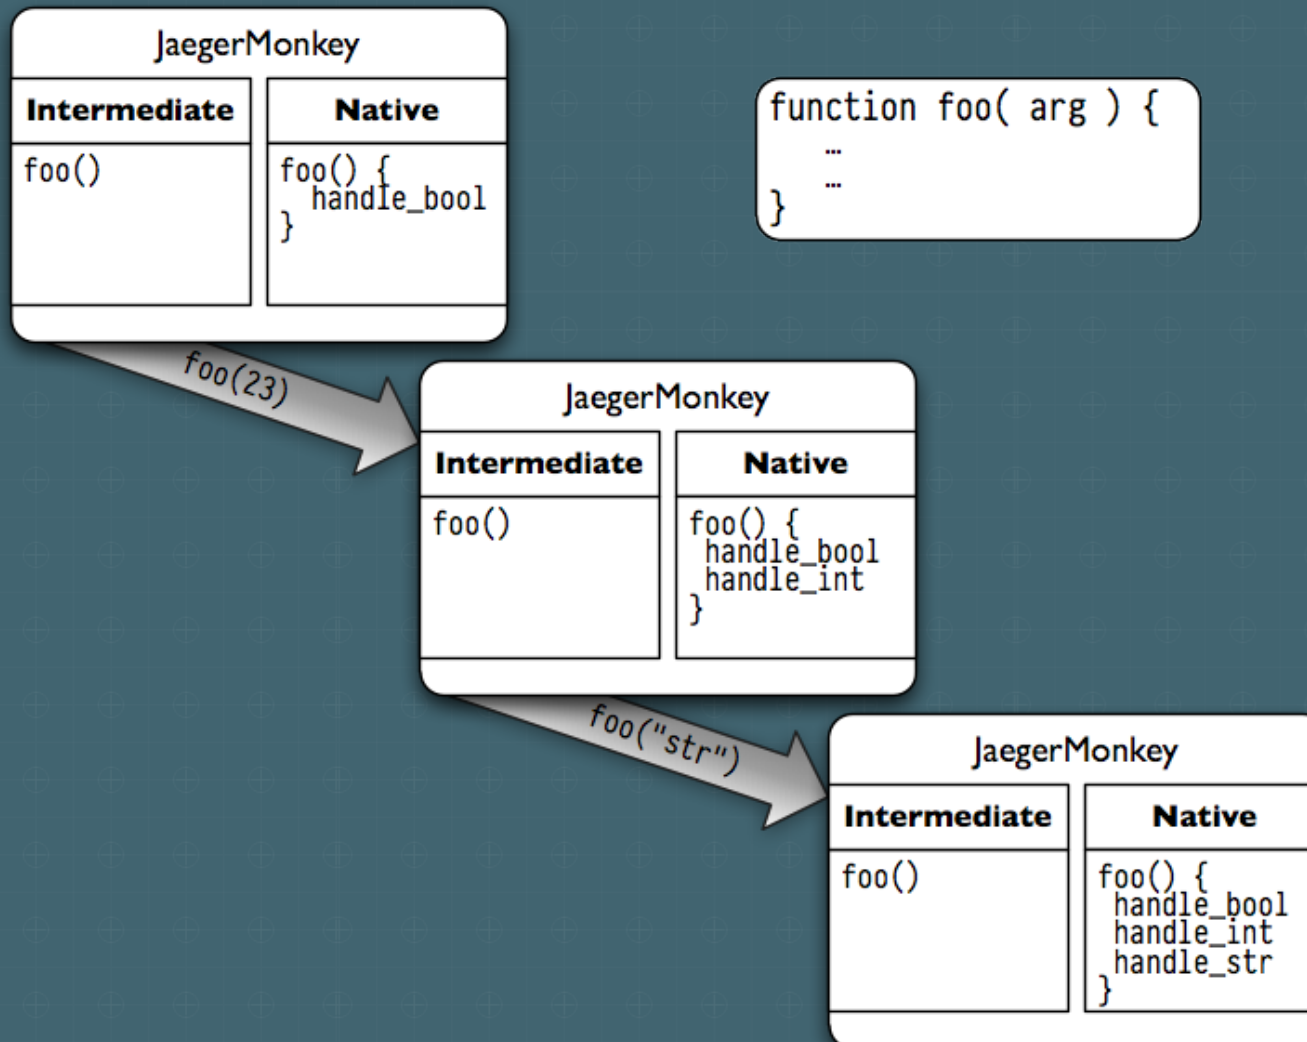
```
function a = blah(var b){  
  for(i=0;i<10;i++) {  
    b += i;  
  }  
}
```

- The JIT emitted code needs to be rewritten for each new type it encounters

```
blah("hello")  
blah([0,1,2,3])
```

- The Inline Cache makes this possible

JaegerMonkey



JaegerMonkey



- Fast paths are native code emitted by the JIT
 - Pure native code emitted by the JIT for predefined operations
- Slow paths are through the execution of bytecode
 - Inline cache hits sometimes have to go back through slow bytecode execution
- Stub calls are calls into C++ code from JIT pages
 - Typically exist to augment a fast path

JaegerMonkey



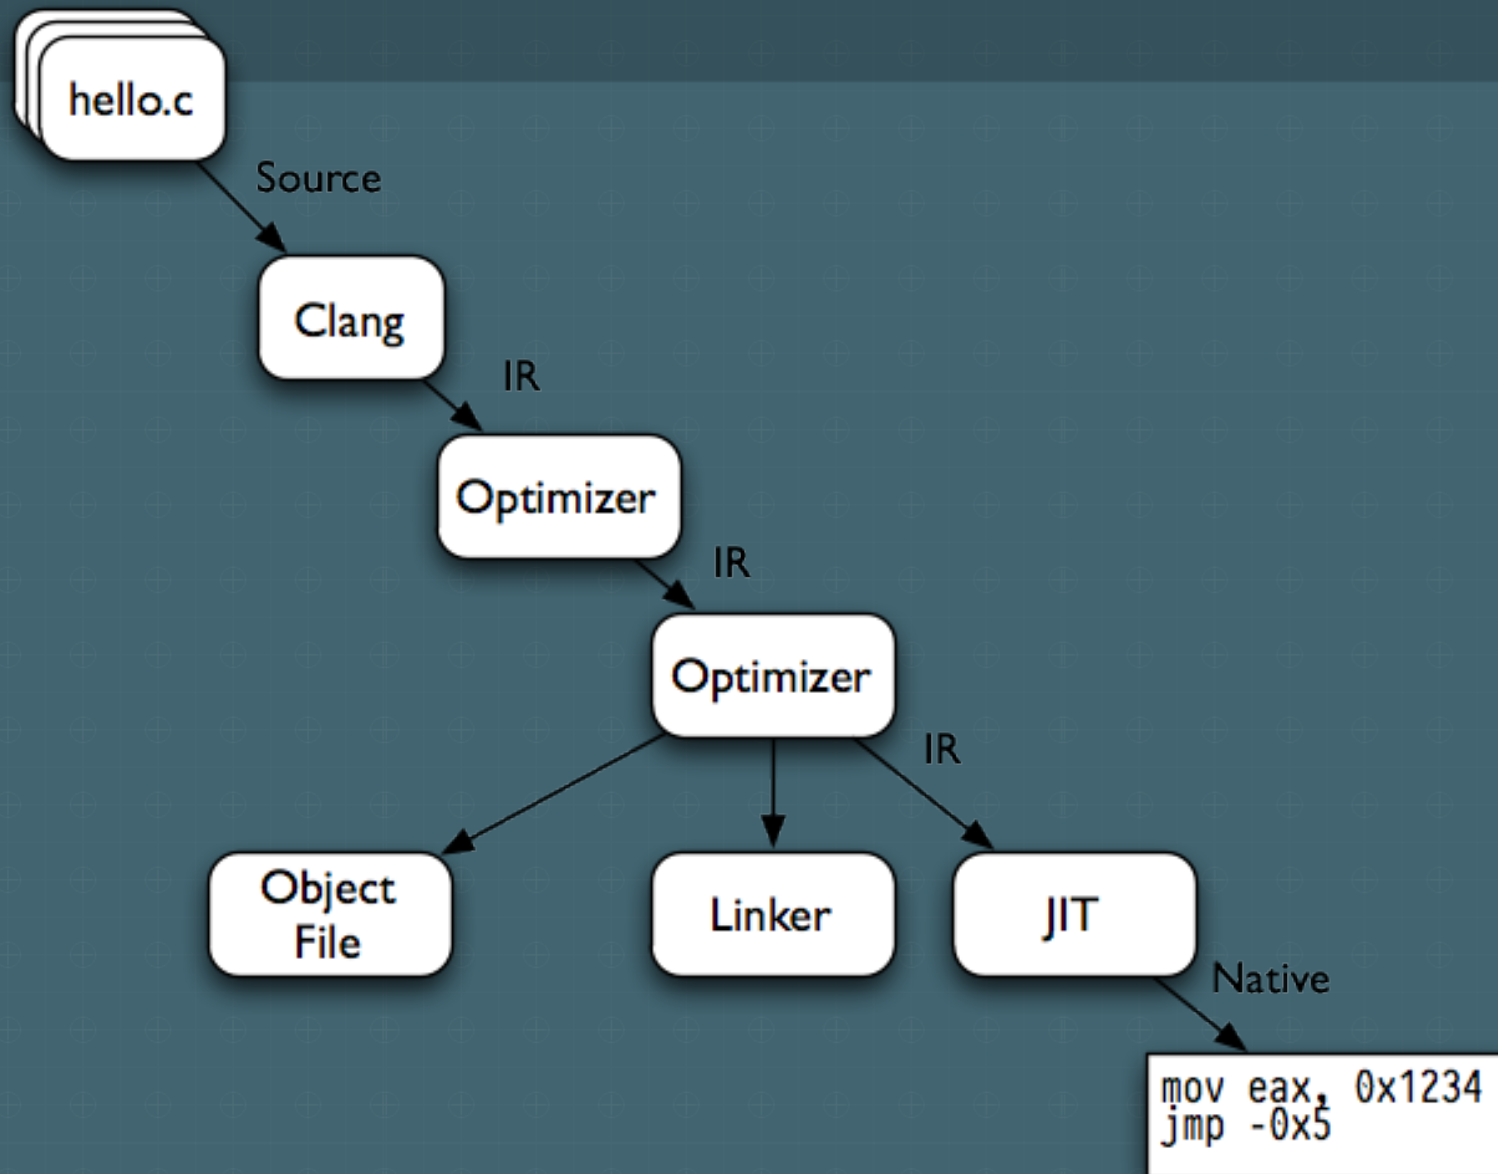
- ExecutableAllocator Class
 - Handles allocating JIT pages to hold code
- ExecutablePool Class
 - Handles managing the larger allocations into `□pools` to hold native code
 - Pools are chosen based on the size of the code that needs to be stored

LLVM

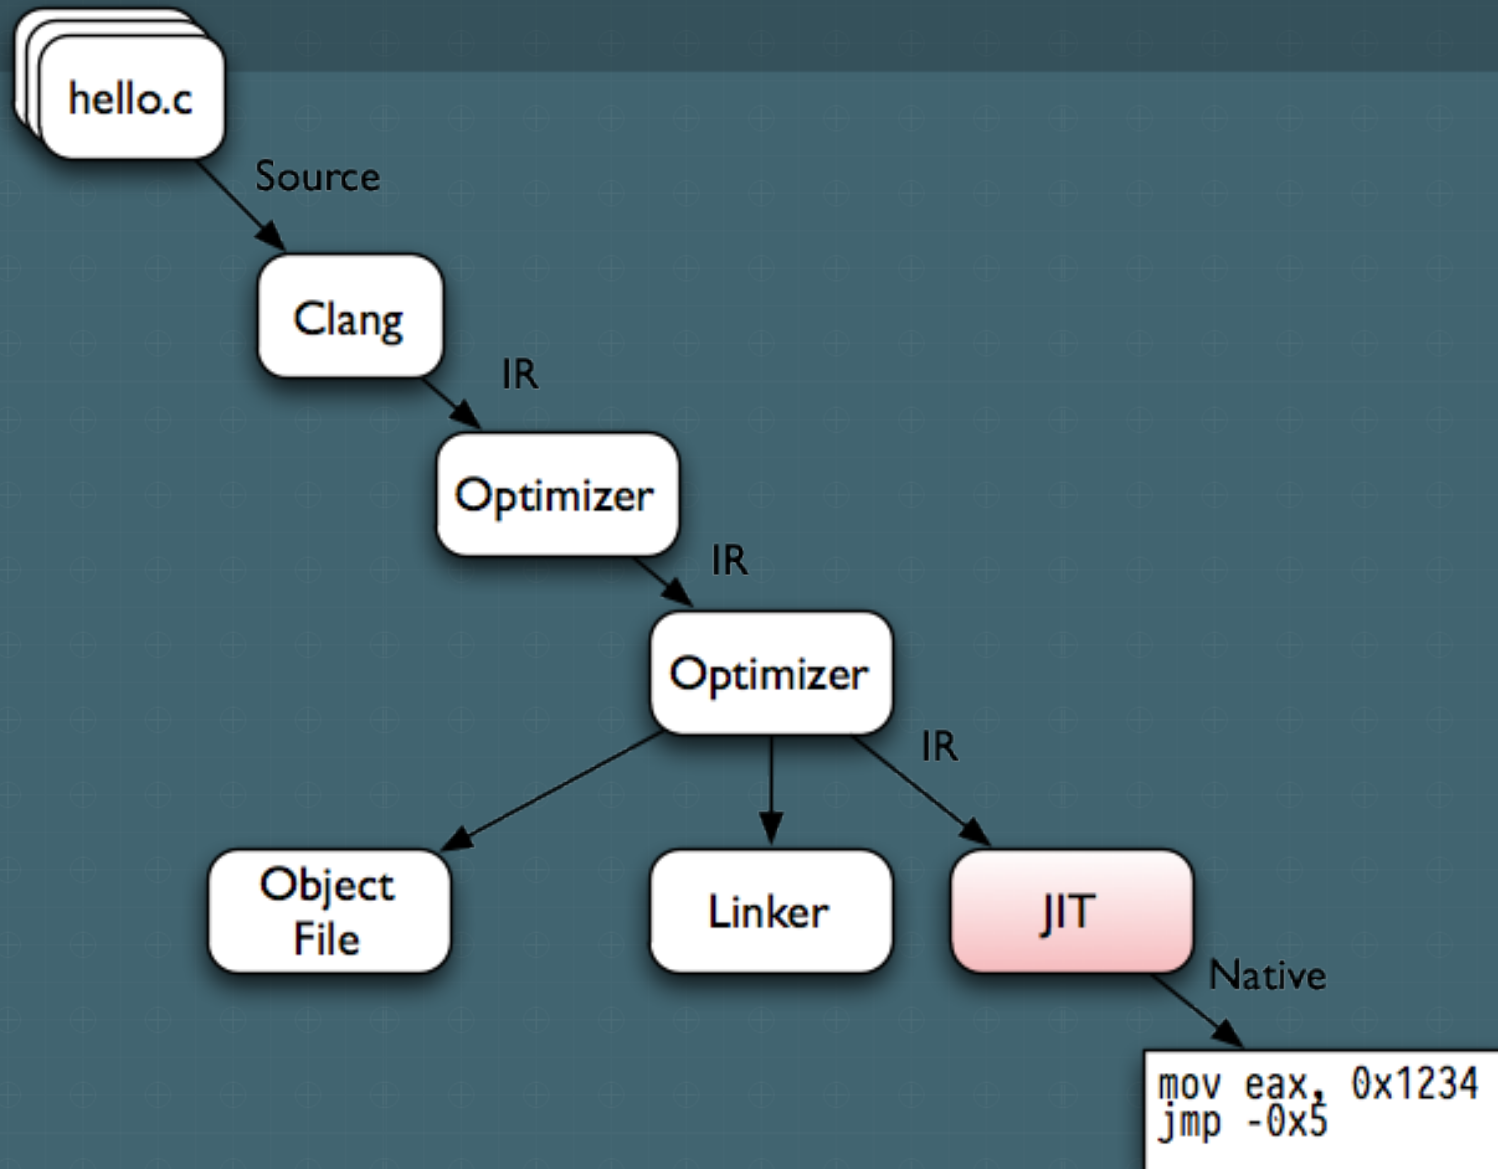


- Originated in UIUC
- A suite of libraries from the start
- Initially used GCC as a front end
- Now supports C, C++ and Objective-C entirely
- Frontends available for
 - Python, Ruby, Haskell, PHP, etc

LLVM



LLVM



LLVM



- Typical progression:
 - I have a project that compiles *something*
 - Need to make it faster
 - Integrate with LLVM!

LLVM Integration

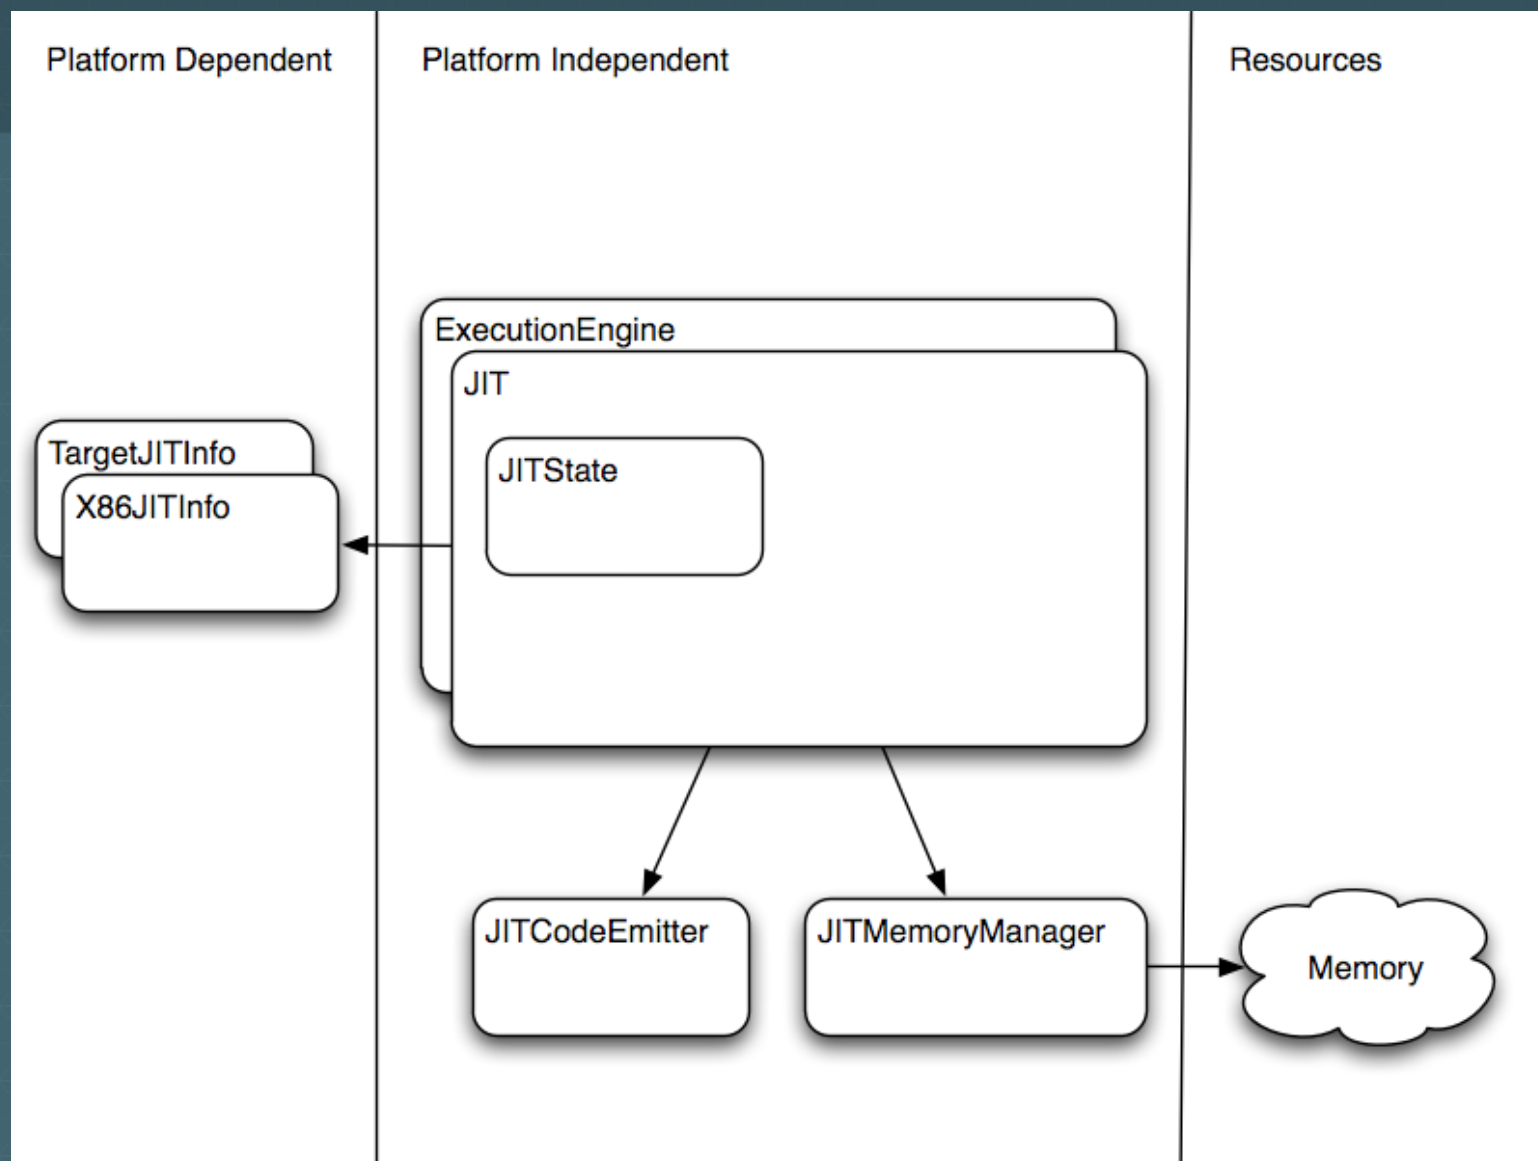
"The LLVM JIT and You"

- Popular integration strategies
- Emit IR directly, create a Module
 - MacRuby, GHC
- Have your own VM instruction set, translate instruction by instruction to LLVM equivalents, then emit
 - Rubinius, ClamAV

LLVM JIT

- Assume a Module is created
- Connect a Module to an ExecutionEngine
- Request a handle to a function, ask the ExecutionEngine for it
- ExecutionEngine emits code for the function, for all data it references and emits stubs for all functions it calls

LLVM JIT



LLVM Attack Scenarios

- Bitcode injection.
 - Portable shellcode!
- Bitcode parser flaws
 - Shouldn't be trusted, but will probably be
- Poor emission
 - Incorrect code being emitted
- Liberal Memory Protections
 - Common vector in general
- Mistranslated instruction set

JITs and Security

- Compiling traditional executables is *typical* for developers
- Code compilation is a trust boundary
 - You've accepted your vendor's code and binary
 - But now your compiling *my untrusted code*



Incorrect Code Emission

- JITs don't always produce perfect code
 - Its impossible
- Compiler bugs are often caught during development and testing
- What can happen when the JIT emits incorrect code?

Incorrect Code Emission

- Java x64 JIT *bug* patched on June 18th, 2011

- Unintended code emission:

`addq %rsp,0xffffffff2b ; shift the stack pointer!`

`popfq ; pop 64 bits from stack+0xffffffff2b ; load the lower 32 bits into RFLAGS`

Intended code emission:

`addq (%rsp),0xffffffff2b ; add 0xffffffff2b to the value at %rsp`

`popfq ; pop 64 bits from stack, load ; the lower 32 bits into RFLAGS`

Incorrect Code Emission

- There are many examples of this
 - Mozilla Bugzilla ID 635295 (Firefox 4.0 Beta)
 - Execution of an invalid branch due to an inline cache that existed for a free'd object
 - MS11-044 Microsoft .Net CLR JIT
 - The JIT produced code that confused an object as NULL or non-NULL
 - This was a great logic bug example!

Incorrect Code Emission

- What usually triggers them?
 - Use after free
 - Integer over/underflows (miscalculation of code paths)
 - Incorrect logic during code emission
- Are incorrect JIT code emissions a new bug class?
 - Depends on the root cause
 - Not for us to decide, but it should be debated by the security research community

JIT Primitives + Traditional Bugs

- JIT engines can be ...
 - the source of vulnerabilities
 - a means to exploit them

Exploitation Primitives

- JITs introduce unique exploitation primitives that would otherwise not be present in an application
 - JIT Spray
 - RWX Page Permissions
 - Reusable code sequences at predictable addresses

JIT Spray

- Publicized by Dion's talk
- Create enough constants to contain native shell code, link together by semantic NOPs
- Transfer execution to mid-instruction, set up a stage 2 and begin executing
- Only 2 JIT engines randomize *VirtualAlloc*
 - V8, IE9
 - Does it matter on 32 bit?

JIT Spray

- JIT Spray in Firefox through JaegerMonkey

- var constants = [0x12424242, 0x23434343, 0x34444444, 0x45454545, 0x56464646, 0x67474747, 0x78484848, /test/]
- 0x40a05e: call 0x82d1820 NewInitArray ; **create an array**
0x40a063: mov %eax,%edi ; **\$edi holds returned array object**
0x40a065: mov 0x24(%edi),%edi ; **load obj->slots in to \$edi**
0x40a068: movl \$0xffff0001,0x4(%edi) ; **JSVAL_TYPE_INT32 into object->slots[1]**
0x40a06f: movl \$0x12424242,(%edi) ; **1st constant into object->slots[0]**
0x40a075: mov %eax,%edi
0x40a077: mov 0x24(%edi),%edi
0x40a07a: movl \$0xffff0001,0xc(%edi)
0x40a081: movl \$0x23434343,0x8(%edi) ; **2nd constant**
0x40a088: mov %eax,%edi
0x40a08a: mov 0x24(%edi),%edi
0x40a08d: movl \$0xffff0001,0x14(%edi)
0x40a094: movl \$0x34444444,0x10(%edi) ; **3rd constant**
0x40a09b: mov %eax,%edi
0x40a09d: mov 0x24(%edi),%edi
0x40a0a0: movl \$0xffff0001,0x1c(%edi)
0x40a0a7: movl \$0x45454545,0x18(%edi) ; **4th constant**

Memory Protections

- Nearly all JITs we surveyed produce RWX pages
 - Breaks DEP / W^X
 - Breaks assumption behind mirror pages
 - Knowledge of both RW / RX pages not required
 - Blind Execution
 - Overwrite RWX JIT page contents
 - Trigger the original JIT'd script
- This isn't going away for Inline Cache designs

Memory Protections

- RWX pages can be repurposed for different things
 - Array index read/write
 - Point into JIT page
 - Write raw shell code, trigger JavaScript
 - Read pointers to a DLL that may be mapped beyond your reach
 - Overflows
 - Heap overflow in adjacent RW page
 - ROP
 - No need to find that VirtualAlloc stub

gaJITs

- ROP Gadgets are small sequences of code found in an existing DLL or .text
 - Combine them to get arbitrary code execution
- gaJITs: Predictable instructions on JIT pages at predictable offsets
- JIT's produce lots of native code
 - You aren't constrained to just one library mapping
 - Does not require controllable constants like JIT Spray

gaJITs

- Finding usable gaJITs depends on the JIT design
 - *ret* or branch based control flow?
 - inline caching
 - (in)frequent calls to C++ stubs
- How does script function A get turned into native code B where native code B contains gaJIT X
 - Requires the right source code inputs to generate them
 - Requires a specific gadget-finding tool

JIT Feng Shui

- Our version of Heap Feng Shui... except for JITs
 - Heap Feng Shui
 - Alex Sotirov 2007
 - Influence the heap layout via JavaScript
 - JIT Feng Shui
 - Untrusted input influences JIT output
 - Specific inputs creates predictable code patterns
- We could have called it jiuJITSu...

JIT Feng Shui

- Controlling register contents with a TraceMonkey gaJIT

gaJIT at offset 0x9e18 (10 matches)

pop esi ; pop edi ; pop ebx ; pop ebp ; ret

- LLVM
 - Portable shellcode!

JIT Feng Shui + gaJITs

- Circumvents constant masking
 - Defeated by NOP padding
 - Defeated by allocation restrictions
- Difficult and noisy
 - Requires a JIT spray to map enough pages
- Not researched on other JITs / architectures yet...

JIT Protections

- The OS provides some basic protections to the process
 - (ASLR) Address Space Layout Randomization
 - (DEP) Data Execution Prevention
 - JITs can negate these by design
- JIT engines have no control over their input
 - ... but completely control their output

Emission Randomization

- Emitting code at runtime allows for randomizing addresses
 - High 22 bits (54 on 64 bit [fact check]) get randomized due to ASLR, or opting into a random address
 - Low ~3 bits get randomized due to NOP insertion
 - The middle bits get optimized within code region


Emission Randomization

- Memory for emission is allocated via mmap or VirtualAlloc
 - VirtualAlloc is *not* randomized by default
 - You can request the address you want mapped
 - V8 and IE9 do this
 - mmap on Linux randomizes anonymous mappings
- Extend ASLR to compiler-allocated memory

Randomization




Randomization



Randomization

- Intra-page offsets (bottom 10 bits) are still predictable
- Since you're emitting code, you can shift each function emitted by inserting NOPs


Randomization



Randomization

- Function emission is still predictable
- If you're batching the functions you're emitting, you can shuffle the order at which they are produced

Randomization



Guard Pages

- Firefox 5.0 adjacent heap and JIT pages

02808000-0280c000 rw-p Read/Write heap memory

0280c000-0281c000 rwxp Read/Write/Execute JIT page

- If an overflow occurs in the first RW heap mapping an attacker can write native code into the RWX page

- Trigger JavaScript to get code execution

- Blind execution

- No need to gain EIP or overwrite a function pointer

02808000-0280c000 rw-p Read/Write heap memory 0280c000-0281c000 r--p Read Only memory

0281c000-0282c000 rwxp Read/Write/Execute JIT page

Constant Folding

- 4 byte constants allows room to insert instructions on x86
- Chained 4 byte chunks allows for a stage 1 payload
- Solution: Fold large constants into 2-byte maximum constants and reassemble at runtime
- **Problem:** If the instructions are predictable an attacker can bypass this by injecting the right constants
- V8 did this for awhile, now they use constant blinding

Constant Blinding

- XOR all untrusted immediate values by a secret cookie
- Generate a random value at startup
 - untrusted immediate \wedge secret cookie
- Emit code that XORs the value at runtime

```
xor eax, 0x00112233
```



```
mov eax, 0x84521310  
xor eax, 0x84433123
```

Allocation Restrictions

- JIT Spray requires mapping a lot of memory
- Capping the number of pages helps mitigate this attack
- For language runtimes some info about code can be known ahead of time
 - code size
 - libraries used
- Unfortunately this protection mechanism makes more sense for browsers than language runtimes

JIT Comparison

	V8	IE9	Jaeger Monkey	Trace Monkey	LLVM	JVM	Flash/Tamarin
Secure Page Permissions	N	Y	N	N	N	N	N
Guard Pages	N	N	N	N	N	N	N
Page Randomization	Y	Y	N	N	N	N	N
Constant Folding	N	N	N	N	N	N	N
Constant Blinding	Y	Y	N	N	N	N	N
Allocation Restrictions	Y	Y	N	N	N	N	N
Random NOP Insertion	Y	Y	N	Y	N	N	Y
Random Code Base Offset	Y	Y	N	Y	N	N	Y

jitter

- jitter is our toolchain for
 - Tracing JIT code emission
 - Tracking JIT memory permissions
 - JIT Fuzzer coverage
 - Searching for gaJITs
- Implemented as a set of Nerve scripts
 - Uses the ragweed debugging framework
 - We also wrote a native Java JIT hook

jitter

- Support for LLVM and Firefox JITs
 - Nerve breakpoint files for specific JIT hook points
 - Interact with the process at each breakpoint with Ruby
 - Extract arguments, code, instructions...
- Generic script for tracking JIT page allocations
 - Just needs a list of call sites
 - Can be used to start support of new JIT engines
- gaJIT finder is built in
 - Receives an array of JIT pages
 - Outputs locations of repeated gaJITs
 - Easily repurposed for other ROP tools

fuzzer(s)

- Fuzzing JIT engines is difficult
 - Testcases must have valid syntax
 - Multiple components before you hit the JIT
- Rubinius Fuzzer (LLVM JIT)
- JavaScript Grammar fuzzer (Firefox JITs)
 - Targets the JIT and interpreter only
 - DOM bugs are boring

rubyfuzz

- Ruby fuzzer for targeting Rubinius
 - Generated Ruby code from a subset of Ruby grammar
 - Avoided Rubinius VM to target other Ruby implementations.
 - MacRuby, JRuby, YARV, MRI, etc
- Fuzzer driver framework Hoke
 - Generated test cases external to the Ruby implementation

rubyfuzz

- Modeled Ruby grammar as Ruby objects
 - Terminals -> arrays
 - Non-terminals -> generators
- Permuted method invocations, block definitions, block invocations, and other Ruby constructs.
- Seeded with common Ruby constructs

fuzzer(s)

- JavaScript Grammar fuzzer for Firefox JITs
- Describe JavaScript in flat text files
 - types, methods, properties, keywords, operators
- Parse text files and serialize into Ruby OpenStruct
- Iterate over the grammar
 - Fast Paths
 - Inline Caches
 - C++ Stubs
- Hundreds of millions of iterations through ./js

fuzzer(s)

- A note on fuzzing for info leaks
 - Fuzzing should be fast
 - Instrumentation is slow
- Differential fuzzing for info leaks
 - Two JavaScript implementations
 - d8 (V8) / js (Mozilla)
 - Feed them the same testcase
 - Record the output
 - What is the expected output type/value?
- Can be generalized to multiple implementations of any language

A bug our fuzzer found

- Our fuzzer found a critical bug in SpiderMonkey

```
a = new Array();  
a.length = 4294967240;  
b = function bf(prev, current, index, array) {  
    document.write(current);  
}  
a.reduceRight(b, 1, 2, 3);
```

- Info Leak: read arbitrary data from *current*
- Code Execution: call a method on *current* object

Demo?

Questions

?

Please complete the Speaker Feedback Surveys.
This will help speakers to improve and for Black
Hat to make better decisions regarding content
and presenters for future events.

