

BLUEHAT

IL 2022

An Analysis of Speculative Type Confusion Vulnerabilities in the Wild

 Adam Morrison 

About the speaker

Associate professor, CS, Tel Aviv University.

Started career doing vulnerability research.




About the speaker

Associate professor, CS, Tel Aviv University.

Started career doing vulnerability research.

Now: Software/hardware interactions.

- From microarchitecture to OS.



About this talk

Joint work with **Ofek Kirzner**.

Full details available in academic paper:

An Analysis of Speculative Type Confusion Vulnerabilities in the Wild

Ofek Kirzner Adam Morrison
Tel Aviv University

Spectre attacks



The New York Times

Researchers Discover Two Major
Flaws in the World's Computers

TECH

Businesses Rush to Contain
Fallout From Major Chip
Flaws

Software patches to plug holes could slow computers, experts say

THE WALL STREET JOURNAL.
Businesses Rush to Contain Fallout From Major Chip

The Washington Post
Democracy Dies in Darkness

Technology

Huge security flaws revealed — and tech
companies can barely keep up

Executive summary



Spectre v1 has no hardware fix; software mitigations required.


Popularized as a bounds-check bypass attack.

We show **other Spectre v1 vectors** in real code.

⇒ **Mitigating Spectre v1 in software requires some rethinking.**

Spectre attacks

Spectre attacks




Goal: Leak data
from the victim
address space

exploit speculative execution

leak secret data

gadget

Spectre variant 1: Bounds Check Bypass



Goal: Leak data
from the victim
address space

exploit branch prediction

leak secret data

gadget

Spectre variant 1: Bounds Check Bypass



Goal: Leak data
from the victim
address space



```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass



Goal: Leak data
from the victim
address space

foo(**x**)

```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass



speculation starts



Goal: Leak data
from the victim
address space

`foo(&secret-array1)`

```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass



speculation starts

speculative access



Goal: Leak data
from the victim
address space

foo(&secret-array1)

array[x]=&secret


```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass



speculation starts

speculative access

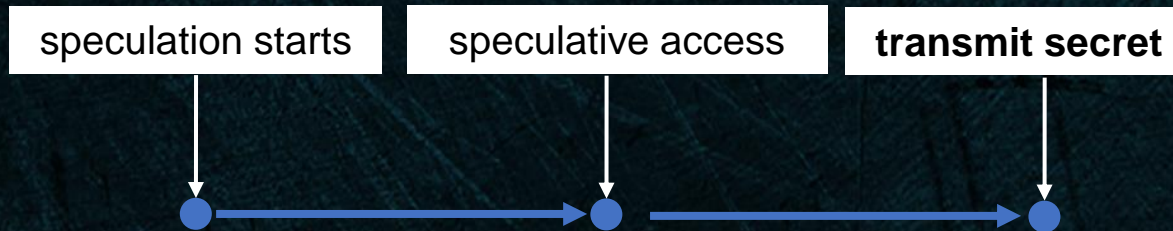


Goal: Leak data from the victim address space

`foo(&secret-array1)`

```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```



Spectre variant 1: Bounds Check Bypass



Goal: Leak data
from the victim
address space

foo(&secret-array1)

Cache state
encodes y




L1 cache

$y=7$

$y=17$


```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass




Goal: Leak data from the victim address space

`foo(&secret-array1)`



L1 cache


```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
          
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass



Goal: Leak data
from the victim
address space

**SIDE
CHANNEL**



L1 cache

```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```



Spectre variant 1: Bounds Check Bypass



SIDE CHANNEL: secret-dependent microarchitectural state.



Goal: Leak data from the victim address space



L1 cache

```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre variant 1: Bounds Check Bypass




SIDE CHANNEL: secret-dependent microarchitectural state.

Spectre enabled **arbitrary data** to be transmitted via a side channel.



Goal: Leak data from the victim address space



L1 cache


```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre v1 is a threat to OS kernels



Attacker:
unprivileged
user

Attacker
Victim: OS
kernel



Goal: Leak data
from the victim
address space

exploit system calls

read any physical memory

```
void foo(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre v1 mitigation

No hardware fix for Spectre v1



While Variant 1 will continue to be addressed via software mitigations,



For all flavors of variant 1, the AMD mitigation recommendation is software only

Spectre v1 mitigation in the Linux kernel

No hardware fix for Spectre v1

⇒ Linux has a special API to ensure bounds checks are respected under speculation

```
void function_called_from_syscall(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array1[x];  
        z = array2[y * 4096];  
    }  
    // ...  
}
```



```
void function_called_from_syscall(long x) {  
    // ...  
    if (x < array1_len) {  
        y = array_index_nospec(array1[x], array1_len);  
        z = array2[y * 4096];  
    }  
    // ...  
}
```


Spectre v1: not only a bounds check bypass

Quoting from the Spectre paper [Kocher et al., 2019]:


Variant 1: Exploiting Conditional Branches. In this variant of Spectre attacks, the attacker mistrains the CPU's branch predictor into mispredicting the direction of a branch, causing the CPU to temporarily violate program semantics by executing code that would not have been executed otherwise.

SPECULATIVE TYPE CONFUSION

Misspeculation makes the victim execute with some variables holding values of the wrong type, and thereby leak memory content


Speculative type confusion: example

Misprediction:
Type confusion



```
void syscall_helper(struct Base* obj) {  
    if (obj->type == TYPE1) {  
        struct Type1* o = (struct Type1*) obj;  
        leak(o->value);  
    }  
    if (obj->type == TYPE2) {  
        ...  
    }  
}
```

```
struct Base {  
    enum Type type;  
};  
  
struct Type1 {  
    struct Base base;  
    ...  
    uint32_t value;  
};
```




Our results

Observation: speculative type confusion may be much more prevalent than previously hypothesized.

We analyzed the Linux kernel, looking for speculative type confusion.

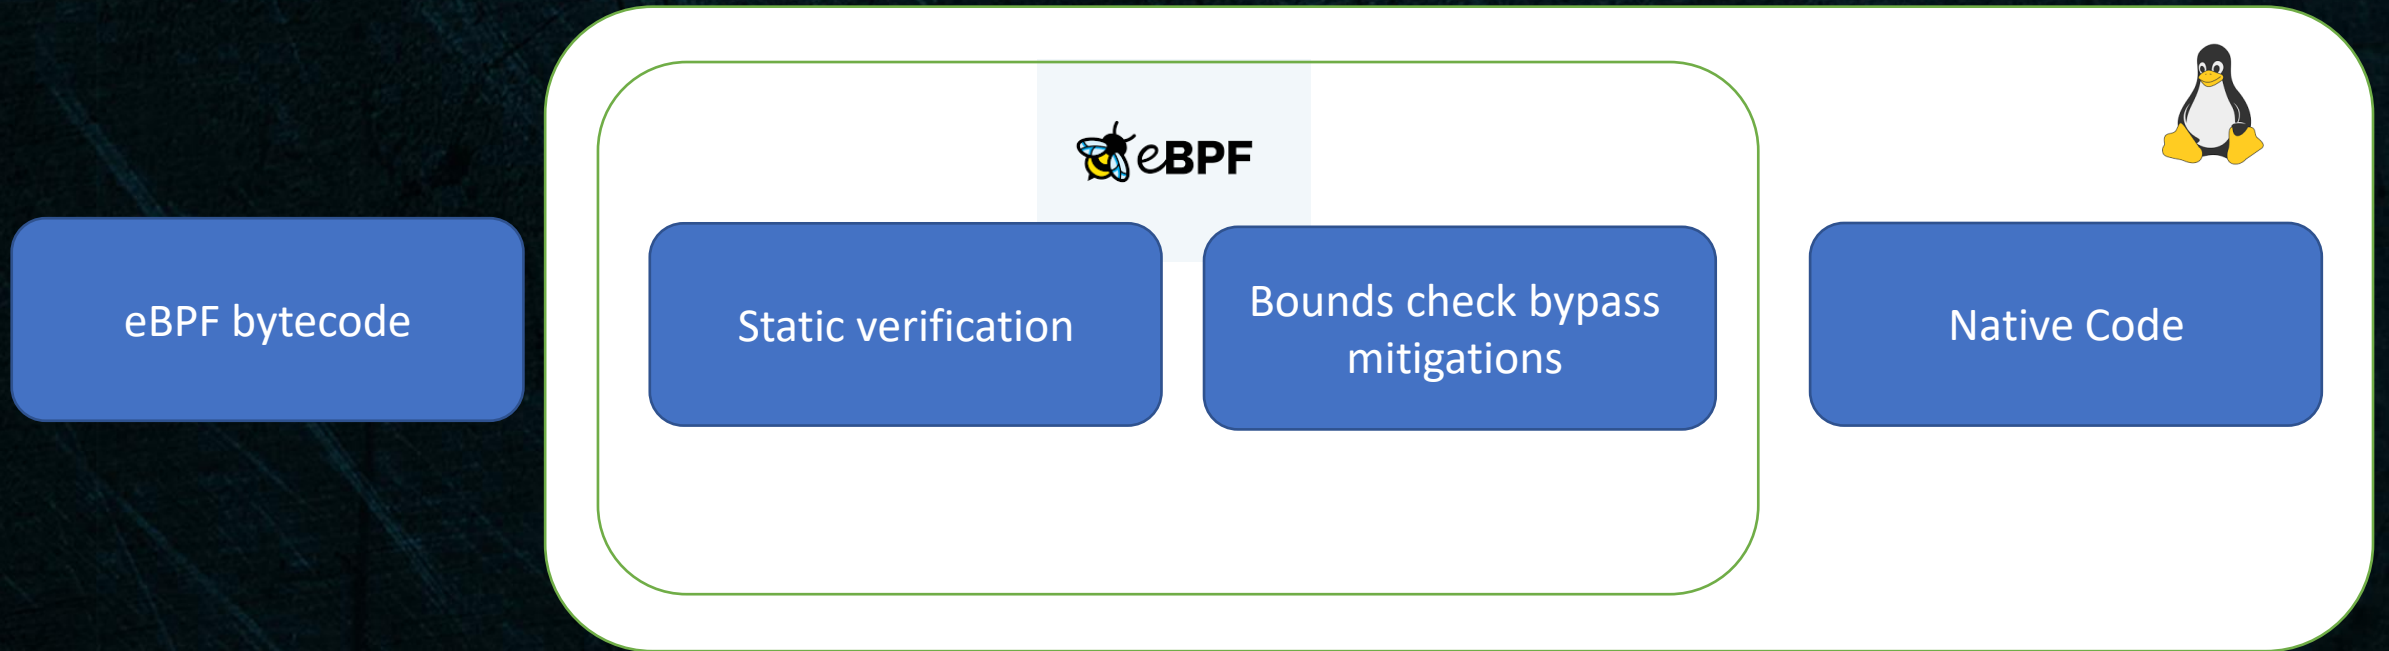
Found new types of speculative type confusion.



eBPF: Speculative Type Confusion

eBPF

Linux subsystem, enabling user-defined programs in kernel.



eBPF verifier vulnerability

```
r0 = *(u64 *)(r0)
A: if r0 != 0x0 goto B
|   r6 = r9
B: if r0 != 0x1 goto D
|   r9 = *(u8 *)(r6)
C:   r1 = M[(r9 & 1) * 512]
D:...
```

Flows considered by eBPF verifier

$r0 == 0$

```
r0 = *(u64 *)(r0)
A: if r0 != 0x0 goto B
|   r6 = r9
B: if r0 != 0x1 goto D
|   r9 = *(u8 *)(r6)
C:   r1 = M[(r9 & 1) * 512]
D:...
```

$r0 == 1$

```
r0 = *(u64 *)(r0)
A: if r0 != 0x0 goto B
|   r6 = r9
B: if r0 != 0x1 goto D
|   r9 = *(u8 *)(r6)
C:   r1 = M[(r9 & 1) * 512]
D:...
```

otherwise

```
r0 = *(u64 *)(r0)
A: if r0 != 0x0 goto B
|   r6 = r9
B: if r0 != 0x1 goto D
|   r9 = *(u8 *)(r6)
C:   r1 = M[(r9 & 1) * 512]
D:...
```


eBPF verifier vulnerability

Speculative flows are not verified

```
r0 = *(u64 *) (r0)
A: if r0 != 0x0 goto B
|   |   r6 = r9
B: if r0 != 0x1 goto D
|   |   r9 = *(u8 *) (r6)
C:   r1 = M[(r9 & 1) * 512]
D:...
```


Predicted “false”

Predicted “false”

```
// r0 = ptr to an array entry (verified != NULL)
// r6 = ptr to stack slot (verified != NULL)
// r9 = scalar value controlled by attacker
```

```
if r0 == 0x0 and r0 == 0x1
  r6 = r9
  r9 = *(u8 *) (r6)
  r1 = M[(r9 & 1) * 512]
```

read
arbitrary
memory




Training mutually exclusive branches



Can both be false

```
A: if r0 != 0x0 goto B
    r6 = r9
B: if r0 != 0x0 goto
    r9 = *(u8 *)
```

Shadow gadget



Mutually exclusive

```
A: if r0 != 0x0 goto B
    r6 = r9
    r0 != 0x1 goto D
    r9 = *(u8 *) (r6)
```

Unprivileged process can read arbitrary memory addresses at a rate of ~6.5 KB/sec

Compiler Introduced Speculative Type Confusion

Compilers might create speculative type confusion

Innocent looking code is compiled in a way that introduces vulnerability

Compiler reasoning:
Branches are mutually exclusive

(trusted) ptr argument held in x86 register %rsi

attacker-controlled

```
void syscall_helper(cmd_t* cmd, char* ptr, long x)
{
    cmd_t c = *cmd;
    if (c == CMD_A)
    {
        ...
    }
    if (c == CMD_B)
    {
        y = *ptr; // y = *%rsi
        z = array[y * 4096];
    }
    // ...
}
```

code during which x moves to %rsi

Can we find it in the wild?

Binary level analysis of Linux.

Focused on system calls, which have well-defined user-controlled interface.

Leakage mechanism is out of scope: aiming at finding speculative attacker-controlled memory dereference.

compiler	flags	# vulnerable syscalls
GCC 9.3.0	-Os	20
GCC 9.3.0	-O3	2
GCC 5.8.2	-Os	0
GCC 5.8.2	-O3	0

Reusing registers for a function call

```
syscall(foo_t* uptr) {  
    foo_t kfoo;  
    // some code  
    if (uptr)  
        copy_from_user(&kfoo,  
                        uptr,  
                        ...);  
    f(uptr ? &kfoo : NULL);  
    // rest of code  
}
```

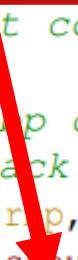
rdi != 0
⇒ rdi = stack var

rdi == 0
⇒ reuse rdi

Reusing registers for a function call

```
syscall(foo_t* uptr) {  
    foo_t kfoo;  
    // some code  
    if (uptr)  
        copy_from_user(&kfoo,  
                        uptr,  
                        ...);  
    f(uptr ? &kfoo : NULL);  
    // rest of code  
}
```

```
# args: uptr in %rdi  
...  
testq %rdi, %rdi  
je L # jump if %rdi == 0  
# set copy_from_user args  
...  
# %rip contains addr of  
# stack buffer  
mov %rip, %rdi  
call copy_from_user  
L:callq f
```



rdi != 0

⇒ rdi = stack var

rdi == 0

⇒ **reuse rdi**

Stack slot reuse

```
long keyctl_instantiate_key_common(key_serial_t id,  
                                struct iov_iter *from,  
                                key_serial_t ringid) {  
    struct key *dest_keyring;  
    // ... code ...  
    ret = get_instantiation_keyring(ringid, rka, &dest_keyring);  
    if (ret < 0)  
        goto error2;  
    ret = key_instantiate_and_link(rka->target_key, payload,  
                                plen, dest_keyring,  
                                instkey);  
    // above call dereferences dest_keyring  
}
```

Out param

Allocated with 1-byte opcode rather than by subtraction (4-bytes), by chance contains user-controlled value

```
# %rcx is a live register from caller  
push %rcx  
# ... code ...  
lea 0x18(%r14),%rsi # rka argument  
mov %rsp,%rdx # &dest_keyring argument  
mov %r15d,%edi # ringid argument  
callq get_instantiation_keyring # returns error  
test %rax,%rax # if (ret < 0)  
mov %rax,%rbx  
js error2 # mispredict no error  
...  
mov (%rsp),%rcx # dest_keyring argument  
# dest_keyring could be old %rcx if not  
# overwritten in get_instantiation_keyring()  
callq key_instantiate_and_link
```

Loads value from stack

Overwrites out param in any case

Small change → exploitable
Hard to reason about

Speculative Polymorphic Type Confusion

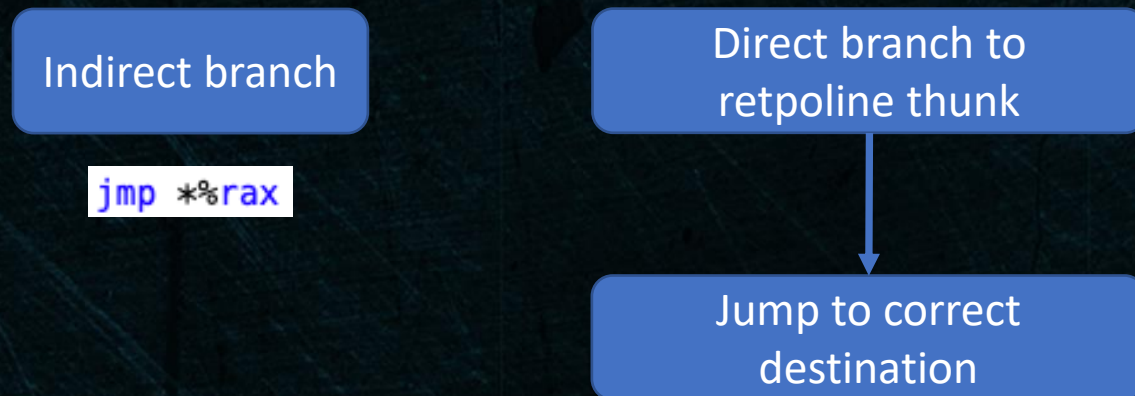
Spectre v2 mitigations

Spectre v2 exploits misprediction of indirect branch target addresses

Retpolines: block indirect branch prediction

Optimization: restrict speculation to **valid targets** [Linux, Amit et al., 2019]

Might create speculative type confusion vulnerabilities



```
# %rax = branch target
cmp $0xFFFFFFFF, %rax # target1?
jz $0xFFFFFFFF
cmp $0xFFFFFFFF, %rax # target2?
jz $0xFFFFFFFF
...
jmp ${fallback} # jmp to retpoline thunk
```


Speculative polymorphic type confusion

```
struct Common { void (*foo) (void*); };  
struct A { struct Common common; char* ptr; };  
struct B { struct Common common; long user_controlled_scalar; };
```

```
void some_code_path(struct Common* common) {  
    /* ... */  
    common->foo(common);  
}
```

```
void foo_A(struct Common* common) {  
    char x = *((struct A*) common)->ptr;  
    leak(x);  
}
```

foo_B()

Speculative polymorphic type confusion

```
struct Common { void (*foo) (void*); };  
struct A { struct Common common; char* ptr; };  
struct B { struct Common common; long user_controlled_scalar; };
```

```
void some_code_path(struct Common* common) {  
    /* ... */  
    common->foo(common);  
}
```

```
# %rax = branch target  
cmp $0xFFFFFFFF, %rax # target1?  
jz $0xFFFFFFFF  
cmp $0xFFFFFFFF, %rax # target2?  
jz $0xFFFFFFFF  
...  
jmp ${fallback} # jmp to retpoline thunk
```

B → user_controlled_scalar

```
void foo_A(struct Common* common) {  
    char x = *((struct A*) common)->ptr;  
    leak(x);  
}
```

misprediction

foo_B()

Analysis

Analysis

- Linux code analysis - looking at ways in which polymorphism can lead to speculative type confusion

Analysis

Analysis

- Linux code analysis - looking at ways in which polymorphism can lead to speculative type confusion

Results

- Flagged potentially vulnerable: 1000s
- "Array indexing" instances: 100s
- All – not exploitable(?) E.g., limited user value control

Analysis

Analysis

- Linux code analysis - looking at ways in which polymorphism can lead to speculative type confusion

Results

- Thousands - flagged potentially vulnerable
- Hundreds - "array indexing" instances
- All - limited speculation window or limited control on user value

Conclusion

- **Were a conditional branch-based mitigation used instead of retpolines, the kernel's security would be on shaky ground**

Limitations / future work

Our analyses are PoC.

Not exhaustive, have false positives, and false negatives.

Much room for improvement!

⇒ More vulnerabilities.

Spectre v1 software mitigation

~~Spot (manual, Linux style)~~




Complete (compiler-based)
E.g.: LLVM SLH




Hardware support might be
required?

SLH overhead on Linux syscalls



Summary

Analysis



Conclusion

Speculative type confusion is prevalent, hard to detect

Takeaways

More bugs

Rethink Spectre v1 mitigations