

Asura: A huge PCAP file analyzer for anomaly packets detection using massive multithreading

DEF CON 26, Aug 12 2018

Ruo Ando

**Center for Cybersecurity Research and Development
National Institute of Informatics**

Outline

- ❑ **“Too many packets, too few resources”**
 - **100,000,000 vs 1000,000,000,000**

- ❑ **ASURA: “Huge PCAP file vs Massive threads”**

Overview

Task based decomposition

Selection of features and containers

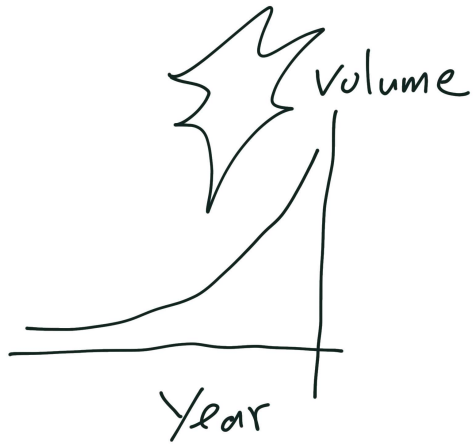
Reduction by massive threads

- ❑ **Demo and Experimental results**

- ❑ **Conclusions**



Story behind Asura



**Traffic explosion.
Internet Traffic continues to
increase at exponential rate,
no end in sight.**

**Too many packets, too few
professionals.
Cyber attack has become more
sophisticated.**



Reference: The Scream @ public domain

100,000,000 vs 1000,000,000,000

❑ **“The universe is not complicated, there's just a lot of it.”**

- Richard Feynman

❑ Unreasonable Effectiveness of Data

If a machine learning program cannot work with a training of a million examples, then the intuitive conclusion follows that it cannot work at all.

However, it has become clear that machine learning using a huge dataset with a trillion items can be highly effective in tasks for which machine learning using a sanitized (clean) dataset with a only million items is NOT useful.

Chen Sun, Abhinav Shrivastava, Saurabh Singh, Abhinav Gupta, “Revisiting Unreasonable Effectiveness of Data in Deep Learning Era”, ICCV 2017

<https://arxiv.org/abs/1707.02968>

Overview of Asura

- ❑ **Portable and reasonable**

GPU, Spark are still high-cost.

- ❑ **Posix Pthreads (explicit parallel programming model)**

**Pthreads represent the assembly languages of parallelism.
Maximum flexibility.**

- ❑ **Parser – full scratch, without libpcap**

Flexible. More scalable compared with tshark (in some cases).

- ❑ **Compact but powerful**

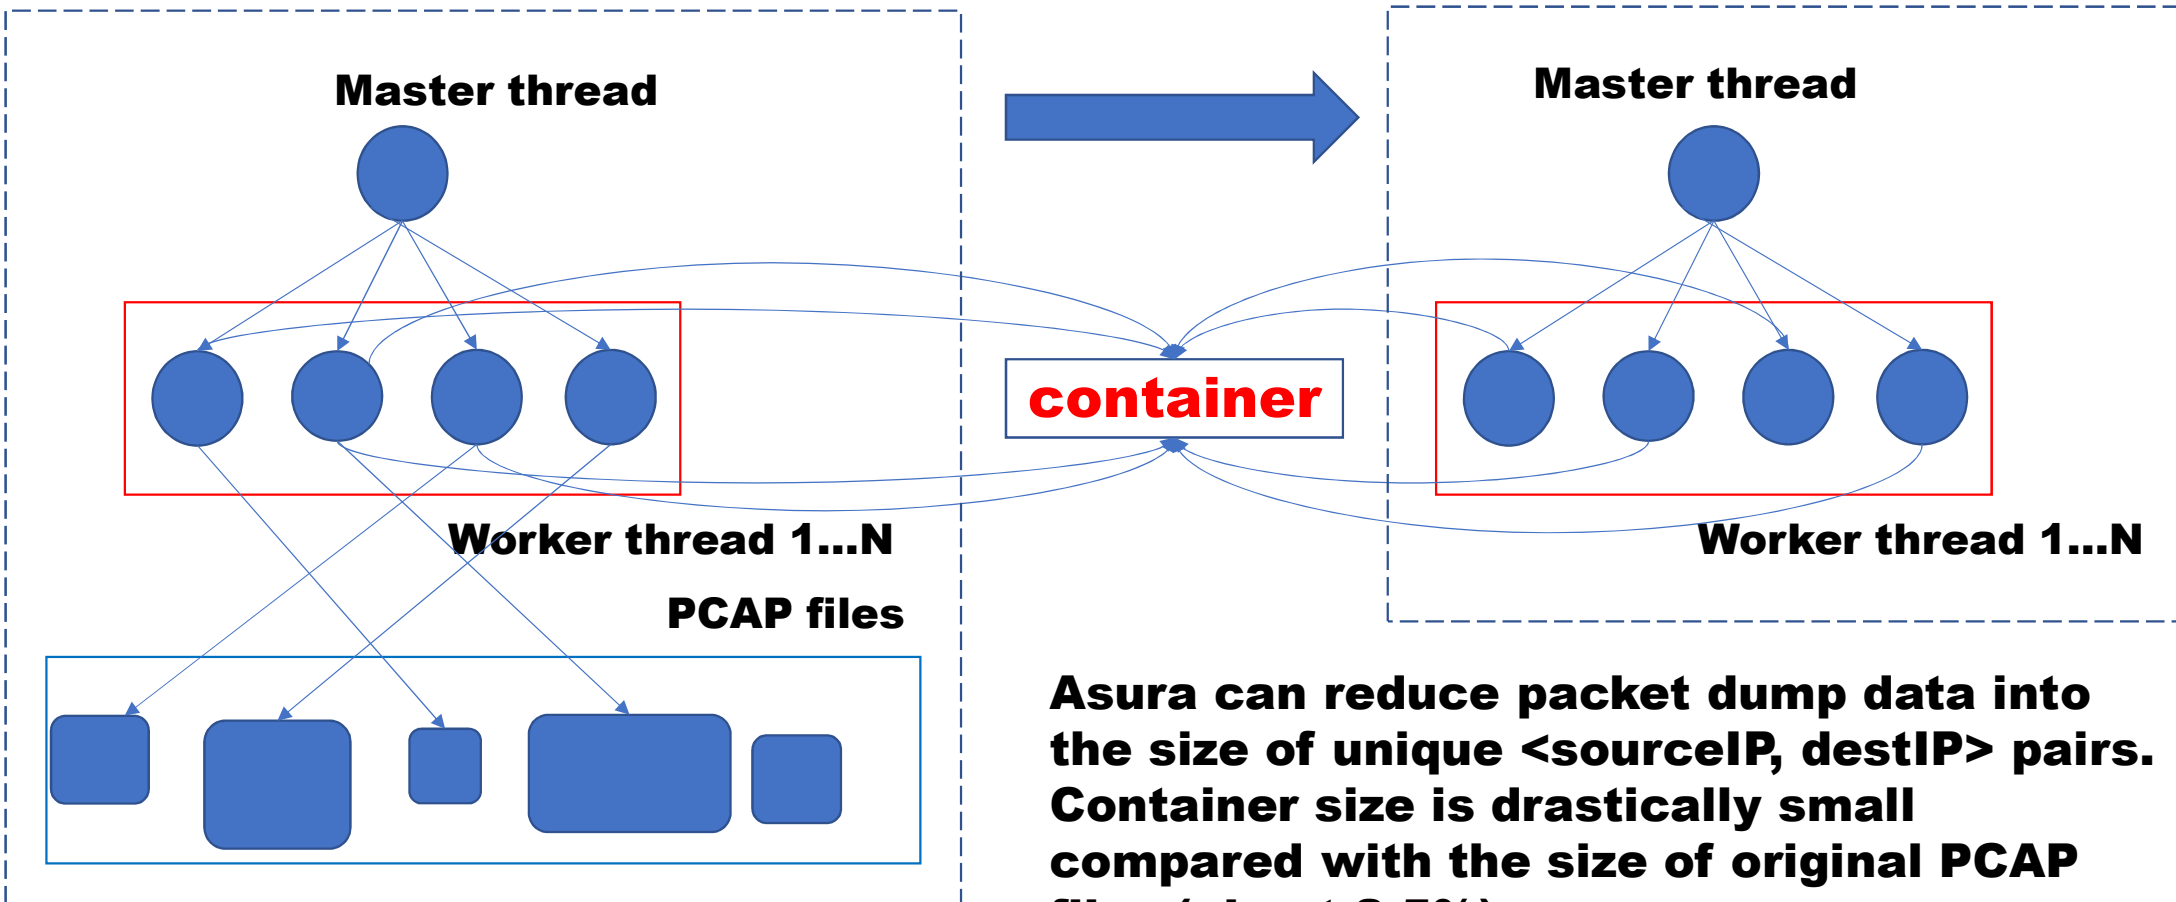
Asura has thousands lines of code.

Asura can process 76,835,550 packets in 200-400 minutes.

Asura: Huge PCAP file vs Massive threads

Reduction (task decomposition)

Clustering (data decomposition)

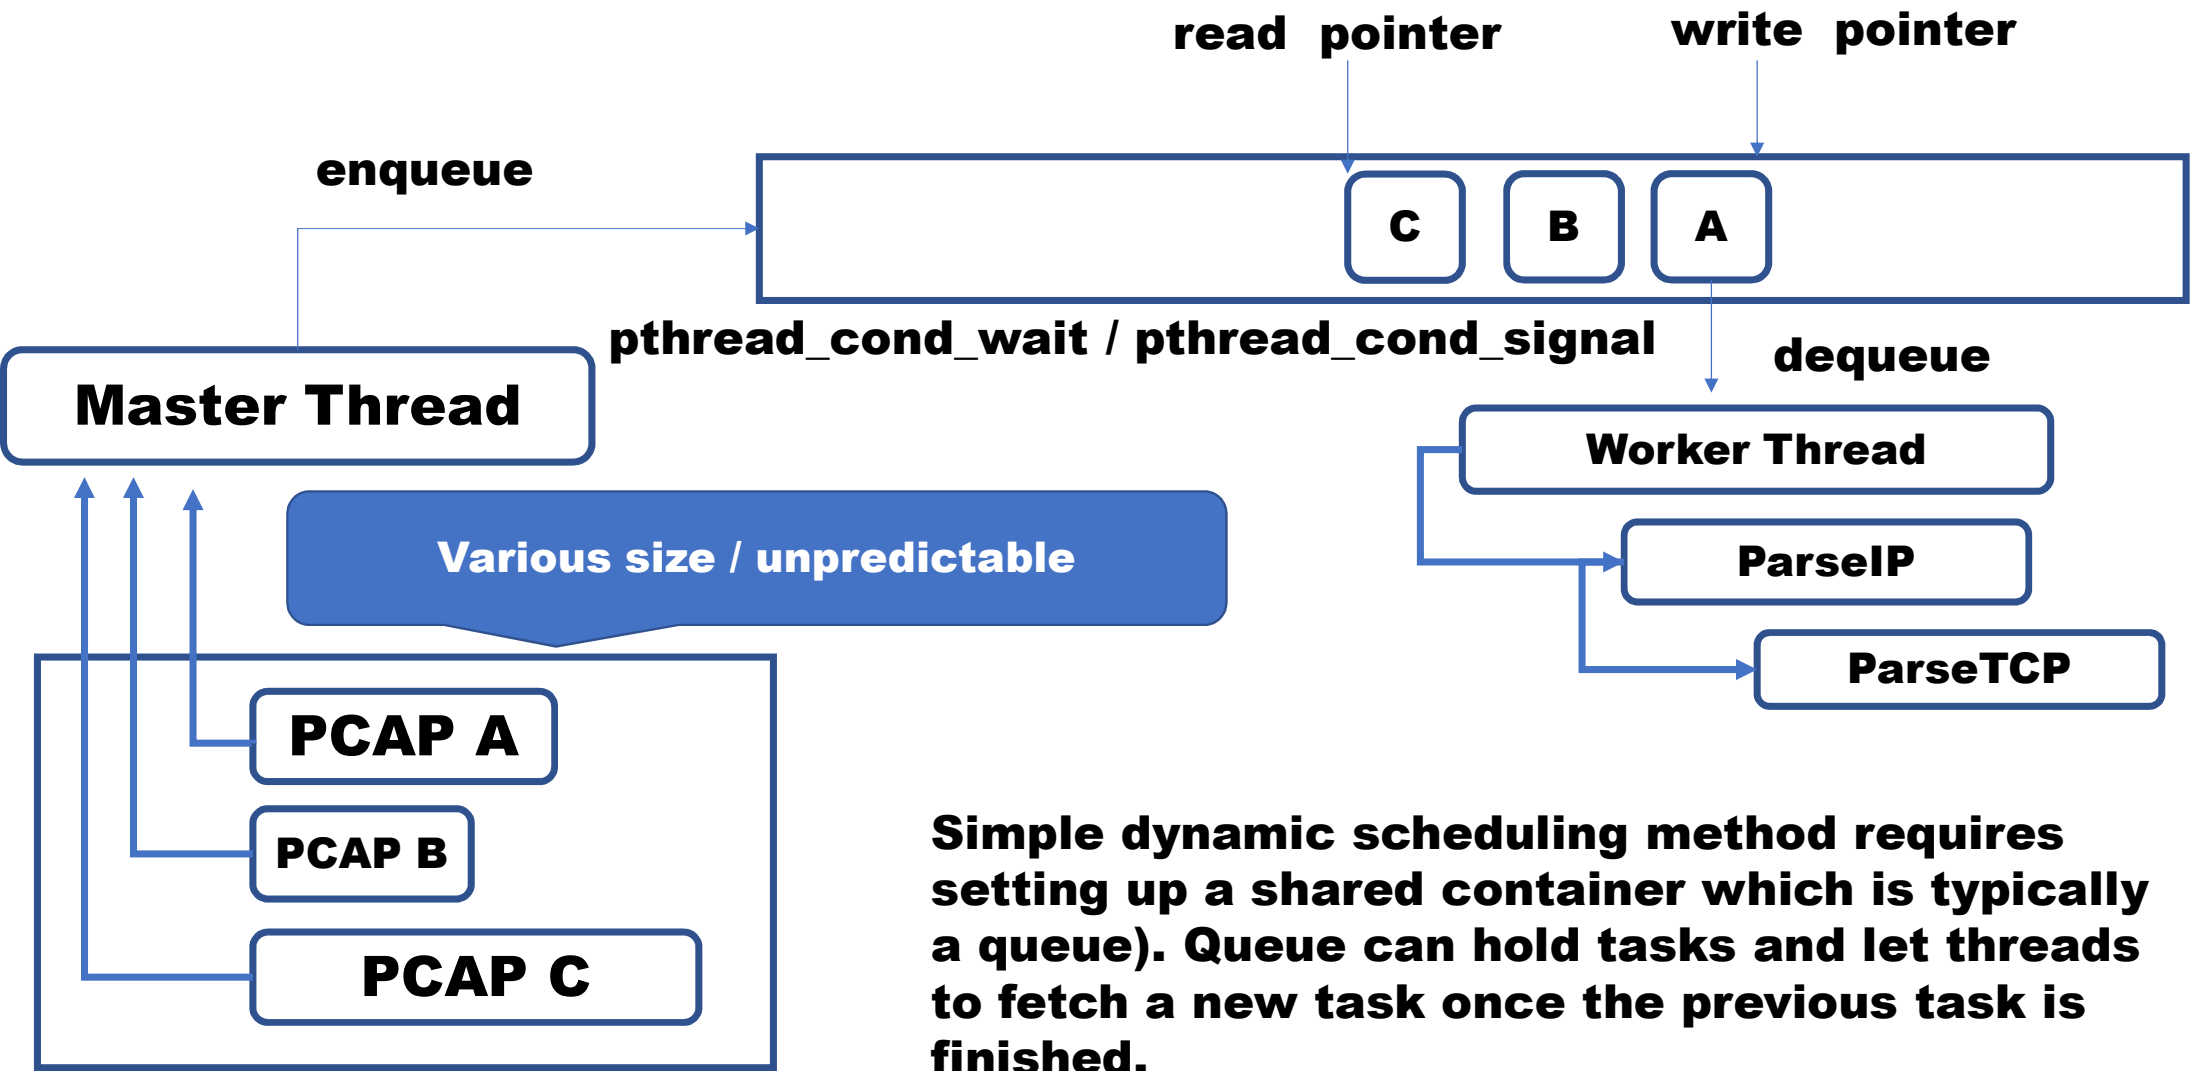


Task decomposition vs Data decomposition

If we want to transform code into a concurrent version, there are two ways.

- ❑ One way is **data decomposition**, in which the program cope with a large collection of data and can process every chunk of the data independently.
- ❑ Another way is **task decomposition**, in which the process is divided into a set of independent task so that threads can run in any order.
- ❑ As with PCAP file parsing, load balance is an important factor to take into consideration, especially when PCAP files are variable sizes. **Real-world PCAP file is NOT organized in a regular pattern and unpredictable.**

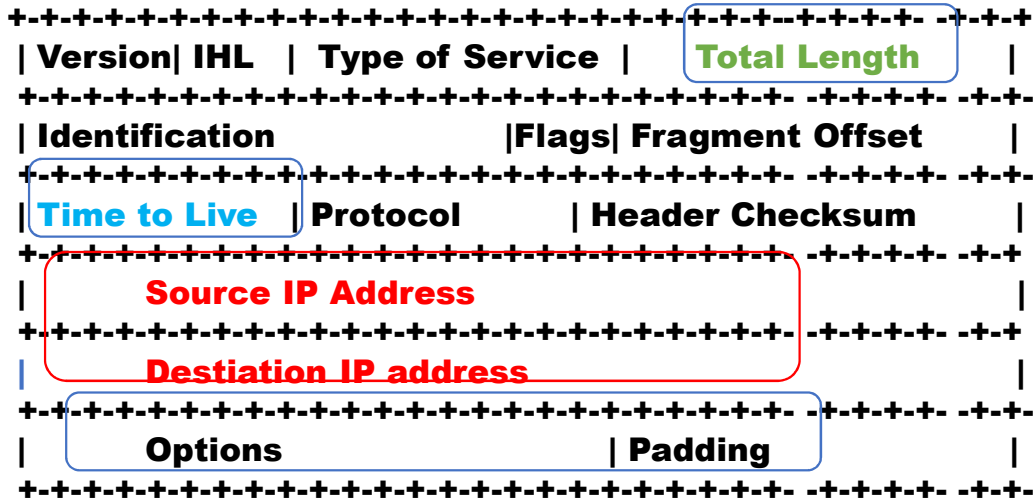
Task based decomposition (dynamic scheduler)




Simple dynamic scheduling method requires setting up a shared container which is typically a queue). Queue can hold tasks and let threads to fetch a new task once the previous task is finished.

Feature selection

IP HEADER



TCP HEADER



Asura can reduce packet dump data into the size of unique <sourceIP, destIP> pairs.


```
1: /* STRUCTURE I: srcIP, destIP */
2: typedef struct _addrpair {
3:   map<string, string> m;
4:   pthread_mutex_t mutex;
5: } addrpair_t;
6: addrpair_t addrpair;
```

Key Value

{<sourceIP, destIP>, V[i]}

```
1: /* STRUCTURE II: reduced */
2: typedef struct _reduced {
3:   map<int, int> count;
4:   map<int, int> tlen;
5:   map<int, int> ttl;
6:   map<int, int> sport;
7:   map<int, int> dport;
8:   pthread_mutex_t mutex;
9: } reduced_t;
10: reduced_t reduced;
```

Reduction by massive threads



❑ Procedures of worker thread are intuitively simple:

1. lock struct _reduced
2. aggregate the members of struct _reduced
3. unlock struct _reduced

Choice of container

- ❑ **STL + Posix Pthreads:** Exposing expose the control of parallelism at its lowest level. Maximum flexibility, but at a high cost in terms of hacker's effort (sometimes painful).
- ❑ **Intel TBB:** An excellent library for task-based parallelism mainly for scientific computation. Therefore Input data should be well structured and organized.
- ❑ **Nvidia Thrust:** C++ template library for CUDA based on the Standard Template Library (STL). Map < > container has not been implemented yet (as far as I know).

PCAP files are NOT organized in a regular pattern or the parsing of PCAP files is different or unpredictable for each element in the stream. So **STL + Posix Pthreads is the choice.**

Experimental results

**18G: 76,835,550 packets
with max queue size 1024
- about 5 to 6 hours**

# of threads	time
1	real 976m46.680s
2	real 474m0.328s
500	real 287m21.413s
1000	real 346m16.257

**18G: 76,835,550 packets
with max queue size 52
- about 7 to 9 hours**

# of threads	time
200	real 877m9.874s
500	real 464m42.022s
1000	real 493m24.110s
4000	real 523m43.533s

Kernel tuning (Ubuntu 16 LTS)


☐ Logical cores N

grep processor /proc/cpuinfo | wc -l

☐ files and procs - /etc/security/limits.conf

It should be more than 1024

☐ posix queue size - /etc/security/limits.d



	sourceIP	destIP	rare rate
MALWARE	*.*.*.*	*.*.*.*	0.73%
FINGERPRINTING	*.*.*.*	*.*.*.*	2.36%
BRUTE FORCE SSH	*.*.*.*	*.*.*.*	3.23%
NORMAL	*.*.*.*	*.*.*.*	91.95%

1 minutes Demo



Conclusions

- ❑ **Asura is full-scratch (and painful) implementation with POSIX Pthreads and C++ STL.**
- ❑ **Leveraging Pthreads which represents assembly language in parallelism provides maximum flexibility.**
- ❑ **Asura adopts task-based implementation for processing huge, heterogeneous and unpredictable PCAP file stream.**
- ❑ **Asura has thousands lines of code and can process 76,835,550 packets in 200-500 minutes.**

Thank you !

<https://github.com/RuoAndo/Asura>