

# **RSAC**Conference2022

San Francisco & Digital | June 6 – 9

SESSION ID: HTA-T02

## **My Fuzzy Driver**

**Eran Shimony**

Senior Security Researcher  
CyberArk  
@EranShimony

**Mark Cherp**

Senior Security Researcher  
CyberArk  
@OcamRazr

**TRANSFORM**



# Disclaimer

Presentations are intended for educational purposes only and do not replace independent professional judgment. Statements of fact and opinions expressed are those of the presenters individually and, unless expressly stated to the contrary, are not the opinion or position of RSA Conference LLC or any other co-sponsors. RSA Conference does not endorse or approve, and assumes no responsibility for, the content, accuracy or completeness of the information presented.

Attendees should note that sessions may be audio- or video-recorded and may be published in various media, including print, audio and video formats without further notice. The presentation template and any media capture are subject to copyright protection.

©2022 RSA Conference LLC or its affiliates. The RSA Conference logo and other trademarks are proprietary. All rights reserved.

# RSA<sup>®</sup>Conference2022

TL;DR

**kAFL vs Drivers == Bugs**



# Our Menu



## Target

What to attack and how?

## Internals

Healthy dose of Windows

## Bugs

Vulnerabilities deep dive

## Automation

Discovery, harness and  
grammar

## Fuzzing

kAFL setup with some  
tweaks



# RSA<sup>®</sup>Conference2022

## Target

**What to attack and why?**



# Don't run away from Windows

- We aim to escalate privilege from a weak point, accessible drivers are a good candidate
- Windows has more than 300 drivers with tons of legacy code written in C
- Windows is “considered” to be closed source, which makes it harder and less explored
- Windows drivers might be challenging but they are also rewarding 😊


# RSA<sup>®</sup>Conference2022

## Internals

**Healthy dose of Windows**




# Windows' privilege level architecture





# Windows Drivers 101




## Part of the kernel

- An extension of the Windows kernel
- Share System's address space
- Essentially a part of the kernel



## Signing

- Must be digitally signed
- Driver signing can be disabled
- Requires Admin+



## Responsibility

- Talking with hardware, devices
- Filtering
- Not blue screening your machine

# Too many Driver models



# WDM Drivers

## PERSONAL DATA

Age	23 Years
Usage	Abundant
Successor	KMDF
Location	Windows OS
Amount	Over 200

## STATS



## PERCENTAGE OF TOTAL



## PURPOSE

- ✓ Talking with devices, power management...
- ✓ Supply kernel support for applications

# RSA<sup>®</sup>Conference2022


## Bugs

**Vulnerabilities deep dive**






# Driver Bug Hunting 101



## Access

Which driver do I have access to?

What are my permissions?



## Communication

How do we communicate with the driver?

How does it parse user data?

# Driver's Anatomy I

## Access

In Windows, you talk to a driver via a device

The driver creates a device, and it should specify which users can access it. It uses SDDL string for it, or defines it in the inf file:

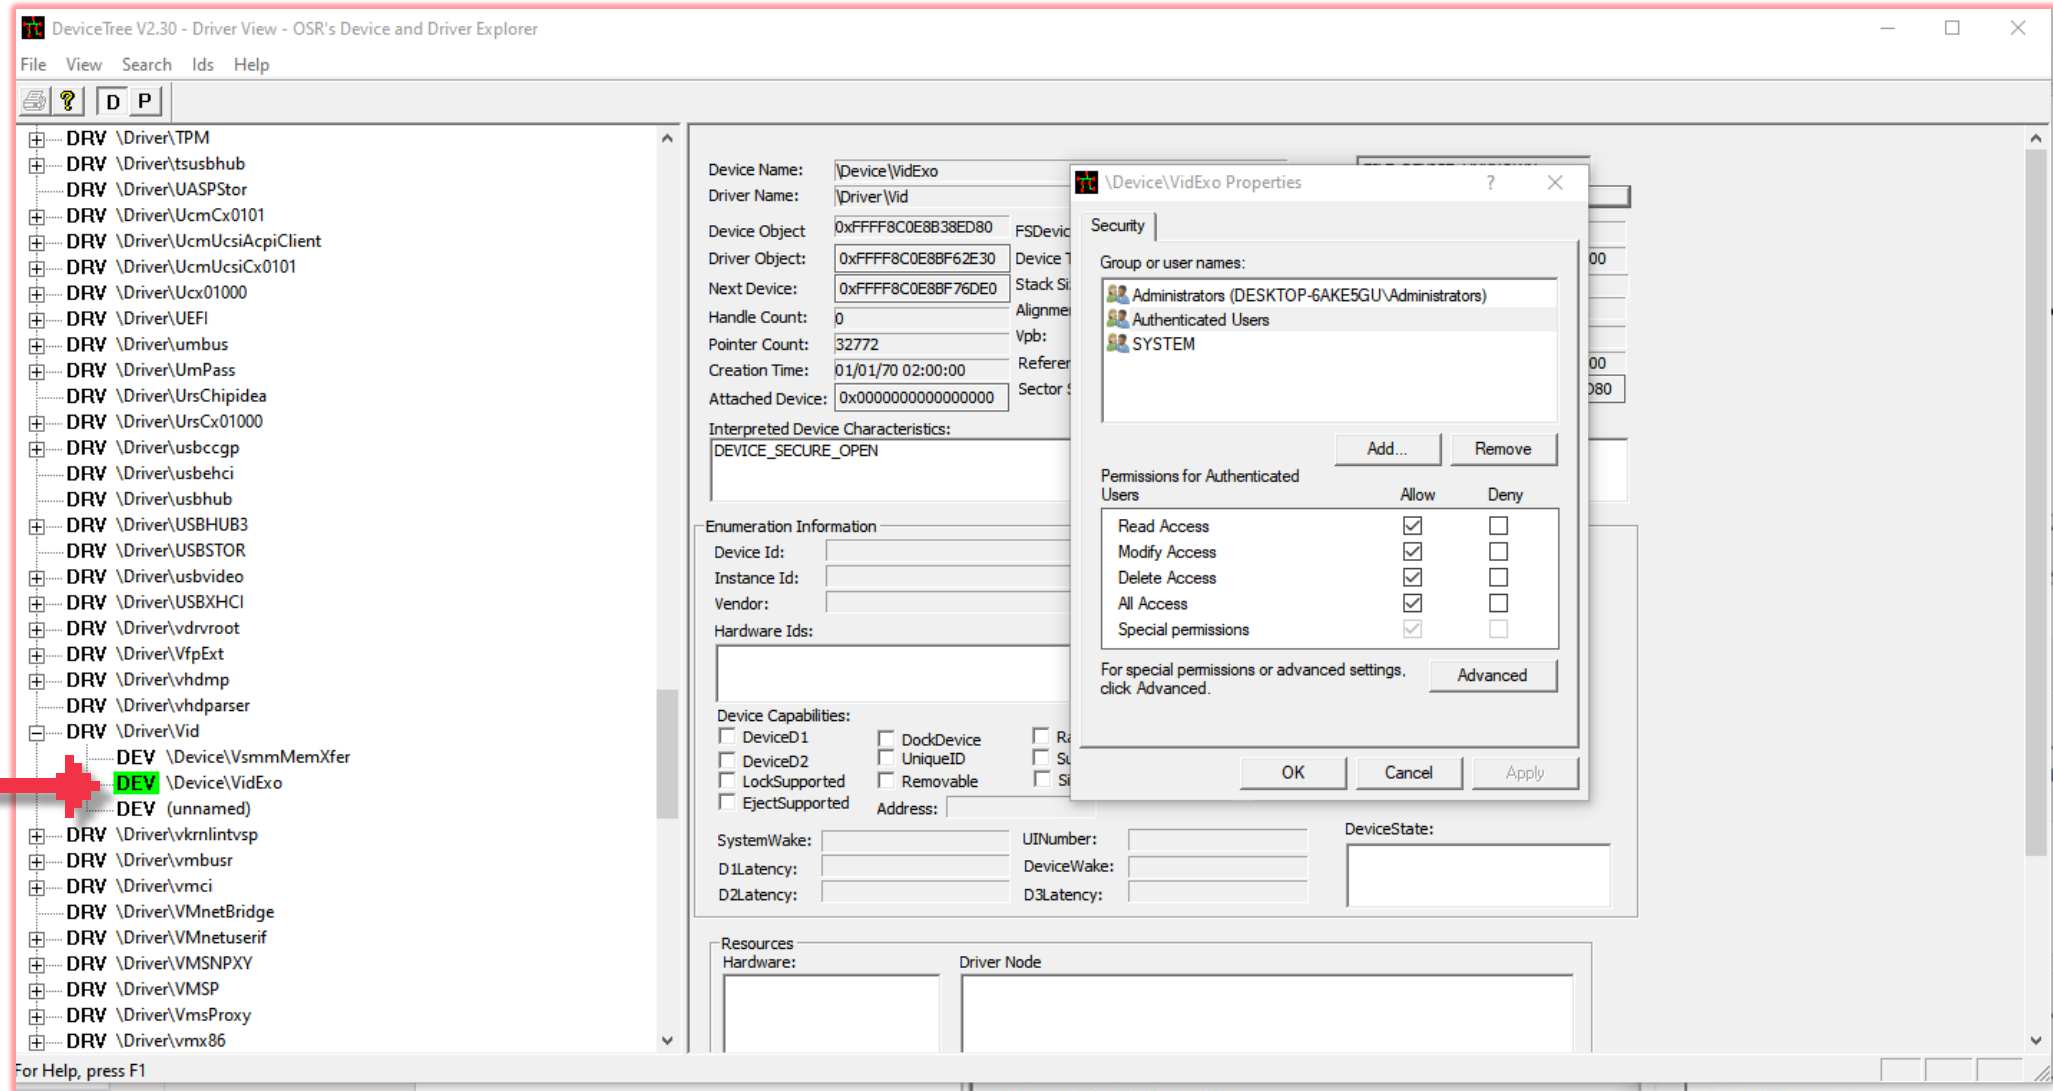
- *IoCreateDevice()* or *IoCreateDeviceSecure()*

Followed by exposing the device to the user via calling:

- *IoCreateSymbolicLink()* or *IoRegisterDeviceInterface()*

The user opens a handle to the device via calling *CreateFile()\NtCreateFile()*

# Show me some permissions



# Driver's Anatomy II



## Communication

The Driver object registers dispatch routines through *DriverObject->MajorFunction[IRP\_MJ\_XXX]*

The Dispatch routines are invoked by the *IoManager* when the following operations are done on the device:

- Create - *NtCreateFile*
- Close - *NtCloseFile*
- Read - *NtReadFile*
- Write - *NtWriteFile*
- Device\_Control – *DeviceIoControl -> FileIoControlDevice*

These dispatch routines our initial go-to places



# Typical DriverEntry

```
extern "C" NTSTATUS DriverEntry(IN PDRIVER_OBJECT DriverObject, IN PUNICODE_STRING RegistryPath)
{
    UNICODE_STRING DeviceName, SymbolicLink, sddlString;
    PDEVICE_OBJECT deviceObject;
    RtlInitUnicodeString(&DeviceName, L"\\Device\\testydrv");
    RtlInitUnicodeString(&SymbolicLink, L"\\DosDevices\\testydrv");
    RtlInitUnicodeString(&sddlString, L"D:P(A;;GA;;;SY)(A;;GA;;;BA)");

    //Create a device
    IoCreateDevice(DriverObject, 65535, &DeviceName, FILE_DEVICE_UNKNOWN, FILE_DEVICE_SECURE_OPEN, FALSE, &deviceObject);
    //IoCreateDeviceSecure(DriverObject, sizeof(65533), &DeviceName, FILE_DEVICE_UNKNOWN, FILE_DEVICE_SECURE_OPEN, FALSE, &sddlString, NULL, &deviceObject);
    //Create a symbolic so the user can access the device
    IoCreateSymbolicLink(&SymbolicLink, &DeviceName);

    //Populating Driver's object dispatch table
    DriverObject->MajorFunction[IRP_MJ_DEVICE_CONTROL] = TestyDispatchIoctl;
    DriverObject->MajorFunction[IRP_MJ_CREATE] = TestyDispatchCreate;
    DriverObject->MajorFunction[IRP_MJ_CLOSE] = TestyDispatchClose;
    DriverObject->MajorFunction[IRP_MJ_READ] = TestyDispatchRead;
    DriverObject->MajorFunction[IRP_MJ_WRITE] = TestyDispatchWrite;
    DriverObject->MajorFunction[IRP_MJ_CLEANUP] = TestyDispatchCleanup;
    DriverObject->DriverUnload = TestyUnloadDriver;
    return STATUS_SUCCESS;
}
```

# Dispatch Routine

```
NTSTATUS TestyDispatchIoctl(PDEVICE_OBJECT DeviceObject, PIRP Irp)
{
    PIO_STACK_LOCATION CurrentStackLocation;
    NTSTATUS ntStatus = STATUS_SUCCESS;
    PVOID SystemBuffer = NULL;
    ULONG InputBufferLength = 0;
    ULONG OutputBufferLength = 0;
    DWORD IoControlCode = 0;

    CurrentStackLocation = IoGetCurrentIrpStackLocation(Irp);
    InputBufferLength = CurrentStackLocation->Parameters.DeviceIoControl.InputBufferLength;
    OutputBufferLength = CurrentStackLocation->Parameters.DeviceIoControl.OutputBufferLength;
    SystemBuffer = Irp->AssociatedIrp.SystemBuffer;
    IoControlCode = CurrentStackLocation->Parameters.DeviceIoControl.IoControlCode;
```

# IRPs

**I/O request packet**, is a struct that has all the parameters any of the dispatch routines would ever need, generated when an action is performed on a device

The driver receives a pointer to the generated **IRP structure** by the *IoManager*, it consists of:

*StackLocation* – holds many important members

*Requestor mode* – Kernel or User

*Buffers* – Depends on the transfer type

*IoStatus.Information* – How many bytes are written to the output buffer

♥ A misuse of any of these fields would probably cause a bug

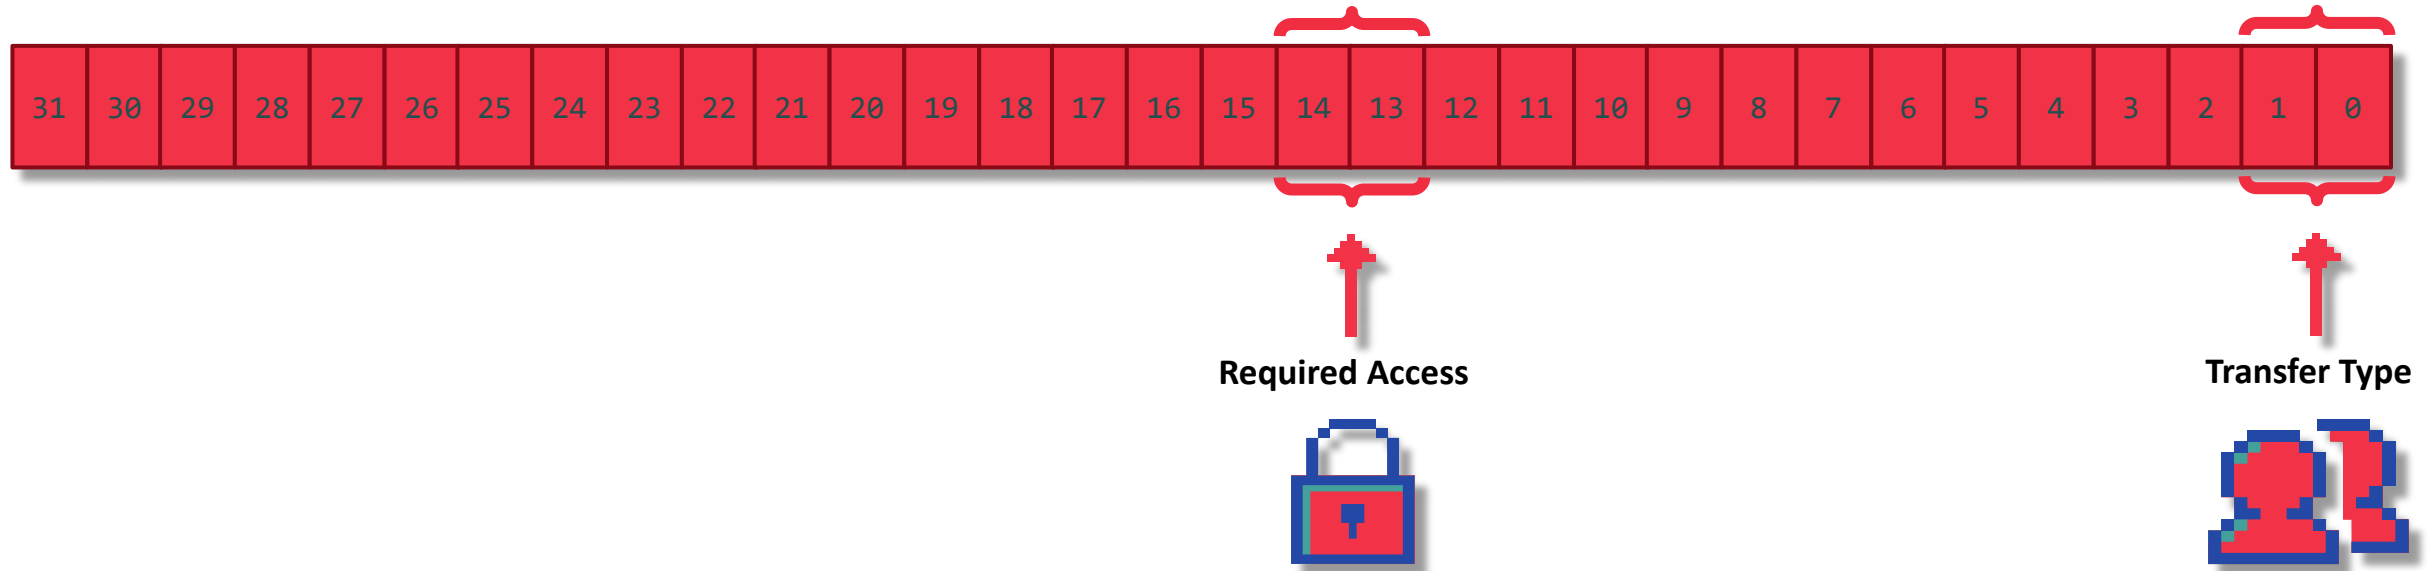
# How to generate IRPs

```
BOOL DeviceIoControl(  
    HANDLE hDevice,  
    DWORD dwIoControlCode,  
    LPVOID lpInBuffer,  
    DWORD nInBufferSize,  
    LPVOID lpOutBuffer,  
    DWORD nOutBufferSize,  
    LPDWORD lpBytesReturned,  
    LPOVERLAPPED lpOverlapped  
);
```



# IOCTL Code

## 32-bit IOCTL



# Required Access



*Two bit number* – describes the required permissions to send the *DeviceIoControl* request, it is based on your permissions in the call to *NtCreateFile*.

- There are four different options:
  - *FILE\_ANY\_ACCESS*
  - *FILE\_READ\_ACCESS*
  - *FILE\_WRITE\_ACCESS*
  - *FILE\_READ\_ACCESS | FILE\_WRITE\_ACCESS*

# Transfer Type



*The two least significant bits* – describe how the *IoManager* treats user's data, there are many nuances the driver developer must know

- There are four different options:
  - *METHOD\_NEITHER*
  - *METHOD\_BUFFERED*
  - *METHOD\_IN\_DIRECT*
  - *METHOD\_OUT\_DIRECT*

# METHOD\_NEITHER = Mother of all evil

- The two bits are on -> ioctl number ends with 11
- The IoManager is lazy, the buffers and their lengths reside in User-Mode, the kernel does not copy them to kernel space, they can be paged-out
- The input and output buffers are:
  - *Irp->CurrentStackLocation.Parameters.DeviceIoControl.Type3InputBuffer*
  - *Irp->UserBuffer*
- The user can allocate and deallocate the buffers making the pages invalid, so you must be cautious when dealing with them
- Every access to the buffers must be in a try except block



# Probing

- Probing validates that an address resides in User-Mode:
  - *ProbeForRead* (\*Address, Length, Alignment)
  - *ProbeForWrite* (\*Address, Length, Alignment)
- But if you don't probe or you do not use the function correctly, then:
  - It throws an exception on invalid address so the call must be inside a try except block
  - If Length is 0 it does nothing, it passes the validation without probing
  - If not probed and the kernel reads from it = **BSoD or info leak**
  - If not probed and the kernel writes to it = **BSoD or arbitrary write**

# Bug Example I

```

if ( IOCTL - 0x220C00 <= 0x27 )
    return sub_140028698(DeviceObject, IRP);
if ( CurrentStackLocation->Parameters.DeviceIoControl.IoControlCode != 0x2F0003 || InputBufferLength < 0x18 )
    return ((__int64 (__fastcall *))(_DEVICE_OBJECT *, IRP *))qword_140065278(DeviceObject, IRP);
SystemBuffer = (BYTE *)IRP->AssociatedIrp.SystemBuffer;
v13 = 0;
InputBuffer = (int64)CurrentStackLocation->Parameters.DeviceIoControl.Type3InputBuffer;
InputBuffer2 = InputBuffer;
OutputBufferLength3 = OutputBufferLength;
if ( SystemBuffer && (*(DWORD *) (InputBuffer + 20) & 1) != 0 )
    OutputBuffer = SystemBuffer;
else
    OutputBuffer = IRP->UserBuffer;
v43 = IRP;
if ( *(QWORD *)InputBuffer == 0x43736C266128A8C4i64 && *(QWORD *) (InputBuffer + 8) == 0x4151AA59370630B6i64 )
    return sub_140027EC8(DeviceObject, IRP);
if ( *(_QWORD *)InputBuffer == 0x475215BAE7F772BCi64 && *(_QWORD *) (InputBuffer + 8) == 0x57462BA4FA7660BFi64 )

```

METHOD\_NEITHER ends with 11

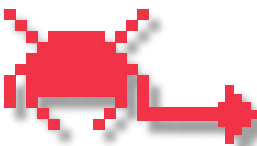
# Method\_Buffered

- The two bits are off -> ioctl number ends with 00
- In *METHOD\_BUFFERED* the IoManager copies the buffers and their lengths to the kernel in a secure manner. Therefore, they don't reside in the user's memory space.
- No need for probing, but if the buffer has embedded fields, like pointer addresses, lengths and so on, they need to be treated properly!
- The IRP buffer is used both for input and output:
  - *Irp->AssociatedIrp.SystemBuffer*

# Method\_Buffered - Continued

- The *Irp->IoStatus.Information* indicates how many bytes are to be copied to user's *OutputBuffer*
- If *Irp->IoStatus.Information* > *InputBufferLen* and *OutputBufferLen* > *InputBufferLen*
  - The rest of the system buffer data copied to (*OutputBuffer*) is uninitialized data
  - $SizeOfData = OutputBufferLen - InputBufferLen$
- Would cause kernel leak unless the buffer is properly initialized
  - Assuming:
    - $OutputBufferLen = 0x1000, InputBufferLen = 0x8$
    - $Irp->IoStatus.Information = OutputBufferLen$
  - Leakage of 0xFF8 bytes to user-mode

# Bug Example II



```

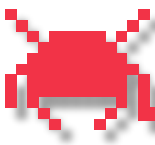
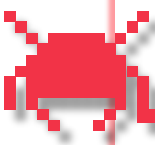
DisptachIoctlFDO (PDEVICE_OBJECT Device_Object, IRP *IRP)
{
case 0x2D2324u:                METHOD_BUFFERED ends with 00
    NTStatus2 = DoSomething(Device_Object, IRP);
    ...

    __int64 __fastcall DoSomething(PDEVICE_OBJECT Device_Object, PIRP irp)
    {
        _IO_STACK_LOCATION *CurrentStackLocation; // rbx
        PVOID Device_Extension; // rbp
        BYTE SystemBuffer; // si
        __int64 InputBufferLength; // r9
        ULONG OutputBufferLength; // [rsp+50h] [rbp+8h] BYREF

        CurrentStackLocation = irp->Tail.Overlay.CurrentStackLocation;
        Device_Extension = Device_Object->DeviceExtension;
        irp->IoStatus.Information = 0i64;
        SystemBuffer = (BYTE *)irp->AssociatedIrp.SystemBuffer;
        InputBufferLength = CurrentStackLocation->Parameters.DeviceIoControl.InputBufferLength;
        OutputBufferLength = CurrentStackLocation->Parameters.DeviceIoControl.OutputBufferLength;

        if ( OutputBufferLength >= 0x107C0 && *( QWORD *)&SystemBuffer )
        {
            FillOutputDataAndRestWithZeros((__int64)Device_Extension, *( _BYTE **)&SystemBuffer, &OutputBufferLength);
            IRP->IoStatus.Information = OutputBufferLength;
            return 0i64;
        }
        else
        {
            IRP->IoStatus.Information = 0x107C0i64;
            return 0x800000005i64;
        }
    }
}

```


# Direct I/O

- The Right bit is on = *METHOD\_IN\_DIRECT* and left is off 01
- The Left bit is on = *METHOD\_OUT\_DIRECT* and right is off 10
- The IRP supply two pointer to buffers, both in kernel:
  - *Irp->AssociatedIrp.SystemBuffer* – The input buffer
  - *Irp->MdlAddress* – A second buffer used or buffered
- The second buffer is the “direct” which always paged = faster access, this buffer can be used for input or output
- The kernel can also create an MDL with *IoCreateMdl()* followed by locking it with *MmProbeAndLockPages()* which would throw an exception on invalid page



# Direct I/O MDL

- “An MDL is a structure that describes the fixed physical memory locations that comprise a contiguous data buffer in virtual memory”
- In simpler words, The MDL describes the data buffer at a fixed position in physical memory, which will always be paged in and locked (you need to do that) in memory. It is a double mapping, one for the user, and another for the kernel



# Bug Example III

- Most of the bugs involve not checking for null, as the user can send a null buffer, thus causing *MmGetSystemAddressForMDLSafe()* to be sad
- Also, if creating an MDL yourself in the kernel, make sure you use the correct virtual address of a buffer not like in here:

```
ControlableVirtualAddress = (void *)*((_QWORD *)SystemBuffer2 + 3);
if ( ControlableVirtualAddress )
{
    AllocatedMdl = IoAllocateMdl(ControlableVirtualAddress, dwMdlSize, 0, 0, Irp);
    IrqlLevel = GetIRQLLevel(KernelGetCurrentIrql());
    LODWORD(Irpb) = *((_DWORD *)SystemBuffer2 + 3);
    // Local pages in physical memory, would crash on a bad address
    MmProbeAndLockPages(AllocatedMdl, 0, IoModifyAccess);
    if ( (AllocatedMdl->MdlFlags & 5) != 0 )
        LocalMappedSystemVA = AllocatedMdl->MappedSystemVa;
    else
        LocalMappedSystemVA = MmMapLockedPagesSpecifyCache(AllocatedMdl, 0, MmCached, 0i64, 0, dword_1400E8D84 | 0x10u);
}
```

# How to Be Lazy



## Enumerate

Enumerate device names and discover IOCTLs



## Generate

Generate harness



## Fuzz

Fuzz with the generated harness

# RSA<sup>®</sup>Conference2022

## Automation

**Discovery, harness and grammar**



# Device Handles and IOCTLs Discovery

## Enumerating All Device Drivers in the System

05/31/2018 • 2 minutes to read • 📄 🗑

The following sample code uses the EnumDeviceDrivers function to enumerate the current device drivers in the system. It passes the load addresses retrieved from this function call to the GetDeviceDriverBaseName function to retrieve a name that can be displayed.

```
C++
#include <windows.h>
#include <psapi.h>
#include <ntchar.h>
#include <stdio.h>

// To ensure correct resolution of symbols, add Psapi.lib to TARGETLIBS
// and compile with -DPSAPI_VERSION=1

#define ARRAY_SIZE 1024

int main( void )
{
    LPVOID drivers[ARRAY_SIZE];
    DWORD cbNeeded;
    int cDrivers, i;

    if ( EnumDeviceDrivers( drivers, sizeof( drivers ), &cbNeeded ) && cbNeeded < sizeof( drivers ) )
    {
        TCHAR szDriver[ARRAY_SIZE];
        cDrivers = cbNeeded / sizeof( drivers[0] );
        _tprintf( TEXT( "There are %d drivers:\n" ), cDrivers );
        for ( i=0; i < cDrivers; i++ )
        {
            if ( GetDeviceDriverBaseName( drivers[i], szDriver, sizeof( szDriver ) / sizeof( szDriver[0] ) ) )
            {
                _tprintf( TEXT( "%d: %s\n" ), i+1, szDriver );
            }
        }
    }
    else
    {
        _tprintf( TEXT( "EnumDeviceDrivers failed; array size needed is %d\n" ), cbNeeded / sizeof( LPVOID ) );
        return 1;
    }

    return 0;
}
```

The SetupDiGetClassDevs function returns a handle to a device information set that contains requested device information elements for a local computer.

## Syntax


```
C++
WINSETUPAPI HDEVINFO SetupDiGetClassDevs(
    const GUID *ClassGuid,
    PCWSTR Enumerator,
    HWND hwndParent,
    DWORD Flags
);
```

## Device Console (DevCon.exe) Commands

04/20/2017 • 4 minutes to read • 📄 🗑

DevCon (DevCon.exe) is a command line tool that can display detailed information about devices on computers running Windows. You can also use DevCon to enable, disable, install, configure, and remove devices. DevCon uses the following syntax.

```
devcon [/m:\computer] [/r] command [arguments]
```




Discover Devices

Enumerate Device Names  
and Discover IOCTLs



# Discovering Accessible Devices





# IOCTLs Discovery




DriverEntry(...)

IOCTL Dispatch Routine

IOCTL Conditioning

Test IOCTLs




Track DriverObject

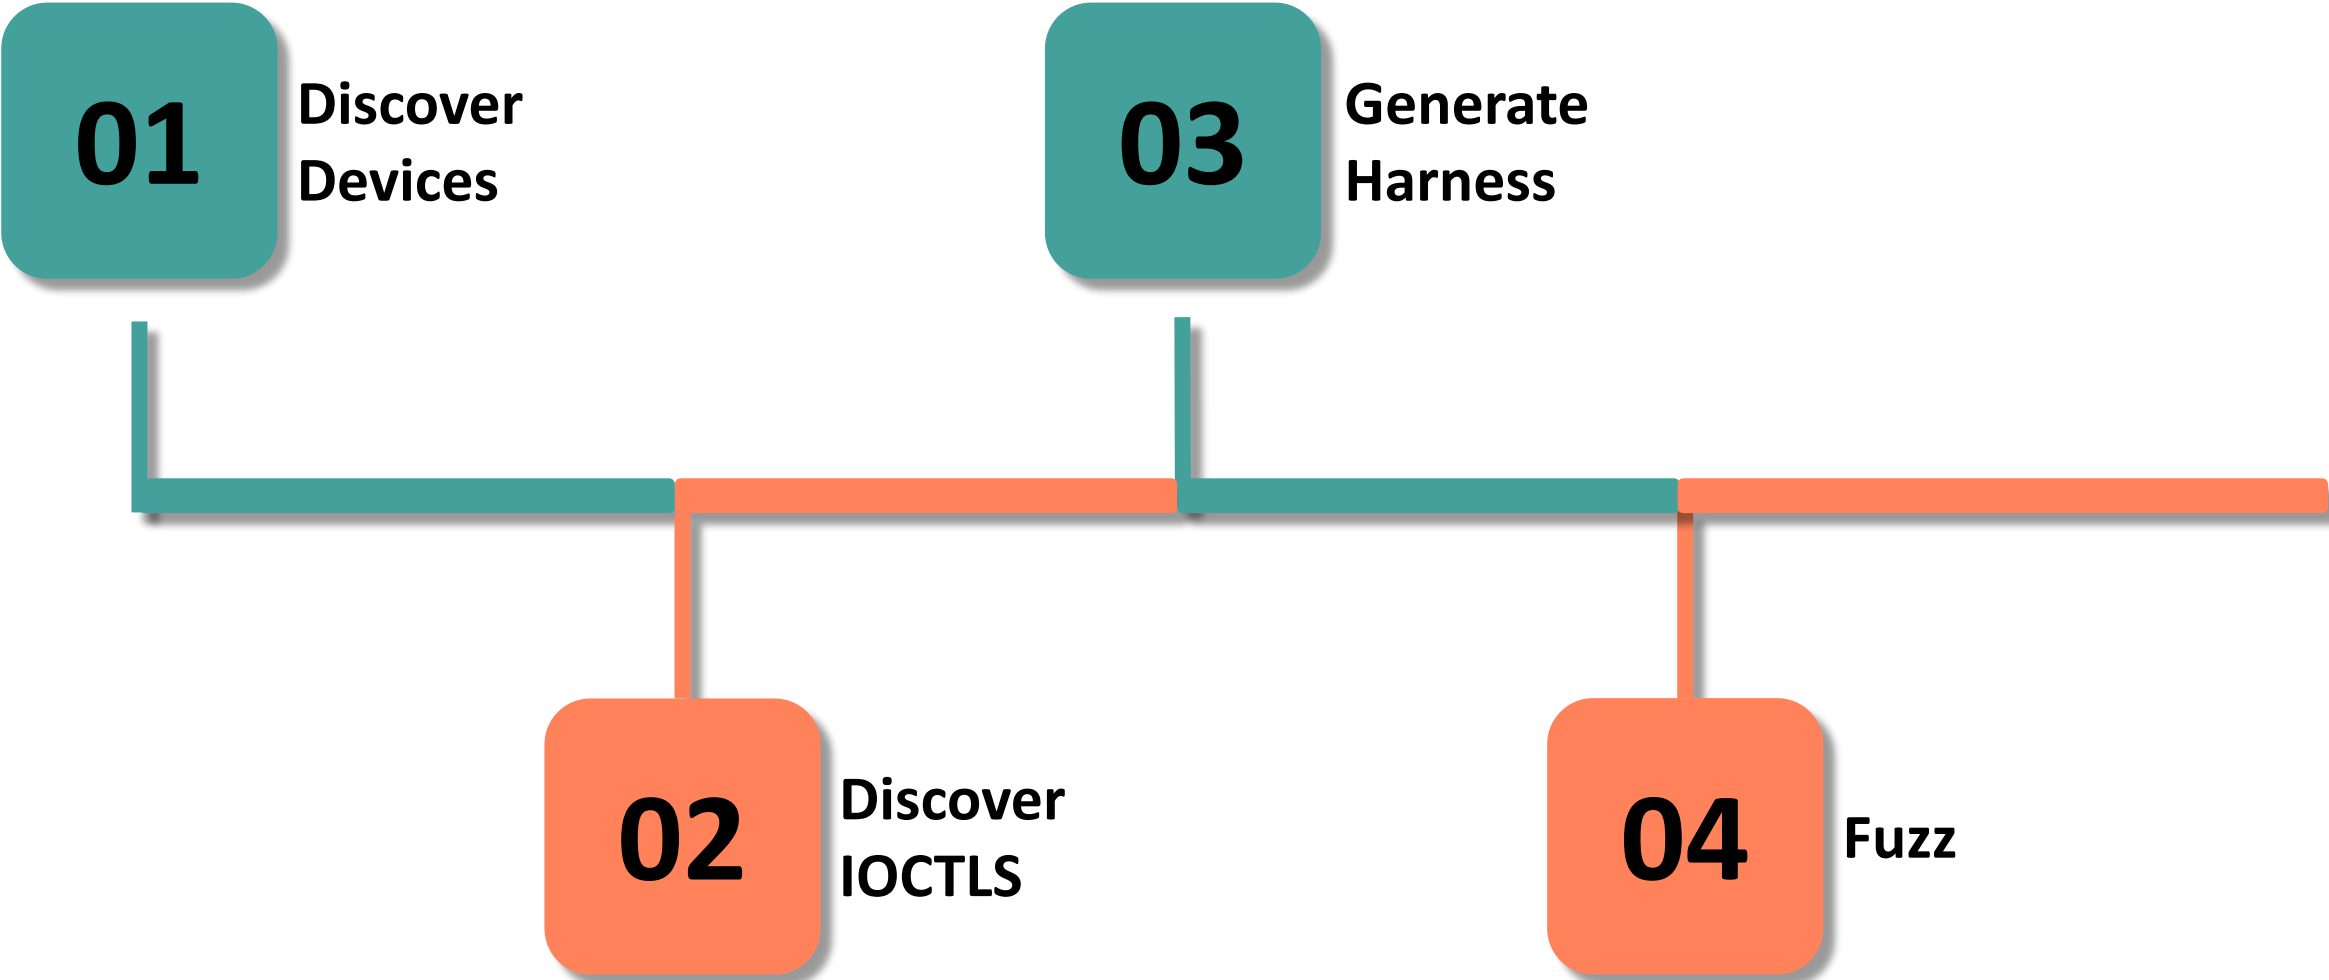
Track IOCTL code

Aggregate IOCTLs

# Harness Generation



# Automation Summary



# Grammar 101

- After discovering some bugs with a simple grammar, we started examining more advanced approaches
- We know the shape of the “high-level” input passed to the IOCTL
- We don’t know how the “deserialized” data inside the driver looks like
- We want to guide our Fuzzer towards relevant inputs and that’s where grammar comes into play
- Grammar is essentially a way to tell the Fuzzer which inputs are more likely to increase coverage

# Grammar Approaches

## Simple template matching (i.e. Regex)

```
// Regex to match phone numbers  
"^([+]*([0-9]{1,4}){0,1}[-\\s\\./0-9]*$"
```



```
// Possible matches  
"+(123) - 456-78-90"  
"+972-548099912"
```

## Input-to-state correspondence (i.e. Redqueen)

```
//Try Input = "SEEDVALUE"  
cmp eax, "ABCD"
```



```
//Observe  
eax == "VALU"
```



```
//Replace  
"SEEDVALUE" with "SEEDABCDE"
```

## Large-scale mutation (i.e. Grimoire)

```
# Original statement  
print ("aabbccdd")
```



```
# Mutated  
print ("aabbccdd")
```



```
# Mutated new paths  
print ("aabbccdd")
```

# RSAC<sup>®</sup>Conference2022


## Kernel Fuzzing

**kAFL setup with some tweaks**



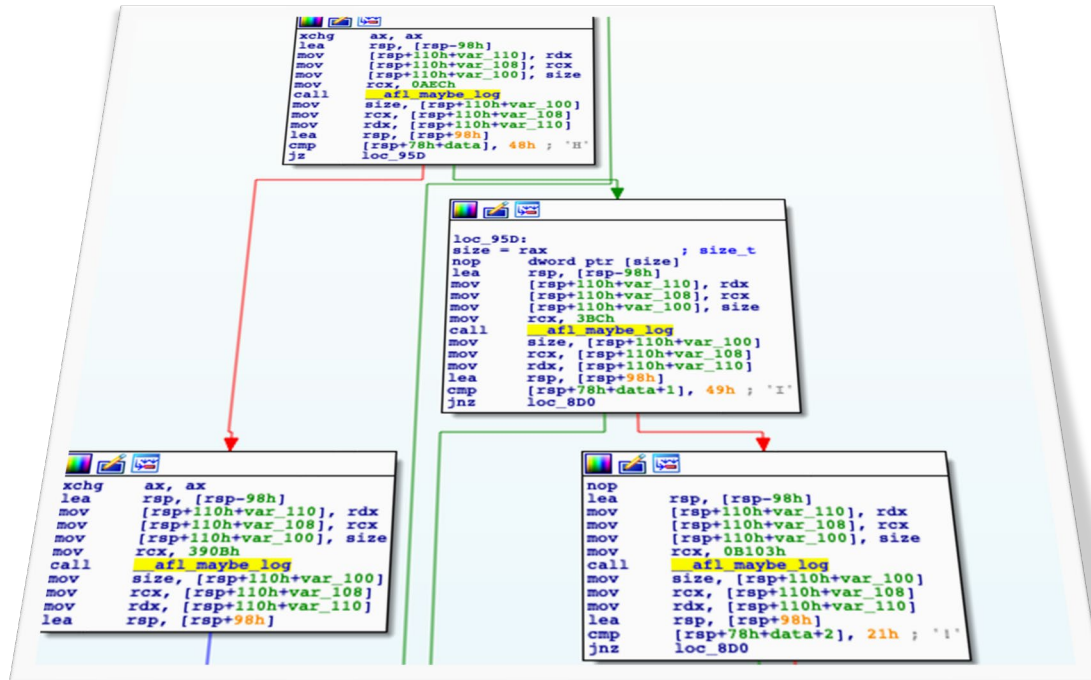
# Fuzzing Concepts

- **Feedback** – Get feedback on target execution state
- **Fitness** – Decide according to some metric (**coverage**) how





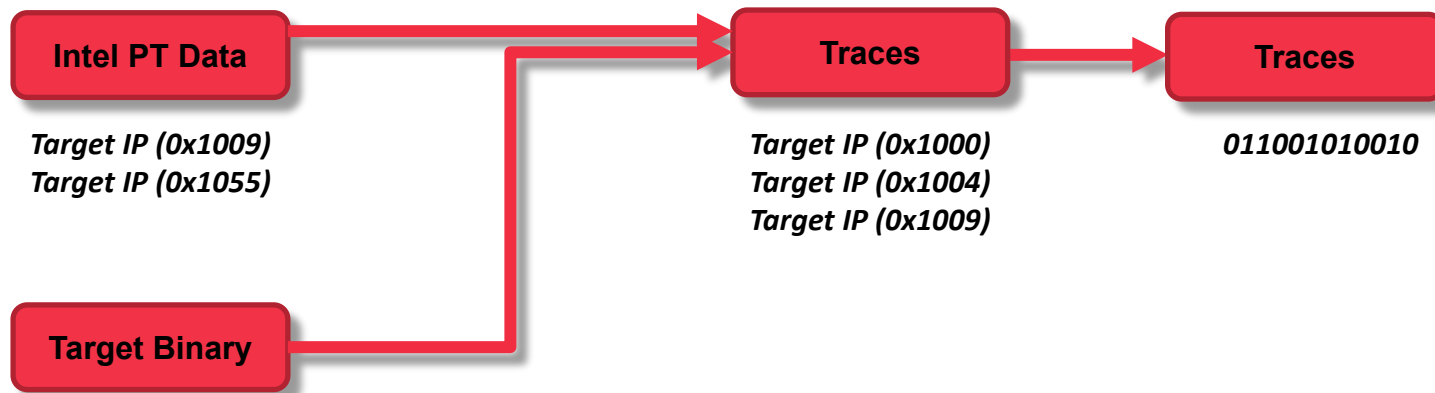
# AFL Feedback



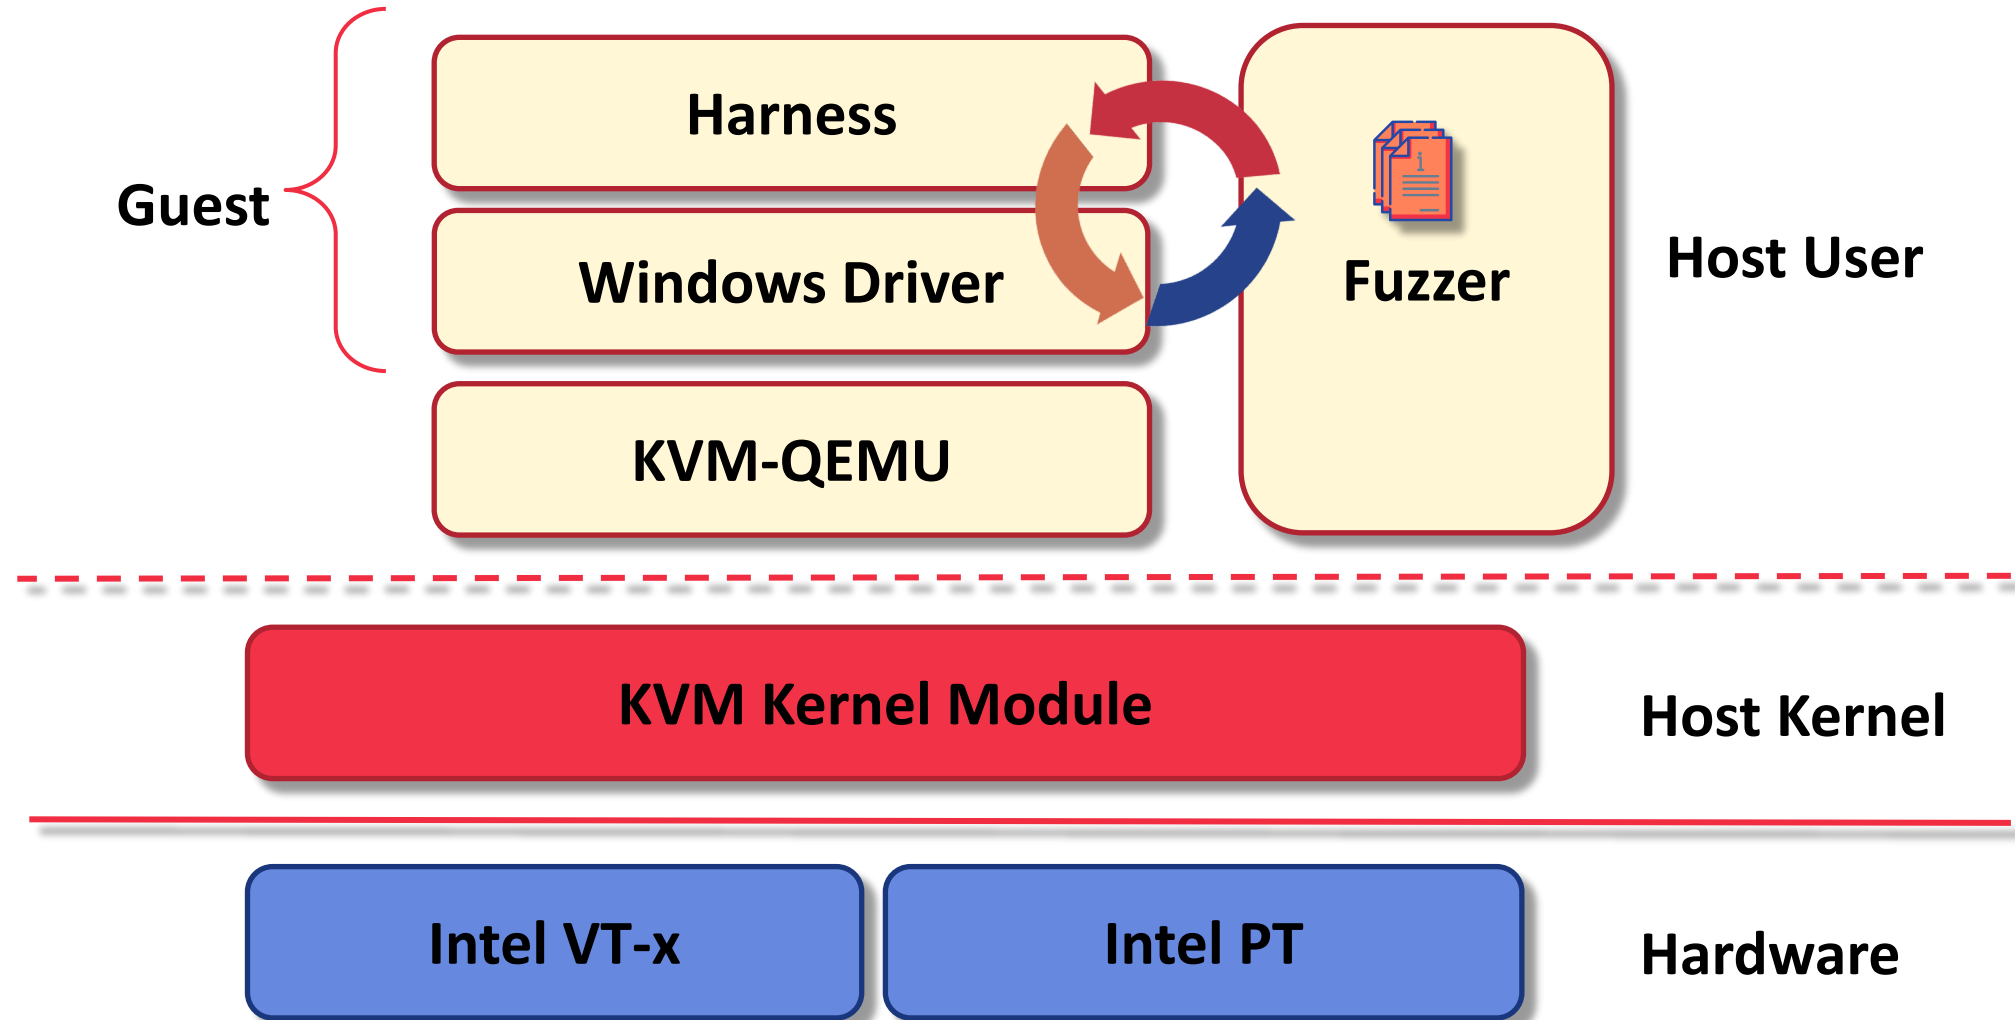
- The feedback is implemented by hooking each code block with a special report snippet
- Upon receiving a cue for entering a new block, the fuzzer updates the bitmap
- Instrumentation can be done during compile time, binary rewrite or dynamically

# kAFL Feedback


- kAFL doesn't modify the target, instead it maps Intel PT to the coverage bitmap
- Intel Processor Tracing is essentially a real-time CPU instructions tracing mechanism
- Intel PT can trace a specific driver at an address range



# kAFL Virtualization Infra



# Why Virtualize?



## Crash Protection

Crashing the target guest won't  
crash the Fuzzer on the host



## Closed Source

VT-x + Intel PT enable efficient  
target agnostic fuzzing

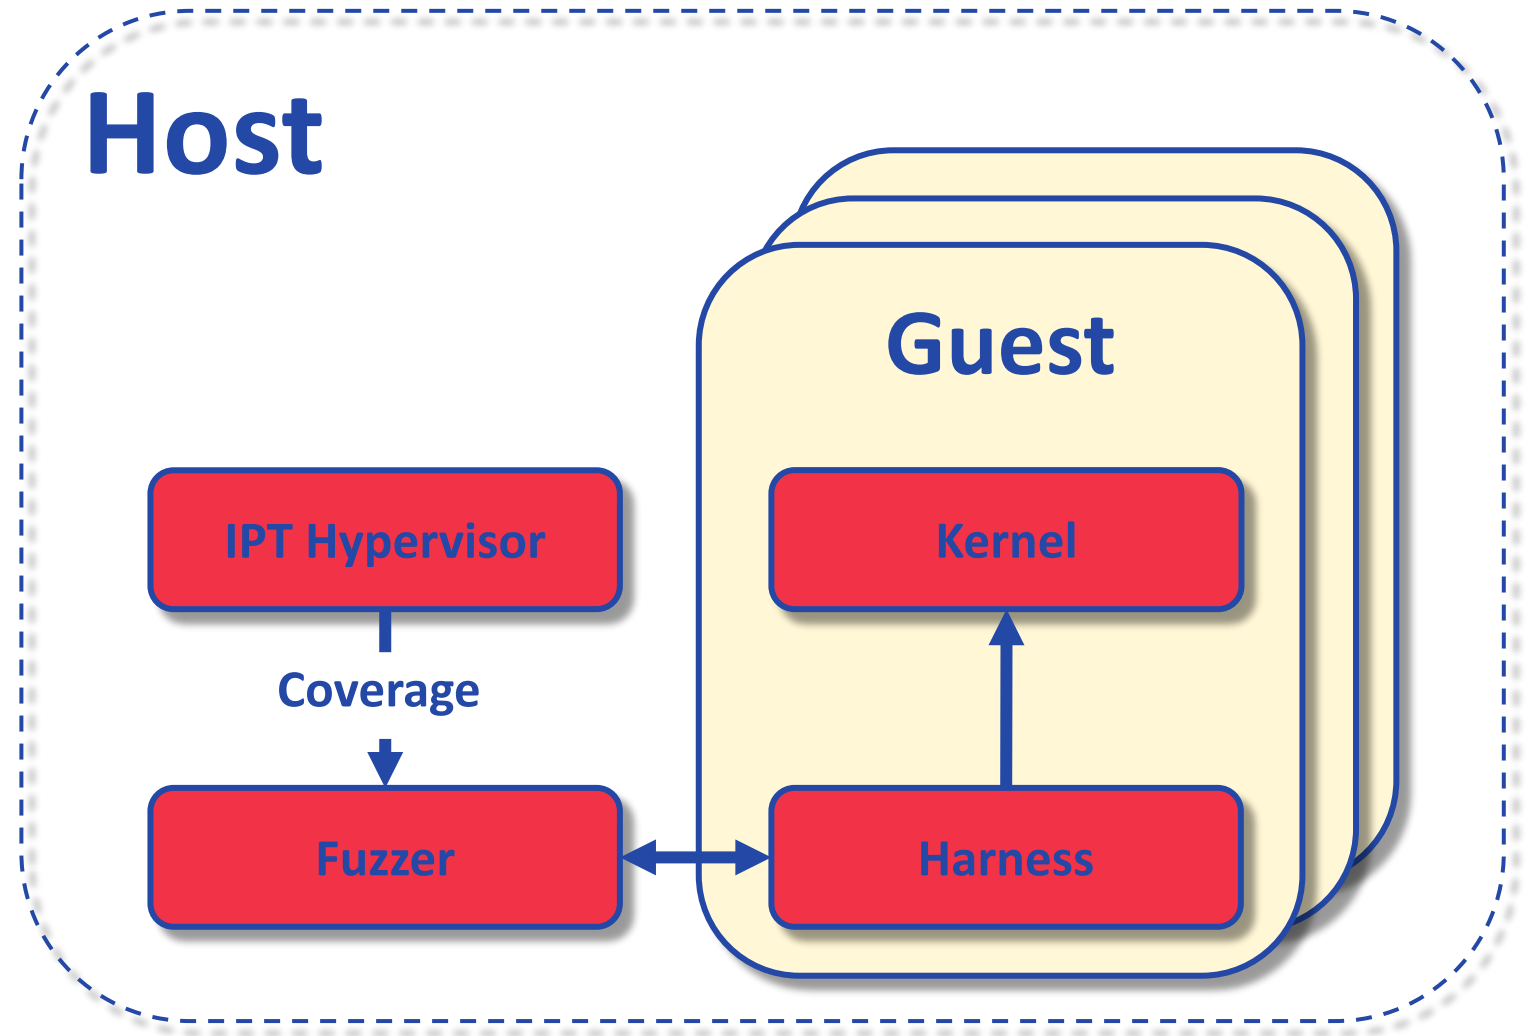
# Host ↔ Guest

## Host to Guest

- Send inputs
- Overwrite panic handler

## Guest to Host

- Request and get next payload
- Disclose CR3 value
- Disclose panic handler address



# Setup Story

- Fuzzer executes from a **VT-x and KVM enabled host** (Latest Ubuntu in our case)
- Fuzzing **target is virtualized** on the Guest (Latest Windows in our case)
- Inputs are requested by the guest via **a Hypercall API**
- Host passes the input via a **shared memory buffer**
- Guest **harness** runs inputs vs kernel
- Host collects **Intel PT traces** and they are converted to a coverage bitmap

# Harness: Initialization



```
#define IOCTL_KAFL_INPUT    0x00010000
```



```
int main(int argc, char** argv)
```

```
{
```

```
    kAFL_payload* payload_buffer = (kAFL_payload*)VirtualAlloc(0, PAYLOAD_SIZE, MEM_COMMIT, PAGE_READWRITE);  
    memset(payload_buffer, 0xff, PAYLOAD_SIZE);
```

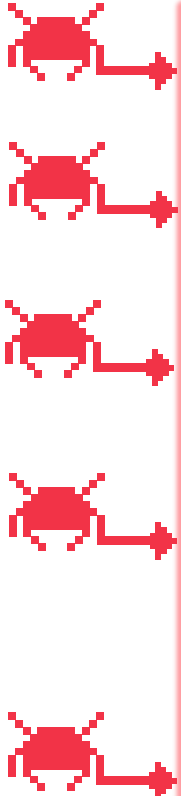
```
    /* open vulnerable driver */  
    HANDLE kafl_vuln_handle = NULL;  
    BOOL status = -1;
```



```
    kafl_vuln_handle = CreateFile((LPCSTR)"\\.\GLOBALROOT\Device\XXXXXXXX",  
        GENERIC_READ | GENERIC_WRITE,  
        FILE_SHARE_READ | FILE_SHARE_WRITE,  
        NULL,  
        OPEN_EXISTING,  
        FILE_ATTRIBUTE_NORMAL | FILE_FLAG_OVERLAPPED,  
        NULL  
    );
```



# Harness: Fuzzing Loop



```
/* this hypercall submits the current CR3 value */
kAFL_hypercall(HYPERCALL_KAFL_SUBMIT_CR3, 0);

/* submit the guest virtual address of the payload buffer */
kAFL_hypercall(HYPERCALL_KAFL_GET_PAYLOAD, (UINT64)payload_buffer);

while(1){
    kAFL_hypercall(HYPERCALL_KAFL_NEXT_PAYLOAD, 0);
    /* request new payload (*blocking*) */
    kAFL_hypercall(HYPERCALL_KAFL_ACQUIRE, 0);

    /* validate input (simple grammar) */
    validate_input(payload_buffer->data, payload_buffer->size)

    /* kernel fuzzing */
    DeviceIoControl(kafl_vuln_handle,
        IOCTL_KAFL_INPUT,
        (LPVOID)(payload_buffer->data),
        (DWORD)payload_buffer->size,
        NULL,
        0,
        NULL,
        NULL
    );

    /* inform fuzzer about finished fuzzing iteration */
    kAFL_hypercall(HYPERCALL_KAFL_RELEASE, 0);
}
```

# Demo

kAFL in action = Vulnerability

# What's Next?



## OTHER OSs



## GRAMMAR

Fine tune drivers' specific grammar requirements



## HYPERVISORS

Attack surfaces with more privilege levels fragmentation

# Takeaways


- What Developers can do:
  - The first step in exploiting drivers is to access them
  - If you restrict that access then the attack surface is nullified
  - You can also fuzz your own drivers as a part of the QA process
- What Attackers can do:
  - Use our automation scripts
  - Start fuzzing the kernel
  - Examine KMDF and NDIS drivers for more bugs

# Takeaways

- What Defenders can do:
  - Naïve: Make sure every driver is up-to-date
  - Advanced:
    - Monitor *DeviceIoControl* requests to devices
    - Block / Alert in case the requesting executable is unsigned

# Credits

- Google 😊
- Ilja van Sprundel
- James Forshaw
- kAFL
- Syzkaller
- Alex Ionescu
- HackSys Team



# Thank You For Listening !

Do you have any questions?

[mark.cherp@gmail.com](mailto:mark.cherp@gmail.com) | @OcamRazr

[eranhelforc@gmail.com](mailto:eranhelforc@gmail.com) | @EranShimony

