

# Machine Learning on Spark

He Yunlong

SSG/DRD/PRC

# Agenda

Machine Learning in Today

Introduction to Spark

Machine Learning on Spark

Large Scale Neural Network

Optimization


Q & A

# Machine Learning in Today


**Old Time:** Before Mobile Internet Booming Era: Data Scale is small :

- Shallow analysis is enough, no need complex algorithm
- Limited data set could not get precise training model.
- Computing Capability is limited by technology


**Today : Big Data laid a foundation of Machine Learning:**




The widespread of smart phones and the development of IoT provides comprehensive data sources for **Big data**.



In the past couple of decades, **Computing Power** is growing exponentially by following the Moore's Law



With the rapid development of science and technology, more **Complex Models** are extracted, built-up and deployed in industry



More and more **Efficient Algorithms** of Machine Learning are researched and developed by scientists and domain experts

**Data are the Greatest Strategic Resources for Internet Companies**

ML: Big data + Computing Power + Complex Model + Efficient Algorithm

- **Create user experiences**
- **Create commercial values**

# Machine learning Phase - 1

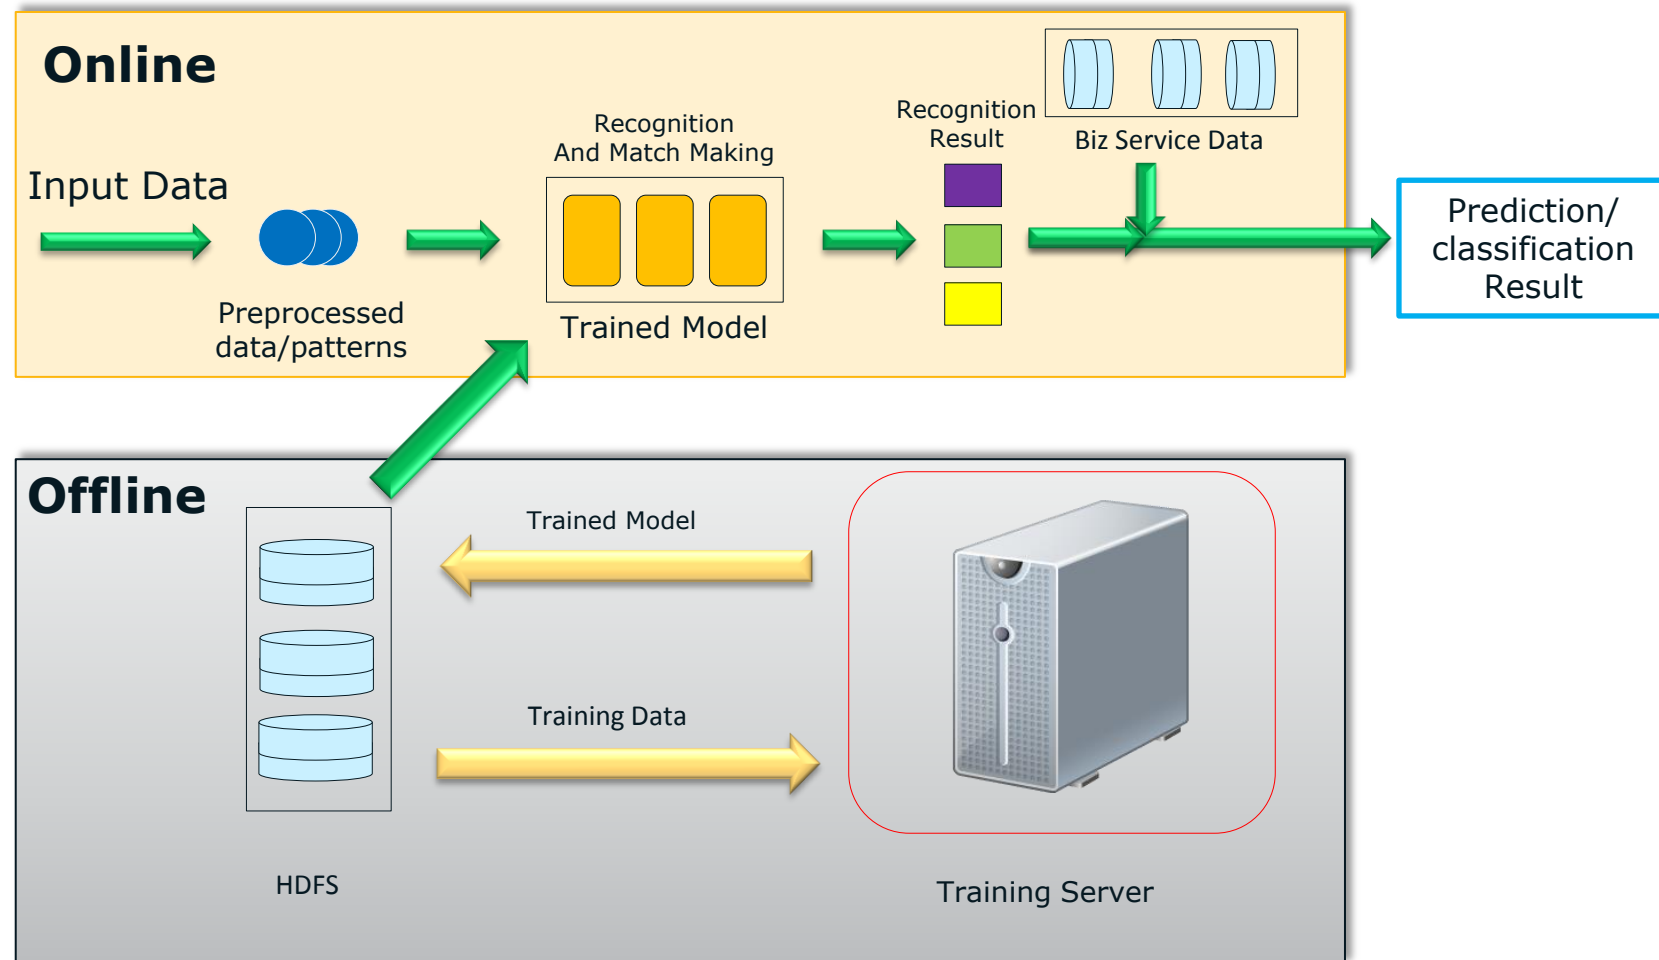
**Model** small

**Dataset** small

**Training time** minutes ~ hours

**Tools** Matlab, R, Python  
...

**Services** Junk Detect, Association



# Machine learning Phase - 2

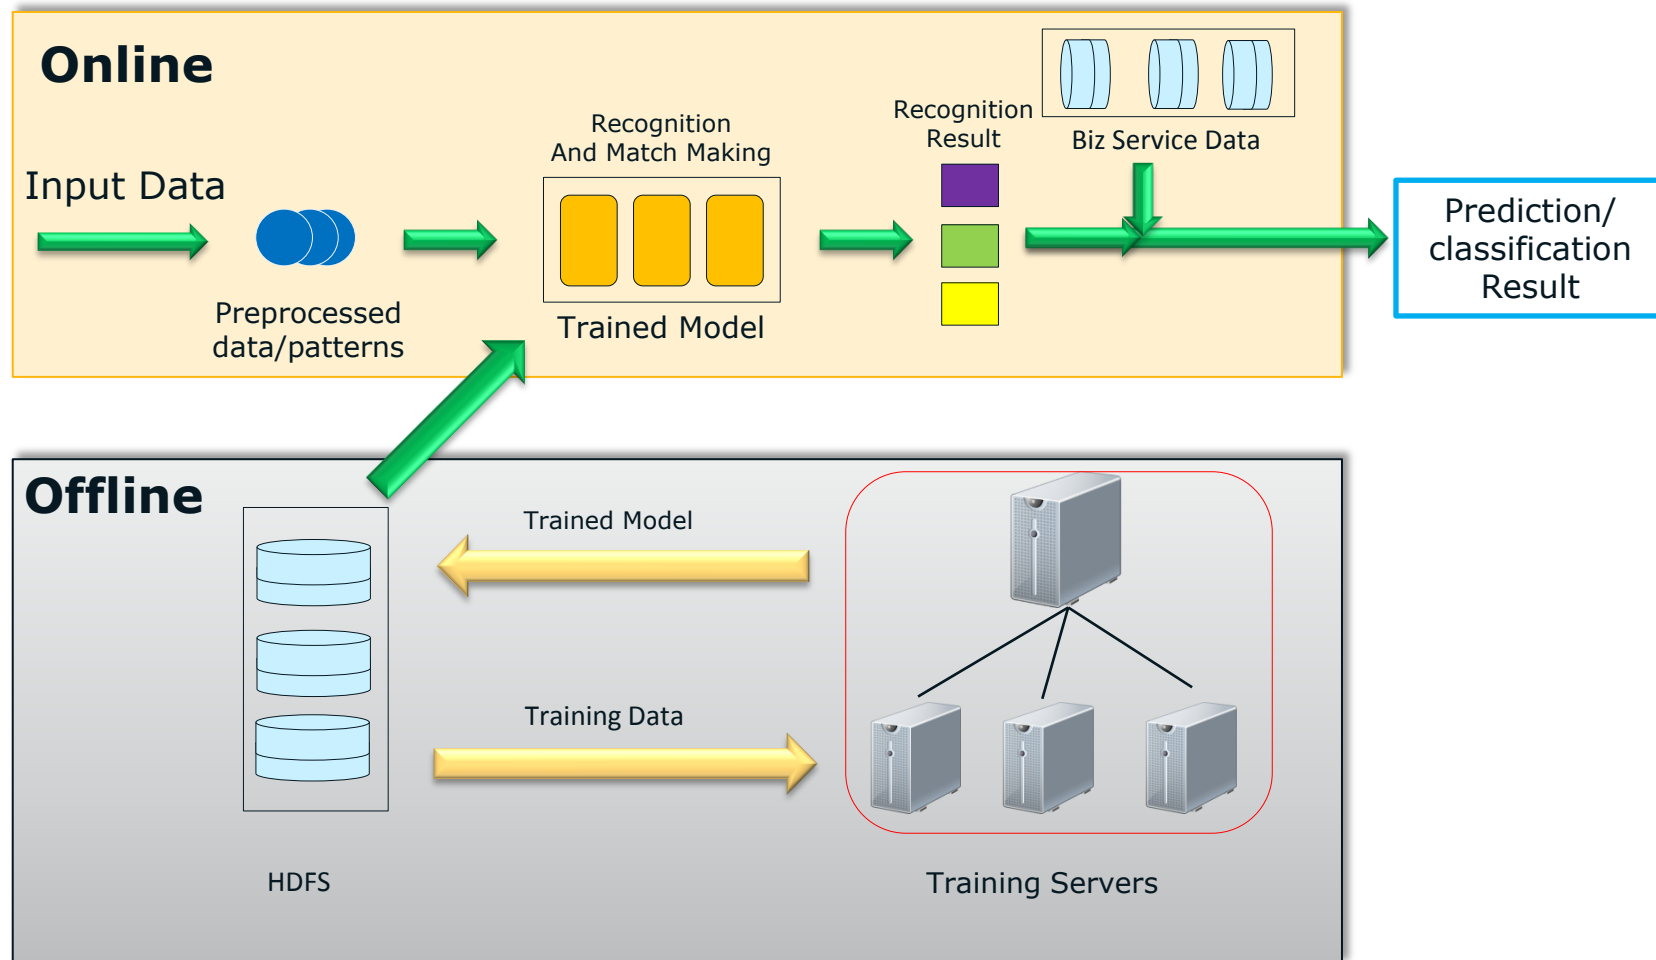
**Model** Small

**Dataset** big

**Training time** minutes ~ days

**Tools** Mahout,  
Mllib,  
...

**Services** CTR,  
Doc Classify



# Machine learning Phase - 3

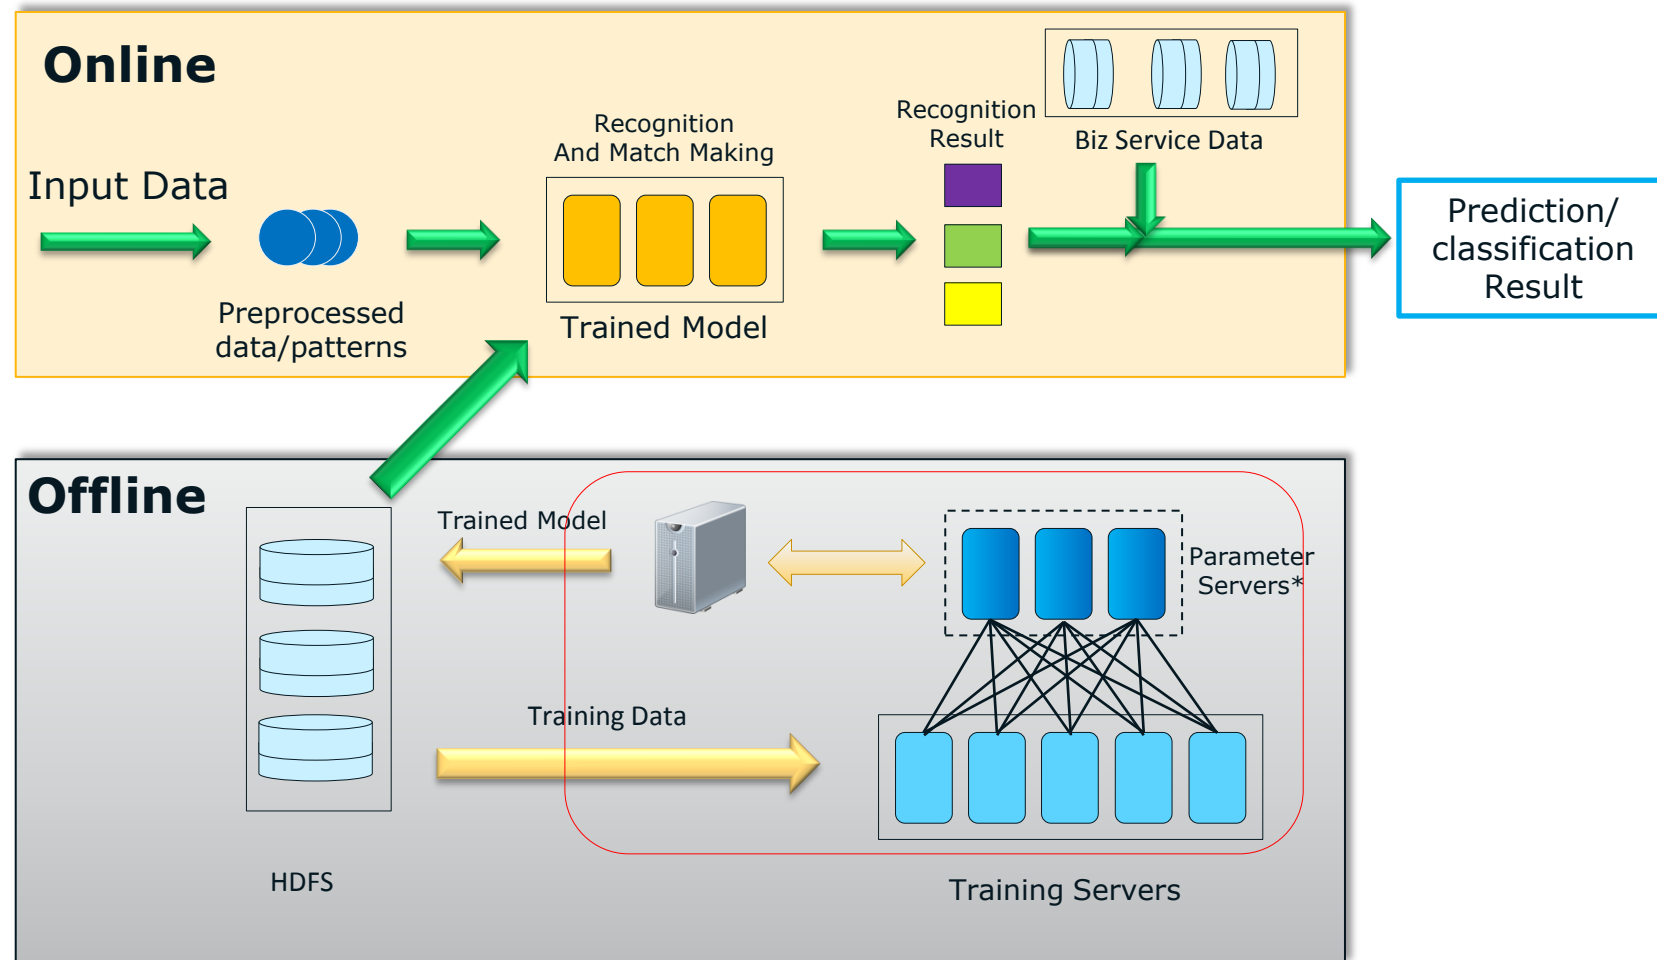
**Model** big

**Dataset** huge

**Training time** minutes ~ days

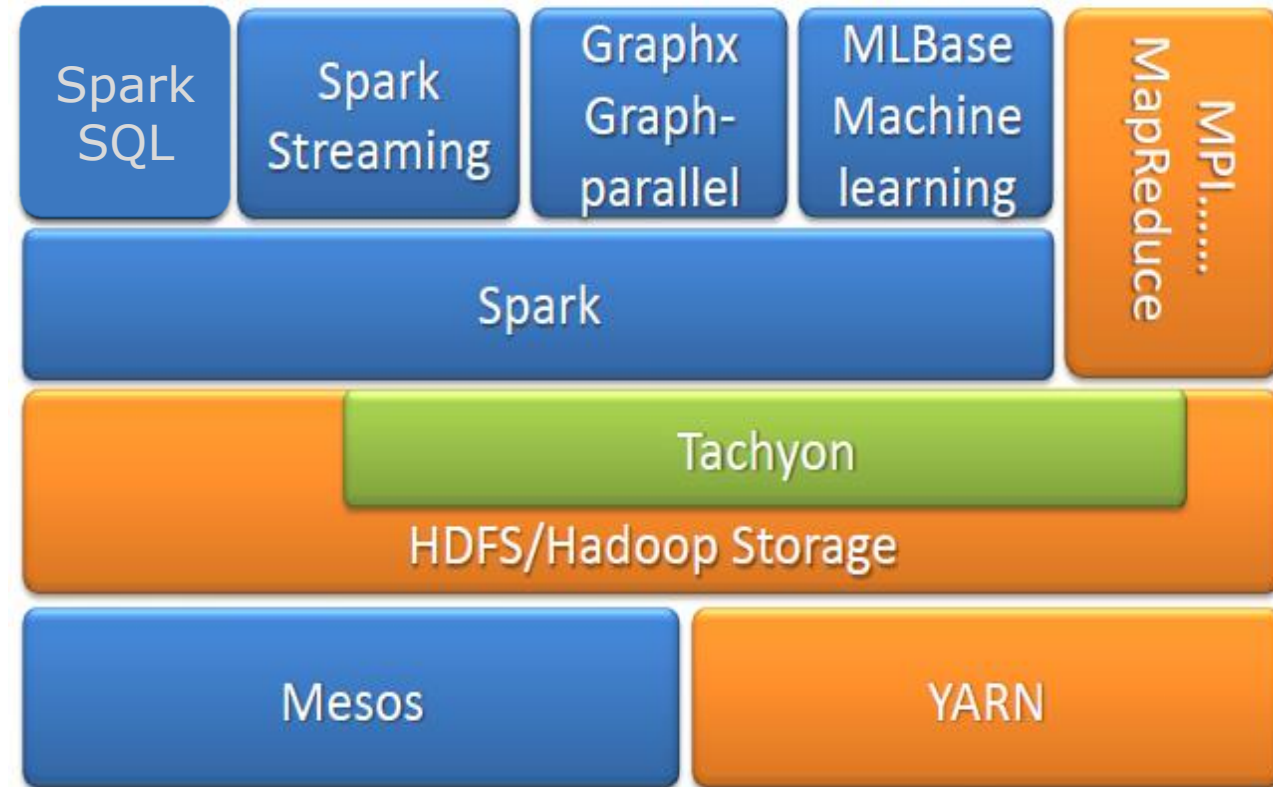
**Tools** ParamServer,  
DistBelief  
...

**Services** Speech Recog,  
Image Search




# About Spark

- Fast In-Memory data analytics cluster computing framework
- Originally developed in the AMPLab, became an Apache Top-Level Project in February 2014
- Suitable for Iterative tasks
- Proven scalability to 2000 nodes in the research lab on EC2 and 1000 nodes in production.



# Spark - Program Model : RDD

- General task graphs
- Automatically pipelines functions
- Data locality aware
- Partitioning aware to avoid shuffles





# Spark - Program Model : RDD

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(_.startsWith("ERROR"))  
messages = errors.map(_.split('\t')(2))  
cachedMsgs = messages.cache()
```


Base RDD

Transformed RDD

Cached RDD

```
cachedMsgs.filter(_.contains("foo")).count  
cachedMsgs.filter(_.contains("bar")).count  
...
```

Parallel operation



# MLlib: Machine Learning on Spark

- **Classification**

- logistic regression, linear support vector machine(SVM), naive Bayes, classification tree

- **Regression**

- generalized linear models (GLMs), regression tree

- **Collaborative filtering**


- alternating least squares (ALS)

- **Clustering**

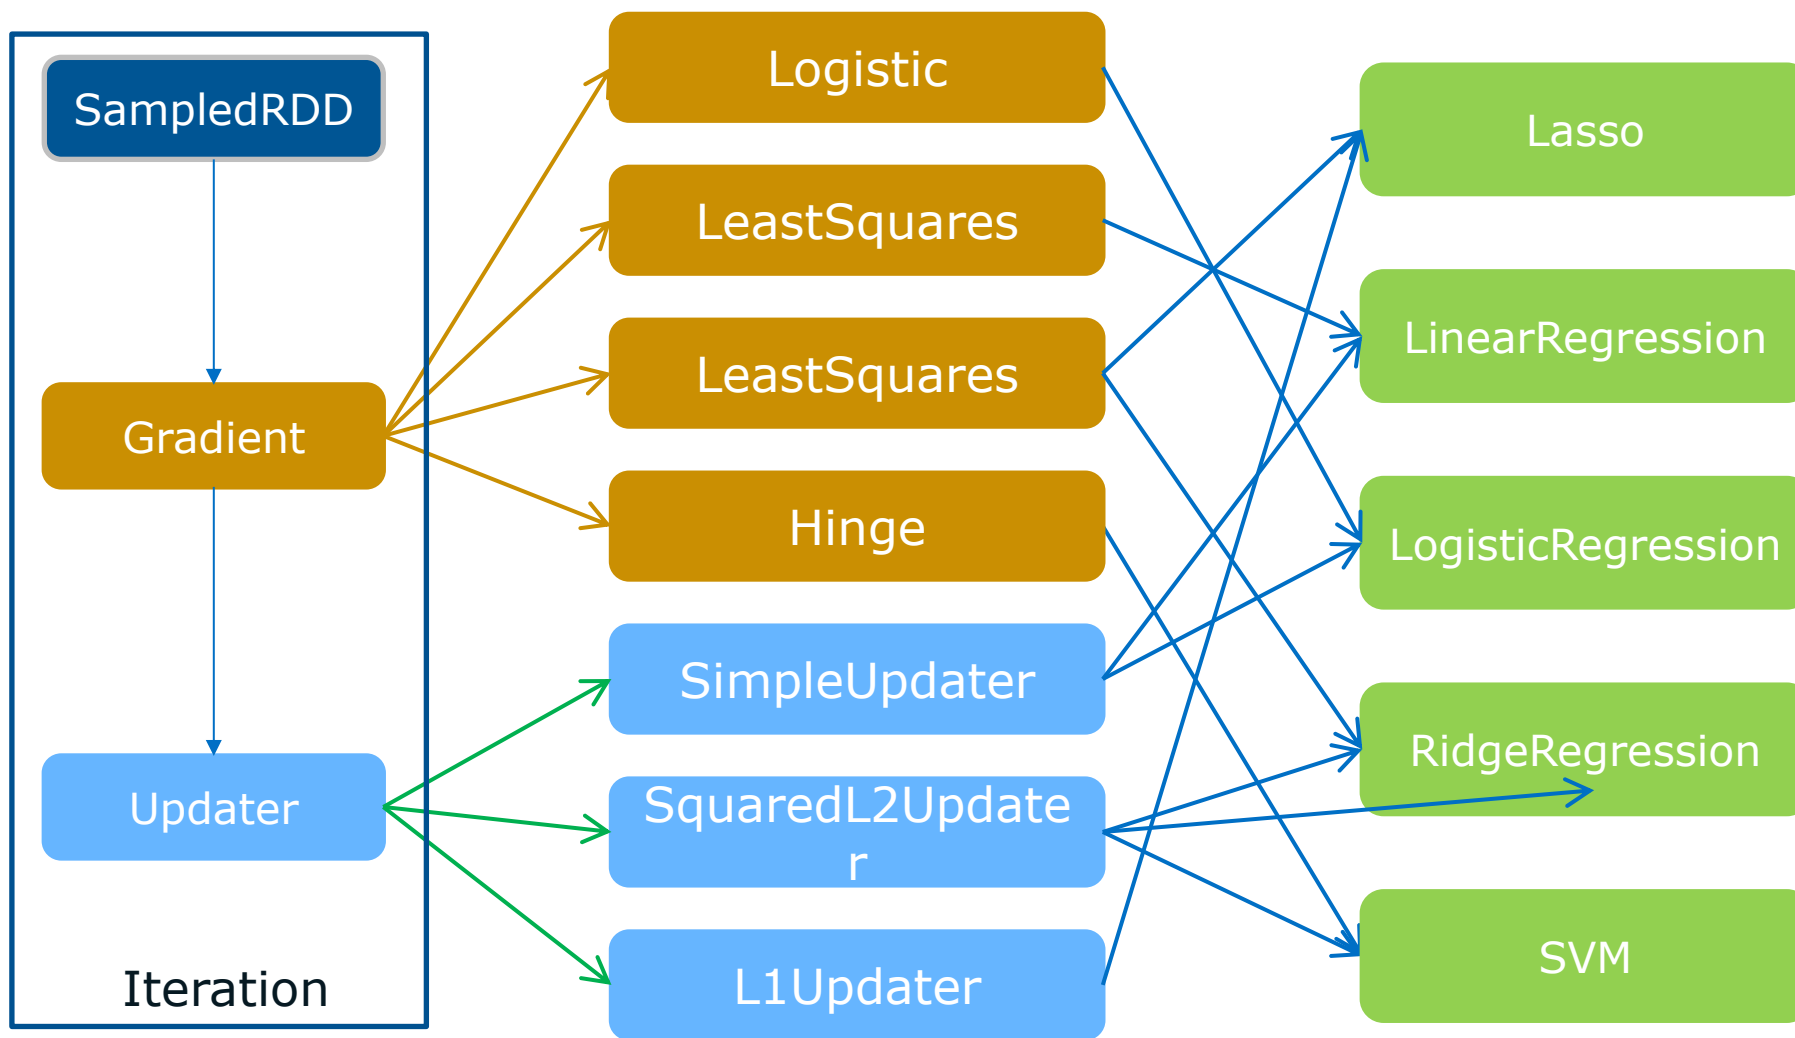
- k-means

- **Decomposition**


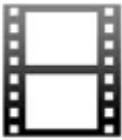







- singular value decomposition (SVD), principal component analysis (PCA)



# Mlib – Regression



# Mlib – Collaboration Filter

			
	★	★★★★	?
	★	★★★	★★
	★★★★	?	★
	★	?	★★
	?	★★★	★★
	★★★★	★★	?

```
// Load and parse the data
```

```
val data = sc.textFile("mllib/data/als/test.data")  
val ratings = data.map(_.split(',') match {  
  case Array(user, item, rate) =>  
    Rating(user.toInt, item.toInt, rate.toDouble)  
})
```


```
// Build the recommendation model using ALS
```

```
val numIterations = 20  
val model = ALS.train(ratings, 1, 20, 0.01)
```

```
// Evaluate the model on rating data
```

```
val usersProducts = ratings.map { case Rating(user, product,  
  rate) =>  
  (user, product)  
}  
val predictions = model.predict(usersProducts)
```

# Mlib - Dimension reduction+ k-means




```
// compute principal components
val points: RDD[Vector] = ...
val mat = RowMatrix(points)
val pc = mat.computePrincipalComponents(20)

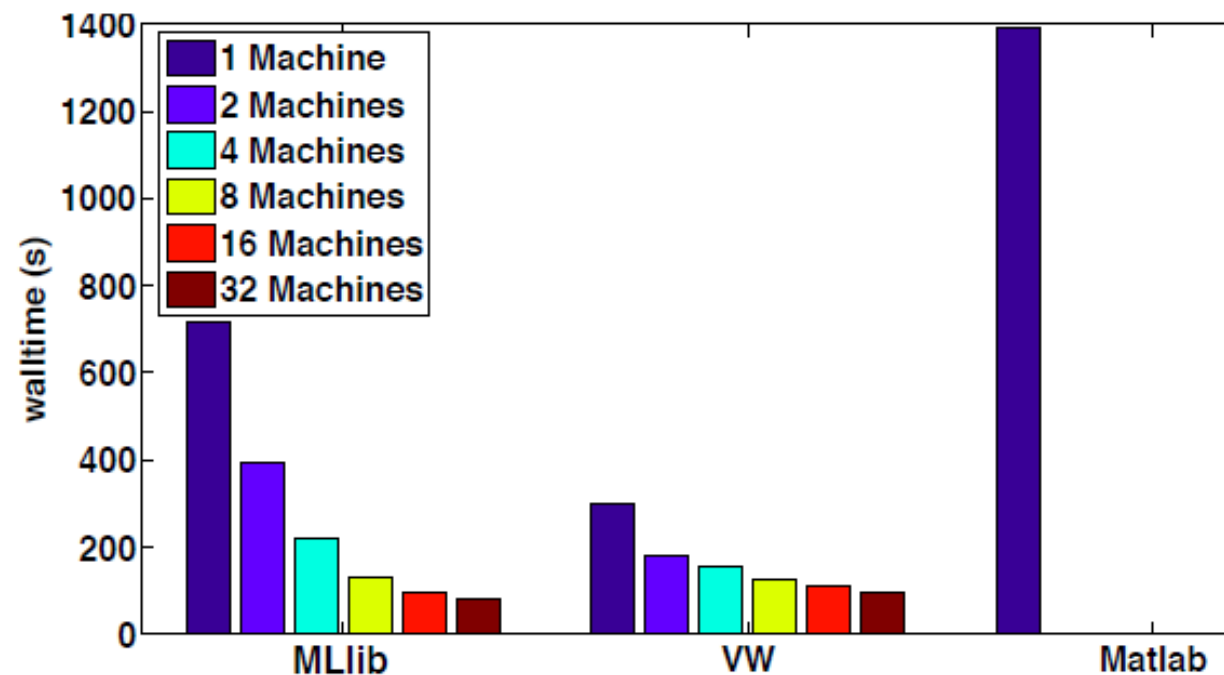
// project points to a low-dimensional space
val projected = mat.multiply(pc).rows

// train a k-means model on the projected data
val model = KMeans.train(projected, 10)
```

# MLlib – Performance



Running times for iterations after the first in Hadoop, HadoopBinMem, and Spark. The jobs all processed 100 GB.



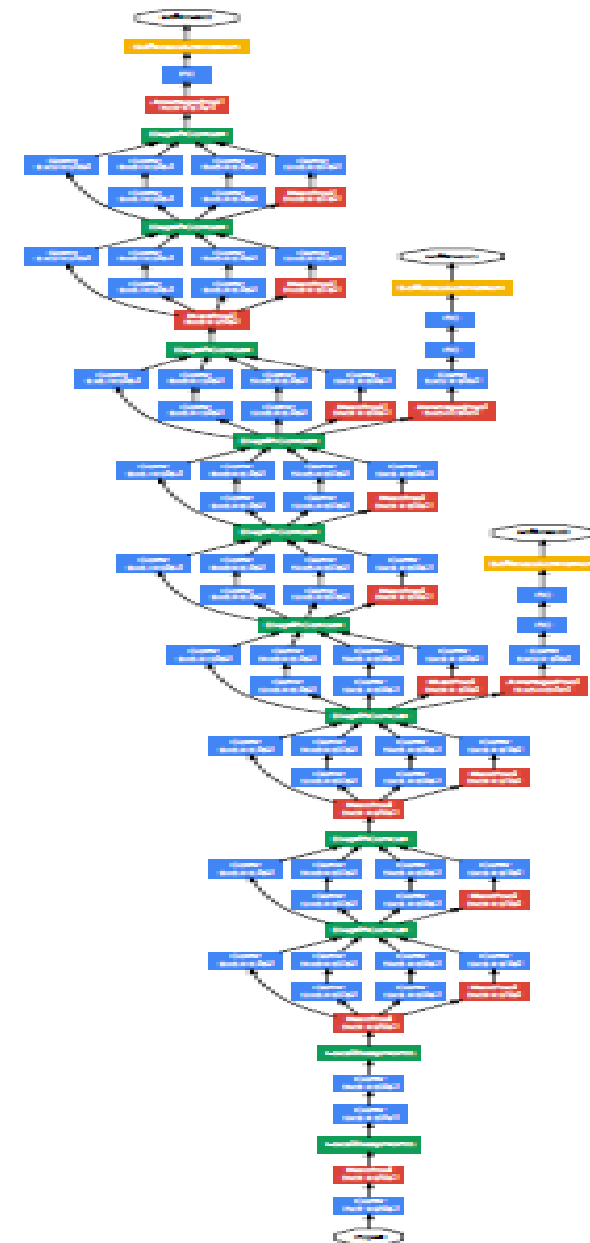
- Fixed Dataset: 50K images, 160K dense features.
- MLlib exhibits better scaling properties.
- MLlib is faster than VW with 16 and 32 machines.

# Mlib 1.1?

- **Model selection!**
  - training multiple models in parallel
  - separating problem/algorithm/parameters/model
- **Learning algorithms!**
  - Latent Dirichlet allocation (LDA)
  - Random Forests
  - Online updates with Spark Streaming
- **Optimization algorithms!**
  - Alternating direction method of multipliers (ADMM)
  - Accelerated gradient descent
- **Neural Network?**


# New Challenge: Deep Learning

- Challenges
  - Very big model
  - Huge training data
- Spark Limitation
  - Only local model supported
  - RDD is read-only, expensive for neural network parameters
  - Broadcast is not feasible for big model
- We need
  - Distributed model
  - High-performance training process





# Distributed Neural Network on Spark



- Distributed parameters
- Configurable server/worker nodes
- Multiple training workers
- Parallel fetching/training/pushing in each worker
- Adaptive learning rate


# Example - Word2Vec

This tool provides an efficient implementation of the continuous bag-of-words and skip-gram architectures for computing vector representations of words. These representations can be subsequently used in many natural language processing applications and for further research

- dog => [0.792 -0.177 0.98 -0.9 ....]
- cat => [0.76 0.12 -0.54 0.9 0.65 ....]

In some cases, word2vec can be used to modelling non-wording services, which makes its model very large


Distributed parameter servers helps to scaling the model size linearly



Word	Cosine distance
France,	0.729900
Italy	0.720465
MORZINE,	0.681200
Germany	0.680331
Spain	0.673912
Russia	0.666366
Poland	0.652955
Spain,	0.648663
France.	0.646427
Germany,	0.642493

# Test Result

- Extensibility
  - Linear extendable model size
  - Huge dataset supported
- Accuracy
  - Tradeoff between accuracy and performance, small batch size can raise accuracy, but hurt performance
  - Adaptive learning rate help to raise accuracy, but enlarge parameters too.
- Performance
  - Network is bottleneck, 10GbE is preferred
  - Multi-worker can obviously speed-up training
  - Optimization with MKL can speed up by nearly 50%



# Q & A